

# Lecture 1

---

Mon. Aug. 26, 2019

## Classic Book (to learn C)

- kernighan & Ritchie, The “K&R” of C

## C is still evolving

- Different standard version
  - K&R C
  - ANSI C
    - C89, C90 (ISO 9899:1990)
  - C99 1999
  - C11 2011
  - C18 2018

## C should be...

- Simple
- easy to compile
- typed [weakly]
- support low-level memory access
- ideal for embedded controller, OS (operating system), ...

## Yet...

- C is powerful
- C is fast

## C is a purely procedural language

- No object-orientation whatsoever

## Central Dogma

- Object of interest: Computations
- Main abstraction tool: Procedures/ functions
  - Caller / Callee
  - Abstracts away “How things are done”
- Programming means
  - Organizing processes as procedures
  - composing processes thru procedure calls

## Procedural Programming in C

- Adheres to the philosophy
- Generates very efficient code
- Exposes as many low-level details you wish to see
- Provides full control over memory management (no Garbage Collection)
- The Programmer is fully in charge

## Resources

### What do you need

- C compiler
- linker
- text editor

### We will use the GNU toolchain

- Compiler: GCC (or CC)
- Linker: ld (invoked by GCC)
- Editor: vim, nano, emacs

## Workflow

- Use text editor to edit source files
- Use compiler to generate “.o” files
- Use linker to link multiple ‘.o’ files into executables

### Text Editor



```
#include <stdio.h>

/* comments */
// single linek comments

int main(void)
{
    printf("Hello, world! \n");
    return 0;
}
```

## Comments

**Compiler ignores everything between “/\* \*/”**

## The ‘main’ function

A special function that defines the entry point for the program

- This is where the OS transfers control once the program starts

## printf

- C library function to print on the standard output for the process
- takes at least a string as argument
  - Between double quotations like “**This is a string**”
- ‘\n’ is a new line character

## Including header files

### #include

- imports a header files with the specification of functions that exists in a library to be linked w/ the program

### Compile

#### What is CC really doing?

- 3 steps
  - preprocesses hello.c
  - compile hello.c to hello.o
  - links hello.o /w libc
  - writes executable file a.out
- **You can and often will separate the 3 steps!**
  - No necessarily that often unless working with big files
  -

## A second program

### Purpose

- Read an integer from the standard input: n
- Compute the sum of all integers between 1 & n
- Print the result on the standard output

### New concepts

- Standard input and output

## Summary

### C program consists of functions

- main() is the entrance of a program

### function consists of a sequence of statements

### Variables can be declared in a function (like main)

- These are local variables
- Variables must be declared

### Statements are terminated with ‘;’

- Empty statements are allowed
  - empty statements

The ‘main function’ taking arguments

**‘main’ function can take 2 arguments**

- argc: the number of arguments
- argv: an array of arguments

If-else statements

## Lecture 3

---

Mon. Sept. 4, 2019

### Operators

- Conventional arithmetic, bitwise, and logical operators
  - + - \* / %
  - & | ~ ^ << >>
  - && || !
- Pre/post increment/decrement (as in Java, C++ etc.)
  - i++ ++i j++ --j
  - c = i++; // c will be (i - 1)
  - c = ++i; // c will be the same as "i"
- Simple & Compound assignment operators
  - more to come!

### Precedence & Associativity

- Precedence determines which operation is done first
  - If operators have the same precedence, use associativity
  - use parentheses

i + j \* 10 - k / 20

(i + (j \* 10)) - (k / 20)

### Assignment Operators

- Assignment Operator
  - LHS (Left Hand Side) is something that can be written to (e.g. to a variable)  
 $LHS = Expression$
  - LHS and Expression have “compatible” types

## CSE 3100 Master Notes

- The value of Expression is assigned to LHGS & becomes the value of the assignment operation
- Compound Assignment operators ( $+=$ ,  $*=$ , ...)
  - $\text{var op} = \text{expr} \Leftrightarrow \text{var} = \text{var op expr}$
- ex.):
  - $a = x + y; \quad b = c = d = 0;$
  - $i += 10; \quad // \quad i = i + 10$
  - $j *= 5; \quad // \quad j = j * 5$

Assignments **ARE NOT** statements

- assignments are **expressions** and “ $=$ ” is the operator
  - you can chain them!
  - you can use them inside larger expressions

Integer Data Types

- char
- short int  $\rightarrow$  short
- int
- long int  $\rightarrow$  long
- long long int  $\rightarrow$  long long

\*\*\*C code can be confusing, people flex their skills in C, but it can be very unreadable  
“Don’t be that guy flexing skills, just because you can do it in C” - Wei We

- Consider x86\_64 (64-bit architecture)

size (in bits)	signed	unsigned
8	char -128 ... 127	unsigned char 0 ... 255
16	short -32786 ... 32767	unsigned short
32	int	unsigned int
32	long - $2^{31}$ .. $2^{32} - 1$	unsigned long
64	long long	unsigned long long

- Consider x86\_64 (64-bit architecture)

size (in bits)	signed	unsigned
8	char -128 ... 127	unsigned char 0 ... 255

## CSE 3100 Master Notes

16	short -32786 ... 32767	unsigned short
32	int	unsigned int
32	long - $2^{31}$ .. $2^{32} - 1$	unsigned long
64	long long	unsigned long long

### How much space?

- How to determine the amount of space for some type?
- Operator size of gives the # of bytes needed for a type of variable
  - You will need this later to dynamically allocate Space

### Character (char) Data Type

- char has 8 bits (a byte)
- ASCII Code (American Standard Code for Information Interchange)
  - Characters are mapped to an integer in 0 ... 127
  - An ASCII character can be stored in char
- Classes in ASCII
  - 0 ... 31: “Control” character (a.k.a non printable)
  - 48 ... 57: Digits
  - 64... 90: Upper case letters
  - 97... 122: Lower case letters

### So...

- The character “h” is none other than ... 72
  - `char h1 = "H", h2 = 72; // h1 & h2 have the same value`
- Observe how...
  - ‘0’ through ‘9’ are consecutive!
  - ‘A’ through ‘Z’ are consecutive!
  - ‘a’ through ‘z’ are consecutive!

`char ch = '8';`  
`int x = ch - '0' // what is the value of x?`

### Non-Printable Characters?

- These are sometimes useful
  - Showing the constant (literal)

'\n'	newline
------	---------

## CSE 3100 Master Notes

'\r'	carriage-return
'\f'	form-feed
'\t'	tabulation
'\b'	backspace
'\x7'	audible bell (x indicates hexadecimal)
'\07'	audible bell (0 indicates octal)

### Basic Data Types: Floating Point

- A few floating point types
  - Consider x86\_64 again

size (in bits)		

### Automatic Type Conversion

- When an operator has operands of different types, the operands are **automatically** converted to a common type by the compiler
  - In general, a lower rank operand is converted into the type of the higher rank one, where
    - *char < short < Int < long < long long < float < double < long double*
    - **long double is of highest rank**
  - E.g. “1” gets converted to double before performing the addition in the expression “1 + 2.5”
- Automatic conversion can also occur across assignments
  - The value of the expression on the right hand side may be widened to the type of the LHS, e.g., “double d = 1”
  - Or **narrowed** (possibly with information loss), e.g., “int i = 2.5”;
    - Read book for more details!

### Type Casting: Explicit Type Conversion

- Useful to convert an operand to another type before doing arithmetic
  - (*<Type>*)*<expression>*

## CSE 3100 Master Notes

- Ex.): integer or double?

<pre>int x = 10; int y = 3;  double z = x / y</pre>	
-----------------------------------------------------	--

What about Booleans?

- K&R and C89/C90 do not have a Boolean data type
  - 0 “means” FALSE and anything else “means” TRUE
  - Common to use int or char to store Boolean values and define convenience macros
- C99 introduced Boolean

**Be Mindful...**

- Sometimes the results may not be as expected!
  - What is the size (in bits) of each operand?
  - Are your operands signed or unsigned?
- Ex.):

<pre>unsigned int x = 3; unsigned int y = 7; unsigned int z = x - y;</pre>	z holds the binary representation of -4 but reading it as an unsigned int yields a very different value!
<pre>_Bool b1; char b2, b3; int i = 256; // 0x100  b1 = i; b2 = i; b3 = i != 0;</pre>	b1 is 1 because i is not 0 b2 is - because the lowest 8 bits in “i” are 0 b3 is 1 because i is not 0  Do you want b2 or b3?

## Lecture 4

---

Mon. Sept. 9, 2019

### Flow of Control

- Statements are normally executed sequentially
- For selective or repeated execution we have all the usual suspects from Java / C++ / Python
  - blocks
  - if & if-else
  - while

## CSE 3100 Master Notes

- for
- switch
- break
- continue

### Blocks (compound statements)

- List of statements enclosed by { and }
  - considered as a single statement
  - No semicolon after closing }
  - can be empty
  - can be nested (block in block)
  - useful for branching/ loop statements
  - can define variables at beginning of blocks
    - can mix declarations and code in c99

### Comparison & Logical Operators

- Comparison operators that compare two expressions
  - Pay attention to types  
== != > < >= <=
- Logical Operators  
    &&   ||   !
- The result is either 0 or 1 (of int type)
  - 0 = false, 1 = true

### Branching: if & else

- “exp” is typically a comparison or logical expression, but can be ANY expression (float,double, pointer, ...)
- The statements can be compound statements (blocks)
- Or other if statements!
  - Beware of the dangling else (“else” matches the nearest preceding “if”, use blocks to disambiguate)

\*\*\*Lot of things are “passable” in C and can compile and run fine, but may not be readable whatsoever by other programmers. Beware of this when using branching - Wei Wei

**Ex.):**

```
int i, j, min;
if (i < j)
    min = i;
```

### Ternary Operator

- Takes 3 expression as operands

exp ? exp2 : exp3

- exp1 is evaluated first
- If exp1 is non-zero (true), exp2 is evaluated and its value is used as the value of the ternary expression
- If exp1 is zero (false), exp3 is evaluated and its value is used as the value of the ternary expressions

- ex.):

min = i < j ? i : j

### While Loop

- Very similar to python,
- ex.): computing sum of 0 .. 99

```
int i = 0, sum = 0
while (i < 100){
    sum = sum + i
    i++;
}
// Same as
while (i < 100) sum += i++;
```

### Do-While Loop

- Checks condition after executing loop body
  - the statement is executed at least once
- ex.): computing sum of 0 .. 99

```
int i = 0, sum - 0;
do{
    sum = sum + 1;
    i++;
} while (i < 100);
```

### For Loop

- Sometimes called “counting” loop
  - more like swiss-army knife!
- Three expression:
  - initialization, condition, increment
- Equivalent to

## CSE 3100 Master Notes

```
exp1;  
while (exp2){  
    <stmt>;  
    exp3;  
}
```

- 4 ways to use for-loops to compute the sum 0..99

```
sum = 0;  
for (i = 0; i < 100; i++) sum = sum +i;
```

### 2nd way - w/ all initializations inside

```
for (sum = i = 0; i < 100; i++) sum += i;
```

### 3rd Way - Empty Body

```
for (sum = i = 0; i < 100; sum += i++);
```

### 4th Way - Comma Operator

```
for (sum = 0, i = 0; i < 100; sum += i, i++);
```

## Comma Operator

- Takes 2 expressions

*exp1, exp2*

- *exp1* is evaluated first, then *exp2* is evaluated
- *exp2* is the result of the whole expression

- Has the lowest precedence
- Associated from left to right

*exp1, exp2, exp3*     $\leftrightarrow$     *(exp1, exp2), exp3*

- Order can make a difference, e.g.,  
*for (sum = 0, i = 0; i < 100; sum += i, i++);*

- is not the same as

*for (sum = 0, i = 0; i < 100; i++, sum += i);*

## Multiway branching using “else if”

- if statement can contain multiple else statements

## Switch example

```
// Assume all variables are defined as int
```

```
switch(i){  
    case 0;  
        n0 ++; break; // Note the break statement  
    case 1;  
    case 2;
```

### Break Statement

- Most commonly used in switch statements
  - Prevents “fall-through” to the next case
- Also works in loops (for, while, do-while)
  - Loop execution terminated immediately, control resumes at statement immediately following the loop

### Continue Statement

- Skip the rest of the current loop iteration and continue to the next one
- Can be used within for, while, and do-while loops
  - Can appear in a nested if / else
  - If used in nested loops, it applies to the “innermost” enclosing loop
  - For “for” loops, go to the evaluation of the “increment” expression

```
{ // begin loop body  
    ...  
continue;  
    ...  
} // end loop body
```

## Lecture 8

---

Weds. Sept. 11, 2019

### Function Definitions

- No nesting (cannot define functions in a function)
- Return type can be “void” (no return value expected)
  - if missing, compiler assumes int
- return statements

```
return;      // terminates execution and return control to caller  
return expr; // terminate and pass value of expr back to caller
```

- Execution also terminates if end of function body reached
  - return value undefined

### Function Declarations (Prototypes)

- Functions can be defined in any order
  - declare a function before first use if definition comes later
  - function prototypes often placed in header files (and reused)

```
#include <stdio.h>
int fahrToCelsius(int);

int fahrToCelsius(int degF) {
    return 5 * (degF - 32) / 9;
}
```

### Example: computing $b^n$

- Power function
  - returns an int
  - 2 parameters: base “b” and exponent “n”, both int
- Functions can declare local variables
  - parameters & local variables can only be accessed inside the function
  - Storage class “auto” by default i.e., discarded when function returns
- Parameters are passed by value
  - “n” is changed in the power function but “i” DOES NOT change

### Static & Global Variables

- Static local variables
  - Not visible outside function but retain value across function calls
- Global variables
  - declared outside functions; retained for entire duration of program
  - Storage class “extern” by default
    - can be accessed from functions in other files
  - Static global variables
    - visible only in functions defined in the same file following variable declaration

### Be Cautious with static and global variables

- “Nice” functions only depend on their inputs
- Static and global variables have “side-effects”
  - retain values across function calls
  - change the meaning of the function at each call
  - You can understand the function without holding all the code in your head
- unless you have really good reasons and really know what you are doing, DO NOT USE STATIC OR GLOBAL VARIABLES

## Function call context

- Function call context includes
  - copies of function arguments (call-by-value\_)
  - auto local variables
  - return address
- Call contexts managed automatically using the **execution stack**
  - a stack frame is created automatically for each function call
  - the frame lasts for the duration of the call
  - discarded automatically when the function terminates
  - NOTHING in the frame survives the call

# Lecture 9

---

Mon. Sept. 16, 2019

## A C Primer (5): Arrays and Pointer Basics

### Arrays

- New Data types
- Arrays represent a linear, contiguous collection of “things”
- Each “thing” in the array has the same fixed type
- ex.):
  - array of characters
  - array of integers
  - array of doubles...

### Array example

```
int x[5]; // define an array of 5 int's
// accessing array elements is similar to accessing list in Python
// the index starts from 0
x[0] = 1; x[1] = 2; x[2] = 3; x[3] = 4; x[4] = x[3] + 1;
// initialize array with a list
int y[5] = {1, 2, 3, 4, 5};
// Number of elements is optional if all elements are listed
int z[] = {1, 2, 3, 4, 5};
// Specify the value of first 2 elements. The rest are set to 0
```

## CSE 3100 Master Notes

```
int a[5] = {1, 2}; // C99. b will have 1, 2, 0, 0, 5.  
int b[5] = {1, 2, [4] = 5};
```

### Array In Memory

- Think about how array elements are stored in memory
- Index starts from 0, the last one is 4 = (5 - 1)

### String Initialization

- A string is a char array that ends with a 0 (null character)
  - Memory that stores 0 is part of the string
- It can be initialized with a list of characters
- or a string (double-quoted literal)

### Arrays as Automatic Variables

- You can declare arrays inside any function or block
  - Destroyed when exiting from the function or block
- Variable length arrays(VLA, C99) The size of your array can depend on function arguments or other known value

```
int foo(int n,int k) {  
    int x[n]; // The value of n is known at this time  
    // Like other auto variables, x is kept on  
    // the stack and is NOT initialized  
  
    for (int i = 0; i < n; i++)  
        x[i] = 0;  
    .....  
    return -1;  
}
```

### Array Assignment

- You CANNOT assign a whole array at once to another array
  - Even when the types match

### Arrays and Functions

- Arrays can be passed to functions!
  - 1 BIG CAVEAT
- Calling convention in C
  - by value for everything
  - except arrays
- Arrays are always passed BY REFERENCE

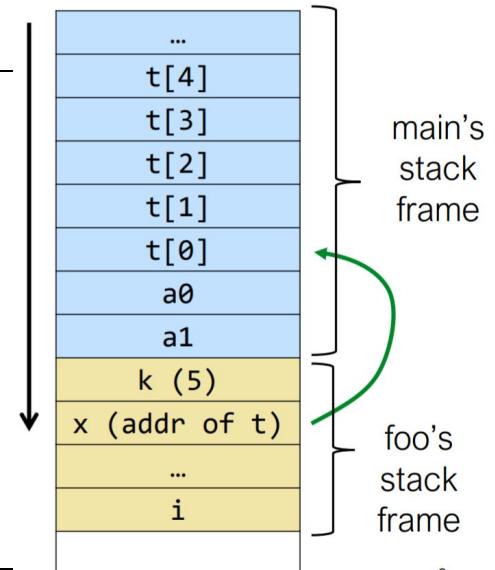
- passed as “pointers” - we’ll look at pointers soon
- Functions cannot return arrays
  - No easy assignments

## Array argument Example

- Address of “t” is passed to foo()
- Modifications to “x” are visible in main

```
int foo(int x[], int k) {
    for (int i = 0; i < k; i++)
        x[i] = i;
    return x[0];
}

int main() {
    int t[5];
    int a0 = foo(t, 5);
    int a1 = t[4];
    printf("%d %d\n", a0, a1);
    // more code
}
```



9

## Multidimensional arrays

```
// declaration and initialization
int h[2][3] = { {0, 1, 2}, {10, 11, 12} };
```

	0	1	2
0	0	1	2
1	10	11	12

## Pointers

- Perhaps the scariest part of C
- Yet...
  - The most useful part of C!
- Pointer is simply
  - a value
  - denoting address of a memory cell

## Variables and Memory

- The memory is an array of bytes
- Every byte in memory is numbered: the address!
  - An address is just an unsigned integer
- Every variable is kept in memory, and is associated with 2 numbers
  - The address
  - The value stored at that address

## Implicit Address Use

- address are used implicitly by the compiler all the time

```
int foo (int v)
{
    int a, b; // allocate storage space for a and b
    a = 1; // store 1 to memory, at a's address
    b = a; // load value from a's address, write to b's
address
    v = v + b; // load value from v's address, add to value at
// b's address, and write result to v's address
    return v;
}
```

	Address	Value
	100012	
b	100008	
a	100004	1

## Explicit Use: Pointers

- A pointer is a variable that holds an address of something
- Declaration

```
// declare p to be a pointer to an int
• the value of p is an address of an int, p itself has an address
```

## Referencing and dereferencing

- Two new operators
  - & Reference - “get” the address of something
  - \*

Ex.):

```
int x = 10, y;
// px is a pointer to int, i.e., the address of an integer
int *px; // px itself has an address

px = &x; // &x is the address of x. Save it to px

// *px: use px as an address to get the value at that location
y = *px; // and save the value in y

// save 20 to location pointed to by px (use px as an address)
*px = 20; // px has x's address, so x becomes 20
```

## Picture It

```
// assume 32 bits in and address
int x = 10, y;
int *px;
px = &x;
y = *px;
*px = 20;
```

	Address	Value
x	1007	
y	1006	
px	1020	10
	1016	?
	1012	1020
	1008	
	1004	
	1000	

## Revisit the Example

```
int foo(int x[], int k) {
    for (int i = 0; i < k; i++)
        x[i] = i;
    return x[0];
}
```

x is the starting address af an array, which is the same as the address of element 0.

```
// x is the address of an int
int foo(int * x, int k) {
    for (int i = 0; i < k; i++)
        x[i] = i; // use the pointer as an array
    return x[0];
}
```

## Array Question 1

```
// how many integers are in a, and in b?
int a[] = {1, 2, 3, 4};
int b[4] = {1, 2, 3};
```

**Both have 4 integers**

## Array Question 2

```
// how many bytes (characters) in c? and in d?
char c[] = {'a', 'b', 'c', 'd'};
char d[] = "abcd";
in "c" is 4, in "d" is 1
```

### Array Question 3

```
int a[10];  
If a[0]'s address is 1000, what is a[4]'s address?  
1016
```

### Array Question 4

```
char t[10][20];  
If t[0][0]'s address is 1000, what is t[1][1]'s address?  
1041
```

## Lecture 10 - C6: Dynamic Memory Allocation

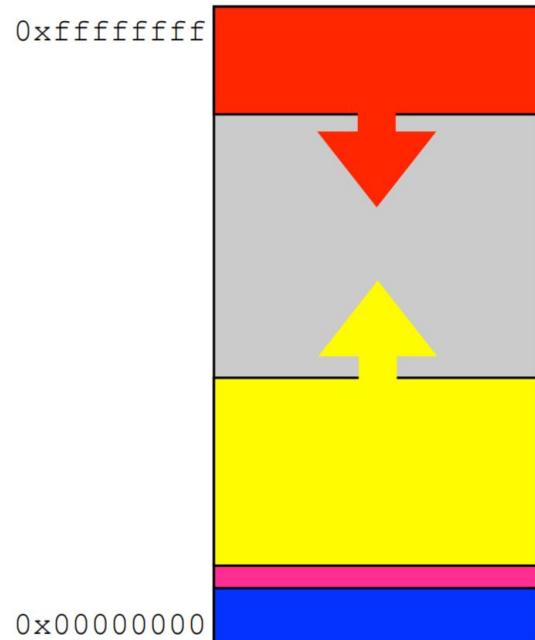
Weds. Sept. 18, 2019

Recall the memory model...

- Three pools of memory
  - Static/global
  - Stack
  - Heap
- Each pool features
  - Different lifetime
  - Different allocation/deallocation policy

Static/global memory pool

- This is where...
  - All constant (including string literals) are held
  - Global variables
  - All variables declared
- Allocated when - the program starts
- Deallocated when - the program terminates
- FIXED size - compiler needs to know the size & make reservations



Stack

- This is where...
  - Memory comes from local variables in functions!
- Easy to manage because it is automatic!

## CSE 3100 Master Notes

- Allocated automatically when entering the function
- De-allocated automatically when you leave the func.
  - Scope is that of func.
  - Should not be used after the func. returns
    - For ex): indirectly used via a pointer ← DO NOT DO

**default stack size is 2 MBs**

**need to increase stack size for large arrays and deep recursion**

### Heap

- This is where...
  - Memory comes from **manual** “on-the-fly” allocations
- Who is in charge?
  - **The programmer** for both allocation / deallocation
- Lifetime of memory blocks?
  - As long as they are not freed

### Requesting memory on the heap

- `#include <stdlib.h>`
- `void* malloc(size_t size);`
  - `size_t` is an unsigned integer data type defined in `<stdlib.h>`
  - used to represent sizes of objects in bytes
- If successful, a call to `malloc(n)` returns a **generic pointer (`void *`)**
  - It points to a memory block of “`n`” bytes on the heap
- If not successful, `NULL` is returned

```
char* p = malloc(100); // request 100 bytes
```

### Generic pointers: `void*`

- Pointer to a memory block whose content is “untyped”
- use for raw memory operations or in generic funcs.
- Automatic casting when assigned to other pointer types
  - `int * pox = malloc(6 * sizeof(int));`
- Requires casting before dereferencing for read / write
  - `*(int *)pv; // use pv as an int *`
- `NULL`, a special pointer value useful for initializations, error handling
  - `#define NULL ((void*) 0)`

### Can't always get what you want

- A call to `malloc()` may fail
  - for ex. when out of memory
- In this case you get back a `NULL` value

## CSE 3100 Master Notes

- Not much to do except report the error (and terminate nicely)
- Idiom

```
char*p = malloc(100); // request 100 bytes
if (p == NULL) {
    // report error and finish
    perror("Not enough memory");
    exit(1);
}
```

### How much Space?

- You need to tell malloc() how many bytes you need
- To request space for an array, need
  - # of elements
  - amount of space for each array element
    - sizeof(t) returns # of bytes needed for a value for type T
- ex.):

```
int* pox = malloc(6 * sizeof(int)); // request space for 6 ints
if (pox == NULL)
    report error and finish;
```

### Another Way

- #include <stdlib.h>
- void\* calloc(size\_t, nmemb, size\_t size);
- calloc() is implemented in terms of malloc()
  - calloc() also initializes the content to 0

```
int* pox = calloc(6, sizeof(int)); // request space for 6 ints
if (pox == NULL)
    report error and finish;
```

### Adjusting the Size

- #include <stdlib.h>
- void \* realloc(void\* ptr, size\_t size);
- What if you change your mind?
  - you requested 100 bytes, but now need 200

```
char* p = malloc(100); // request 100 bytes
...
p = realloc(p, 200); // p may change!
```

- Before a call to realloc(p, size), p must be
  - A pointer returned by a previous malloc/calloc/realloc

## CSE 3100 Master Notes

- o or NULL, in which case the call is equivalent to malloc(size)

### Deallocation

- #include <stdlib.h>
- void free(void \*ptr);
- Straightforward
  - o simply call the lib. func. "free"
  - o takes a pointer to the block to free

```
free(ptr);
```

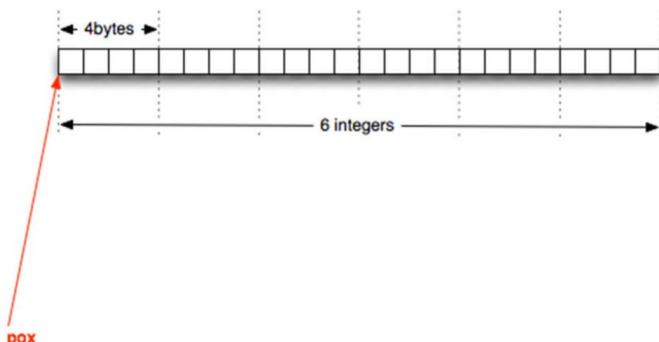
- Do not free a pointer twice!
  - o After freeing a pointer, set it to NULL!

### Two key rules

- Rule 1: Everything you request should be freed, eventually
- Rule 2: Only free what is allocated via malloc/calloc/realloc/
- Consequences of not following the rules
  - o Memory "leaks"
    - Your program will eventually run out of memory
  - o Undefined behavior and horrible crashes
    - Freeing unallocated mem. or already freed mem.
    - May cause a memory error and program crash
    - Worse, may corrupt heap and cause crash later
    - Even worse, program may keep running, totally corrupting your data, without you knowing

### Pointers and arrays

- after `pox = malloc(6 * sizeof(int))`
- That looks like an array
  - o and you can use `pox` as if it is



### Example: Use pointer as array

```
// programmers have to manage memory themselves before C99
void doSomething(int n){
    int * pox;
    pox = malloc(sizeof(int)*n);      // request mem from heap...
    *pox = 0;                      // set the int at the address pox to 0
    pox[0] = 0;                     // same thing
```

```

    pox[1];           // the int after pox[0]
    // more lines...
    free(pox);       // remember to free!
}

```

## Array of pointers

- Pointers, like integers, can be placed in an array

```

int a0;
int a1;
int a2;

char *p0 = malloc(10);
char *p1 = malloc(10);
char *p2 = malloc(10);

```

```

// array of int's
int a[3];

// array of pointers
char * p[3];

// Can also do with a loop
p[0] = malloc(10);
p[1] = malloc(10);
p[2] = malloc(10);

```

## Example: Command line argument

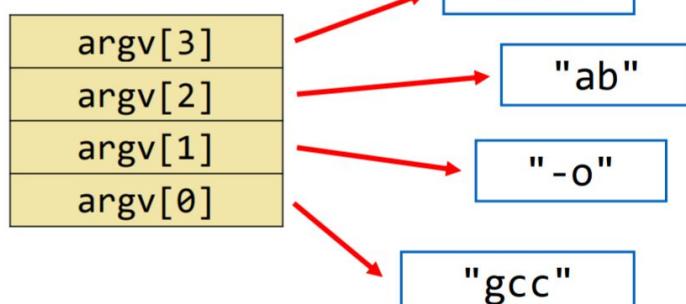
```
int main (int argc, char *argv[]);
```

- argc: the # of arguments on the command line
- argv: array of pointers to characters
  - the # of elements is argc
  - each element in an array points to null-terminated strings

- Example:**

```
$gcc -o ab ab.c
```

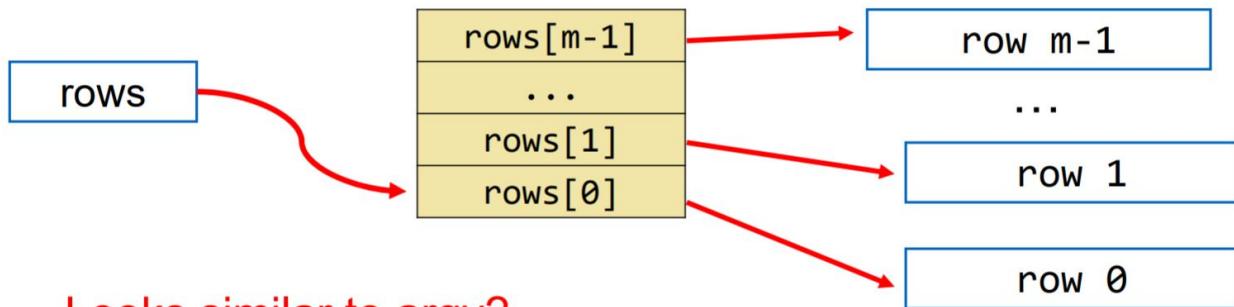
argv



What is argv[3][1]?

Example: allocating 2d dynamical array

```
void doSomething(int m, int n)
{ int **rows;
  rows = malloc(sizeof(int *) * m); // array of pointers
  for (int i = 0; i < m; i++)
    rows[i] = malloc(sizeof(int)*n); // one int array for each row
  ...
  for (int i = 0; i < m; i++)
    free(rows[i]);
  free(rows);
}
```

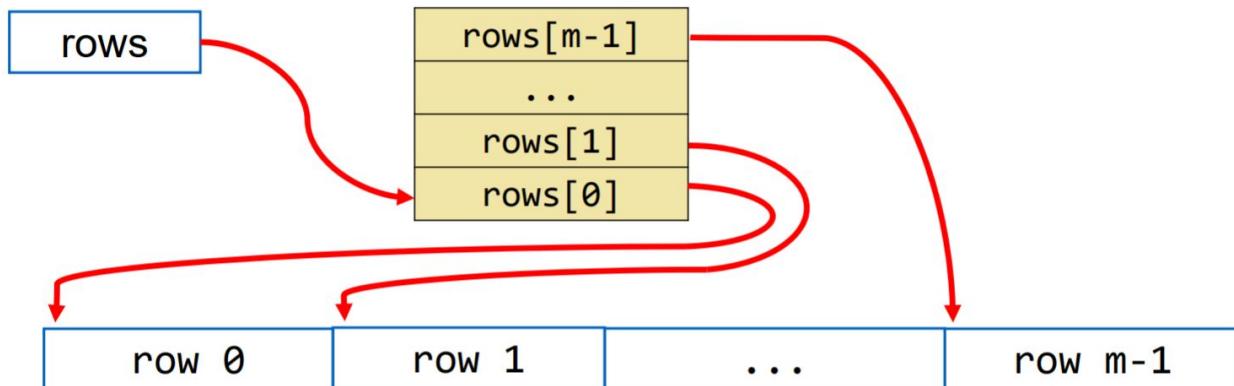


Looks similar to argv?

18

Example: allocating 2d dynamical array (take 2)

- Requesting all the memory space needed by data with one malloc() all
  - Instead of call malloc() m times, for each row
- Calculate rows[1], rows[2], and so on (pointer arithmetic! next lecture)
  - You can even calculate the address of each element (e.g. rows[10][5])



Example: allocating 2d dynamical array (take 3)

```

void doSomething(int m, int n)
{
    int(* arr)[n]; // arr is a pointer to an array of n int's
    arr = malloc(m * n * sizeof(int)); // one malloc() for data
to access
    arr[1][2] = 1;
    arr[2][3] = arr[1][2] + 10;
    ...
    free(arr); // free memory
}

```

n must be constant if Variable Length  
Array(VLA) is not supported



Pointers taking the address of...

- A static?
  - The address is never going to go “bad”
  - The static lives as long as the program!
- A stack [automatic] variable?
  - The address is valid as long as the variable!
  - When the function returns... The address is bogus
- A heap variable?
  - The address is valid as long as the variable is!
  - The variable disappears when explicitly de-allocated (freed)

# Lecture 11 - C7: Pointer Arithmetic & Structures

---

Mon. Sept. 23, 2019

## Pointers are addresses

- The value of a pointer is a byte address
  - Unsigned integer used to number bytes in memory
  - Range is between
    - 0x00000000 and 0xFFFFFFFF [32-bit]
    - 0x0000000000000000 and 0xFFFFFFFFFFFFFF [64-bit]
- Corollary
  - If a pointer is an integer, you can do arithmetic
  - To computer addresses

## Pointer Addition Example

- Suppose p is a pointer to an int, and its value is 10000
- $p + 1$  is not the next byte address
- It is the address of next item (of type int)

## Adding a pointer and an integer

- Add an integer to a pointer, the result is a pointer of the same type
  - It is dif. from reg. integer addition
  - The integer is **automatically scaled** by the size of the type pointed to
- Suppose p is a pointer to type T, and k is an integer
  - Then both " $p + k$ " and " $k + p$ "
  - Are valid expressions that evaluate to a pointer to type T
  - Have a byte address equal to

**(unsigned long)(address stored in p) + k \* sizeof(T)**

C standard does not allow arithmetic on void \*

gcc has an extension, treating sizeof(void) as 1

Address	Value		To access values
1020		p + 5	*( $p+5$ ) OR $p[5]$
1016		p + 4	*( $p+4$ ) OR $p[4]$
1012		p + 3	*( $p+3$ ) OR $p[3]$
1008		p + 2	*( $p+2$ ) OR $p[2]$
1004		p + 1	*( $p+1$ ) OR $p[1]$
1000		<b>p = 1000</b>	*p OR $p[0]$
996		p - 1	*( $p-1$ ) OR $p[-1]$
992		p - 2	*( $p-2$ ) OR $p[-2]$

## Pointers Subtraction

- Subtract one pointer from another: both must have the same type
- The result is the number of data items between the two pointers!
  - Not the number of bytes

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *p = malloc(sizeof(int)*10);
    int *last = p + 9;
    int dist = last - p; // both are int *
    printf("Distance is %d\n", dist);
    free(p);
    return 0;
}
```

## Output

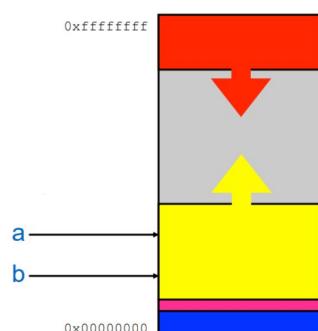
```
$ gcc ptrsub.c
$ ./a.out
Distance is 9
$
```

## Pointer comparison

- You can also compare two pointers
- Purpose
  - Check boundary conditions in arrays
  - Manually manage memory blocks
- Semantics
  - Simply based on memory layout!
  - Compare bits in pointers as unsigned integers!

Check if you are done:

```
while (b < a) {
    // do something
}
```



## Effects of casting types?

- If you cast a pointer type...
  - Any subsequent pointer arithmetic will use the type you choose
- Do not want scaling? Casting a pointer to (char \*)
  - Because sizeof(char) is 1

```
int * t;
char * p = (char *) t + 8;    // 8 is not scaled
char * q = (char *) (t + 8); // 8 is scaled
```

## Arrays and pointers

- Arrays and pointers can often be used interchangeably

```
int a[10];
int *p = a;

// all of the following evaluate to the value of a[0]
*p           p[0]           *a           a[0]

// all of the following evaluate to the value of a[1]
*(p+1)       p[1]           *(a+1)       a[1]

// all of the following evaluate to the address of a[0]
// type is int *
p             &p[0]          a           &a[0]
```

↑

"array of int" becomes "pointer to int" (array decay)

## Example: arrays and pointers

- Equivalent ways of initializing an array

## CSE 3100 Master Notes

```
int a[10], *p = a; // not *p = a; it is int *p; p = a;  
  
for(int i=0; i<10; i++) a[i] = i; // array indexing  
  
for(int i=0; i<10; i++) p[i] = i; // indexing via pointer  
  
for(int i=0; i<10; i++) *(p+i) = i; // explicit pointer arithmetic  
  
for(i=0; i<10; i++) i[p] = i; // obfuscated but valid C!  
  
for(i=0; i<10; i++) *p++ = i; // common pointer use idiom
```

```
int * tp = p;  
p++;  
*tp = i;
```

Arrays and pointers are NOT the same

```
int a[10];  
int *p = a; // a is converted to int *  
  
// a is still an array after &  
&a // pointer to array of 10 int's int (*)[10]  
&p // pointer to a pointer to int int **  
  
// a is still an array after sizeof  
sizeof(a) // 40 because a is an array of 10 int's  
sizeof(p) // 8 because p is a pointer  
  
p++; // can increment p  
a++; // cannot increment a; this will not compile  
// Similar to n++ vs 2++
```

Example: Arrays and pointers are NOT the same

```
#include <stdio.h>

void foo(int *x)
{
    printf("%lu\n", sizeof(x));
}

int main()
{
    int a[10], *p;

    p = a; // a is converted to int *
    // a is still an array in sizeof
    printf("%lu %lu\n", sizeof(a), sizeof(p));
    foo(a); // a is converted to int *
}
```

## Output

```
% gcc array.c
% ./a.out
40 8
8
```

## Structures

- Mechanism to define new types
  - Also known as “compound types”
  - Use to aggregate related variables of different types
- Structures
  - Structures can have a type name
  - Can have “members” of any types
    - Basic types
    - Pointer
    - Arrays
    - Other structures
- Structure variable definition
  - Specifies variable name
  -

```
struct student_grade {
    char *name;
    int id;
    char grade[3];
};

...
struct student_grade student1;
struct student_grade
    student2, student3;
struct student_grade all[200];
```

## Structure example

```
struct Person {
    int     age;
    char   gender;
};

int main(){
    struct Person p;

    p.age = 44;
    p.gender = 'M';

    struct Person q = {44, 'M'};

    return 0;
}
```

### Structure type declaration

### Structure variable definition

Syntax for field access  
similar to Java and Python

### Structure variable definition and initialization

## Example: Array of Structures

- Member name is a char array
- Cabeats
  - Names cannot be more than 31 characters long
  - Four person in family
  - Indexed 0..3
- Array of structures for the whole family
- Nested initializers

## typedef

- Struct. names can be long
- C provide the ability to define type abbreviations
  - typedef declaration
    - Give existing type new type name
- Make code more readable
- Structure and typedef declarations often combined

## Operations on struct

- Assignment
  - All struct members copied
- Can be passed to funcs.
- If pass by value, cannot change members in funcs.
- Passing or returning large structures can be costly

```
#include <stdlib.h>

struct Person {
    char   name[32];
    int    age;
    char   gender;
};

int main()
{
    struct Person family[4] = {
        {"Alice",34,'F'},
        {"Bob",40,'M'},
        {"Charles",15,'M'},
        {"David",13,'M'}
    };
    int juniorAge = family[3].age;
    return 0;
}
```

```
struct Person {
    char   name[32];
    int    age;
    char   gender;
};
typedef struct Person TPerson;
int main()
{
    TPerson family[4];
    ...
    return 0;
}
```

```
typedef struct Person {
    char   name[32];
    int    age;
    char   gender;
} TPerson;
```

\*\*\*use pointers to structures!

Pass Structure by reference

```
typedef struct Person{
    char name[32];
    int age;
    char gender;
}
```

```
TPerson * init_Person(TPerson *p, char * name, int age, char
gender) {
    strcpy(p->name, name);           // (*p).name
    p->age = age;                  // (*p).age
    p->gender = gender;            // (*p).gender
    return p;
}
// Study the demo code is in the demo repo
// especially the lines that copy name
```

```
TPerson a, b, c;
...
a = b;
c = searchPerson("name");
```

## Structure Alignment

- Structure members are aligned for the natural types

Alignment requirements on x64 architecture	
char	1
short	2
int	4
long	8
float	4
double	8

```
struct struct_random {
    char     x[5]; // bytes 0-4
    int      y;    // bytes 8-11
    double   z;    // bytes 16-23
    char     c;    // byte 24
}; // Total size 32
```

```
struct struct_sorted {
    double   z;    // bytes 0 - 7
    int      y;    // bytes 8 - 11
    char     x[5]; // bytes 12 - 16
    char     c;    // byte 17
}; // Total size 24
```

## Arrays and pointers

- Copying arrays

```
// using array indexing
void copy_array0(int source[], int target[], int n)
{
    for(int i = 0; i < n; i++)
        target[i] = source[i];
}

// using pointers
void copy_array1(int source[], int target[], int n)
{
    for(int i = 0; i < n; i++)
        *target++ = *source++;
```



```
int * tp = source;
source++;
*target = *tp;
target++;
```

## Typecasting Pointers

- C lets you cast pointers in any way you like
- You can “forge” pointers to point wherever you wish...

```
float f;
char p = * (char *) &f;           // 1st byte representing f
char ch = * (char *) 1000; // access any byte in memory
```

- THAT'S WHAT MAKES C VERY ATTRACTIVE FOR LOW-LEVEL PROGRAMMING
- THAT IS ALSO VERY POWERFUL AND THUS DANGEROUS

## Returning more than one value from functions

- Use references (caller prepares the storage)

```
long int strtol (const char* str, char** endptr, int base);
```

- Use arrays (caller prepares the storage)

```
int pipe(int pipefd[2]);
```

- Return a structure (costly if structure is large)
- Return a pointer (to array or structure)

- Must be dynamically allocated or static. RTM (read the manual)!

```
char strdup(const char *str1);
```

## CSE 3100 Master Notes

```
struct tm *localtime( const time_t *time );
• Use global variables (DON'T)!  
errno
```

### Typedef

```
// Example of typedef. Think about how you would define a variable
typedef    int    BOOL;
typedef    char   name_t[100];
typedef    char *Pointer;
```

### Self-referential structures

```
struct Person{
    int age;
    char gender;
    char name[32];
    struct Person * parents;      // A pointer to this type of
struct
} person1, person2;           // Can define variables here
```

### Self-referential structures - 2

```
struct student{
    char name[128];
    // Can have a pointer to a struct defined later.
    // However, you cannot define an array of book here (e.g.
books[8])
    struct book * books;
};

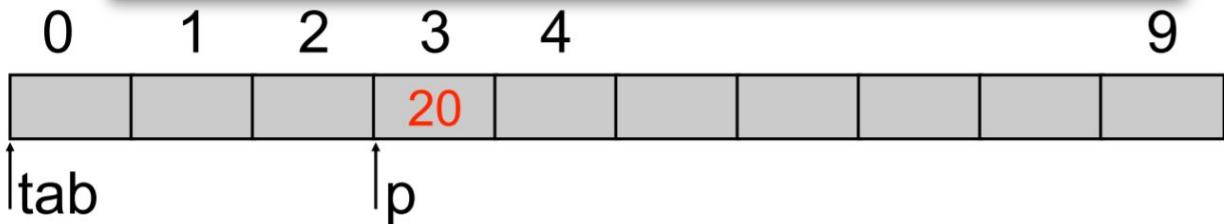
struct book {
    char title[128];
    struct student * owner;
    struct book * next; // A pointer to this type of struct
};
```

### Example of Pointer Arithmetic

- Simple illustration

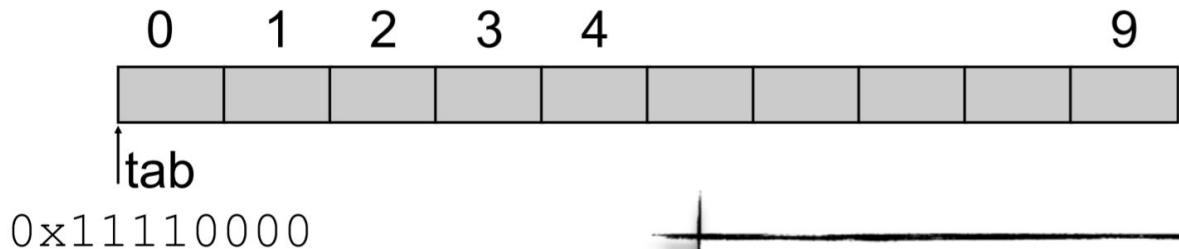
```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int)*10);
    tab[3] = 10;
    int *p = tab + 3;
    printf("What is at tab+3? = %d\n", *p);
    *p = 20;
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```



But what about memory addresses?

- Same story...!



$\text{tab} + 1 = ?$

~~0x11110001~~ ?

0x11110004    ?

### WHY?

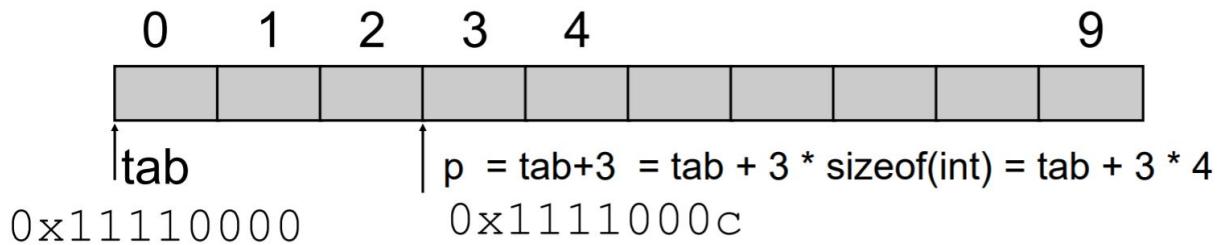
Simply because  $\text{tab}$  is a pointer to an int and an int is 4-bytes wide!

## Bottom Line

```
#include <stdlib.h>

int main()
{
    int *tab = (int*)malloc(sizeof(int) * 10);
    tab[3] = 10;
    int *p = tab + 3;
    printf("What is at tab[3]? = %d\n", *p);
    printf("What is at tab[3]? = %d\n", tab[3]);
    return 0;
}
```

The offset **3** is scaled by the compiler with the size of the type to get an address in bytes



## Memory Alignment Requirements

- Memory used to store a value of type X MUST
  - be lined-up on a multiple of natural alignment for X
- Why?
  - Performance!
- If you do not respect alignment requirements...
  - BUS ERROR (sigbus)
  - The O.S. will kill your program

## Good News and Bad News

- The C compiler handles alignment 99% of the time
- Programmers have to handle the rest - when you do pointer arithmetic of course!
  - Do not assume the location of the fields
- When you call sys. routines w/ specific alignment needs
  - Your arguments must comply
  - Use compiler annotations to force specific alignment (beyond our scope, simply remember that this exists!)

## Example

```
#include <stdio.h>
struct Person {
    char name[32];
    int age;
    char gender;
};
typedef struct Person TPerson;

TPerson init(char name[], int age, char gender) {
    TPerson p;
    int i;
    for(i = 0; name[i]>0; i++)
        p.name[i] = name[i];
    p.name[i] = '\0';
    p.age = age;
    p.gender = gender;
    return p;
}

void print_info(TPerson p) {
    printf("name: %s, age: %d, gender: %c\n",
           p.name, p.age, p.gender);
}
```

```
int main() {
    TPerson family[4];
    family[0] = init("Alice",34,'F');
    family[1] = init("Bob",40,'M');
    family[2] = init("Charles",15,'M');
    family[3] = init("David",13,'M');
    print_info(family[0]);
    print_info(family[1]);
    print_info(family[2]);
    print_info(family[3]);
    family[1] = family[0];
    print_info(family[1]);
    return 0;
}
```

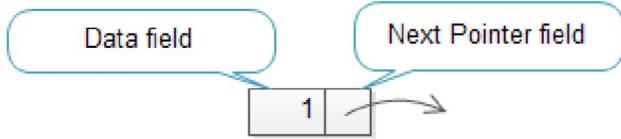
## Output

```
./a.out
name: Alice, age: 34, gender: F
name: Bob, age: 40, gender: M
name: Charles, age: 15, gender: M
name: David, age: 13, gender: M
name: Alice, age: 34, gender: F
```

## Lecture 12 - C8: Linked Lists, Enums, Func. Ptrs.

Weds. Sept. 25, 2019

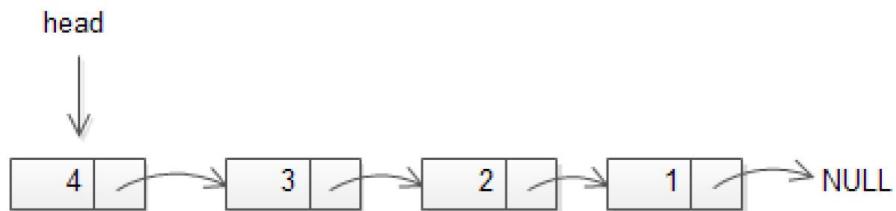
## Example: Linked List



## Head

```
node * head;           // head is a pointer, not a node!
head = NULL;           // at beginning, it is empty
```

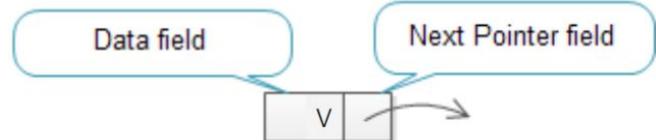
- After adding nodes into the list,



## Create a node

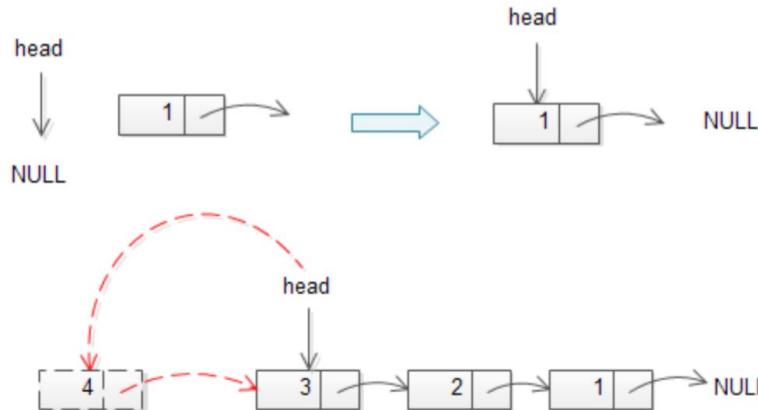
```
node* new_node(int v) // create a node for value v
{
    node * p = malloc(sizeof(node)); // allocate memory
    assert(p != NULL); // you can be nicer

    // Set the value in the node.
    p->v = v; // you could do (*p).v
    p->next = NULL;
    return p; // return
}
// is it similar to creating objects using "new"?
```



## Prepend

```
node * prepend(node * head, node * newnode)
{
    // how?
}
```



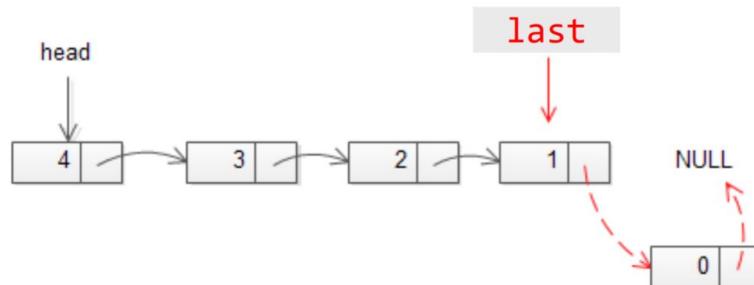
- cannot perform the following or else the head will point to the 1st node, and the 1st node points back to the head

Find the last one

```
node * find_last(node * head)
{
    if (head != NULL) ;
}
```

Append

```
node * append(node * head, node * newnode)
{
    node * last = find_last(head); // find the last one
    if(last == NULL) // if the list is empty, new node is the
    head
        return newnode;
    last -> next = newnode;
    newnode -> next = NULL;
    return head; // return the (unchanged) head
}
```



## Enumeration types (ABC 7.5)

- User-defined integer-like types:
- Names look like C identifiers
  - are listed (enumerated) in definition
  - treated as integer

## Enumeration types

```
// enum start from 0 by default
enum week {Sun, Mon, Tue, Wed, Thur, Fri, Sat};
enum week dow = Mon;

// But can be initialized; Warning is 2, Error is 3, etc.
enum status {OK = 1, Warning, Error, Fatal};
```

## Type qualifier: const

```
// constant int
const int a = 10;           // cannot change a
// a pointer to constant int
const int *pa = &a;    // can change pa, but not *pa
// a constant pointer to an int
int * const pb = &b      // can change *pb, but not pb
// constant pointer to a constant int
const int * const pc = &a; // cannot change *pc or pc

// cannot change the source string
char * strcpy(char * dest, const char* src);
```

## Function pointers

```
/* function returning */
int func();
/* function returning to integer */
int * func();
/* pointer to function returning integer */
int (*func)();
/* pointer to function returning pointer to int */
int * (*func)();
```

## Pointer to function example

```

int mymax(int a,int b)
{
    return (a > b) ? a : b
}
// a pointer to function
int (*pf)(int a, int b);

// assign a value to the pointer
pf = mymax;           // C99 style. Note that it is not mymax()
pf(3, 5);
pf = & mymax;
(*pf)(3,5);

```

## Use of function pointers

- Call-back mechanism
  - Generic functions (ex. coming next)
  - pthread\_create()
  - Dynamic signal handlers, ....
- You can store function pointers in arrays
  - And arrays stored in structures!
  - And you can simulate objects in Object Oriented Languages!

## Example: quicksort in C library

- The prototype (in <stdlib.h>)

```

void qsort(void * base
           size_t nel,
           size_t width,
           int (*compare)(const void *, const void *));

```

- qsort takes...
  - base:
  - nel:
  - width:
  - compare: a pointer to a function that compares two values

*\*\*\*sort only knows*

## Why passing a function to qsort?

- Need to tell qsort() how to compare items in the array

## CSE 3100 Master Notes

- we have a generic quickSort implementation
  - Do not want to implement one for each type of data
- The qsort() implementation calls the comparator to rank elements

```
int (*comprae)(const void *a, const void *b);
```
- the function takes the address of 2 items to be compared,
- and returns:
  - 0 if \*a EQUALS \*b
  - a positive value if \*a is GREATER THAN \*b
  - A negative value if \*a is LESS THAN \*b

### Example of compare() function

- when qsort() compares items, it provides the address of two elements to be compared.

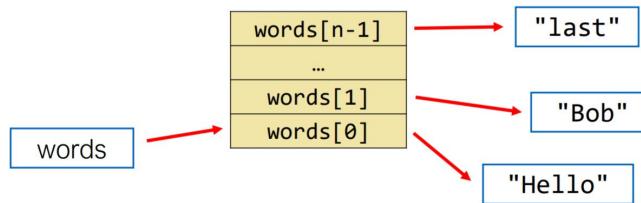
```
int compare_int(const void *a, const void *b)
{
    // qsort() does not know the type, but you know
    return *(int *a) - *(int *)b;
}

int compare_double(const void *a, const void *b)
{
    // qsort() does not know the type, but you know
    return
}
```

### Example: sort array of strings

- Example are pointers to strings
  - Need to compare string, instead of pointer

```
int compare_string(const void *a, const void *b){
// how to compare *a and *b?
// for example, a is &words[0] and b is &word[1]
```



### Compare string pointers

- An element in array words is (char \*).
- a is the address of an element of type (char \*). So, a's type is (char\*)\*

```

int compare_string(const void * a,const void * b)
{
    char *s1, *s2;
    s1 = *(char **)a;      s2 = *(char **)b;

    return strcmp(s1, s2); // use library function to compare
}

// or on one line
int compare_string(const void * a,const void * b)
{
    return strcmp(*(char**)a,*(char**)b);
}

```

## Calling quicksort()

- See complete demo code in the demo repo

```

int compare_string(const void* a,const void* b)
{
    return strcmp(*(char**)a,*(char**)b);
}

```

Type casting to char \*\* before dereferencing

```

int some_function(void)
{
    ...
    char** words = malloc(sizeof(char*)*n);
    ...
    qsort(words,n,sizeof(char*),compare_string);
    ...
}

```

## Lecture 13 C9: I/O and Files

### errno

- Most C library functions can “fail”
  - When they do, they return a flag reporting failure... (-1)
  - Some set of global variable to report the exact error code  
*errno*

// to use errno, include <errno.h>

- Check manual page to interpret the error code
- Print a more descriptive message with perror()
- Avoid functions that set errno in multithreaded code
  - Prefer thread-safe version when available

```
void perror(const char *str);
```

### Files and directories

- A file is an object that stores information, data, etc.
- 
- Example: files you create with an editor (.c, .h, Makefile, readme, etc.) executable generated by the compiler, and gcc itself other devices, like screen, keyboard, ...
- In Linux, files are organized in directories
  - A directory can have subdirectories and files
  - The top directory is /
- A path specifies the location of file/directory in the file system

/home/john

- In Unix/Linux, everything is a file

### The stdio library

- ```
#include <stdio.h>
```
- Declares FILE type and function prototypes
    - FILE is an opaque type (system dependent) for operating on files
      - It is a structure, but do not try to change it directly!
    - Use library functions to access FILE objects, via pointers (FILE \*)
  - Defines “standard” streams stdin, stdout, stderr
    - Created automatically when program starts
    - They are files!
  - The library is linked automatically by the compiler

### Files and I/O API

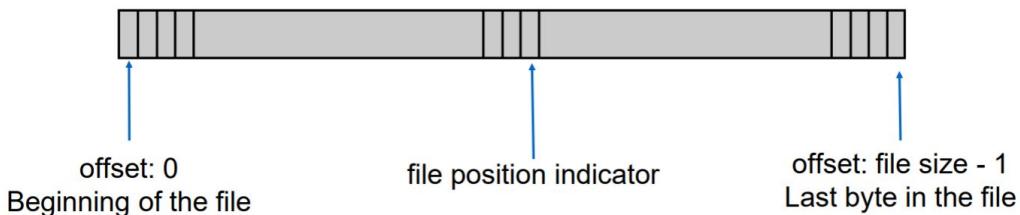
- In C, a file is simply a sequential stream of bytes
- The “f” family of functions (fopen, fclose, fread, fgetc, fscanf, fprintf,...) are C library functions to operate on files

## CSE 3100 Master Notes

- All these use a FILE\* abstraction to represent a file
- The C library provides buffering
  - That's why sometimes you do not see output of printf immediately We will learn another set of functions provided by OS

### File as stream of bytes

- Before use a file must be “open”
  - This sets a position indicator for reading and/or writing
- Each read/write starts from current position, and moves the indicator
  - Writing after last byte increases the file size
- Position indicator can also be changed with fseek
- All open files are closed when program ends
  - Good practice to close explicitly when no longer needed



### Opening Streams

FILE \*fopen( const char \*filename, const char \*mode)

- Open the file filename in mode as a stream of bytes
- Returns a pointer to FILE (FILE \*) or NULL (and errno is set)
- Mode
  - “r” : Reading mode
  - “r+” : Read and write
  - “w” : Writing mode, file is created or truncated to zero length
  - “w+” : Read and write, but the file is created or truncated
  - “a” : Append mode, the file is created if it does not exist
  - “a+” : Read and append, the file is created if it does not exist. Reading starts at the beginning, but writing done at the end

### Closing Streams

int fclose (FILE \*stream)

- Close a stream
- Returns
  - 0 if it worked
  - EOF if there was a problem (and errno is set)

### fgetc / fputc (one byte at a time)

```
int fgetc( FILE *stream);  
int fputc(int c, FILE *stream):  
    • Read or write one (ASCII_ character (8-bits) at a time  
        ○ Can be slow for large files  
    • fgetc reads a character from the stream and returns the character just read in (as  
        unsigned char extended to int)  
        ○ Returns EOF when at the end of file or on error  
    • fputc writes the character received as argument to the stream and returns the character  
        that was just written out  
        ○ Returns EOF on error
```

### getc / putc and ungetc

```
int getc(FILE *stream);  
int putc(int c, FILE *stream);  
    • Same as fgetc/fputc except they may be implemented as macros  
    • Use fgetc/fputc unless you have strong reasons not to  
int ungetc(int c, FILE *stream);  
    • Pushes last read char back to stream, where it is available for subsequent read  
        operations  
    • Only one pushback guaranteed
```

### getchar / putchar

```
int getchar(void)  
// same as fgetc(stdin)  
    • Reads a character from stdin  
    • Returns the character just read in, or EOF on end-of-file errors  
int putchar(int c)  
// same as fputc(c, stdout)  
    • Writes the character received as argument on stdout  
    • Returns the character that was just written out, or EOF on errors
```

### More than one byte: get a line

```
char *fgets(char *buf, in size, FILE *in)  
    • Reads the next line from in into buffer buf  
    • Halts at '\n' or after size-1 characters have been read  
        ○ NUL is placed at the end  
    • Returns pointer to buf if ok, NULL otherwise  
    • Do not use gets(char *)! – buffer overflow
```

## CSE 3100 Master Notes

```
int fputs(const char *str, FILE *out)
• Writes the string str to out, stopping at '\0'
• Returns number of characters written or EOF
```

### Formatted output

```
int fscanf(FILE *stream, const char *format, ...);
int fprintf(FILE *stream, const char *format, ...):
• Formatted input from file and output to file
• Like scanf() / printf(), but not from stdin or to stdout
```

### For binary data

```
size_t fread (void *ptr, size_t sz, size_t n, FILE *stream);
size_t fwrite(void *ptr, size_t sz, size_t n, FILE *stream);
• Read / write a sequence of byte from / to a stream
• Return the number of items read or written
○ If smaller than n, EOF or error
```

**Example:**

```
int A[10][20];
size_t n = 10 * 20;
if (fwrite(A, sizeof(int), n, fp) != n) {
    // error
```

### Moving file position indicator

```
long ftell(FILE *stream);
• Read file position indicator
• Return -1 on error
int fseek(FILE * stream, off_t offset, int whence);
• Set the file position indicator
• Return 0 on success and -1 on error
```

**Example:**

```
fseek(fp, 0, SEEK_SET);           // move to the beginning
fseek(fp, 200, SEEK_CUR);         // move forward 200 bytes
fseek(fp, -1, SEEK_END);          // move to the last byte
```

### More useful stdio functions

```
//Check if end-of-file is set (after a read attempt!)
int feof(FILE * stream);

//Force write of buffered data
```

```
int fflush(FILE * stream);
```

- Read the manual pages!
- Check the return values!

## Lecture 14

---

Mon. Sept. 30, 2019

- forgot what we did lol

## Lecture 15

---

Weds. Oct. 2, 2019

Practice Exam 1

# Lecture 16 - P1: Intro to Processes

Mon. Oct. 07, 2019

## Process Basics

- A process is an instance of a program being executed
  - Core operating system (OS) concept
- In a multiprocessing OS
  - Multiple programs can be executed at the same time
  - Multiple instances of a program can be executed at the same time
- Executing multiple programs
  - Single-core: time-sharing
  - Multi-core: true parallelism + time-sharing

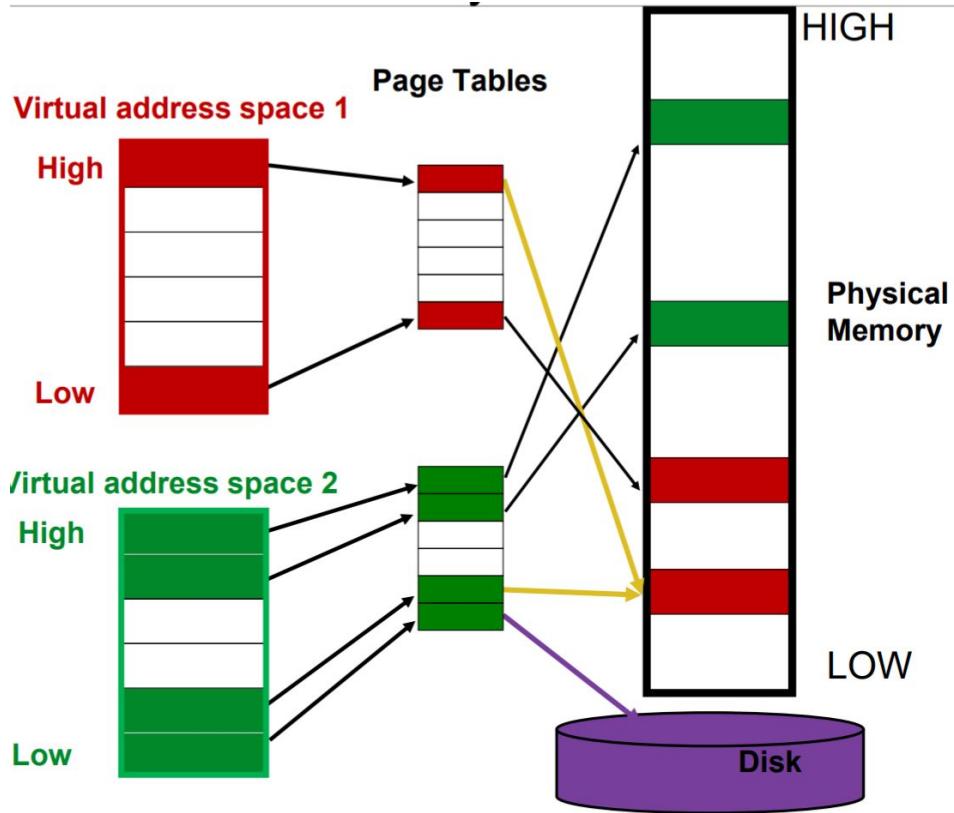
## Process Management: OS View

- OS maintains a process table
  - Each process has a table entry, called process control block (PCB)
  - Typical PCB info
- OS scheduler picks processes to be executed at any given time
  - When a process is suspended, its state is saved in PCB
  - What about the process memory?

| Process management        | Memory management             | File management   |
|---------------------------|-------------------------------|-------------------|
| Registers                 | Pointer to text segment info  | Root directory    |
| Program counter           | Pointer to data segment info  | Working directory |
| Program status word       | Pointer to stack segment info | File descriptors  |
| Stack pointer             |                               | User ID           |
| Process state             |                               | Group ID          |
| Priority                  |                               |                   |
| Scheduling parameters     |                               |                   |
| Process ID                |                               |                   |
| Parent process            |                               |                   |
| Process group             |                               |                   |
| Signals                   |                               |                   |
| Time when process started |                               |                   |
| CPU time used             |                               |                   |
| Children's CPU time       |                               |                   |
| Time of next alarm        |                               |                   |

## Paged Virtual Memory: How Processes Share Memory

- Physical memory is shared by all processes
- Page table maps virtual address to physical address
- Multiple virtual pages can be mapped to the same physical pages



## Process Management: User's View

- Events which cause process creation
  - System initialization
  - User request to create a new process (e.g., shell command)
  - Executing a shell script, which may create many processes
- Events which cause process termination
  - Normal program exit
  - Error exit
  - Fatal error, e.g., segmentation fault
  - Killed by user command or signal (Ctrl-C)

## Useful Commands

- **ps** - List running processes
- **pstree** - Display the tree of processes
- **top** - Dynamic view of memory & CPU usage + processes that use most resources (to exit top, press q)
- **kill** - Kill a process given its process ID
  - Try **-9** option if simple kill does not work
- Additional functions/options in man page of each command

### Process Management: Programmer's View



- Process birth
  - Processes are created by other processes!
  - A process always starts as a clone of its parent process
  - Then the process may upgrade itself to run a different executable
    - Child process retains access to the files open in the parent
- Process life
  - Child process can create its own children processes
- Process death
  - Eventually calls exit or abort to commit “suicide”
  - Or gets killed

### Birth via Cloning

- The function to create a new process in your code

```
#include
pid_t fork(void);
```
- Child is an exact copy of the parent
  - Both return from fork()
- **Only difference is the returned value**
  - In the **parent** process:
    - fork() returns the process identifier of the child (> 0)
    - If a failure occurred, it returns -1 (and sets errno)
  - In the **child** process: fork() returns 0 (zero)

### Concurrency

- Parent and child processes return from fork() concurrently
  - They may return at the same time (on a multicore machine) or one after the other
  - Cannot assume that they return at the same time or which one “returns first” (even on a uni-core)
    - Order is chosen by OS scheduler

### Cloning Effect

- On memory
  - The parent and child memory 100% identical
  - But are viewed as distinct by OS (“copy-on-write”)
  - Any memory change (stack/heap) affects only that copy
  - Thus the parent and child can quickly diverge
- On files
  - *All files open in the parent are accessible in the child!*

## CSE 3100 Master Notes

- I/O operations in either one move the file position indicator
- In particular
  - *stdin*, *stdout*, and *stderr* of the parent are accessible in the child

### What can the parent do?

- Depends on application!
  - It could wait until the child is done (dies!)
    - Typical of a shell like bash/ksh/zsh/csh/....
  - It could run concurrently and check back on the child later
  - It could run concurrently and ignore the child
    - If child dies it enters a zombie state

### Waiting on a child

```
#include <sys/wait.h>

pid_t wait(int * status);
pid_t (waitpid(pid_t pid, int * status, int options);
```

- Purpose
  - Block the calling process until a child is terminated
    - Or other state changes specified by options
  - Report status in \*status (which is ignored if NULL is passed)
    - The cause of death
    - The exit status of the child (what he returned from main)
  - Return value identifies the child process (or -1 on error)
  - Run “man -S2 wait” for full details

### Zombies

- A dead process, waiting to be 'reaped' (checked by its parent)
  - You cannot kill it, because it is already dead
  - Most resources released, but still uses an entry in the process table
- Parents should check their kids
  - On some systems, parents can say they do not want to check
- When a parent dies, ‘init’ becomes the new parent
  - Then the zombie child is reaped

### System calls

- APIs used to request services from the OS kernel
  - Example: fork()
  - System calls are more expensive than normal function calls
  - Manuals for system calls are in section 2

## CSE 3100 Master Notes

*man -S2 intro ; man -S2 syscalls*

### Summary

- Clone a process with fork()
  - The child is exactly the same as the parent
  - Check the return value
- Parents wait for child processes
  - Reap the zombies!

## Lecture 17 - P2: Exec and Low-Level I/O

---

Mon. Oct. 09, 2019

### Process upgrades

- Usually...
  - A fresh clone wants to run different code
- This is done by
  - Loading another executable3 into the process address space
    - [picked up from the file system of course]
- Note: **open files are NOT AFFECTED** by the upgrade operation

### The exec family

- The act of “upgrading” is done by the child with a system call
  - Many variants ‘many -S3 exec’ for all details

```
#include <unistd.h>
int execl(const char *path, const char *arg0, ...
           /*, (char *) NULL */ );
```

- The path to the executable to load inside our own address space
- A list of arguments to be passed to the new executable
- A final NULL pointer to give the “end of argument list”
- If successful, execl() does not return! Started a new process

### Exec example

- We will turn the child process into the following executable

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc,char* argv[]) {
    int i, sum=0;
    for(i=1;i<argc;i++)
        sum += atoi(argv[i]);
    printf("sum is: %d\n",sum);
    return 0;
}
```

This is a simple “adder” program that computes the sum of its integer arguments

## Parent Program

```
int main() { // complete code is in demo/padder
    char *cmd1 = "./adder", *cmd2 = "addder";
    pid_t child = fork();
    if (child == 0) {
        printf("In child!\n");
        execl(cmd1,cmd1,"1","2","3","10",NULL);
        printf("Oops.... something went really wrong!\n");
        perror(cmd1);
        return -1;
    } else {
        printf("In parent!\n");
        execl(cmd2,cmd2,"100","200","300",NULL);
        printf("Oops.... something went really wrong!\n");
        perror(cmd2);
        return -1;
    }
}
```

## How is executable found?

- Specify a path, like /bin/ls
- Specify a file, and the system searches in directories listed in PATH
  - echo \$PATH in bash to see directories separated by ‘:’

in `execl(const char *path, const char *arg0, ...)`

```
/*, (char *) NULL */ );

// execvp() searches paths for file
int execvp(const char *file, const char *arg0, ...
           /*, (char *) NULL */ );
```

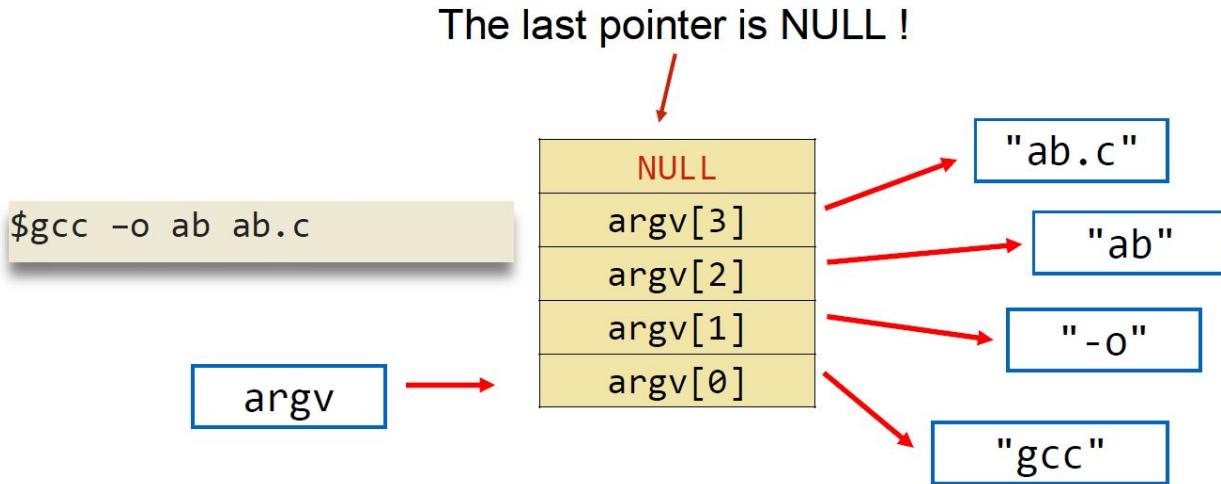
## execv family

```
// If the number of arguments is unknown at compile time
#include <unistd.h>
int execv(const char *path, char *const argv[]);
int execvp(const char *file, char *const argv[]);
```

- The arguments in execv are placed in an array
  - argv is the argv you see in the main function!
- execv needs a path while execvp can search file in PATH
- Start a new process if successful. Similar to exec

## argv to execv and execvp

- Note the NULL pointer at the end
- Why?



## File APIs

- Remember the (C standard library) IO APIs
  - The “f” family (fopen, fclose, fread, fgetc, fscanf, fprintf,...)
  - All these use a FILE\* abstraction to represent a file
    - Additional features: user-space buffering, line-ending translation, formatted I/O, etc.
- UNIX has lower-level APIs for file handling
  - Directly mapped to system calls

## CSE 3100 Master Notes

- Open, close, read, ...
  - Use file descriptors [which are just integers]
  - Deal with bytes only

Some low level file APIs

- Read the man pages (man -s2 ...) for more functions

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int open(const char *path, int oflag);
int close(int fd);

ssize_t read(int fd, void *buf, size_t nbytes);
ssize_t write(int fd, const void *buf, size_t nbytes);

off_t lseek(int fd, off_t offset, int whence);
```

Open a file

```
#include <fcntl.h>
#include <unistd.h>

int open(const char *path, int oflag);
```

- Parameters
  - path: the path to the file to be open/created
  - oflag: read, write, or read and write, and more (on the next slide)
- The function returns a file descriptor, a small, nonnegative integer
  - Return -1 on error

Flags in open()

- Must include one of the following:  
O\_RDONLY (read only), O\_WRONLY (write only), or O\_RDWR (read and write)
- And or=ed (|) with many optional flags, for example,
  - O\_TRUNC: Truncate the file (remove existing contents) if opening a file for write
  - O\_CREAT: Create a file if it does not exist

Example:

```
// remember open() returns -1 on error
```

## CSE 3100 Master Notes

```
fd1 = open("a.txt", O_RDONLY); // open for read
fd1 = open("a.txt", O_RDWR); // open for read and write
fd1 = open("a.txt", O_RDWR | O_TRUNC); read, write, truncate the file
```

Create a file with open()

```
// a mode must be provided if O_CREAT or O_TMPFILE is set
int open(const char *path, int oflag, int mode);
• mode: specify permissions when a new, or temporary, file is created

open("b.txt", O_WRONLY|O_TRUNC|O_CREAT, 0600);

// open b.txt for write. If the file exists, clear (truncate) the
contents.

// if the file does not exist, create one, and set the permission so
that the owner of the file can read and write, but other people
cannot.
```

File descriptor vs stream

```
#include <stdio.h>
int fileno(FILE *stream);
// returns a file descriptor for a stream
```

| FD | FILE * |
|----|--------|
| 0  | stdin  |
| 1  | stdout |
| 2  | stderr |

File descriptors after fork and exec

- Opened files are NOT AFFECTED by the upgrade operation

```
pid_t pid = fork();
assert(*pid >= 0);
if (pid == 0){
    // Child process can access FDs 0, 1, and 2
    // if execl() is successful, gcc can access FDs 0, 1, and 2
    execlp("gcc", "gcc", "a.c", NULL);
    // If control gets here, execlp() failed.
    // Remember to terminate the child process!
```

```
    return 1;  
}
```

## Lecture 18 - P3: Redirection

---

Mon. Oct. 14, 2019

### Review

- Function open() returns a file descriptor, a non-negative integer
- the file descriptor is used later in funcs. like read() and close()
- Every opened file has a file descriptor
  - stdin: 0
  - stdout: 1
  - stderr: 2
- Files opened in a process remain open after fork() and exec()

### Shell redirections

- Available when executing commands in your shell (e.g. bash)
  - Implemented with the close/open/dup technique

```
$ command <infile >outfile
```

    - < infile: Take input from file infile
    - > outfile: Send output to file outfile
- Other variants
  - >> outfile: Append output to file outfile
  - 2> outfile: Send errors to file outfile
  - &> outfile: Send both output and errors to file outfile
- Read the manual for more variants like 2>>, 2>&1, etc.

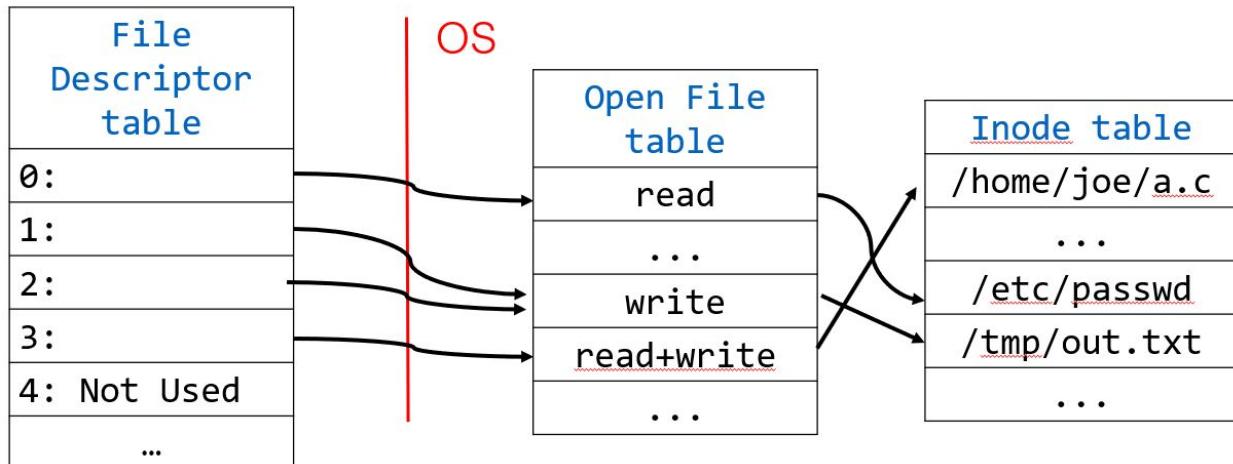
### Shell redirection examples

```
$ sort < file.txt > sorted.txt
```

- sort will read lines from file.txt, instead of terminal
- The output of sort will be saved in sorted.txt
  - You cannot see it on screen
- The statements in sort are not changed
- They read from stdin (0), and print to stdout (1)

## File Descriptor Table

- Each process has a **file descriptor table**
  - Holds indices of entries into the **Open File Table** managed by OS
- The system-wide **Open File Table**
  - Records the mode of the opened files (e.g., reading, writing, appending)
  - Holds index into the **Inode Table** that has the actual file name and location on disk



## Duplicating File Descriptors

- Do not change file descriptor table directly
- Used open() and close() and two new functions

```
#include <unistd.h>

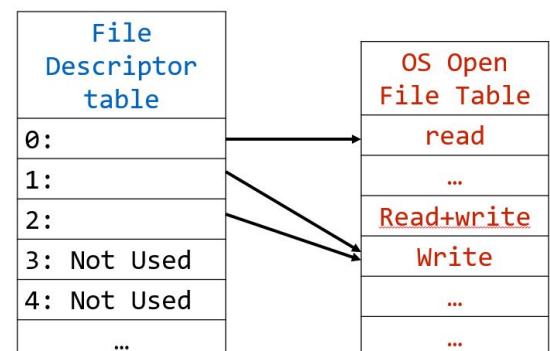
int dup(int oldfd);
int dup2(int oldfd, int newfd);
```

- dup() copies oldfd to the **first available entry** (in FD table)
- dup2() copies oldfd to newfd
  - Closes newfd if it is in use

\*\*\*There is dup3(), but it is not in POSIX. We should not use it in this course.

## Example: stdout redirect

- A program can change its standard output
- How?



## Steps for redirecting stdout

1. open(). open file (and save the file descriptor in fd)
2. dup2(). copy fd to 1 (so that the file descriptor 1 points to the file just opened)
3. close(fd)

## Example: stdout redirect

| <ul style="list-style-type: none"> <li>• Open the new file for writing; 3 is the returned fd             <ul style="list-style-type: none"> <li>◦ We will use 3 instead of a variable in this example</li> </ul> </li> </ul> | <table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th style="text-align: left;">File Descriptor table</th> <th style="text-align: left;">OS Open File Table</th> </tr> </thead> <tbody> <tr><td>0:</td><td>read</td></tr> <tr><td>1:</td><td>...</td></tr> <tr><td>2:</td><td>Read+write</td></tr> <tr><td>3:</td><td>Write</td></tr> <tr><td>4: Not used</td><td>Write</td></tr> <tr><td>...</td><td>...</td></tr> </tbody> </table>    | File Descriptor table | OS Open File Table | 0: | read | 1: | ... | 2: | Read+write | 3: | Write    | 4: Not used | Write | ... | ... |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------|--------------------|----|------|----|-----|----|------------|----|----------|-------------|-------|-----|-----|
| File Descriptor table                                                                                                                                                                                                        | OS Open File Table                                                                                                                                                                                                                                                                                                                                                                                                                                                     |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 0:                                                                                                                                                                                                                           | read                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 1:                                                                                                                                                                                                                           | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 2:                                                                                                                                                                                                                           | Read+write                                                                                                                                                                                                                                                                                                                                                                                                                                                             |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 3:                                                                                                                                                                                                                           | Write                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 4: Not used                                                                                                                                                                                                                  | Write                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| ...                                                                                                                                                                                                                          | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| <pre>// Method 1: two functions. not atomic close(1); dup(3); // Method: better. dup2() closes newfd first dup2(3, 1);</pre>                                                                                                 | <table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th style="text-align: left;">File Descriptor table</th> <th style="text-align: left;">OS Open File Table</th> </tr> </thead> <tbody> <tr><td>0:</td><td>read</td></tr> <tr><td>1:</td><td>...</td></tr> <tr><td>2:</td><td>Read+write</td></tr> <tr><td>3:</td><td>Not used</td></tr> <tr><td>4: Not used</td><td>Write</td></tr> <tr><td>...</td><td>...</td></tr> </tbody> </table> | File Descriptor table | OS Open File Table | 0: | read | 1: | ... | 2: | Read+write | 3: | Not used | 4: Not used | Write | ... | ... |
| File Descriptor table                                                                                                                                                                                                        | OS Open File Table                                                                                                                                                                                                                                                                                                                                                                                                                                                     |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 0:                                                                                                                                                                                                                           | read                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 1:                                                                                                                                                                                                                           | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 2:                                                                                                                                                                                                                           | Read+write                                                                                                                                                                                                                                                                                                                                                                                                                                                             |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 3:                                                                                                                                                                                                                           | Not used                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 4: Not used                                                                                                                                                                                                                  | Write                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| ...                                                                                                                                                                                                                          | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| <pre>close(3);</pre>                                                                                                                                                                                                         | <table border="1" style="border-collapse: collapse; width: 100%;"> <thead> <tr> <th style="text-align: left;">File Descriptor table</th> <th style="text-align: left;">OS Open File Table</th> </tr> </thead> <tbody> <tr><td>0:</td><td>read</td></tr> <tr><td>1:</td><td>...</td></tr> <tr><td>2:</td><td>Read+write</td></tr> <tr><td>3:</td><td>Not Used</td></tr> <tr><td>4: Not Used</td><td>Write</td></tr> <tr><td>...</td><td>...</td></tr> </tbody> </table> | File Descriptor table | OS Open File Table | 0: | read | 1: | ... | 2: | Read+write | 3: | Not Used | 4: Not Used | Write | ... | ... |
| File Descriptor table                                                                                                                                                                                                        | OS Open File Table                                                                                                                                                                                                                                                                                                                                                                                                                                                     |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 0:                                                                                                                                                                                                                           | read                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 1:                                                                                                                                                                                                                           | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 2:                                                                                                                                                                                                                           | Read+write                                                                                                                                                                                                                                                                                                                                                                                                                                                             |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 3:                                                                                                                                                                                                                           | Not Used                                                                                                                                                                                                                                                                                                                                                                                                                                                               |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| 4: Not Used                                                                                                                                                                                                                  | Write                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |
| ...                                                                                                                                                                                                                          | ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                       |                    |    |      |    |     |    |            |    |          |             |       |     |     |

## Implementing redirections

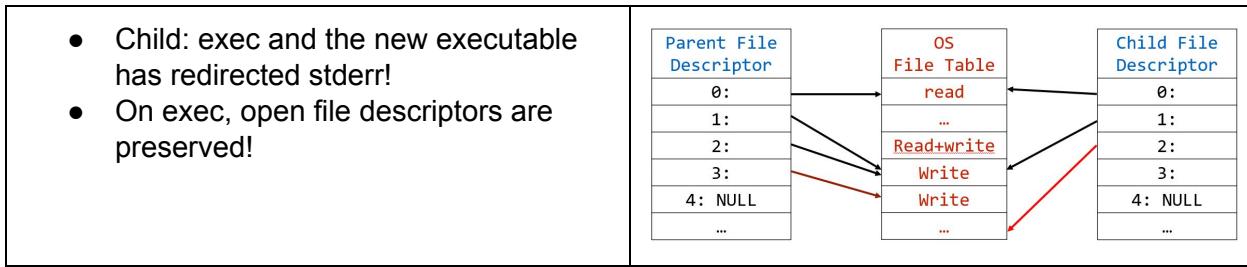
- How does bash do redirection for other processes?
- When bash starts a new process, it does fork() and exec()
- Recall that the file descriptor table **is preserved during fork & exec**

Idea:

- In child process, set up proper file descriptors before upgrading
  - Simply change the files corresponding to stdin, stdout, or stderr

## Example: redirecting stderr for another program

| <ul style="list-style-type: none"> <li>• Assume the parent uses FD 3</li> <li>• After fork(), the child has the same file descriptors as the parent                     <ul style="list-style-type: none"> <li>◦ Close FDs that are not needed!</li> </ul> </li> </ul> | <table border="1"> <tr><th colspan="2">Parent File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> <table border="1"> <tr><th colspan="2">OS File Table</th></tr> <tr><td>read</td><td>← 0:</td></tr> <tr><td>...</td><td>← 1:</td></tr> <tr><td>Read+write</td><td>← 2:</td></tr> <tr><td>Write</td><td>← 3:</td></tr> <tr><td>Write</td><td>← 4:</td></tr> <tr><td>...</td><td>← ...</td></tr> </table> <table border="1"> <tr><th colspan="2">Child File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> | Parent File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  | OS File Table |  | read | ← 0: | ... | ← 1: | Read+write | ← 2: | Write | ← 3: | Write | ← 4: | ... | ← ... | Child File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------|--|----|------------------------|----|--|----|--|----|------------------------------|---------|--|-----|--|---------------|--|------|------|-----|------|------------|------|-------|------|-------|------|-----|-------|-----------------------|--|----|------------------------|----|--|----|--|----|------------------------------|---------|--|-----|--|
| Parent File Descriptor                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| OS File Table                                                                                                                                                                                                                                                          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| read                                                                                                                                                                                                                                                                   | ← 0:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← 1:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Read+write                                                                                                                                                                                                                                                             | ← 2:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 3:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 4:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Child File Descriptor                                                                                                                                                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| <p><b>How do you close FD 3 in child process?</b></p> <ul style="list-style-type: none"> <li>• Child: open the file (to save error output)</li> </ul>                                                                                                                  | <table border="1"> <tr><th colspan="2">Parent File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> <table border="1"> <tr><th colspan="2">OS File Table</th></tr> <tr><td>read</td><td>← 0:</td></tr> <tr><td>...</td><td>← 1:</td></tr> <tr><td>Read+write</td><td>← 2:</td></tr> <tr><td>Write</td><td>← 3:</td></tr> <tr><td>Write</td><td>← 4:</td></tr> <tr><td>...</td><td>← ...</td></tr> </table> <table border="1"> <tr><th colspan="2">Child File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> | Parent File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  | OS File Table |  | read | ← 0: | ... | ← 1: | Read+write | ← 2: | Write | ← 3: | Write | ← 4: | ... | ← ... | Child File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  |
| Parent File Descriptor                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| OS File Table                                                                                                                                                                                                                                                          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| read                                                                                                                                                                                                                                                                   | ← 0:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← 1:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Read+write                                                                                                                                                                                                                                                             | ← 2:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 3:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 4:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Child File Descriptor                                                                                                                                                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| <ul style="list-style-type: none"> <li>• Child:<br/> <pre>dup2(2, 3); close(3); // not yet run</pre> </li> </ul>                                                                                                                                                       | <table border="1"> <tr><th colspan="2">Parent File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> <table border="1"> <tr><th colspan="2">OS File Table</th></tr> <tr><td>read</td><td>← 0:</td></tr> <tr><td>...</td><td>← 1:</td></tr> <tr><td>Read+write</td><td>← 2:</td></tr> <tr><td>Write</td><td>← 3:</td></tr> <tr><td>Write</td><td>← 4:</td></tr> <tr><td>...</td><td>← ...</td></tr> </table> <table border="1"> <tr><th colspan="2">Child File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> | Parent File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  | OS File Table |  | read | ← 0: | ... | ← 1: | Read+write | ← 2: | Write | ← 3: | Write | ← 4: | ... | ← ... | Child File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  |
| Parent File Descriptor                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| OS File Table                                                                                                                                                                                                                                                          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| read                                                                                                                                                                                                                                                                   | ← 0:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← 1:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Read+write                                                                                                                                                                                                                                                             | ← 2:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 3:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 4:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Child File Descriptor                                                                                                                                                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| <ul style="list-style-type: none"> <li>• Child:<br/> <pre>dup2(2, 3); close(3); // now running</pre> </li> </ul>                                                                                                                                                       | <table border="1"> <tr><th colspan="2">Parent File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> <table border="1"> <tr><th colspan="2">OS File Table</th></tr> <tr><td>read</td><td>← 0:</td></tr> <tr><td>...</td><td>← 1:</td></tr> <tr><td>Read+write</td><td>← 2:</td></tr> <tr><td>Write</td><td>← 3:</td></tr> <tr><td>Write</td><td>← 4:</td></tr> <tr><td>...</td><td>← ...</td></tr> </table> <table border="1"> <tr><th colspan="2">Child File Descriptor</th></tr> <tr><td>0:</td><td>→ OS File Table (read)</td></tr> <tr><td>1:</td><td></td></tr> <tr><td>2:</td><td></td></tr> <tr><td>3:</td><td>→ OS File Table (Read+write)</td></tr> <tr><td>4: NULL</td><td></td></tr> <tr><td>...</td><td></td></tr> </table> | Parent File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  | OS File Table |  | read | ← 0: | ... | ← 1: | Read+write | ← 2: | Write | ← 3: | Write | ← 4: | ... | ← ... | Child File Descriptor |  | 0: | → OS File Table (read) | 1: |  | 2: |  | 3: | → OS File Table (Read+write) | 4: NULL |  | ... |  |
| Parent File Descriptor                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| OS File Table                                                                                                                                                                                                                                                          |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| read                                                                                                                                                                                                                                                                   | ← 0:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← 1:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Read+write                                                                                                                                                                                                                                                             | ← 2:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 3:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Write                                                                                                                                                                                                                                                                  | ← 4:                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    | ← ...                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| Child File Descriptor                                                                                                                                                                                                                                                  |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 0:                                                                                                                                                                                                                                                                     | → OS File Table (read)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 1:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 2:                                                                                                                                                                                                                                                                     |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 3:                                                                                                                                                                                                                                                                     | → OS File Table (Read+write)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| 4: NULL                                                                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |
| ...                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |                        |  |    |                        |    |  |    |  |    |                              |         |  |     |  |               |  |      |      |     |      |            |      |       |      |       |      |     |       |                       |  |    |                        |    |  |    |  |    |                              |         |  |     |  |



### Redirecting stdout for a child process

- A process would like to start a new program, and redirect stdout of the new process to a file
  - Where & when should the file be opened?
  - Select the best options.
- Before fork(), in parent
  - After fork() in parent
  - Before exec in child (after fork())
  - After exec in child (after fork())
  - None of the above

### Summary of Steps in Child Process

- Close FDs that are opened in parent and not needed in child
- open(). Open a file (and save the file descriptor in fd)
- dup2(). Copy FD to the right place
- close(fd)
- Exec

### More questions on redirecting stdout?

- Where & when should the new file be opened?
- Where & when should you call close(1)/
- Where & when should you call dup?
- Where & when should close (newfd) be called?

### Summary

- A program can direct input/output
  - It is done with open(), close(), dup(), or dup2()
- FDs are preserved on exec
  - Close file descriptors that are not needed

# Lecture 19 - P4: Inter-Process Comm. w/ Pipes

Weds. Oct. 16, 2019

## Inter-process communication (IPC)

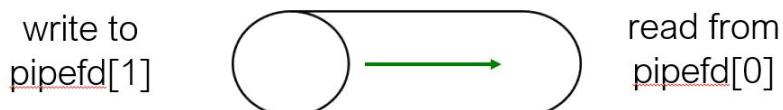
- Files
- Pipes
- Named pipes
- Sockets
- Message queues
- Shared memory
- Synchronization primitives
  - Semaphores, Signal, etc

## pipe()

```
#include <unistd.h>
```

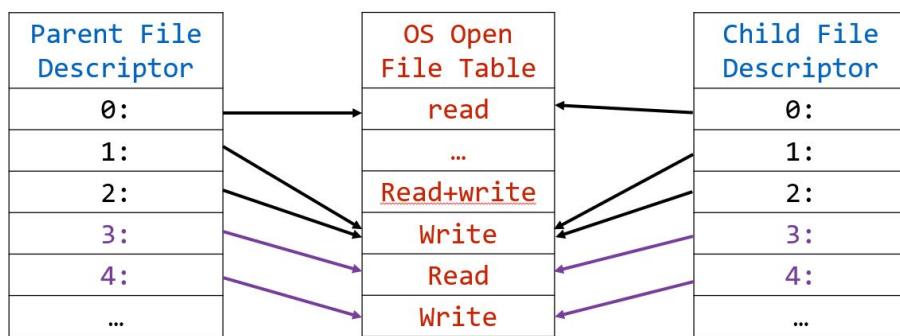
```
int pipe(int pipefd[2])
```

- Create a one-way pipe (a buffer to store a byte stream)
- Two FDs in pipefd. pipefd[0] is the read end, pipefd[1] is the write end
- Return 0 if successful
- Pipes allow IPC. One process writes and the other one reads



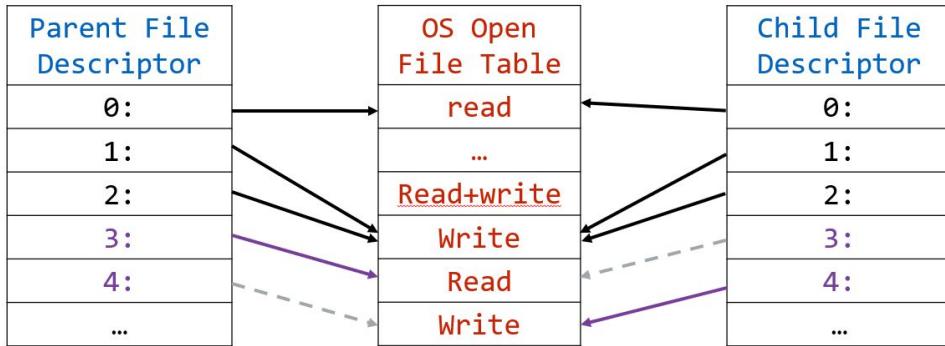
## Connecting two processes

- Parent creates a pipe and gets two FDs (e.g. 3 and 4)
- After fork(), the child has 3 and 4, too
- One process can write to FD 3, and the other can read from FD 4
  - Close unused FD!



## Closing FDs not in use

- If the pipe is for parent to read and for child to write
  - Parent: `close(4)`
  - Child: `close(3);`
- Then the child can write to and parent can read from the pipe. See demo code!



## Questions

- What would you do if you need two-way communication between parent and child?
- After exec, the new program gets the file descriptors for the pipe, too
- How can the new program use the pipe?
  - A program is aware of FDs 0, 1, and 2, but not 3 or 4

## Pipeline in shell

- Shell supports pipelines

```
cmd | cmd2 arg 21 arg22 | cmd3 arg31
```

- stdout of a command is connected to stdin of the next command
  - Done with pipes on Linux/Unix
  - cmd1 writes to a pipe and cmd2 reads from it
- Example:

```
ls | tr a-z A-Z | wc
```

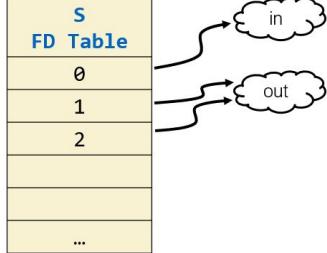
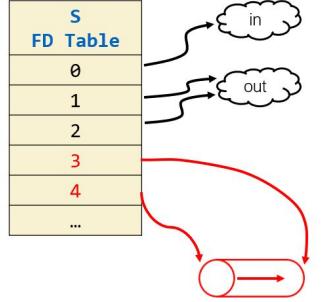
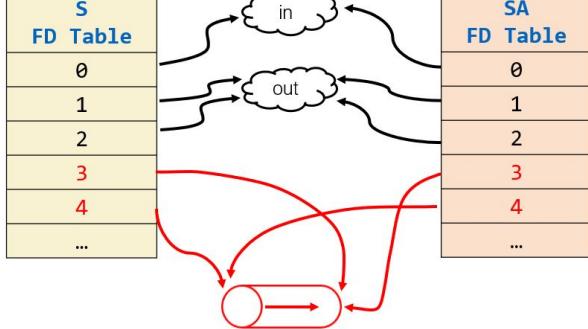
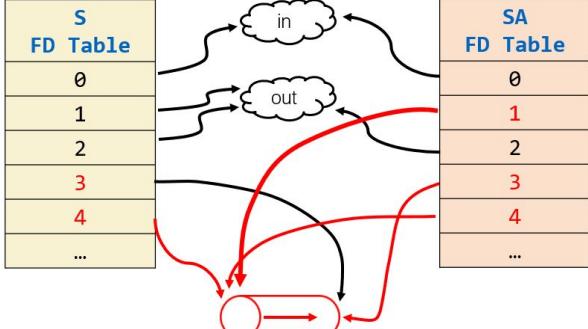
Example: connect two programs with a pipe

**Start a pipeline in program S (aka, the shell): A | B**

- High-level strategy (missing clean up!)
  - Create a pipe
  - Fork #1
    - In child process
      - Redirect stdout to the write end of the pipe
      - Start A, by calling exec
  - Fork #2

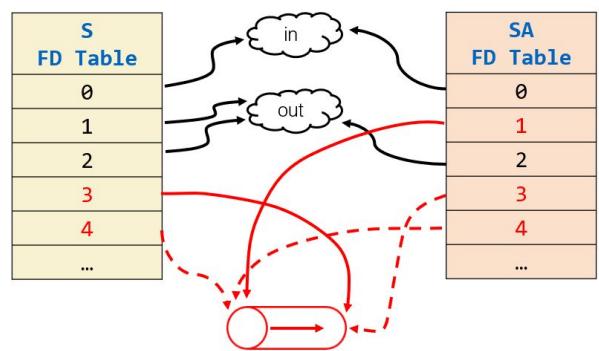
## CSE 3100 Master Notes

- In child process
  - Redirect stdin to the read end of the pipe
  - Start B, by calling exec

|                                                                                                                                                                                                                |                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>At the beginning</b> <ul style="list-style-type: none"> <li>• S has only 0, 1, and 2 open</li> </ul>                                                                                                        |  <p>A diagram showing the initial state of the FD Table for process S. The table has columns for FD number (0, 1, 2, ...) and file descriptor type (in, out). FD 0 points to 'in', FD 1 points to 'out', and FD 2 points to another 'out' entry. A red oval at the bottom indicates the current position.</p>                                                            |
| <b>Pipe Creation</b> <ul style="list-style-type: none"> <li>• S creates a pipe by calling pipe()                     <ul style="list-style-type: none"> <li>◦ A pair of FDs is returned</li> </ul> </li> </ul> |  <p>A diagram showing the state of the FD Table after creating a pipe. Process S now has FD 3 (read end) and FD 4 (write end) in its table. The pipe is represented by a red oval with two arrows pointing between FD 3 and FD 4. A red arrow also points from FD 4 back to the current position indicator.</p>                                                          |
| <b>Fork #1</b> <ul style="list-style-type: none"> <li>• S: fork()                     <ul style="list-style-type: none"> <li>◦ FD table is duplicated</li> </ul> </li> </ul>                                   |  <p>A diagram showing the state of both the parent (S) and child (SA) processes after a fork. Both have identical FD tables with FD 0 (in), FD 1 (out), FD 2 (out), FD 3 (read end), and FD 4 (write end). The pipe connection is shown with red arrows between the corresponding FD pairs in both tables. Red ovals indicate the current positions in each table.</p> |
| <b>Redirect in first child process</b> <ul style="list-style-type: none"> <li>• SA: dup2(4, 1)                     <ul style="list-style-type: none"> <li>◦ Or close(1); dup(4);</li> </ul> </li> </ul>        |  <p>A diagram showing the state of the parent (S) and child (SA) processes after redirecting FD 1. In the child's FD table, FD 1 now points to the pipe's read end (FD 3). The pipe connection is shown with red arrows between FD 4 (write end) in the parent and FD 3 (read end) in the child. Red ovals indicate the current positions.</p>                         |

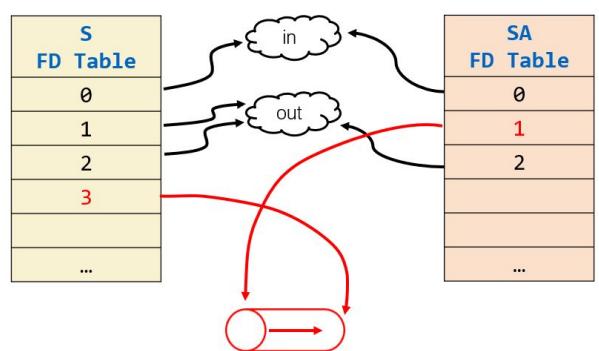
### Clean up #1

- S: close(4)
- SA: close(4), close(3)
- SA can then exec into A



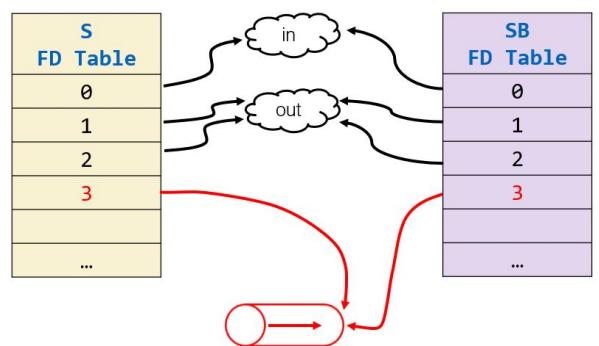
### After clean up #1

- S: close(4)
- SA: close(4), close(3)
- SA can then exec into A



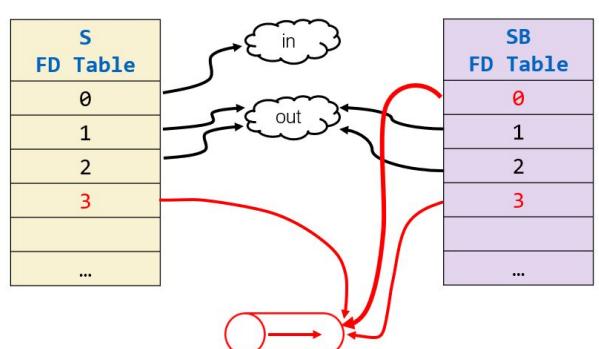
### Fork #2

- S: fork()
  - Note that 4 has been closed in S



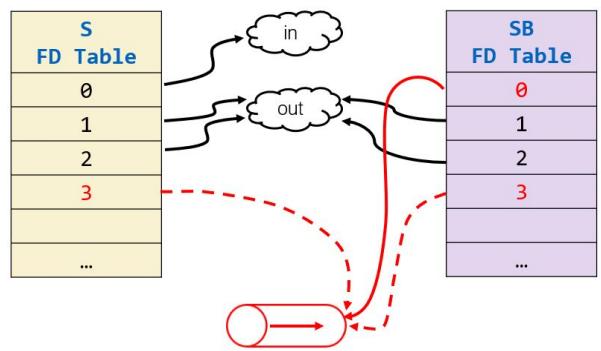
### Redirect in second child process

- SB: dup2(3,0)
  - or close(0); dup(3);



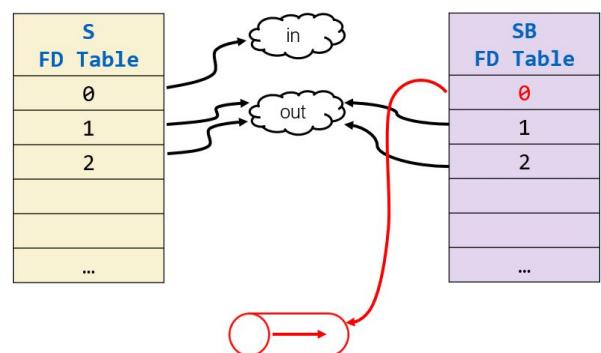
**Clean Up #2**

- S: close(3)
- SB: close(3)
- SB can then exec into B



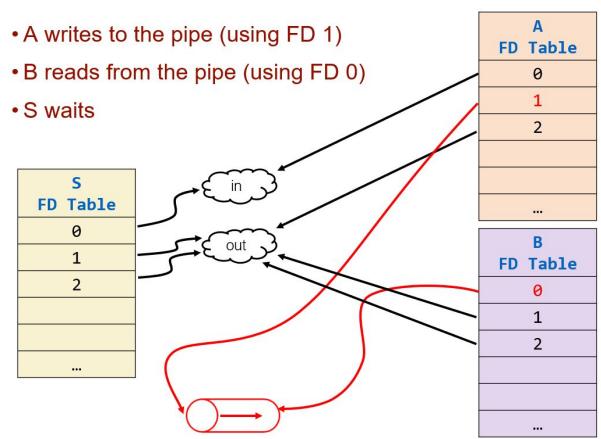
**After clean Up #2**

- S: close(3)
- SB: close(3)
- SB can then exec into B



**Final set up**

- A writes to the pipe (using FD 1)
- B reads from the pipe (using FD 0)
- S waits



FDs of a dying process

- When a process ends, all its open FDs are automatically closed
- What happens to the processes on the other end of the pipe?

**Example:**

**Assume S does not read or write, but have FDs of the pipe**

- If both A and S die, B gets EOF when all buffered data are consumed
- If A dies, B will wait for more data (assuming S may write)
- If both B and S die, A gets an error (SIGPIPE) when writing
- If B dies, A will wait if the pipe is full (assuming S will read)



Going further...

- You can repeat this to create a long pipeline
  - e.g., connect B's stdout to stdin of another process C
- Draw pictures to find how pipes are used
  - And what FDs need to be closed

### Remember

- Processes are running in parallel once they are created
  - Although we showed the operations in sequence
- All processes in the pipeline are running concurrently on Linux
  - As soon as data are sent in the pipe...
  - The next process can pick them up on the work

Atomicity of read() and write()

```
n_r = read(fd, buf, N);
n_w = write(fd, buf, N);
    • write() and read() returns the # of bytes actually read/written
    • The number maybe less than the requested
```

**Atomicity** - The degree a program is guaranteed to be isolated from other operations that may be happening at the same time

- Atomicity of write() is guaranteed if the # of bytes < PIPE\_BUF
  - The bytes will be consecutive
  - The default value of PIPE\_BUF is 4096 on Linx
- For read(), it is fine if all writes and reads are of the same size
  - Otherwise, need special handling

Starting a 2-stage pipeline - 1

```
// A | B

pipe(pipefd)           // pipefd is an array of 2 int's
pid_a = fork()          // for A
if (pid_a == 0){        // child process for A
    dup2();             // setup stdout for A
    close both FDs in pipefd
    exec to start A     // remember to exit from child on error
```

```

}

close(pipefd[WR_END]);      // No need to keep it open in parent

```

## Starting a 2-stage pipeline - 2

```

pid_b = fork();           // for B
if (pid_b == 0){          // child process for B
    dup2();               // setup stdin for B
    close(pipefd[RD_END]);
    exec to start B      // remember to exit from child on error
}
close(pipefd[RD_END]);    // No need to keep it open in parent

```

## Using Pipes to Sum Matrix Rows Concurrently

- See the complete code in the demo repo.

```

int main(void)
{
    int i, row_sum, sum = 0, pd[2], a[N][N] = {{1, 1, 1}, {2, 2, 2}, {3, 3, 3}};

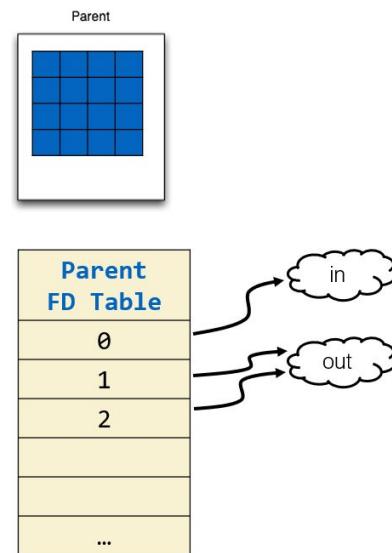
    if (pipe(pd) == -1) error_exit("pipe() failed"); /* create pipe */

    for (i = 0; i < N; ++i)
        if (fork() == 0) { /* create a child process for each row */
            row_sum = add_vector(a[i]); /* compute the sum of a row */
            if (write(pd[1], &row_sum, sizeof(int)) == -1) /* write to pipe */
                error_exit("write() failed");
            return 0;                                /* exit from child */
        }
    /* better to close the write end in the parent */
    for (i = 0; i < N; ++i) {
        if (read(pd[0], &row_sum, sizeof(int)) == -1) /* read from pipe */
            error_exit("read() failed");
        sum += row_sum;                            /* calculate the total */
    }
    printf("Sum of the array = %d\n", sum);
    /* wait for child processes*/

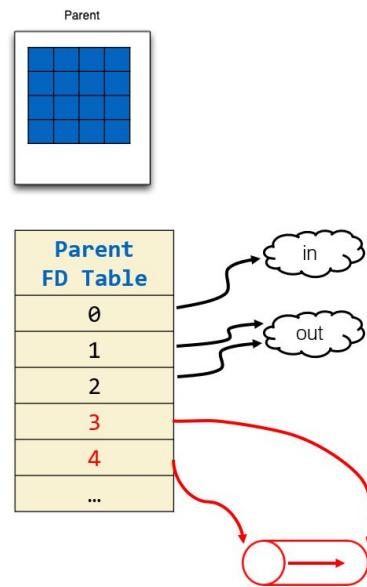
```

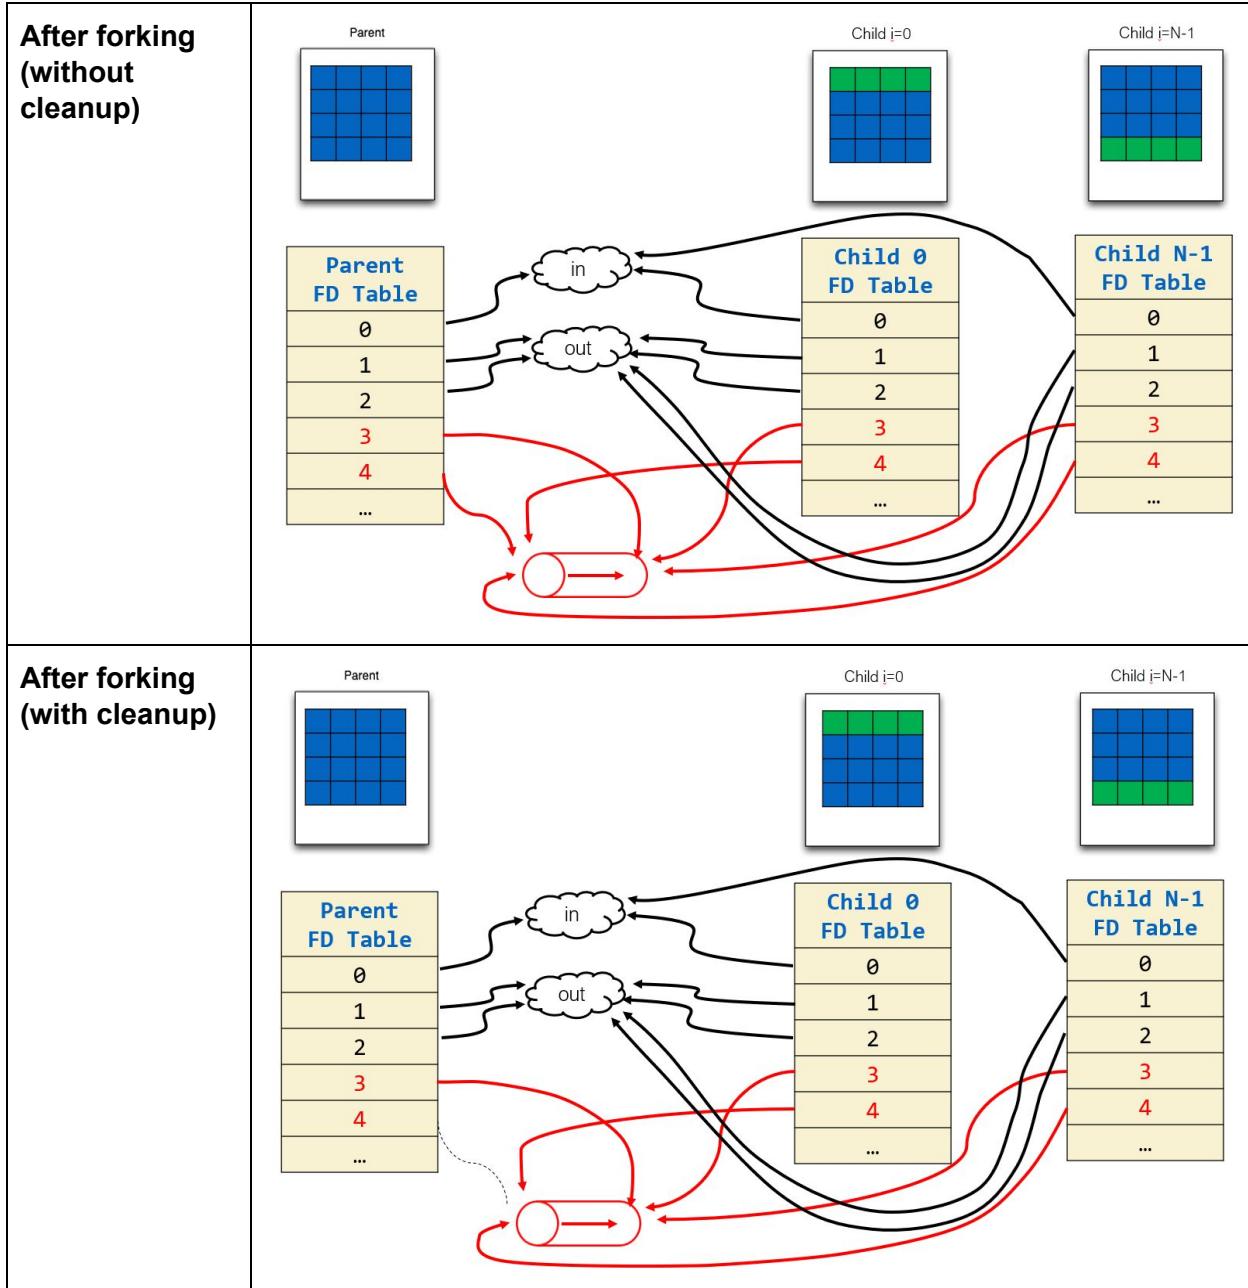
## CSE 3100 Master Notes

### Parent Process



### Pipe Creation





## Lecture 20 - T1: Instruction & Basic Management

Mon. Oct. 21, 2019

### Overview

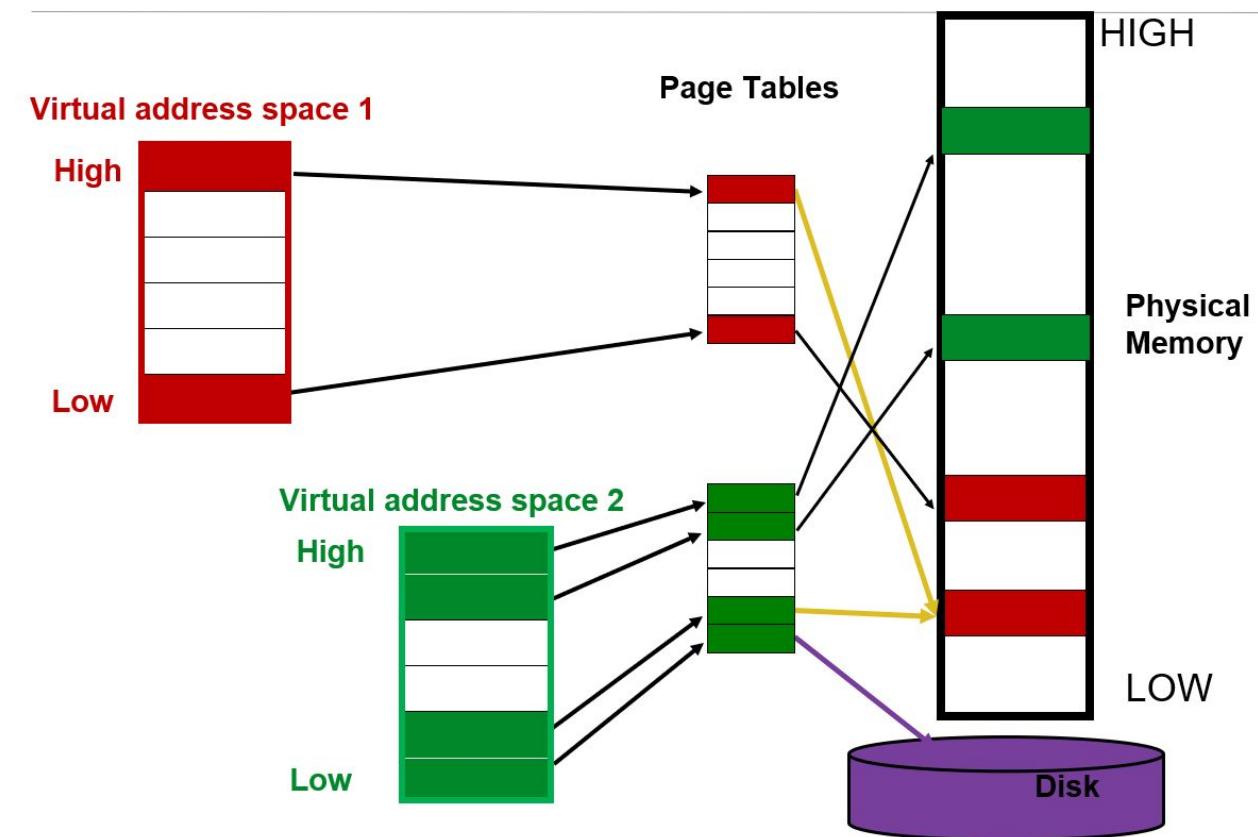
- Motivation & Concurrency

- Thread abstraction
- Thread creation & termination

## Motivation

- Processes for concurrency
  - Process do run concurrently on O.S.
    - On multi-core or single-core processors, through time-sharing
- Processes also provide protection

Each process has its own virtual address space



## Why?

- Well...
  - Nobody can interfere
  - Nobody can move his “stuff”
  - Nobody can play with his “stuff”
  - Nobody can break his toys
- It is very very safe!

### Downsides?

- Nobody to play with
  - You can get bored quite quickly
- Nobody to do your chores
  - Do everything yourself!
- You can't leave the castle
  - It is also your prison
- Communication with the outside is tricky and very limited (but safe)
  - Banging on a pipe...
  - Smoke signals...
  - Sockets...

### Where are her buddies (e.g., processes)?

- In other castles!
  - Equally alone / isolated
  - You cannot get together easily
  - Sharing is limited

### Threads?

- Inviting other ‘living creatures’ inside your castle
  - These “creatures” can
    - Move around independently of you
    - Do work / chores on your behalf
    - Communicate on your behalf
    - Play with you
    - Use/share your toys
- Essentially
  - No limits on their abilities in their address space
  - As powerful as the “Lord of the Castle”

### Refining our definitions

- Process is a bundle grouping
  - A virtual address [memory]
  - A collection of files / sockets [IO]
  - A collection of concurrent threads [execution units]
- Thread is a light-weight entity
  - Can run concurrently w/ other threads
  - Share resources in the process [having same rights]

## CSE 3100 Master Notes

- Confined to a single process [cannot “move” to another]
- Can be created and destroyed [different life cycles]

**Lightweight ‘processes’ that share the address space (and other resources in a process)**

### The Circle of Life

- How is a thread created
  - By another thread! `pthread_create()`
  - It is given a stack to execute and a function to run
- What about the 1st thread?
  - When a process is created, there is a single thread [starts alone]
- How does a thread die?
  - Voluntarily, after completing its task `pthread_exit()`
  - Requested by another thread `pthread_cancel()`
  - When the process dies (along with all threads in it)
    - Any thread calls `exit()`
      - Note that `exit()` is called when `main()` returns!

**!!! Any thread can bring down the whole castle !!!**

### The pthreads API

- ANSI/IEEE POSIX 1003.1 - 1995 standard
- These types of routines:
  - Thread management: create, terminate, join, and detach
  - Mutexes: mutual exclusion, creating, destroying, locking, and unlocking mutexes
  - Condition variables: event driven synchronization

### The Pthreads API naming convention

| Routine Prefix                  | Function                        |
|---------------------------------|---------------------------------|
| <code>pthread_</code>           | General pthread                 |
| <code>pthread_attr_</code>      | Thread attributes               |
| <code>pthread_mutex_</code>     | mutex                           |
| <code>pthread_mutexattr_</code> | Mutex attributes                |
| <code>pthread_cond_</code>      | Condition variables             |
| <code>pthread_condattr_</code>  | Conditional variable attributes |
| <code>pthread_key_</code>       | Thread specific data keys       |

### The Pthreads API

```
# include <pthread.h>
• Add '-pthread' option to compile on Linux
    cc -pthread a.c
• PThread functions do not set errno on errors
• Many types are defined in pthread library. They are opaque objects
    ○ Cannot make assumptions on the representation/implementation
```

For ex.): **should use pthreads\_equal() to compare two thread IDs**

### main() function of thread

```
void * thread_main(void * arg)
• The "main" function of a thread
    ○ Can be any name you like
• Takes a pointer as the only parameter
    ○ It can point to anything, int, char, string, or a structure
• Return a pointer
    ○ It can point to anything, int, char, string, and a structure
    ○ However, do not point to local variables on stack -
```

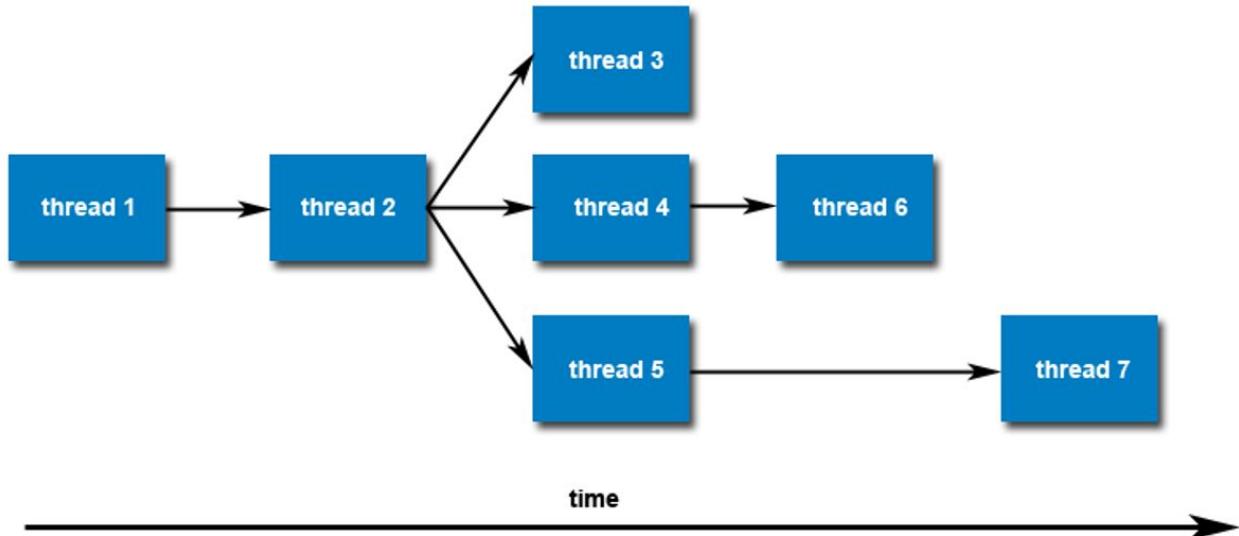
### Threads creation

```
int pthread_create(
    pthread_t * thread,
    pthread_attr_t * attr,
    void * (*start_routine)(void*),
    void * arg);
```

- Returns 0 if successful, and non-zero (> 0) if error
- \*thread is the returned thread ID, if successful
- attr specifies the attribute for the thread. NULL for default
- start\_routine()
- void \* arg is passed to start\_routine()
- Thread equivalent of fork(), but it does create (not "clone")

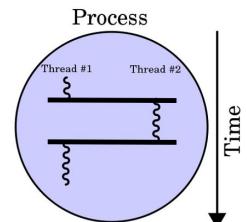
### Thread creation

- Once created, threads are peers, and may create other threads
- There is no implied hierarchy or dependency
  - All threads are equal!



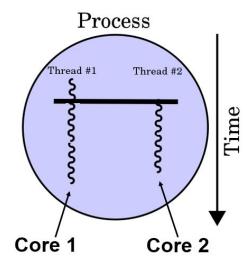
What about concurrency?

- Single-core CPU: Timesharing
  - When one thread is waiting for an IO to complete...
  - ... another thread can use the CPU



What about thread concurrency?

- Multi core CPU: true concurrency
  - MIMD architecture: multiple instructions, multiple data
  - Threads can execute in parallel, one on each core
  - OS can still preempt threads
  - Useful when #threads >> #cores → timesharing is still used!



Thread Termination

- Return from the `start_routine` function, or
- Call `pthread_exit()`

```
void pthread_exit(void* status)
```
- The function always succeeds and does not return
- Status can be obtained by other threads
- Similar to process termination
  - `main()` returns, or `exit()` called by any thread

Joining a thread

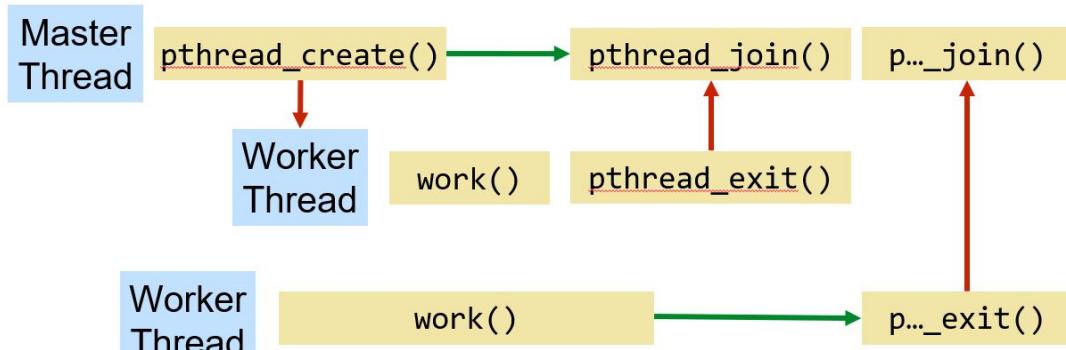
```
int pthread_join(pthread_t tid, void** status)
```

- Wait for a thread to complete
  - Blocks the calling thread until thread id terminates

- Can obtain the exit status of the thread
  - Pass NULL to ignore the return value
  - Why is the type of status void (\*\*)? - A pointer to a pointer that can be changed
- Equivalent of `waitpid()` for processes

### Joining a thread - 2

- Joining is a simple way to accomplish synchronization
  - The calling thread can obtain the target thread's termination return status if it was specified in the target thread's call to `pthread_exit()`



### Reading

- Book "Programming with POSIX Threads"
  - Chapter 1, 2, 3, 4,
  - Chapter 5, Section 1- 4
- Nice to read too
  - Chapter 6, Sectoins 1-5
  - Chapter 8 [debugging]
  - Chapter 9 [reference]
  -

### Passing Arguments and Getting Results Back

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

#define NUM_THREADS 8

struct thread_data
{
    int thread_num;
    char* message;
    int len;
};

void* PrintHello(void* threadarg)
{
    struct thread_data* my_data = (struct thread_data*) threadarg;
    sleep(1 + 5*(my_data->thread_num % 2) );
    my_data->len = strlen(my_data->message);
    printf("Thread #%-d: %s length=%d\n", my_data->thread_num, my_data->message,
        my_data->len);
    pthread_exit(NULL);
}
```

```

int main(int argc, char* argv[])
{
    pthread_t threads[NUM_THREADS];
    struct thread_data thread_data_array[NUM_THREADS];
    char* messages[NUM_THREADS];
    int rc, t;

    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvuyte, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";

    for( t=0; t<NUM_THREADS; t++ ) {
        thread_data_array[t].thread_num = t;
        thread_data_array[t].message = messages[t];
        printf("Creating thread #%-d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, &thread_data_array[t]);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    int grand_total = 0;
    for( t=0; t<NUM_THREADS; t++ ) {
        printf("Joining thread #%-d\n", t);
        rc = pthread_join( threads[t], NULL );
        if( rc ){
            printf("ERROR; return code from pthread_join() is %d\n", rc);
            exit(-1);
        }
        grand_total += thread_data_array[t].len;
    }
    printf("Grand total = %d\n", grand_total);
    pthread_exit(NULL);
}

```

28

## Common ways to use threads

- Pipeline
  - A task is broken into a series of sub-operations, each of which is handled in series, but concurrently, by a different thread (think automobile assembly line)
- Manager/worker

## CSE 3100 Master Notes

- A single thread, the manager assigns work to other threads, the workers. The manager handles all input and parcels out work to workers
- Two common forms: static worker pool and dynamic worker pool
- Peer
  - Similar to the manager/worker model, but after the main thread creates other threads, it participates in the work

### Applications of threads

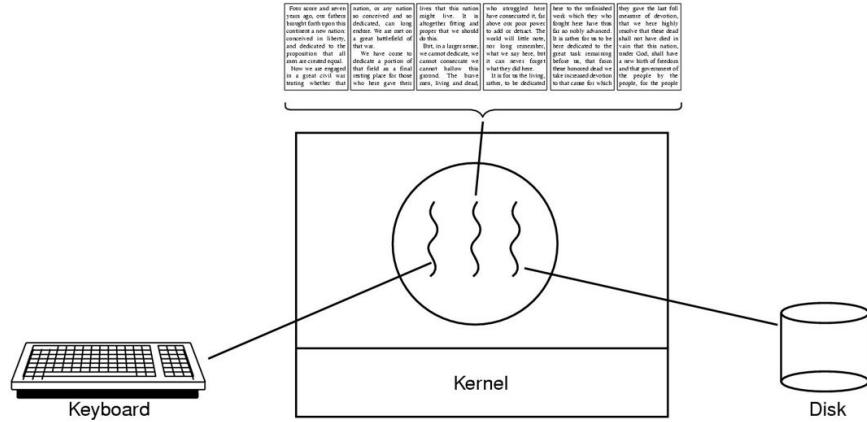
- Parallel computing
  - On multi-core machines, threads can be executed at the same time
- Overlap CPU work with I/O
  - While one thread is waiting for an I/O, others can perform CPU work
- Asynchronous event handling
  - e.g., a web server can both transfer data from previous requests and manage the arrival of new requests
- Priority/real-time scheduling
  - Important tasks can be scheduled with higher priority
- Computer games
  - Each thread controls the movement of an object

### Why is it hard for the princess to do everything by herself?

- Imagine an application like Microsoft Word
- When you are typing, the process needs to
  - Listen for keystrokes and display the text on screen
  - Save data periodically to disk so you do not lose data in case of crash
  - Reform the document while edit the text
- How can a single process (i.e., one princess) can do all these concurrently?

**The answer is she cannot without affecting application performance**

We solve it by creating multiple threads!



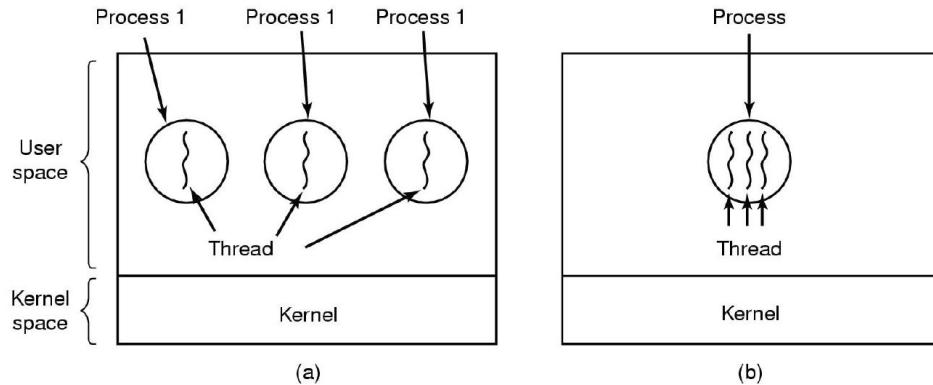
A word processor with three threads.

## Usage Examples?

- We will create multiple threads, one thread for each sub tasks
  - One thread controls a GUI
  - One thread reformats the document
  - One thread does background tasks (e.g., savings!)
- Other examples
  - Worker threads do parallel computations
    - e.g., in matrix-vector multiplication: do all the rows in parallel
    - e.g., simulate agents in parallel [think agents in games!]

## The Classical Thread Model(1)

- a. Three processes each with one thread
- b. One process with threads



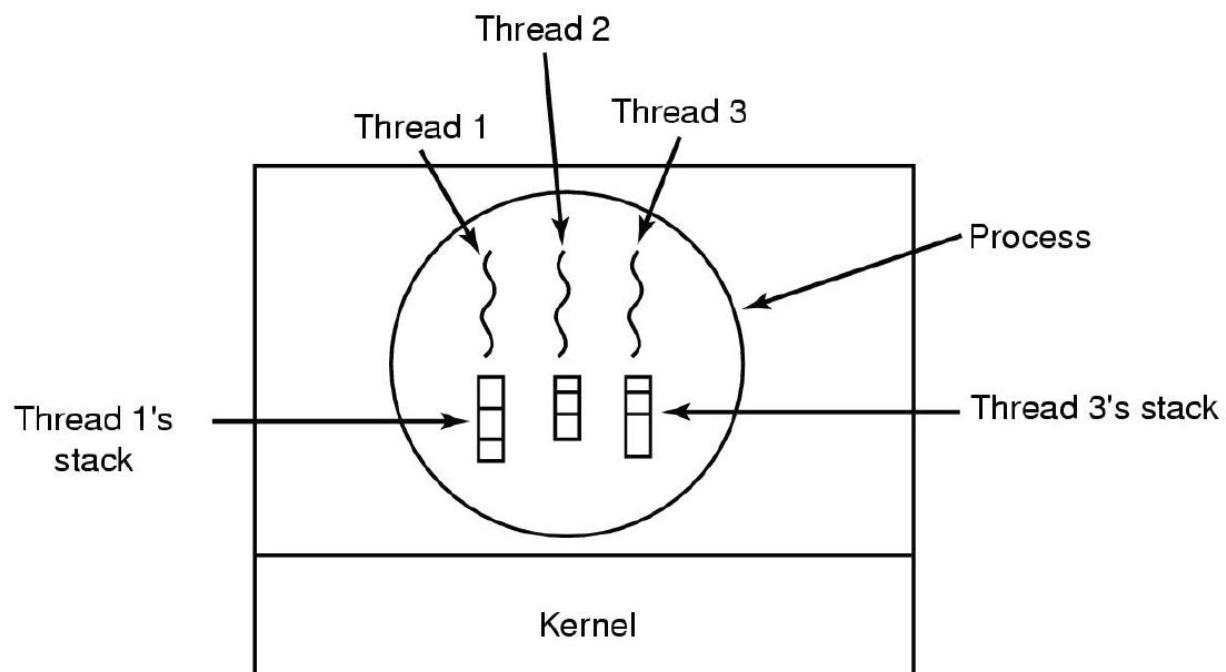
## Classical Thread Model (2)

| Per process items           | Per thread items |
|-----------------------------|------------------|
| Address space               | Program counter  |
| Global variables            | Registers        |
| Open files                  | Stack            |
| Child processes             | State            |
| Pending alarms              |                  |
| Signals and signal handlers |                  |
| Accounting information      |                  |

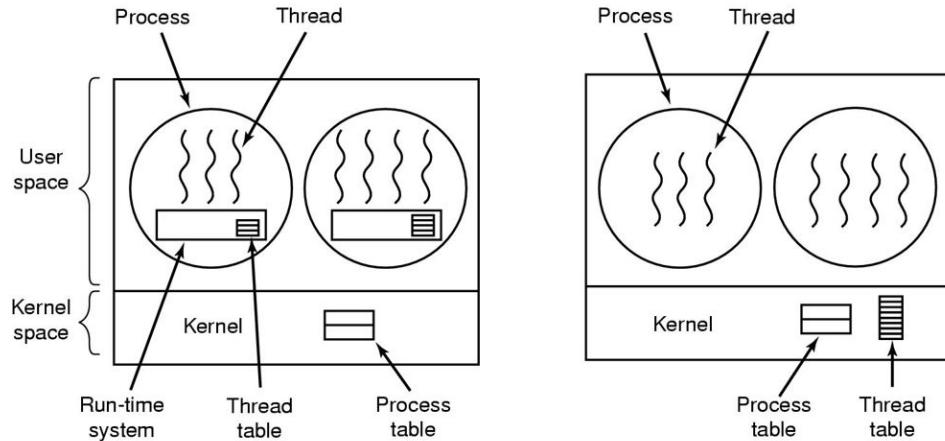
Shared by all threads  
in a process.

Private to each thread.

## Classical Thread Model (3)



## Implementing Threads



- a. A user-level threads package
- b. A threads package managed by the kernel

## The Plot

- What you wish to share (the toy) is
  - The variable  $x$
  - it exists somewhere in the shared virtual address space
  - It has no special status (compared to other regions!)
- The protagonist who both want to play with the toy are
  - thread 1
  - thread 2

## Your Objective

- Protect the toy( $x$ )
  - Make sure nothing “bad” happens to it
  - In particular, make sure it has the correct value at the end
- Recall
  - $x$  exists in memory
  - So you wish to protect memory (an integer)

## Idea

- You cannot protect the memory hold “ $x$ ”
- But...
- You can specify a protocol for everyone using “ $x$ ”
  - If a thread wishes to increase “ $x$ ” ...
    - It must grab a specific “lock”

[lock]

## CSE 3100 Master Notes

- If it has the “lock”, it can do what it wants to x [critical section]
- When does with x it must release the “lock” [unlock]
- If it does not have the lock, it must wait

You do not protect x directly

Instead, you discipline the code that touches x

Idea 1 - Is this slide necessary? - Maifi

- Put x in a “mini-prison”
  - If thread 1 wishes to increment “x”
    - It must take “x” out of its prison, increment it, and put it back
  - If thread 2 wishes to increment “x” and the cell is empty...
    - It must wait until “x” is back in its cell!
- Does this work? [conceptually]
- Does this work? [practically]

32-bit Implication

- Virtual address Space size on 32-bit OS
  - Linux: 2G
  - Windows: 1G
- Typical stack size per thread
  - 8 Megs
- Memory usage goes to
  - Executable: ~ 1 to 50 megs
  - Heap: ~ 1 to 200 megs
  - Stacks: ~ # of threads \* 8 → 100 threads yield 800 megs
- Total near the limit of the address space size.

Circumventing the limit?

- Several “ways” to hop along
  - Make smaller stacks! [but beware of recursion]
  - Separate tasks in several process that
    - Communicate via pipes
    - Communicate via shared virtual memory
  - Use a 64-bit OS!
    - Remember 8 megs =  $2^{23}$
    - Address space size =  $2^{64}$ 
      - What's left is still:  $2^{63}$
    - How many stacks can you have?  $\sim 2^{40}$

### Threads and MIMD Architectures

- **MIMD** - Multiple Instruction Multiple Data
- Process has
  - Multiple execution units
  - All executing independently
  - All executing diff. instructions
  - All operating on diff. pieces of data

### Threads and MIMD Architecture

- Threads
  - Are OS abstractions that capture computation streams
  - Can be scheduled on a MIMD processor
  - All concurrent threads execute diff. instructions on diff. data

### Timesharing

- Idea is simple
  - Diff. pattern of interaction [work / IO] for threads
  - When one thread is “waiting” for an IO to complete...
  - ... another thread can use the CPU for some computing.
  - At any one point in time, CPU is used by only ONE thread
- True concurrency
  - There are > 1 CPUs
  - Threads are executing truly in parallel, one on each CPU
  - OS can still preempt threads
  - Useful when # threads >> # cpus! → time sharing is also used!

### Threads?

- **definition [wikipedia]** - In computer science, a thread of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler (typically as part of an operating system).

### User Level Thread

- Can be implemented on top of OS that does not support threading
  - Each process maintains thread table
  - Per process run time system for thread switching
  - Each process can have its own customized algorithm
- 
- But
    - Blocking system call in one thread can be a problem

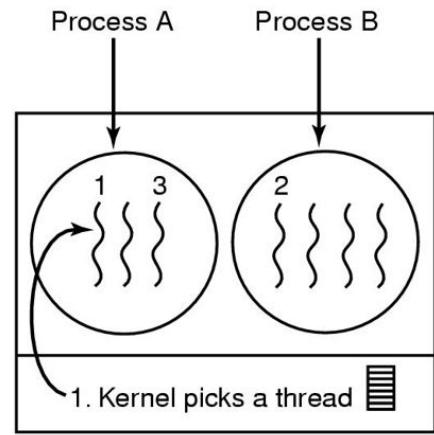
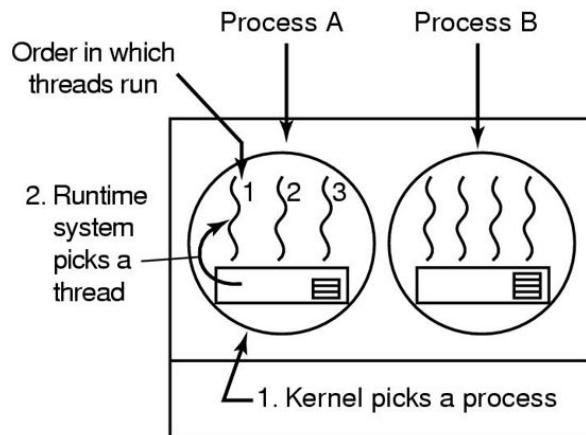
## CSE 3100 Master Notes

- Change all block to non-block call by changing OS
- Within a single process, there is no interrupt. If a thread does not give up voluntarily, other threads in the same will not get a chance
- Threads are mainly for processes that often blocks for I/O
  - User level thread does not make much sense as the goal was to avoid expensive kernel switch. Once switch is done, not much work to switch between threads

### Kernel Level Thread

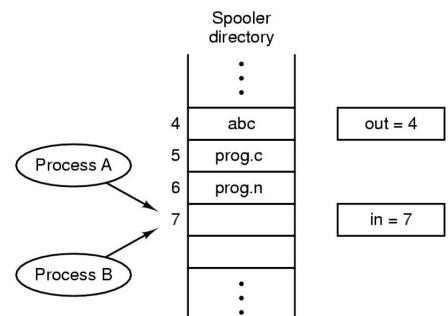
- Kernel has a thread table for all threads
  - Hold thread specific register, stack, state
- Thread recycling
  - Instead of deleting thread data structure, preserve it and reuse it
- But
  - What happens when multithreaded process forks?
  - Signals are per process. Which thread should receive the signal?

Let's check...



### Threads & Castles

- In reality
  - When a process is created there is only 1 living creature (1 thread)
  - The first thread no diff. from other threads
  - All threads have equal rights (unless otherwise specified)



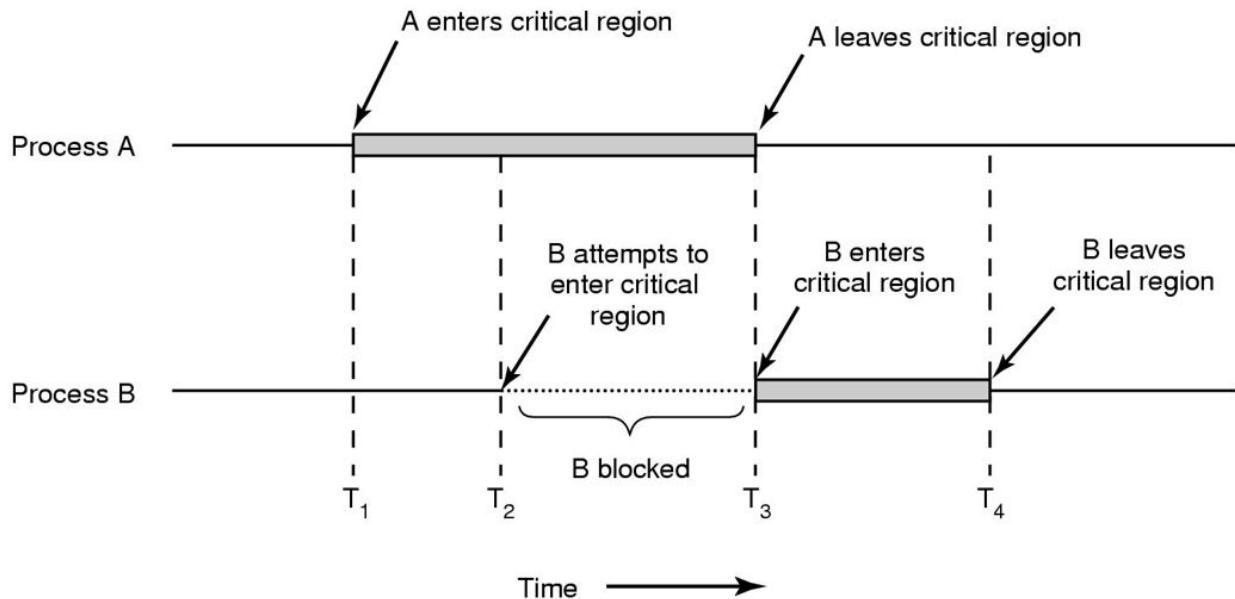
## Race Conditions

- Two processes want to access shared memory at the same time

## Critical Regions (1)

### Conditions required to avoid race condition:

- No two threads may be simultaneously inside their critical regions
- No assumptions may be made about speed or the # of CPUs
- No threads running outside its critical region may block other processes
- No threads should have to wait forever to enter its critical region



## Critical Regions (2)

- Mutual exclusion using critical regions.

## Lecture 21 - T2: Resource Sharing

Weds. Oct. 23, 2019

## Review

```
#include <pthread.h>
// Compile and link with '-pthread'
```

```

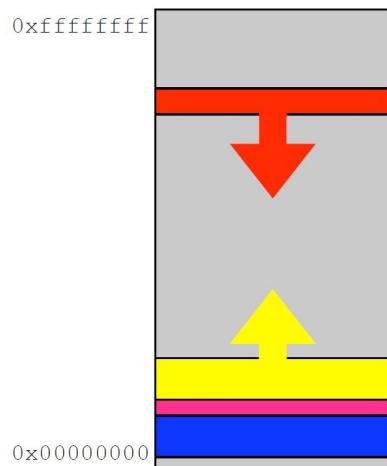
int pthread_create(  pthread_t* thread,
                    pthread_attr_t* attr,
                    void* (*start_routine) (void*) ,
                    void* arg);
// Terminating itself or return from start_routine()
void pthread_exit(void *retval);

// avoid zombie threads
int pthread_join(pthread_t thread, void ** retval);

```

## Virtual Address Space

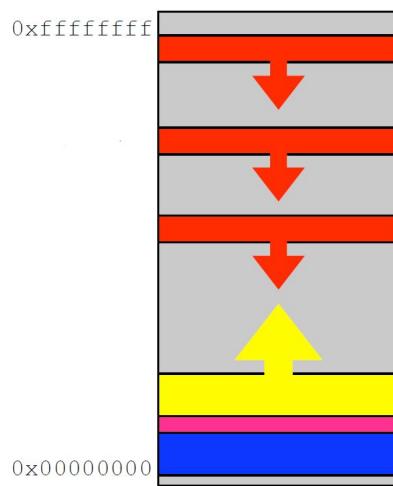
### Process with one thread



**Each thread gets a stack** - Stacks have a max size -- keeps stacks separated to avoid accidental overlap

### Process with 3 threads

- Global data is shared



### 32-bit Implication

- Virtual address Space size on 32-bit OS
  - Linux: 2G
  - Windows: 1G
- Typical stack size per thread
  - 8 Megs
- Memory usage goes to
  - Executable: ~ 1 to 50 megs
  - Heap: ~ 1 to 200 megs
  - Stacks: ~ # of threads \* 8 → 100 threads yield 800 megs
- Total near the limit of the address space size.

### Circumventing the limit?

- Several “ways” to hop along
  - Make smaller stacks! [but beware of recursion]
  - Separate tasks in several process that
    - Communicate via pipes
    - Communicate via shared virtual memory
  - Use a 64-bit OS!
    - Remember 8 megs =  $2^{23}$
    - Address space size =  $2^{64}$ 
      - What's left is still:  $2^{63}$
    - How many stacks can you have?  $\sim 2^{40}$

### Advantages of Threads

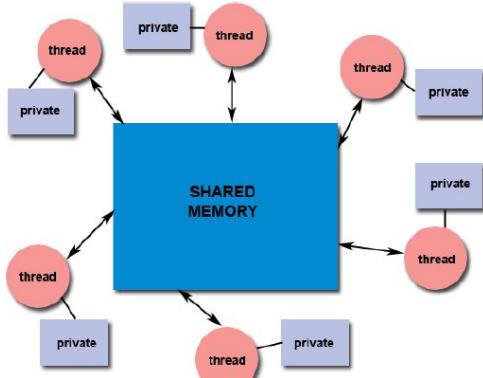
- Simpler programming model
- Easier to coordinate (shared address space, and data)
- Lighter weight than processes
  - Takes less resources than process to manage
- In case of substantial CPU and I/O, thread improves performance
- Light-weight
  - Lower overhead for thread creation
  - Lower context switching overhead
  - Fewer OS resources

### Time (sec.) for creating 50,000 processes/threads

| Platform                               | fork() |      |      | pthread_create() |      |     |
|----------------------------------------|--------|------|------|------------------|------|-----|
|                                        |        |      |      | real             | user | sys |
|                                        | real   | user | sys  | real             | user | sys |
| AMD 2.4 GHz Opteron (8cpus/node)       | 17.6   | 2.2  | 15.7 | 1.4              | 0.3  | 1.3 |
| IBM 1.9 GHz POWER5 p5-575 (8cpus/node) | 64.2   | 30.8 | 27.7 | 1.8              | 0.7  | 1.1 |
| IBM 1.5 GHz POWER4 (8cpus/node)        | 104.1  | 48.6 | 47.2 | 2.0              | 1.0  | 1.5 |
| INTEL 2.4 GHz Xeon (2 cpus/node)       | 55.0   | 1.5  | 20.8 | 1.6              | 0.7  | 0.9 |
| INTEL 1.4 GHz Itanium2 (4 cpus/node)   | 54.5   | 1.1  | 22.2 | 2.0              | 1.3  | 0.7 |

<https://computing.llnl.gov/tutorials/pthreads>

- Shared State
  - Simpler programming model
  - Don't need IPC-like mechanism to communicate between threads



The larger, the better.

| Platform                   | MPI Shared Memory Bandwidth (GB/sec) | Pthreads Worst Case Memory-to-CPU Bandwidth (GB/sec) |
|----------------------------|--------------------------------------|------------------------------------------------------|
| Intel 2.6 GHz Xeon E5-2670 | 4.5                                  | 51.2                                                 |
| Intel 2.8 GHz Xeon 5660    | 5.6                                  | 32                                                   |
| AMD 2.3 GHz Opteron        | 1.8                                  | 5.3                                                  |
| AMD 2.4 GHz Opteron        | 1.2                                  | 5.3                                                  |
| IBM 1.9 GHz POWER5 p5-575  | 4.1                                  | 16                                                   |
| IBM 1.5 GHz POWER4         | 2.1                                  | 4                                                    |
| Intel 2.4 GHz Xeon         | 0.3                                  | 4.3                                                  |
| Intel 1.4 GHz Itanium 2    | 1.8                                  | 6.4                                                  |

### Example: array sum

- Use two threads to compute the sum of integers in an array

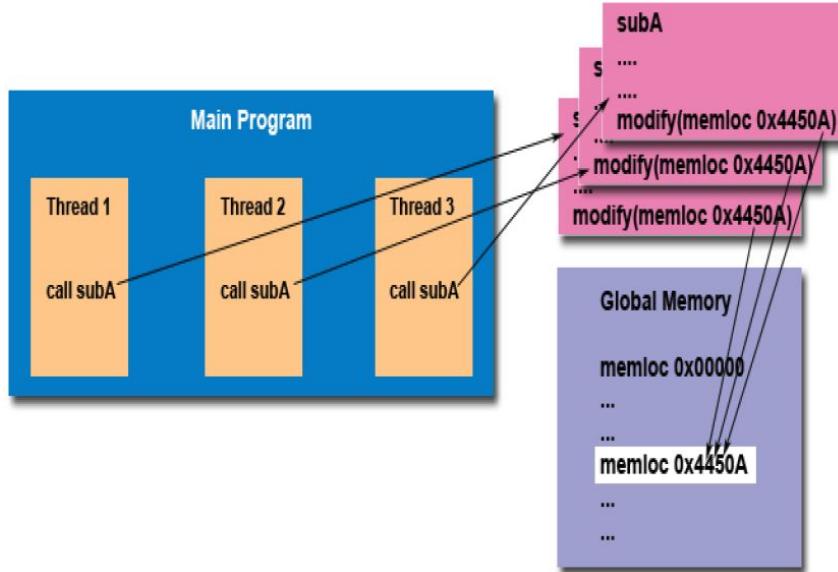
### Disadvantages of threads: What if:

- A thread something stupid like
  - A division by zero?
  - Dereferencing a null pointer?
  - Corrupting a block of memory?

- Access a bad file descriptor?

**The entire process crashes and burns!**

Disadvantages of threads: Shared State.

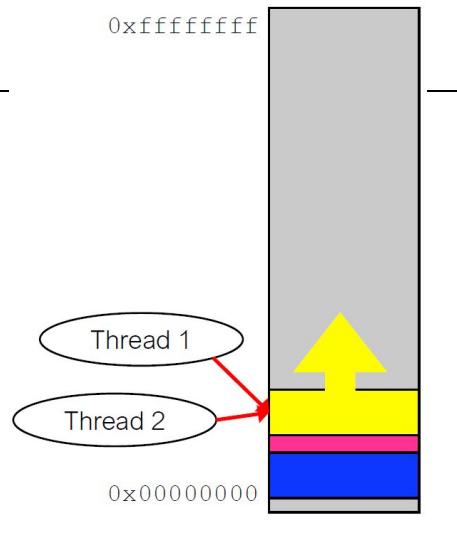


**You may not be aware of shared resources!**

Shared data: global, heap, and even local

```
int a[100];
// a, defined outside of functions,
// can be accessed in all threads.

int main(void)
{
    int i;      // stack
    char * p = malloc(1000); // heap
    void * arg = &i;
    // pass p or arg to threads
}
```



However, Sharing is Unsafe - A simple counting program...

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

long count = 0;

void* increase(void *arg) {
    long i, inc = *(long *)arg;
    for (i=0; i<inc; i++)
        count++;
    pthread_exit(NULL);
}

int main(int argc, char* argv[]){
    pthread_t tid1, tid2;
    long inc = atol(argc >= 2 ? argv[1] : "100");
    pthread_create(&tid1, NULL, increase, &inc);
    pthread_create(&tid2, NULL, increase, &inc);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("counter is %ld\n", count);
    return 0;
}
```

## Example

- Consider the two threads each doing the following

int x = 0;

```
void increase(int cnt) {
    int i;
    for(i=0;i<cnt;i++)
        x = x + 1;
}
```



```
void increase(int cnt) {
    int i;
    for(i=0;i<cnt;i++)
        x = x + 1;
}
```



**What will happen?**

What is happening?

- The addition
  - Becomes more complex in assembly
  - Something like (pseudo-code)

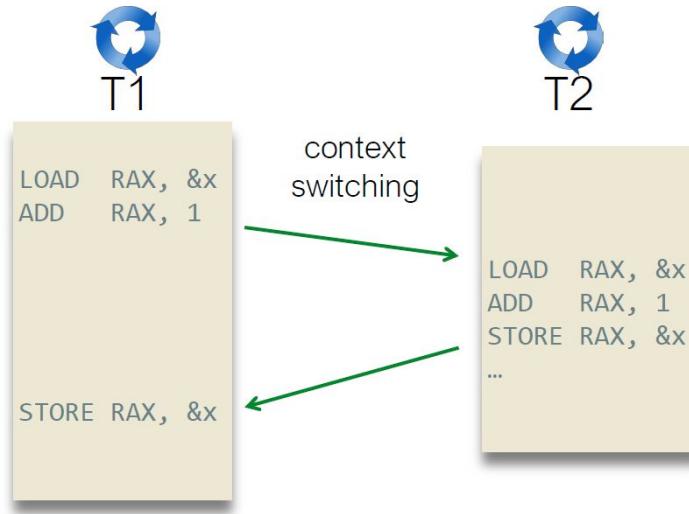
```
void increase(int cnt)
{
    int i;
    for(i=0;i<cnt;i++)
        x = x + 1;
}
```

## CSE 3100 Master Notes

```
// x = x + 1;
```

```
LOAD RAX, &x  
ADD RAX, 1  
STORE RAX, &x
```

- Threads execute this concurrently
- Even on a single core, the execution of a thread can be interrupted



### Lesson

- Even sharing a single integer can go wrong!
- What to do?
- We need coordination!
  - RULES and PROTOCOLS
  - To establish how share data safely and keep everyone happy

### The Road Ahead

- What we will do
  - Define PROTOCOLS and DATA STRUCTURES to safely share
- Examples
  - Mutexes / Spinlocks
  - Conditions
  - Semaphores
  - Barriers
  - Producer / Consumer
  - Reader / Writer
  - ...

### How unsafe can this be?

- Lots of subtle issues
  - It is very easy to get it wrong
- Good multi-threading programming must be disciplined

### Detaching a thread

```
int pthread_detach(pthread_t tid)
```

- The “parent thread doesn’t need to wait
- A thread can detach another thread
- When detached thread terminates, its resources are automatically released
- A thread can detach itself:  

```
    pthread_detach(pthread_self());
```
- Only threads created as joinable and not detached can be joined

### Passing Arguments to Threads -1

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

#define NUM_THREADS 8

struct thread_data {
    int thread_num;
    char* message;
};

void* PrintHello(void* threadarg) {
    struct thread_data* my_data = (struct thread_data*) threadarg;

    sleep(1 + 5*(my_data->thread_num % 2) );

    printf("Thread #%-d: %s length=%zd\n", my_data->thread_num, my_data->message,
           strlen(my_data->message));

    pthread_exit(NULL);
}
```

## Passing Arguments to Threads -2

```

int main(int argc, char* argv[])
{
    pthread_t threads[NUM_THREADS];
    static struct thread_data thread_data_array[NUM_THREADS];
    char* messages[NUM_THREADS];
    int rc, t;
    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";
    messages[5] = "Russian: Zdravstvuyte, mir!";
    messages[6] = "Japan: Sekai e konnichiwa!";
    messages[7] = "Latin: Orbis, te saluto!";

    // continue on the next slide
}

```

## Passing Arguments to Threads -3

```

for( t=0; t<NUM_THREADS; t++ ) {
    //set up struct for thread t
    thread_data_array[t].thread_num = t;
    thread_data_array[t].message = messages[t];

    printf("Creating thread # %d\n", t);
    rc = pthread_create(&threads[t], NULL, PrintHello,
                        (void*) &thread_data_array[t]);
    if (rc) {
        printf("ERROR; return code from pthread_create() is %d\n", rc);
        exit(-1);
    }
    printf("Detaching thread # %d\n", t);
    rc = pthread_detach( threads[t] ); // detach a thread
    if( rc ) {
        printf("ERROR; return code from pthread_detach() is %d\n", rc);
        exit(-1);
    }
}
pthread_exit(NULL);
}

```

### Create a thread as detached

- When calling `pthread_create()`, one of the attributes defines whether the thread is joinable or detached

- By default (NULL attribute) threads are created as joinable

```
pthread_t tid; void * arg      // thread id and argument
pthread_attr_t attr;           // an attribute variable

pthread_attr_init(&attr);      // initialize with default attributes
// set detach state
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);
pthread_create(&tid, &attr, start_routine, arg);    // create thread
pthread_attr_destroy(&attr);           // destroy the attribute
```

## Practice Exam 2

---

Mon. Oct. 28, 2019

### Practice Exam 2 Starter Code

```
//In this practice exam, we use multiple processes to simulate a word game named hangman.
//In the original hangman game, the number of guesses are limited and some graphical
display of a hangman
//is involved to indicate the progress of the game.
//We simplify the game to not to draw a hangman, and not limit the number of guesses.
//The simplified game works as follows.
//Player one chooses a word and indicates the length of the word
// and let the second player to guess the word. For example, the first player
// shows the following string to indicate that the word has 4 letters.
// ----
//The player two suggests a letter, for example, the letter 'e'.
//The player one responses by displaying the correct guess at the right places in the word.
//For example, player one tells player two the following string
// --ee
//Player two suggests another letter, a letter 'f' this time.
//The player one displays the same string
// --ee
//This is because the letter 'f' is not in the word.
//Player two suggests another letter 't'.
//The first player displays
// t--ee
//Then the second player suggests another letter 'r'.
//Player one displays
// tree
//Now since all the hidden letters are displayed. The game is over.
//In this practice exam, we need to use two processes to simulate the two players.
//Also, we need to use pipes for the communications between the two players.
//To be more specific, the child process is player one; the parent process is player two.
//A user will type guesses from the standard input to play the game.
```

## CSE 3100 Master Notes

```
// Search TODO to find the location where the code needs to be completed.

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <assert.h>
#include <sys/wait.h>
#include <errno.h>
#include <ctype.h>

#define PFD_READ 0
#define PFD_WRITE 1

#define MAX_WORD_COUNT 60000          //we have less than 60000 words
#define MAX_WORD_LENGTH 80           //each word is less than 80 letters

void die(char *s)
{
    if (errno)
        perror(s);
    else
        fprintf(stderr, "Error: %s\n", s);
    exit(EXIT_FAILURE);
}

char words[MAX_WORD_COUNT][MAX_WORD_LENGTH];      //2-d array to hold all the
words
int count = 0;           //number of words, initialized to 0

//read words from the file to the array words declared above
//also update the number of words (update variable count)
//We could have avoided using global variables. Try to revise it yourself.
void read_file_to_array(char *filename)
{
    FILE *fp;

    //open the file for reading
    fp = fopen(filename, "r");
    if(fp==NULL)
        die("Cannot open the word list file.");

    // TODO
    // make sure when each word is saved in the array words,
    // There is no white space in words, we can use fscanf().
    // We could also use fgets(). Need to remove '\n' at the end.
    fclose(fp);
```

## CSE 3100 Master Notes

```
}

// write a character to a pipe
void write_char(int pd, char value)
{
    if (write(pd, &value, sizeof(char)) != sizeof(char))
        die("write()");
}

// write a string to FD pd , add '\n' at the end
void write_word(int pd, char *word)
{
    size_t len = strlen(word);

    if (write(pd, word, len) != len)
        die("write()");
    write_char(pd, '\n');
}

// read a char from FD pd and save the result in *pc
// return the return value from read()
int read_char(int pd, char *pc)
{
    return read(pd, pc, sizeof(char));
}

// read a line from FD pd
//
void read_word(int pd, char buffer[], int sz)
{
    char c;
    int count = 0;

    while (read_char(pd, &c) > 0)
    {
        if (count >= sz)
            die("line is too long in read_word().");
        if (c == '\n') {
            buffer[count] = 0;
            return;
        }
        buffer[count ++] = c;
    }
    // could handle error better
    die("read() failed in read_word()");
}
```

## CSE 3100 Master Notes

```
//check if the character guess is in the word
//if it is in the word, update the string so_far in the right places
//for example, if guess is 'e', so_far is "----", and word is 'tree'
//then so_far will be updated to be '--ee'
//if guess is 't', so_far is '--ee', and word is 'tree'
//then so_far will be updated to be 't-ee'
// Return value:
// 0: guess is not in the word
// 1: guess is in the word
int check_guess(char guess, char *so_far, const char *word)
{
    // TODO
    return 0;
}

int main(int argc, char* argv[])
{
    if(argc!= 2)
    {
        printf("Usage: %s seed\n", argv[0]);
        return -1;
    }
    int seed = atoi(argv[1]);
    assert(seed > 0);

    int pdp[2];
    //pipe creation
    if(pipe(pdp) == -1)
    {
        perror("Error.");
        return -1;
    }

    int pdc[2];
    //pipe creation
    if(pipe(pdc) == -1)
    {
        perror("Error.");
        return -1;
    }

    pid_t pid;
    pid = fork();
    if(pid == 0)
    {
        // TODO
    }
}
```

## CSE 3100 Master Notes

```
// close some file descriptors

// read the list in child process
read_file_to_array("dict.txt");

char *my_word;
char so_far[MAX_WORD_LENGTH];

srand(seed);
my_word = words[rand() % count];
fprintf(stderr, "Child: debugging: the word is %s\n", my_word);

// Note that so_far should have enough space
size_t len = strlen(my_word);
for(int i = 0; i<len; i++)
    so_far[i] = '-';
so_far[len] = 0;

//TODO
//repeatedly doing the following
//    send so_far to parent
//    receive a guess (a character) from parent
//    exit from the loop if it was not successful
//    check_guess
//Do some clean up before exit from the process
return 0;
}

else
{
    char guess;
    char so_far[MAX_WORD_LENGTH];

    // TODO
    // close some file descriptors
    // Then do the following in a loop:
    //    read a word from child
    //    print it to stdout
    //    if there is no '-', exit from the loop
    //    read a character from stdin until a letter is found
    //    report error if EOF found.
    //    turn the character to lower case
    //    send it to child
    // close file descriptors before exit from the process

}

//wait for the child process to finish
```

```
    waitpid(pid, NULL, 0);
    return 0;
}

//below is a sample output
//it can also be found in the file sample-output.txt
/*
./hangman 8
To help debugging: the word is pride
-----
e
---e
a
---e
t
---e
i
--i-e
r
-ri-e
p
pri-e
d
pride
The word is pride.
*/
```

## Lecture 23 - T2: Mutual Exclusion (mutex)

---

Weds. Oct. 30, 2019

### Review

- So far
  - You learned how to create threads (pthread\_create)
  - You learned how to terminate thread execution (pthread\_exit)
  - You learned how to wait on threads (pthread\_join)
- You can carry out independent computations with threads
- However, sharing is a problem
  - Remember the shared counter?



What if the shared counter is the balance on your bank account?

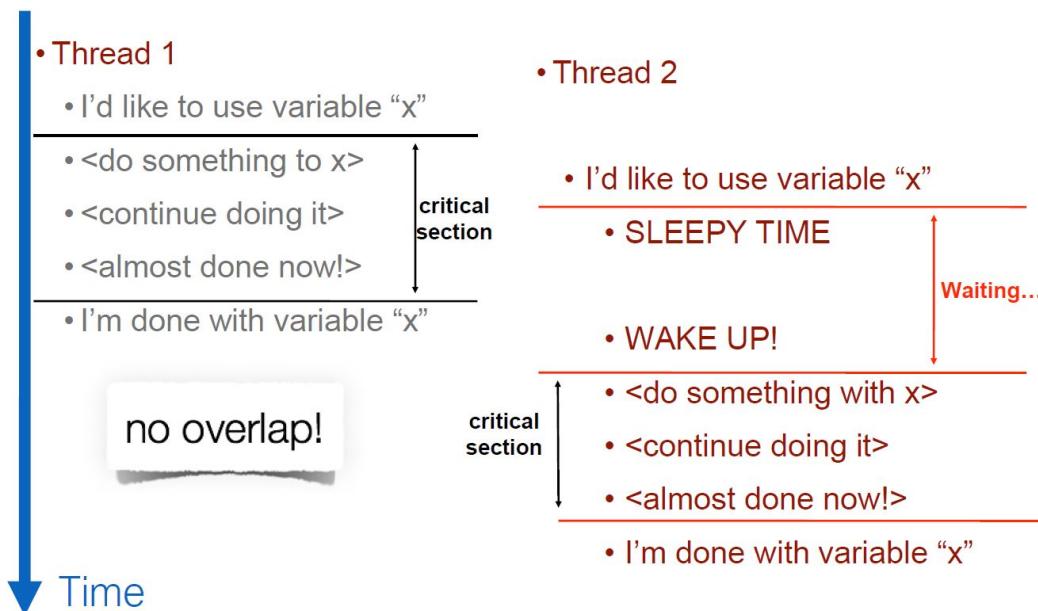
## Overview

- Basic Sharing
  - Mutual exclusion
  - Critical selection
  - Mutex
- Application
  - Concurrent access to shared data structures
  - Counter arrays
  - ...

## Mutual Exclusion

- Objective
  - Protect shared resources
  - Only one thread can access the resources
- Protocol
  - A set procedures for accessing shared resources
- Example:
  - Lock the resource while using it. Cannot lock if it is already locked
  - Wait if the resource is already lock
- Fact
  - Ask everyone nicely, and expect that everyone behaves

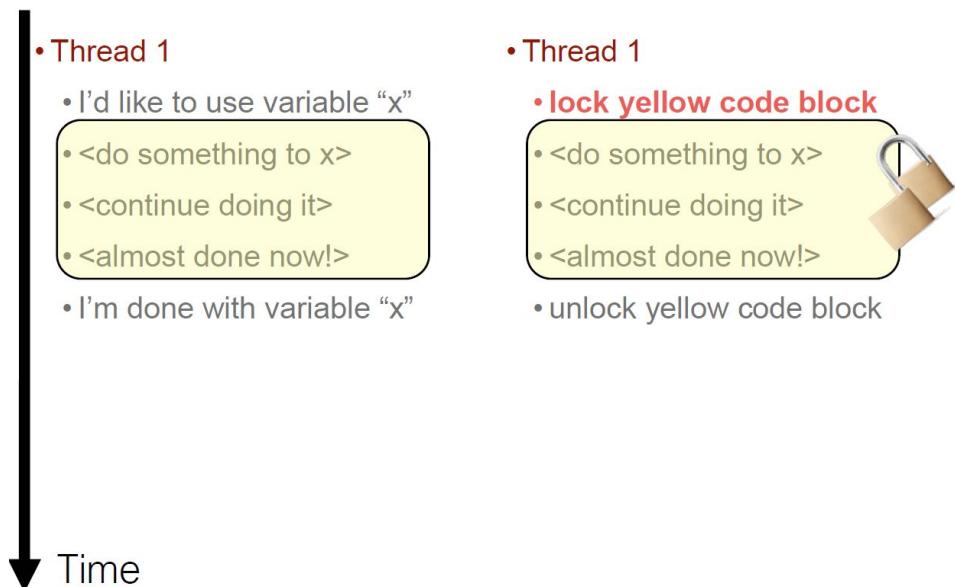
## Protocol



### Race to the critical section

- Multiple threads race each other to get to critical section
  - Critical section is a code segment that accesses shared resources
- One of them “wins” the race
  - Winner
    - acquire a “lock” first
    - executes critical section
    - release the lock (unlock)
  - Loser
    - waits as the resource is locked
    - wakes up when it is its turn to get the lock
    - executes critical section
    - releases the lock

### Protocol



### Mutex

- The Lock is a POSIX abstraction
  - Called a **Mutex** (for **MUTual Exclusion**)
  - Provided by the operating system
- Semantics
  - At most one thread can acquire the lock at any time
  - If a thread “loses” a race, it **falls asleep**
  - Sleepy threads **wake up** when the lock is released

## pthread Mutex types and API

```
pthread_mutex_t; // define a mutex
```

- Functions
  - Initialize the Mutex
  - Destroy a Mutex we no longer need
  - Lock a Mutex
  - Unlock a Mutex
  - [Attempt to Lock a Mutex]

## Creation

```
# include <pthread.h>
int pthread_mutex_init(pthread_mutex_t * mutex,
                      const pthread_mutexattr_t * attr);
```

- Initialize a mutex (i.e., allocate OS resources)
- Return 0 on success
- Example

```
typedef struct MyRecord {
    pthread_mutex_t myLock;
    int           myValue;
} MyRec;

MyRec* makeARecord() {
    MyRec* rec = (MyRec*)malloc(sizeof(MyRec));
    pthread_mutex_init(&rec->myLock, NULL);
    rec->myValue = 0;
    return rec;
}
```

mutex attribute

## Destruction

```
# include <pthread.h>
```

```
int pthread_mutex_destroy(pthread_mutex_t * mutex);
```

- Release resources for lock
- Return 0 on successes
- Example

```

typedef struct MyRecord {
    pthread_mutex_t myLock;
    int             myValue;
} MyRec;

void freeARecord(MyRec* rec) {
    pthread_mutex_destroy(&rec->myLock);
    free(rec);
}

```

### Lock / Unlock

```

#include <pthread.h>
int pthread_mutex_lock(pthread_mutex_t * mutex);
int pthread_mutex_unlock(pthread_mutex_t * mutex);

```

- Enter/leave the critical section. Wait (asleep!) if mutex is locked
- Return Value 0 on success (=!= 0? something went horribly wrong)
- Example

```

typedef struct MyRecord {
    pthread_mutex_t myLock;
    int             myValue;
} MyRec;

void incrementRecordValue(MyRec* rec) {
    pthread_mutex_lock(&rec->myLock);
    rec->myValue = rec->myValue + 1;
    pthread_mutex_unlock(&rec->myLock);
}

```

### Key Observations

- Lock and unlock some pairs
- You need ONE MUTEX per “thing” you wish to protect
  - That’s why we package the integer & the mutex in a struct
  - The mutex protects just that integer
  - And nothing else!
- Lock is **BLOCKING**
- Critical sections should be **short**
  - Why?
- Locking incurs a **big** cost
  - Why?

### Example 1

- Fix the example we had issue

### Example 2: Counter ADT

- A shared counter
  - That was the driving example
  - Make an ADT pairing mutex and counter
  - Have functions to manipulate the ADT
    - increase
    - decrease
    - reset to Zero

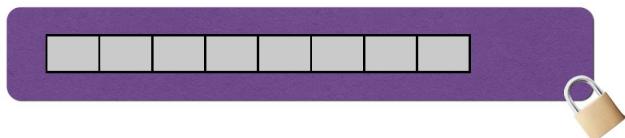
See the Demo Code under t3.counter-adt

### Example 3: array

- Sharing an array
- Key decision: Locking Granularity
  - Lock the entire structure?
  - Lock individual parts?

#### Option 1

- Idea
  - Wrap up the array in a structure
  - Have a single lock for the whole thing
- Issues?



#### Option 2

- Idea
  - Wrap up each value in a structure (value + lock)
  - Make an array of structures
- Issues?



mutex\_trylock()

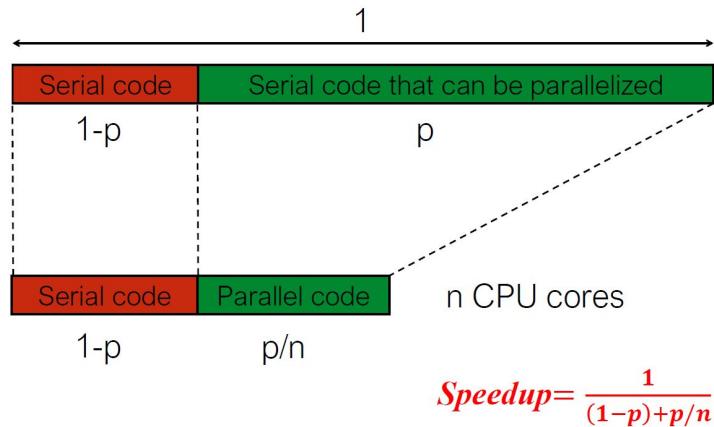
```
#include <pthread.h>
int pthread_mutex_trylock(pthread_mutex_t *mutex);
```

- If the mutex is not locked, it will succeed and lock

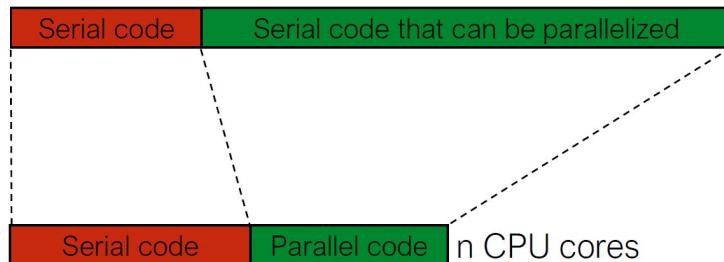
## CSE 3100 Master Notes

- If the mutex is already locked
  - It will NOT block
  - It will return an error code
- Purpose
  - Mix breed. It allows polling before locking
    - sometimes, it is fine to use the shared resource later

### Amdahl's Law



Practical speedup lower due to overhead



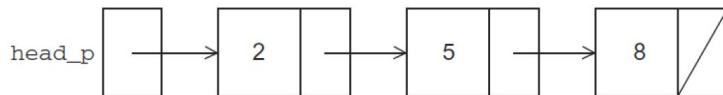
### Mutex Attributes?

- A value of type `pthread_mutexattr_t`
  - An optimal argument to create create a mutex
- APIs to
  - Initialize attribute record (`pthread_mutex_attr_init()`)
  - destroy attribute record (`pthread_mutexattr_destroy()`)
  - modify an attribute in attribute record (set type, get type, etc.)
- Key property: mutex TYPE
  - How to deal with recursive lock?
  - What if other threads try to unlock?

## Mutex Type (from the DOC)

- PTHREAD\_MUTEX\_NORMAL
  - This type of mutex does not detect deadlock. A thread attempting to relock this mutex without first unlocking it will deadlock. Attempting to unlock a mutex locked by a different thread results in undefined behaviour. Attempting to unlock an unlocked mutex results in undefined behaviour.
- PTHREAD\_MUTEX\_ERRORCHECK
  - This type of mutex provides error checking. A thread attempting to relock this mutex without first unlocking it will return with an error. A thread attempting to unlock a mutex which another thread has locked will return with an error. A thread attempting to unlock an unlocked mutex will return with an error.
- PTHREAD\_MUTEX\_RECURSIVE
  - A thread attempting to relock this mutex without first unlocking it will succeed in locking the mutex. The relocking deadlock which can occur with mutexes of type PTHREAD\_MUTEX\_NORMAL cannot occur with this type of mutex. Multiple locks of this mutex require the same number of unlocks to release the mutex before another thread can acquire the mutex. A thread attempting to unlock a mutex which another thread has locked will return with an error. A thread attempting to unlock an unlocked mutex will return with an error.
- PTHREAD\_MUTEX\_DEFAULT
  - Attempting to recursively lock a mutex of this type results in undefined behaviour. Attempting to unlock a mutex of this type which was not locked by the calling thread results in undefined behaviour. Attempting to unlock a mutex of this type which is not locked results in undefined behaviour. An implementation is allowed to map this mutex to one of the other mutex types.

## Sorted Linked List Example



```
typedef struct list_node_s {
    int data;
    struct list_node_s* next;
} LNode;
```

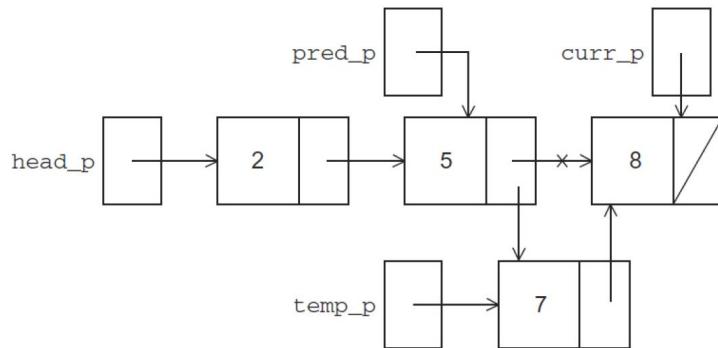
## Testing Membership

```
int member( int value, LNode* head_p) {
    LNode* curr_p = head_p;

    while( curr_p != NULL && curr_p->data < value )
        curr_p = curr_p->next;

    if( curr_p == NULL || curr_p->data > value )
        return 0;
    else
        return 1;
}
```

## Inserting a new value

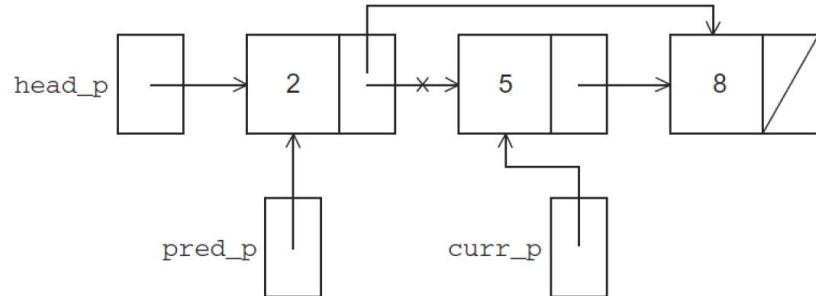


## Inserting a new value

```
int insert( int value, LNode** head_pp) {
    LNode* curr_p = *head_pp;
    LNode* pred_p = NULL;
    LNode* temp_p;

    while( curr_p != NULL && curr_p->data < value ) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }
    if( curr_p == NULL || curr_p->data > value ) {
        temp_p = (LNode*)malloc( sizeof(LNode) );
        temp_p = value;
        temp_p = curr_p;
        if(pred_p == NULL) /* new first node */
            *head_pp = temp_p;
        else
            pred_p->next = temp_p;
        return 1;
    } else /* value already in list */
        return 0;
}
```

## Deleting a value



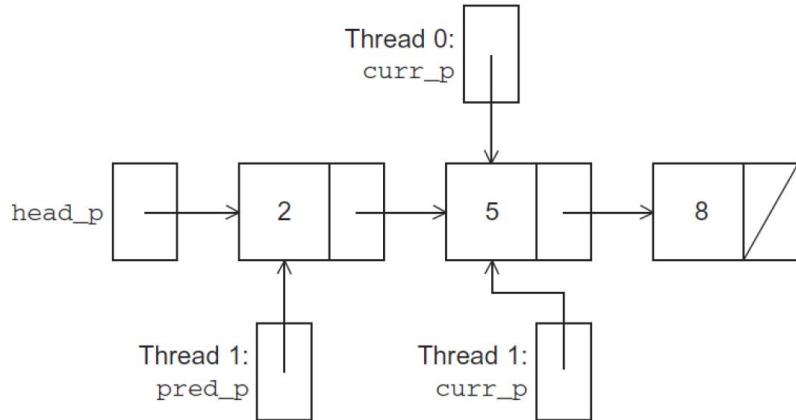
```

int delete( int value, LNode** head_pp) {
    LNode* curr_p = *head_pp;
    LNode* pred_p = NULL;

    while( curr_p != NULL && curr_p->data < value ) {
        pred_p = curr_p;
        curr_p = curr_p->next;
    }
    if( curr_p == NULL || curr_p->data > value ) {
        if(pred_p == NULL) { /* deleting first node in list */
            *head_pp = curr_p->next;
        } else {
            pred_p->next = curr_p->next;
        }
        free(curr_p);
        return 1;
    } else /* value not in list */
        return 0;
}

```

Simultaneous access by two threads?



Solution #1

- Lock the entire list any time attempts to access it
  - A call to each of the three functions protected by a mutex e.g.,

```

pthread_mutex_lock(&list_mutex);
member( value, head_p );
pthread_mutex_unlock(&list_mutex);
  
```

- Issues?

Solution #2

- Instead of locking the entire list, we could try to lock individual nodes
  - A “finer-grained” approach

```

typedef struct list_node_s {
    int data;
    struct list_node_s* next;
    pthread_mutex_t mutex;
} LNode;
  
```

### Implementation of member with one mutex per node

```
int member( int value, LNode* head_p) {
    LNode* curr_p;

    pthread_mutex_lock( &head_p_mutex );
    curr_p = head_p;
    while( curr_p != NULL && curr_p->data < value ) {
        if( curr_p->next != NULL )
            pthread_mutex_lock( &(curr_p->next->mutex) );
        if( curr_p == head_p )
            pthread_mutex_unlock( &head_p_mutex );
        pthread_mutex_unlock( &(curr_p->mutex) );
        curr_p = curr_p->next;
    }
}
```

```
if( curr_p == NULL || curr_p->data > data ) {
    if( curr_p == head_p )
        pthread_mutex_unlock( &head_p_mutex );
    if( curr_p != NULL )
        pthread_mutex_unlock( &(curr_p->mutex) );
    return 0;
} else{
    if( curr_p == head_p )
        pthread_mutex_unlock( &head_p_mutex );
    pthread_mutex_unlock( &(curr_p->mutex) );
    return 1;
}
}
```

### Issues

- More complex to implement
- More memory to store the list
  - One mutex per node
- Slower
  - Each time a node is accessed, a mutex must be locked and unlocked
- Using more than one lock creates opportunities for **DEADLOCK!**

# Lecture 24 - T3: POSIX Condition

Mon. Nov. 4, 2019

## Mutex Review

- Mutex
  - One has to get the lock before entering critical section
  - Yields exclusive access to a resource
  - If a thread cannot get the lock, it waits ...
- Example of two threads sharing data
  - Thread 1 computes a result
  - Thread 2 displays the result on UI when computation done
- Question
  - How does thread 2 know the result is ready?

## Producer-Consumer with Bounded Buffer

- Classic problem
  - Producer(s) put things into a shared buffer
  - Consumer(s) take them out!



## Problem Constants

- Mutual exclusion
  - Only one thread can manipulate the buffer at any time
    - Buffer is shared!
- Synchronization
  - Consumer(s) must WAIT if buffer is empty
  - Producer(s) must WAIT if buffer is full

**Can it be done with only a lock ?**

Well... Sorta... Consumer in pseudo-code

```
do {
    Lock the buffer// Buffer is shared ! May have to wait
    if buffer is not empty
        Fetch data in the buffer
```

```

        Unlock the buffer
} while (data is not fetched) // May fail, so try in a loop
Process the data

```

Consumer in C-like code

```

pthread_mutex_t buffer_lock;
// Consumer fetches data from a buffer, one each time
    fetch = 0
    do {
        pthread_mutex_lock(&buffer_lock);           // Get the lock for
the buffer
        if (nElements > 0){
            data = fetch_from_buffer();
            fetched = 1;
        }
        pthread_mutex_unlock(&buffer_lock);
    } while (fetched == 0);
// Continue to process data

```

What happens if the buffer is empty?

Condition variable?

- A handshake mechanism to say: “There is data for you to use”



Consumer waits on condition

```

pthread_mutex_tbuffer_lock;
// Consumer fetches data from a buffer, one each time.
pthread_mutex_lock(& buffer_lock);// Get the lock for the buffer

while (! nElements> 0) { // keep waiting if buffer is empty
    wait for someone telling me data is ready!
}

data = fetch_from_buffer(); // fetch data from the buffer
pthread_mutex_unlock(& buffer_lock);
// Continue to process data

```

## Consumer waits on condition - 2

It is a **while**, not if

A predicate

```

pthread_mutex_lock(& buffer_lock); // Get the lock for the buffer
while (! nElements > 0) {           // keep waiting if buffer is empty
    wait for someone telling me data is ready!
}
data = fetch_from_buffer(); // fetch data from the buffer
pthread_mutex_unlock(& buffer_lock);

```

Check predicate when mutex is locked  
 mutex is unlocked while waiting,  
 and locked when waiting is over

- Predicate: Logical expression describing the property of an object (or the state) the program needs
  - Examples: the buffer is full, or the buffer is not empty

## POSIX condition variable

- Second and last core synchronization primitive
  - All others (barriers, read-write locks, ...) can be implemented using mutexes and condition variables
- Always paired with a mutex
  - Its access needs to be **mutex** protected!
- If a predicate is not true, a thread can wait on a condition variable
  - Other threads can **signal** on the condition variable when the **predicate has changed**
    - Threads must still check if the predicate is true due to spurious wake-ups

## pthread\_cond API

```

#include <pthread.h>

pthread_cond_t cond;           // Define a condition variable
int pthread_cond_init(pthread_cond_t * cond,
                      const pthread_condattr_t * attr);
int pthread_cond_destroy(pthread_cond_t * cond);
int pthread_cond_wait(pthread_cond_t * cond, pthread_mutex_t *mutex);
int pthread_cond_signal

```

```
int pthread_cond_broadcast
```

### Tips

- The mutex protects
  - The shared and the state (predicate P)
- The condition variable
  - lock the mutex before waiting
    - the mutex is unlocked automatically while waiting
    - Automatically and atomically re-lock the mutex when waiting up
- The predicate P is checked in a while loop!
  - The while loop returns to waiting when P is false
    - Why? (You will see in a moment)

### Typical Structure of using mutex and condition

```
pthread_mutex_t mutex;
pthread_cond_t cond;

pthread_mutex_lock(&mutex);
while (! predicate){
    pthread_cond_wait(&cond, &mutex);
    do_something();
}
access_shared_resources();                                // safe access the share
resources
pthread_mutex_unlock(&mutex);

// continue to perform other tasks
```

### Example setup

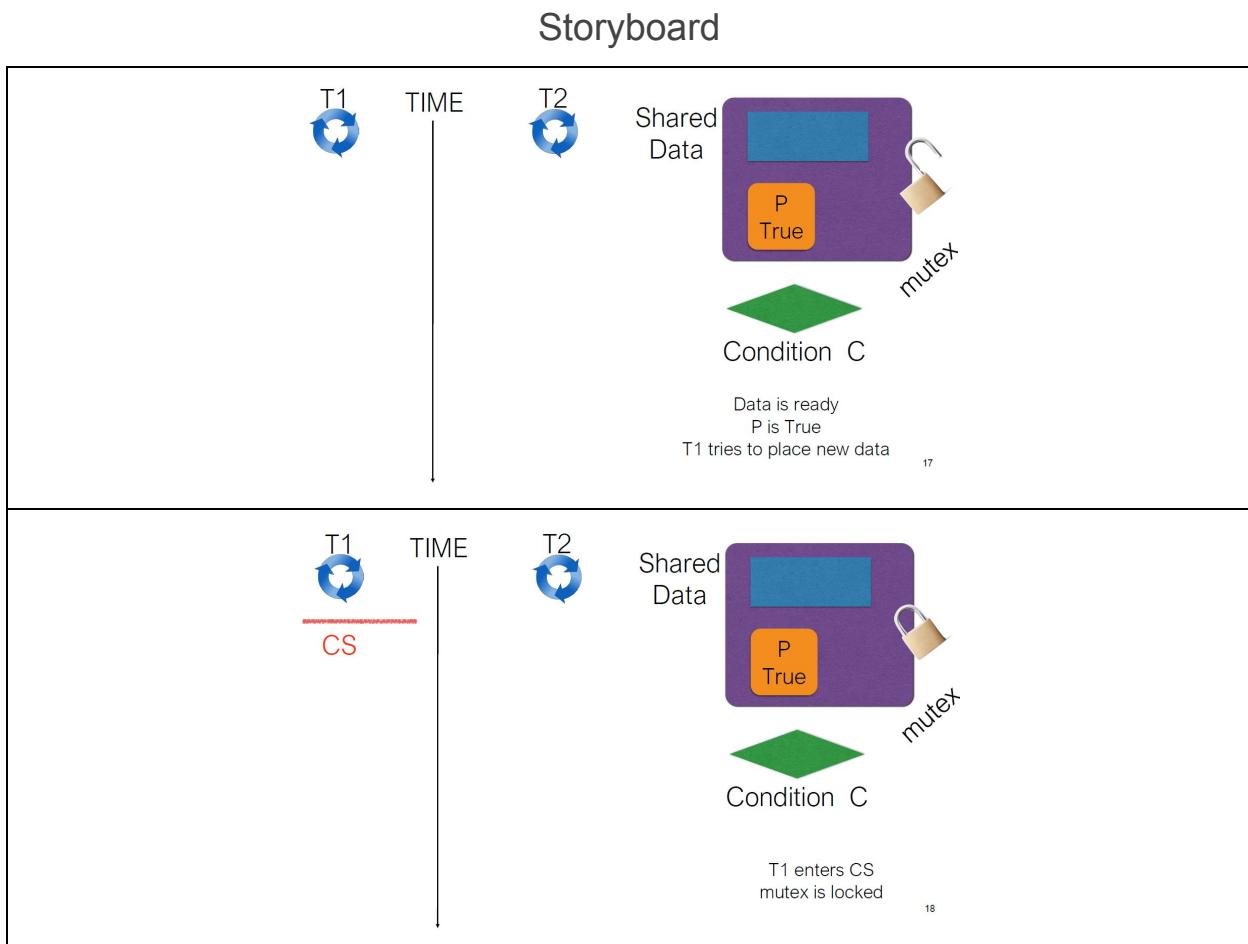
- Assumptions
  - A buffer that can hold only one item
  - A predicate P, indicating if the data is ready for consumer
  - A POSIX condition variable
  - A POSIX mutex
- Threads
  - T1: producer T2: consumer
  - Can be many threads: 1 .. k

## Two scenarios

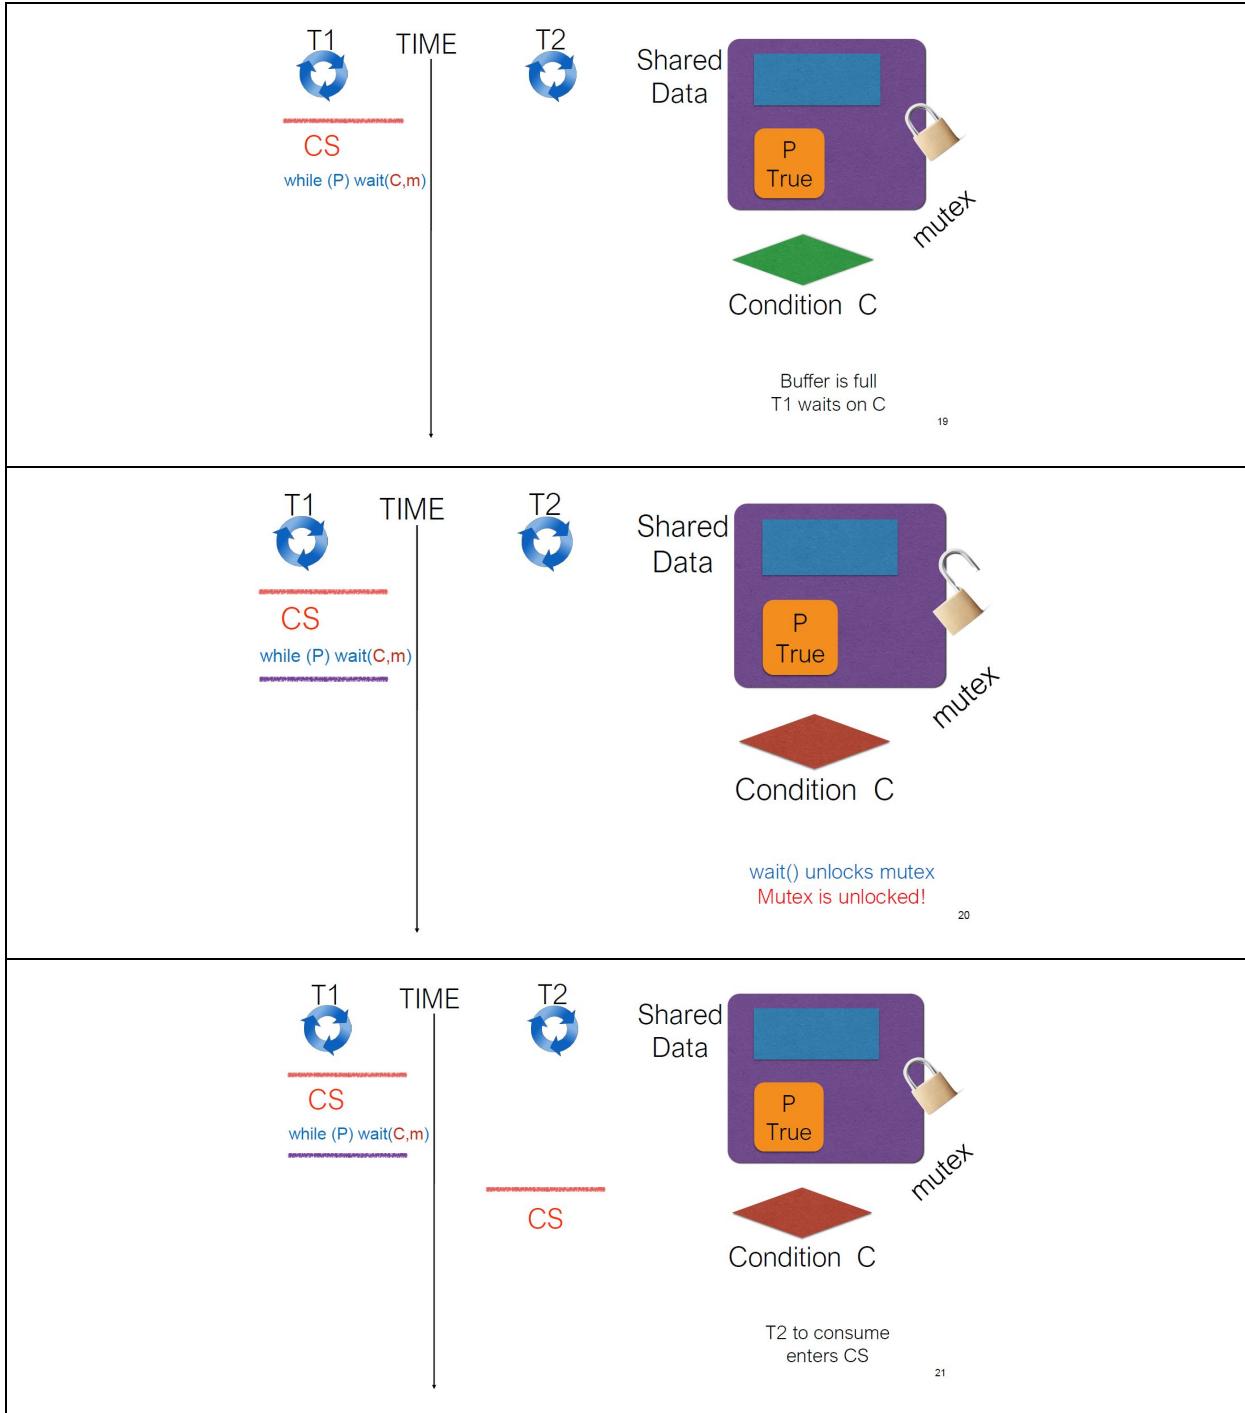
- Scenario 1
  - The producer gets to produce...
  - ...before the consumer consumes the previous one
- Scenario 2
  - The consumer is eager and tries to consume...
  - ...before the producer gets a chance to produce

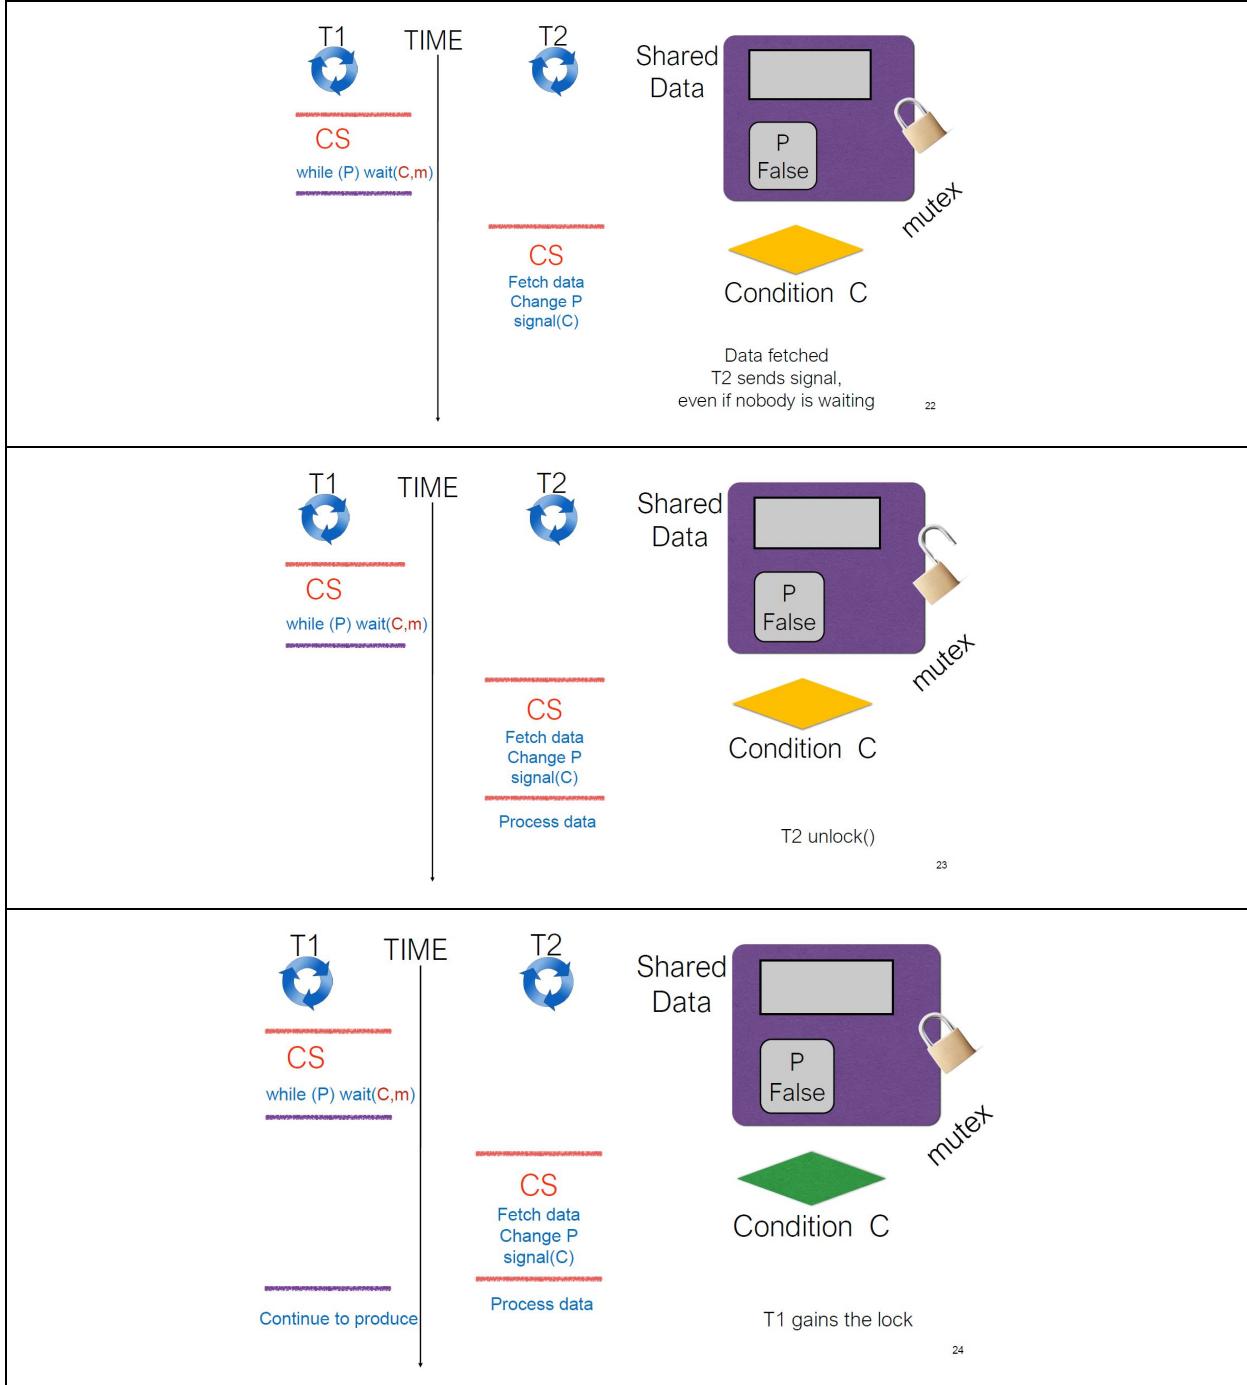
## Scenario 1: eager producer

- Producer tries to produce more, but the buffer is full



## CSE 3100 Master Notes



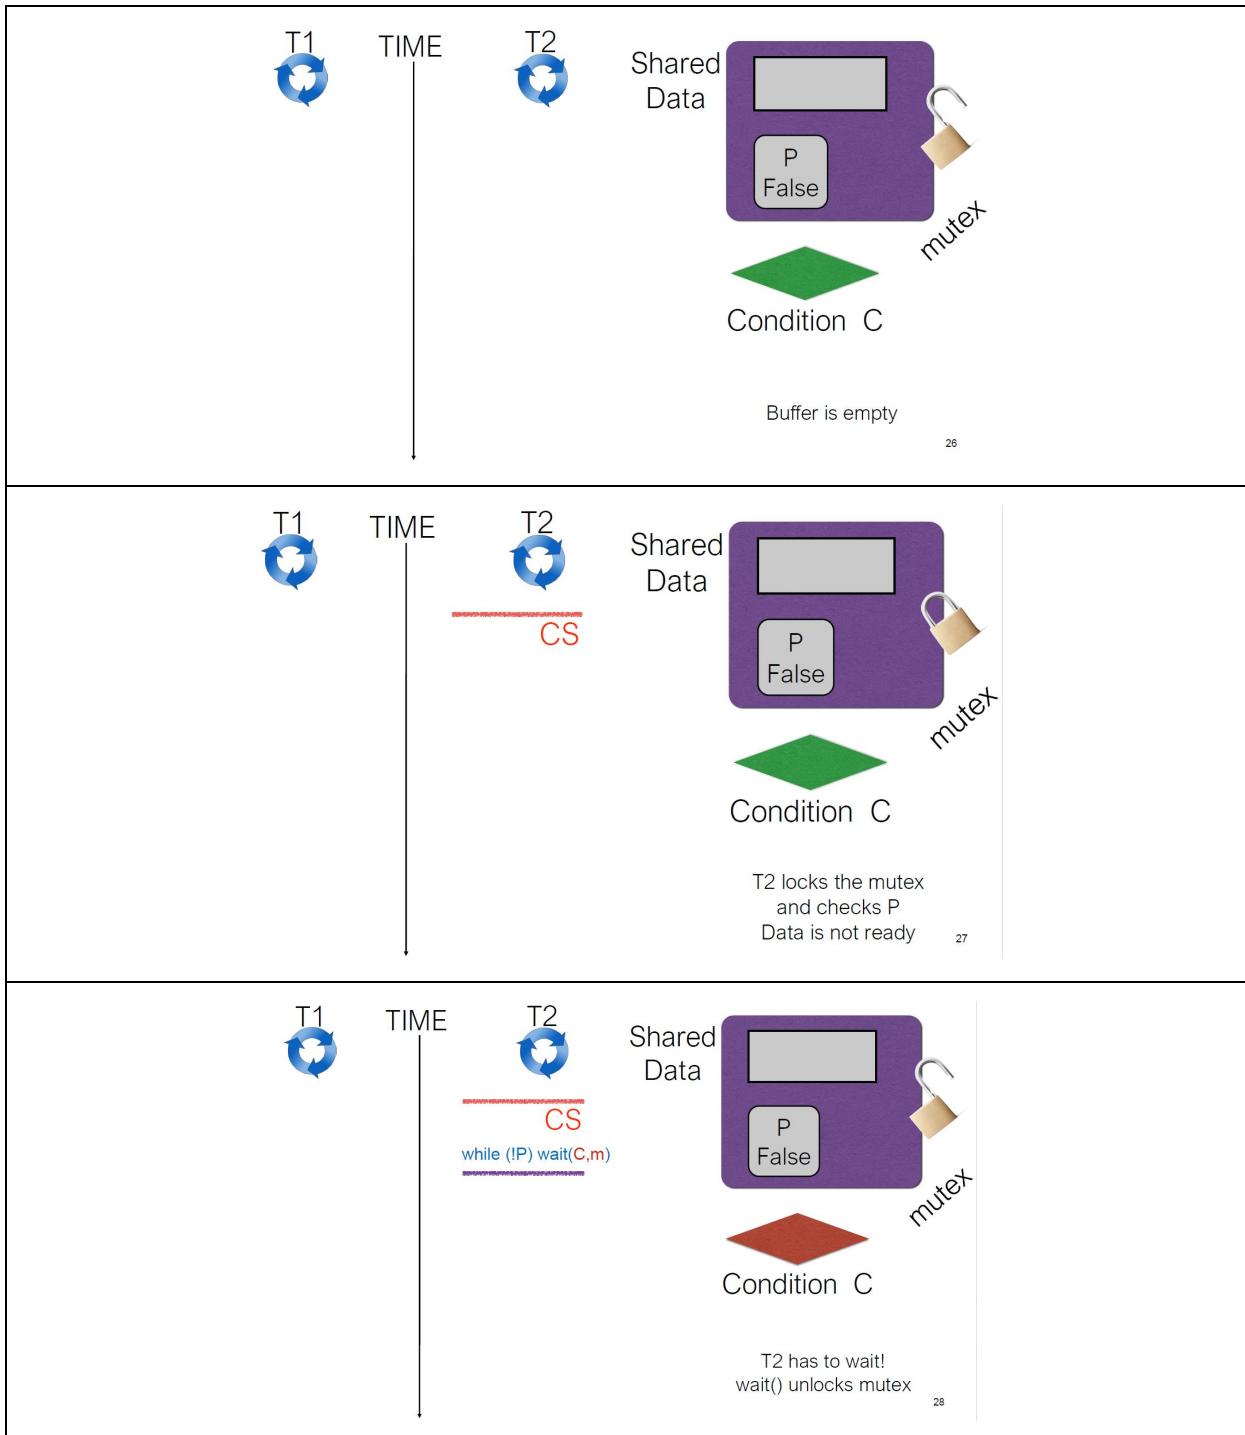


## Scenario 2 (Eager consumer)

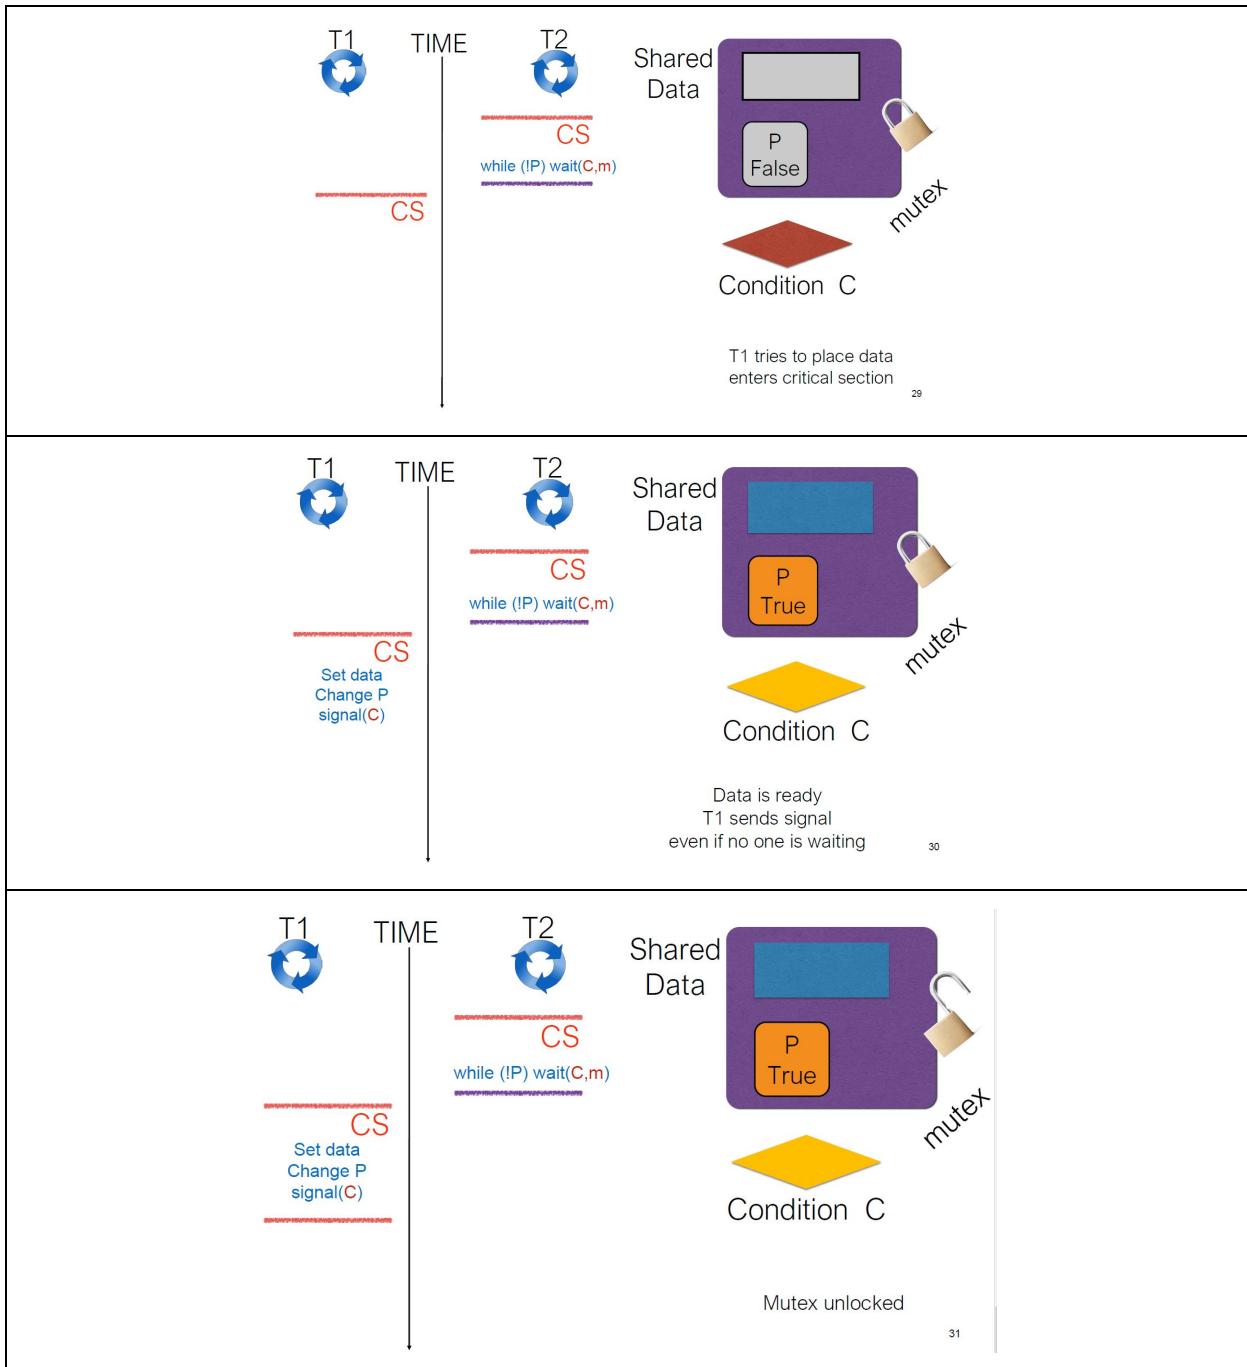
- Consumer tries to fetch data before they are produced

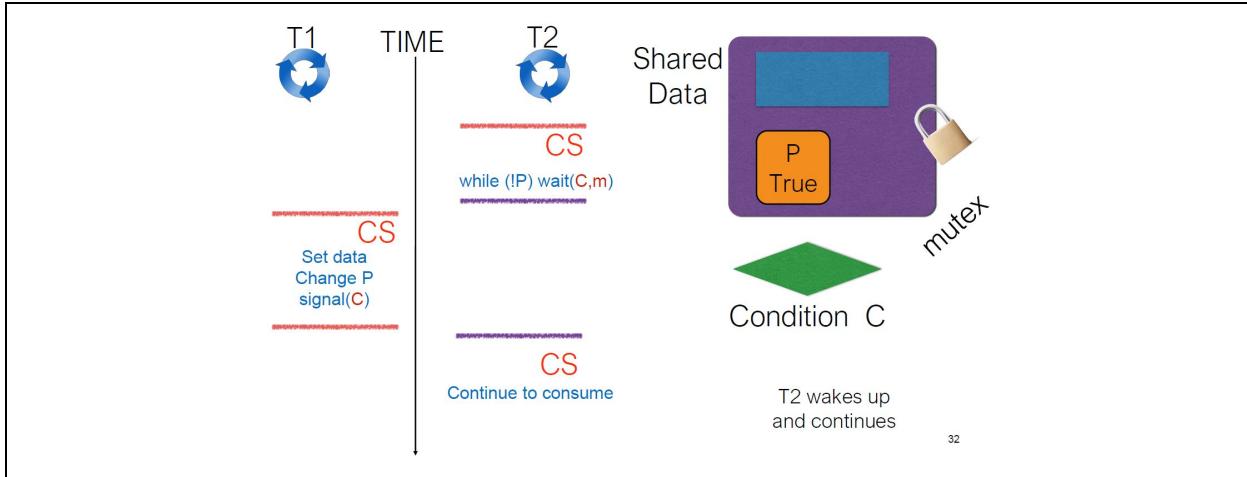
## Storyboard

## CSE 3100 Master Notes



## CSE 3100 Master Notes





32

## Signal vs. Broadcast

```
pthread_cond_signal(pthread_cond_t*cond);
```

- Wakes up **one waiting thread** in the condition

```
pthread_cond_broadcast(pthread_cond_t*cond);
```

- Wakes up **all waiting threads** in the condition
- Yet, only one of the waiting threads can grab the mutex
  - Go back to sleep if it fails

Caveat: a thread calling **wait after the broadcast** will not wake up (even if others in the process of waking up are still stuck in the condition variable —e.g., waiting on getting the mutex lock —)

## Producer in pseudo-code

```

Prepare data
do {
    Lock the buffer// May have to wait
    if buffer is not full// For bounded buffer
        Put data in the buffer
    Unlock the buffer
} while (data is not placed in the buffer)

```

## Producer in C-like code

```

pthread_mutex_tbuffer_lock;
// Producer adds data in a buffer, one each time.
..... // Prepare data here
added = 0;
do {
    pthread_mutex_lock(& buffer_lock);// Get the lock for the
    buffer

```

```
    if (nElements < BUF_SIZE) { // If the buffer is not full
        add_to_buffer(data); // Add data into the buffer
        added = 1;
    }
    pthread_mutex_unlock(& buffer_lock);
} while (added == 0);
```

**What happens if the buffer is full ?**

### Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;

/* other data structures here */

main() {
    /* declarations and initializations */

    task_available = 0;
    pthread_init();
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
}
```

```
void *producer(void *producer_thread_data)
{
    int inserted;
    while (!done())
    {
        create_task();
        pthread_mutex_lock(&task_queue_cond_lock);
        while (task_available == 1)
            pthread_cond_wait(&cond_queue_empty,&task_queue_cond_lock);
        insert_into_queue();
        task_available = 1;
        pthread_cond_signal(&cond_queue_full);
        pthread_mutex_unlock(&task_queue_cond_lock);
    }
}
```

```
void *consumer(void *consumer_thread_data)
{
    while (!done())
    {
        pthread_mutex_lock(&task_queue_cond_lock);
        while (task_available == 0)
            pthread_cond_wait(&cond_queue_full,&task_queue_cond_lock);
        my_task = extract_from_queue();
        task_available = 0;
        pthread_cond_signal(&cond_queue_empty);
        pthread_mutex_unlock(&task_queue_cond_lock);
        process_task(my_task);
    }
}
```

## Lecture 25 - T5: More Threads Synchronization

### More Synchronization Primitives

- So Far
  - Mutexes
  - Condition variables
- Today
  - Read-write blocks
  - Barriers
- Mutexes and cond. variables are primitive constructs
  - Read-write locks & barriers can be built using mutexes and condition variables

### Deadlock

- Deadlock is a state in which no member can make progress
  - All are waiting for other members to take actions
  - A member can be a thread, a process, a computer, etc.
- Example: 2 Mutexes A & B
  - Thread 1, locked A, try to get B
  - Thread 2, locked B, try to get A

### Common Solutions

- Fixed order
  - All the threads lock mutexes in the same order
  - m<sub>0</sub>, m<sub>1</sub>, and so on
- Try and back off
  - Try to lock. If it fails, release all locks & try again

### Readers-Writers Program

- In many applications, a data structure is read frequently but written infrequently
- Constraints:
  - Multiple readers can read simultaneously
    - A reader can start even if other readers are reading
  - Only one writer can write at a time
  - Read & write cannot happen at the same time
    - A reader has to wait if a writer is writing
    - A writer has to wait for other readers & writers to finish

### Solutions?

- Use Pthreads **read-write blocks**
  - Like a mutex w/ **2 diff. lock funcs.** (read/write)
- Use mutex & cond. variables

## CSE 3100 Master Notes

- Keep track of the number of readers and writers
  - Active & waiting
- Choose proper predicate
  - If the predicate is true, proceed to read/write
  - If the predicate not true, wait
- Notify others when a thread is done w/ reading/writing

### Pthread Read-Write Lock

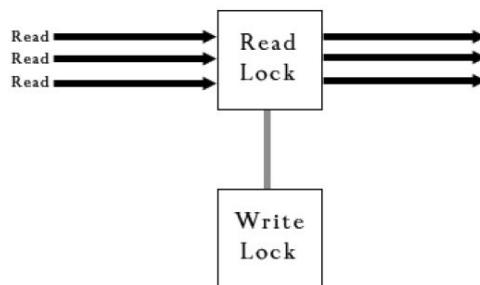
```
#include <pthread.h>
pthread_rwlock_t rwlock; // define a read-write lock

// initialize and destroy the a read-write lock
int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock,
    const pthread_rwlockattr_t *restrict attr);
int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);

// lock functions for readers and writers, and unlock
int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);
int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);
int pthread_rwlock_unlock(pthread_rwlock_t *rwlock);
```

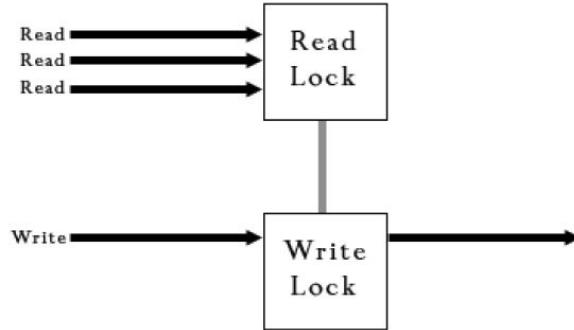
### Pthread Read-Write Locks For Reading

- First lock function locks for **reading**
  - A thread gets read lock if and only if no writer has the lock and no writer is blocked on the lock
  - More than one thread can get the read lock



### Pthread Read-Write Locks For Writing

- Second lock function locks for **writing**
  - A thread gets write lock if and only if no reader or writer has the lock
    - Wait for read lock, too!
  - Only one writer is allowed at any time



## Read-Write Lock Issues

- Writer starvation
  - Too many readers. Writers do not have a chance to start
  - Solution?
- “write locks shall take precedence over read locks”
  - which leads to reader starvation
- Reader starvation
  - Too many writers. Readers do not have a chance
    - Another writer starts when a writer unlocks

## Pthread read-write lock

- Pthread rwlock prefers writer
  - From the manual page

The `pthread_rwlock_rdlock()` function shall apply a read lock to the read-write lock referenced by `rwlock`. The calling thread acquires the read lock if a writer does not hold the lock and there are no writers blocked on the lock.

## Use mutex and condition to implement rwlock

- Using mutex and condition is more flexible, we can adjust our strategy
  - There are many ways
  - Ex.): alternate reader writer action, every other
- Best for us: One mutex and two condition variables (one for rd, one for wr)
  - Keep track of the numbers needed by the predicate

<http://heather.cs.ucdavis.edu/~matloff/158/PLN/RWLock.c>

## Pseudocode for rwlock lock

```

pthread_mutex_lock(&mutex);
// increment/decrement waiting counter and the whole loop
  
```

```
// can be placed in an if branch: if (! predicate)
increment waiting counter
// predicate depends on the policy!
while (! predicate) // check if this thread can lock
    pthread_cond_wait(&cond); // using either rd or wr condition
decrement waiting counter
increment active counter
pthread_mutex_unlock(&mutex);
// rwlock is locked. Can start to read/write
// Note that mutex is unlocked!
```

### Pseudocode rwlock unlock

```
Perform read/write operations // Depending on the lock type
// unlock rwlock when the operation is done
pthread_mutex_lock(&mutex);
decrement active counter
// inform waiting threads
// policy decides checking writer or reader first
if there is a writer waiting
    signal write cond
if there is a reader waiting // only needed for write unlock
    broadcast read cond
pthread_mutex_unlock(&mutex);
```

### Example: rwlock\_writeunlock()

**Does this lock prefer reader or writer?**

- A: reader
- B: writer
- C: don't know

```
status = pthread_mutex_lock(&rw1->mutex);
rw1->w_active = 0
if (rw1->r_wait > 0){
    status = pthread_cond_broadcast (&rw1->read);
    if(status != 0){
        pthread_mutex_unlock (&rw1->mutex);
        return status;
    }
}
else if (rw1->w_wait > 0){
    status = pthread_cond_signal (&rw1->write
```

```

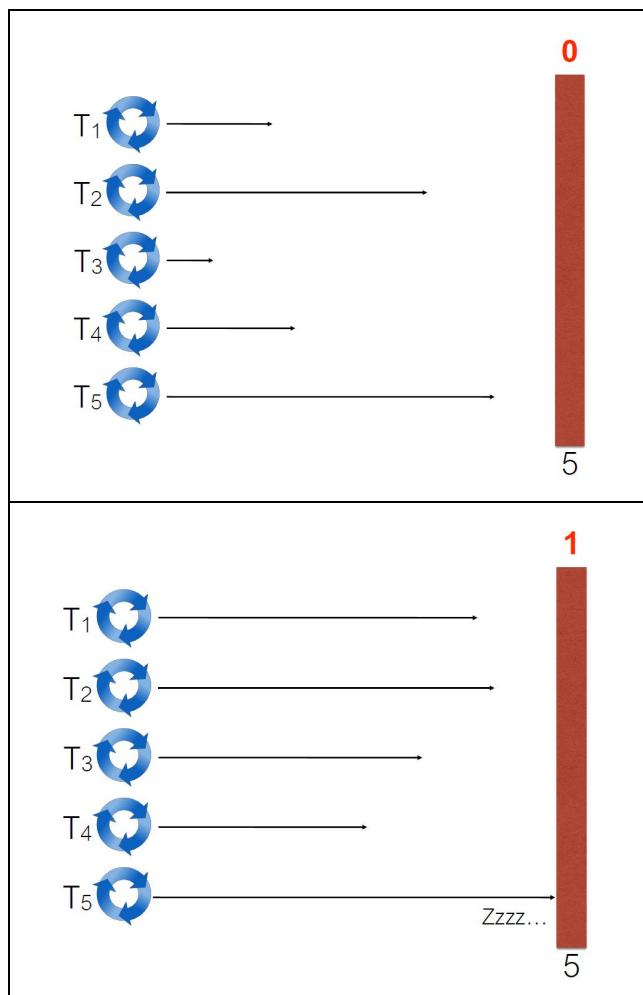
if (status != 0){
    pthread_mutex_unlock (&rw1->mutex);
    return status;
}
pthread_mutex_unlock(&rw1->mutex);

```

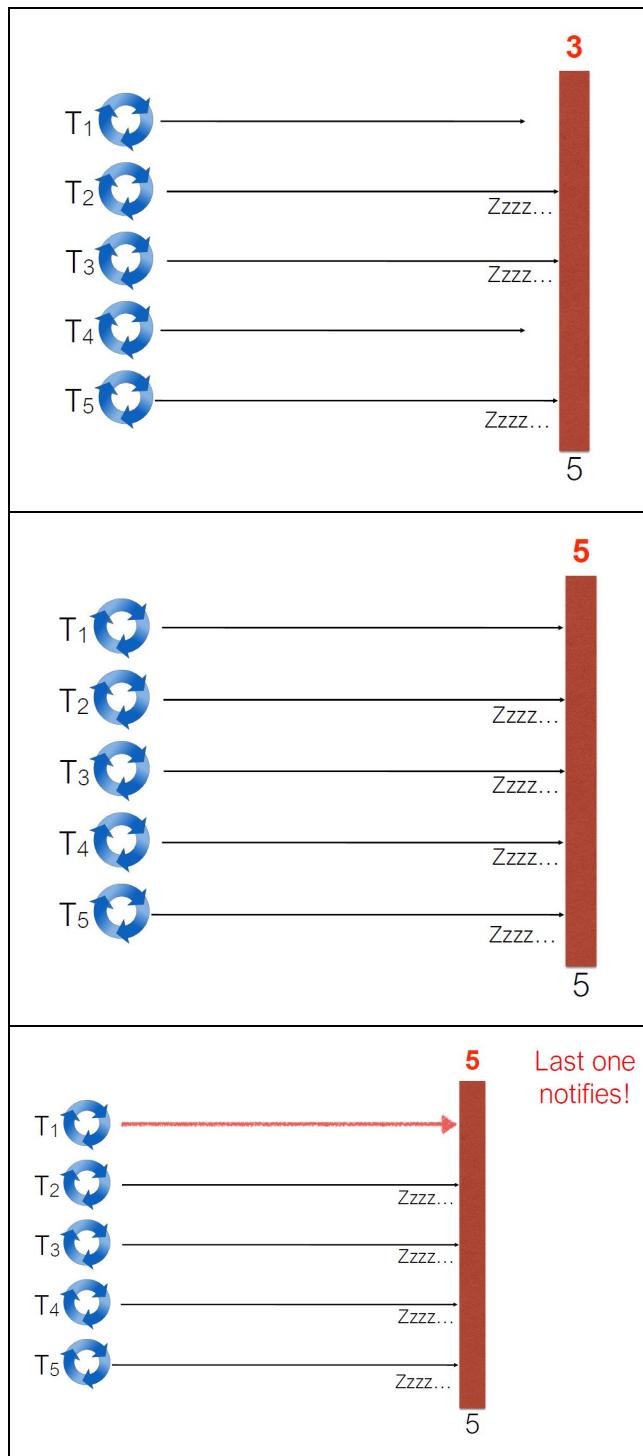
## Barriers

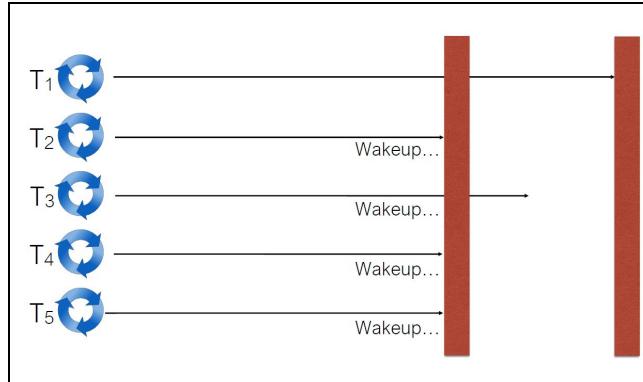
- Purpose
  - For applications where work is done in “phases”
  - Must have “worker” threads wait for the entire “group” to be done before proceeding to next phase
  - Number of workers known a priority

## Visually



## CSE 3100 Master Notes





### POSIX Barriers Support

```
#include <pthread.h>
// create barrier. note the count argument
int pthread_barrier_init(pthread_barrier_t* restrict barrier,
    const pthread_barrierattr_t* restrict attr, unsigned count);

// destroy a barrier
int pthread_barrier_destroy(pthread_barrier_t* barrier);

// wait on a barrier. When the function returns,
// one thread gets PTHREAD_BARRIER_SERIAL_THREAD, others get 0
int pthread_barrier_wait(pthread_barrier_t* barrier);
```

### Sync mechanisms in this course

- We can use mutex, condition, read/write lock, barriers, and only call blocking functions!
- The following mechanisms are not allowed for synchronization!
  - Busy waiting, like spin lock
  - Sleep functions

\*\*\* important for HWs and exams

## Lecture 26 - S1 - Intro to Internet Sockets

---

Mon. Nov. 11, 2019

### Sockets

- Allows both inter-machine and inter-process communications
  - Simple APIs for IPC, including communications over network

## CSE 3100 Master Notes

- Communication channels created without forking • Greater interoperability
- Plan
  - Introduction to networking and sockets
  - Socket-based communications
    - Client and Server
  - Protocol

### Network Sockets

- We will focus on network sockets
  - Networking service is provided by OS
  - Two processes on same/different machines can talk to each other
- TCP/IP provides end-to-end communication layer in the Internet
  - IP: addressing and routing
  - TCP: other services such as error correction and retransmission

### Network Protocol Layers

- Network are complex, with many “pieces”:
  - hosts
  - routers
  - links to various media
  - applications
  - protocols
  - hardware, software

**Question: is there any hope of organizing structure of network? ... or at least our discussion of networks?**

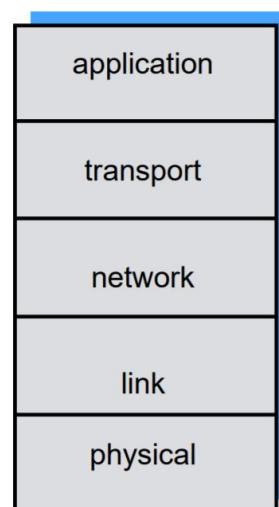
### Why layering?

#### Dealing with complex systems!

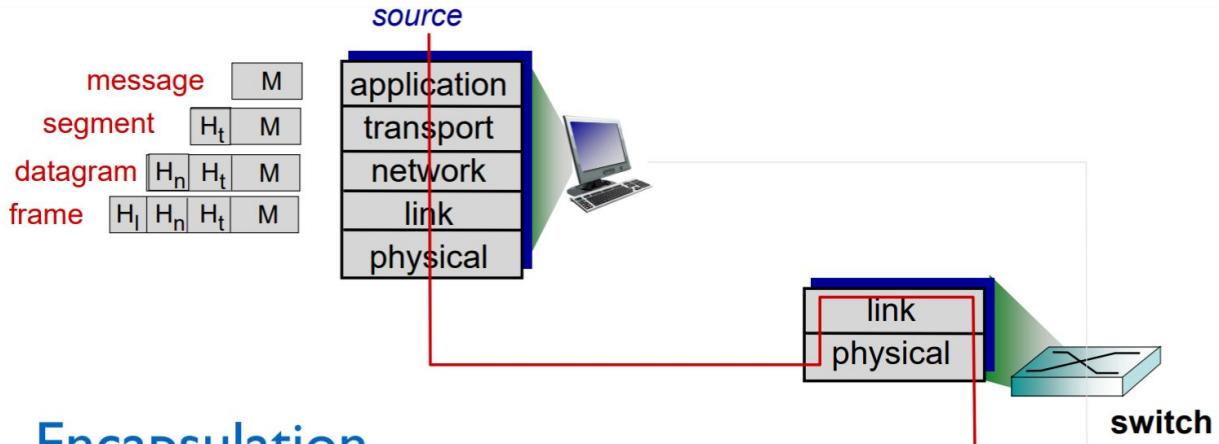
- Explicit structure allows identification, relationship of complex system's pieces
  - layered reference model for discussion
- Modularization eases maintenance, updating of system
  - change of implementation of layer's service transparent to rest of system
- layering considered harmful?

### Internal protocol stack

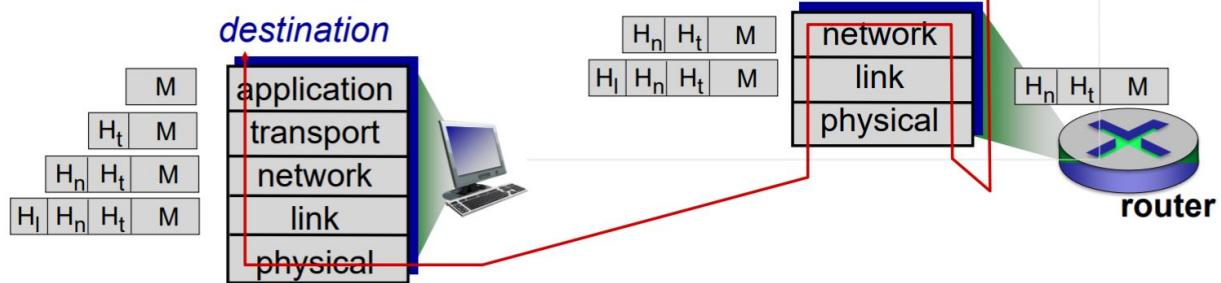
- application: supporting network applications
  - FTP, SMTP, HTTP
    - transport: process-process data transfer
  - TCP, UDP



- network: routing of datagrams from source to destination
  - IP, routing protocols
- link: data transfer between neighboring network elements
  - Ethernet, 802.111 (WiFi), PPP
- physical: bits “on the wire”



## Encapsulation



## IP (Internal Protocol)

- A datagram service at the network layer
  - Data are divided into packets
  - Every packet has a destination “label”: the IP address!
  - Packets are handled independently
  - And delivered with best-effort
    - However, you can see loss, out of order, duplication,..

## IP Address

- IPv4: 4 bytes
  - Each device normally gets one (but can have more)
    - Each interface(boundary between host and physical link) has one IP address
- IPv6: 16 bytes

## CSE 3100 Master Notes

- 128 bit addresses
  - Make it feasible to be very wasteful with address allocations
- Lots of other new features
  - Built-in auto configuration, security options, ...
- Not compatible with IPv4

### IPv4 examples

#### A Google server:

67.218.93.15

#### A gateway at UConn:

137.99.10.1

#### Private addresses:

10.0.0.0 - 10.255.255.255

172.16.0.0 - 172.31.255.255

192.168.0.0 - 192.168.255.255

### Domain name for dummies

- A domain name is an identification string
  - Formed by rules and procedures in DNS (Domain name system)
  - Hierarchical
- Fully qualified domain names (FQDNs) is a complete domain name for a specific computer on the Internet
- Domain name servers (also a DNS) translate domain names to IP address
  - Phonebook for the Internet

#### Example:

www.uconn.edu is a FQDN.

www.google.com is a FQDN. Its IP address is 172.217.7.228

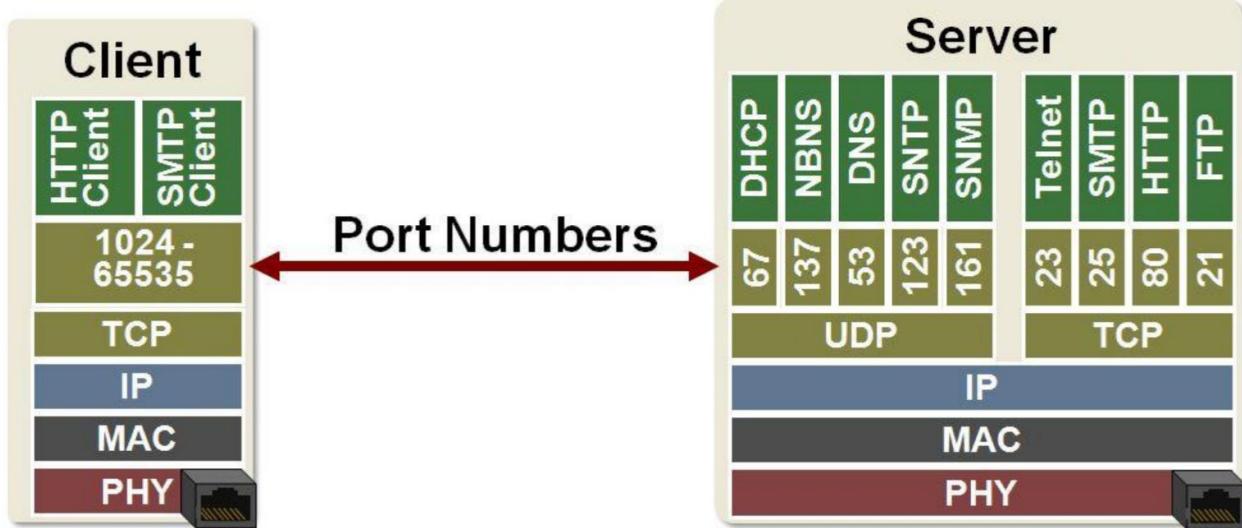
### The Internet Transport Layer

#### Two protocols above IP

- TCP (Transmission Control Protocol) TCP/IP
  - **Connection-oriented** like POTS (plain old telephone service)
  - Flow-control and bi-directional
    - Reliable
- UDP (User Datagram Protocol) UDP/IP
  - **Connection-less** like (unregistered) snail mail
  - No acknowledgments or retransmissions
  - Packets may be delivered out of order and have duplicates
    - Retransmission can be handled in applications

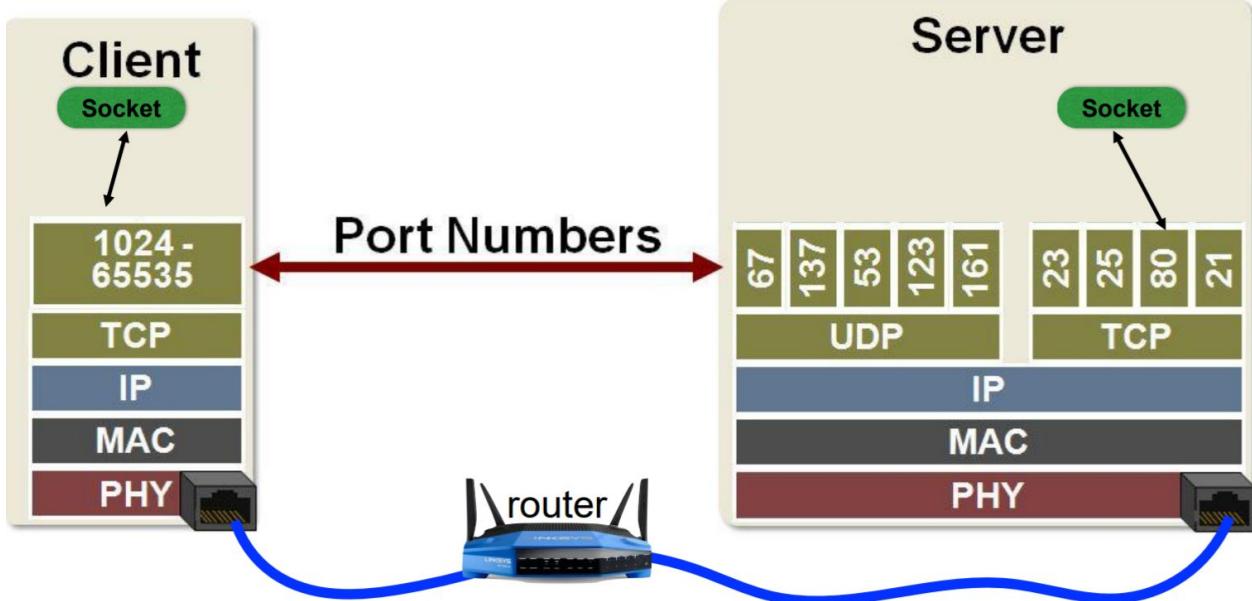
## Ports

- A server can provide many services
- Ports Identify services and applications for both TCP and UDP



## Sockets

- The programmatic abstraction to access ports



### Commonly used TCP port numbers

| Port number | Service/ Protocol |
|-------------|-------------------|
| 21          | FTP               |
| 22          | SSH               |
| 25          | SMTP              |
| 80          | HTTP              |
| 137-139     | NetBios           |
| 443         | HTTPS             |

### Sockets API

- API to create, close, read, and write
  - Sockets are like files!
  - Can use the normal I/O functions
- Two types of sockets
  - Stream sockets for TCP
  - Datagram sockets for UDP
    - Single payload oriented
- Need to associate a socket with a port (coming up next time)
  - An Internet address (IPV4 or IPV6)
  - A port number

### Create a socket

```
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
// Use close() to close a socket
```

Returns a file descriptor for the given domain, type, and protocol.  
Returns -1 on error.

domain: AF\_INET, AF\_INET6, etc.  
type: SOCK\_STREAM, SOCK\_DGRAM, etc.  
protocol: 0.

### Send and recv for TCP

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t send(int sfd,const void *buf,size_t len,int flags);
```

## CSE 3100 Master Notes

```
ssize_t recv(int sfd, void *buf, size_t len, int flags);
```

- Return the number of bytes that have been sent/received
  - It can be less than the number of bytes requested!
  - sent does not mean delivered!
- Both returns -1 on errors
  - Similar to read() and write()
    - You can use read and write if you do not need flags
  - Working with TCP (not UDP)

Example: send

**After a connection is established!**

```
// To send n bytes stored in buf
sent = 0;
while(sent < n){
    int r = send(sid, buf +sent, n - sent, 0);
    if (r < 0) report_error();
    sent += r;
}
```

Example: recv

**After a connection is established!**

```
received = 0;
while(received < n):|{
    int r = recv(sid, buf + received, n - received, 0);
    if (r < 0) report_error();
    received += r;
}
```

Study this slide yourself

**Useful commands. Read the manual for details**

```
#scan ports
nmap localhost

#list FDs for the Internet Need sudo for system files.
lsof - i

#show information about networking subsystem
netstat -l
```

```
netstat -l -4
```

## Lecture 27 - S2: Sockets Programming: Client & Server

Weds. Nov. 13, 2019

### Review

- TCP/IP
  - Name/IP and port number
- Socket
  - We get a file descriptor and we can send and recv

### A Client-Server World

#### Establish a connection between a server and client

- Server
  - Listens on a port for inbound “let’s talk” requests
- Client
  - Initiates the comm. w/ the server
  - Must know
    - Server name (which is used to find the IP address)
    - Protocol to use (TCP/UDP)
    - Port number

### Server side (TCP)

- After creating a socket
  - bind() Server binds (reserves) a port for chatting
  - listen() Server listens on port for “Let’s talk requests”
  - accept() When a request comes in ...
    - Server allocates a separate socket to accept private comm. w/ the client (same host, diff. port)
      - The original socket remains available for future “Let’s talk”
      - The new socket is bound to an ephemeral port
    - Comm. goes on on separate sockets with
      - send() / recv()
    - close() Conversation ends via a close

### Analogy Between Sockets and Telephone

|  |         |           |
|--|---------|-----------|
|  | Sockets | Telephone |
|--|---------|-----------|

## CSE 3100 Master Notes

|                 |                         |                        |
|-----------------|-------------------------|------------------------|
| socket()        | Create a socket         | Buy a telephone        |
| bind()          | Bind to a port          | Plug in / Insert a SIM |
| listen()        | Get ready to respond    | Get ready for ringtone |
| accept()        | Accept requests         | Pick up the phone      |
| send() / recv() | Communication           | Listen and talk        |
| close()         | Close the communication | Hang up                |

### Client side (TCP)

- After creating a socket
  - connect() Client initiates a session
    - Need server name, type, and port
  - Then, comm. goes on with send() / recv()
  - close() Client closes a session

**Note that no bind() is needed on the client side OS allocates a port**

Getting back to server...

```
#include <sys/types.h>
#include <sys/socket.h>

int bind(int sockfd,
         const struct sockaddr *addr,
         socklen_t addrlen);
```

- sockfd is the socket file descriptor. It does not have an address!
- What is struct sockaddr?
  - It holds address information for many types of socket

Socket address types

```
// generic socket address
struct sockaddr{
    unsigned short sa_family; // address family, AF_XXX
    char sa_data[14];           // 14 bytes of protocol address
};

// we are using Internet socket address
struct sockaddr_insa; // socket address IPv4
struct sockaddr_in6 sa6; // socket address IPv6
// struct sockaddr_in* can be cast to struct sockaddr*
```

```
// Do I really need to learn how to set the fields in saand sa6?
```

### addrinfo structure

```
// A recent invention for preparing the socket address
struct addrinfo {
    int    ai_flags;           // AI_PASSIVE, AI_CANONNAME, etc.
    int    ai_family;          // AF_INET, AF_INET6, AF_UNSPEC
    int    ai_socktype;        // SOCK_STREAM, SOCK_DGRAM
    int    ai_protocol;        // use 0 for "any"
    size_t   ai_addrlen;       // size of ai_addr in bytes
    struct   sockaddr * ai_addr; // struct sockaddr_in or _in6
    char   *ai_canonname;      // full canonical hostname
    struct   addrinfo *ai_next; // linked list, next node
}
```

### getaddrinfo()

```
int getaddrinfo(const char *node // "www.example.com" or IP
                const char *service, // "http" or port #
                const struct addrinfo *hints,
                struct addrinfo **res
```

- this function finds addrinfo for you!
  - Provide you a linked-list, res. A list of results!
- **node** is the host name to connect to, or an IP address
- **service** can be a port number, like “80”, or the name of a particular service like “HTTP”
- **hints** points to a struct addrinfo that you’ve already filled out with relevant information.

### Usage of getaddrinfo() for server

```
struct addrinfo hints, *servinfo;           // res will point to the
results

memset(&hints, 0, sizeof(hints));           // make sure the struct
is empty
hints.ai_family = AF_UNSPEC;               // don't care IPv4 or IPv6
hints.ai_socktype = SOCK_STREAM;           // TCP stream sockets
hints.ai_flags = AI_PASSIVE;               // fill in my IP for me
if (getaddrinfo(NULL, "3490", &hints, &servinfo) != 0)
    die("Error");
```

```
// servinfo now points to a linked list of 1 or more struct addrinfos  
// continue on the next page
```

Now, let us create socket and then ibnd!

```
int sockfd, yes = 1;  
struct addrinfo *p      ;  
for (p = servinfo; p != NULL; p = p->ai_next){  
//socket(); setsockopt(); bind();  
// should check return value of errosr!  
sockfd = socket(p->ai_family, p->ai_socktype, p->ai_protocol);  
setsockopt(sockfd, SOL_SOCKET, SO_REUSEADDR, &yes, sizeof(int));  
// bind it to the port we passed in to getaddrinfo();  
bind(sockfd, p->ai_addr, p->ai_addrlen);  
break;  
}  
freeaddrinfo(servinfo);           // free the linked-list
```

\*\*\* should understand it at least once, afterwards, just copy, paste, and use

After bind()

```
// listen on that socket for "Let's talk" message.  
status = listen(sockfd,10);  
checkError(status);  
// client's address information, if you are interested in  
struct sockaddr_storage client_addr;  
while(1) {  
    socklen_t clientSize = sizeof(client_addr);  
    int chatSocket = accept(sockfd,  
                           (struct sockaddr*)&client_addr,  
                           &clientSize);  
    checkError(chatSocket);  
    printf("We accepted a socket: %d\n",chatSocket);  
    close(chatSocket);  
}
```

## Usage of getaddrinfo() for client

```

struct addrinfohints, *servinfo; // will point to the results

memset(&hints, 0, sizeofhints); // make sure the struct is empty
hints.ai_family= AF_UNSPEC; // don't care IPv4 or IPv6
hints.ai_socktype= SOCK_STREAM; // TCP stream sockets
// get ready to connect
if (getaddrinfo("www.uconn.edu", "3490", &hints, &servinfo) != 0)
    die("Error");
// servinfo now points to a linked list of 1 or more struct addrinfos
// continue on the next page

```

But, how can i do bind()?

```

int sockfd, yes = 1;
struct ddrinfo *p;
for (p = servinfo; p != NULL; p = p->ai_next);
// socket();
// should check return value for errors!
sockfd = socket(p->ai_family, p->ai_socktype, p->ai_protocol);
// connect
connect(sockfd, p->ai_addr, p->ai_addrlen);
break;
}
...
freeaddrinfo(servinfo); // free the linked-list

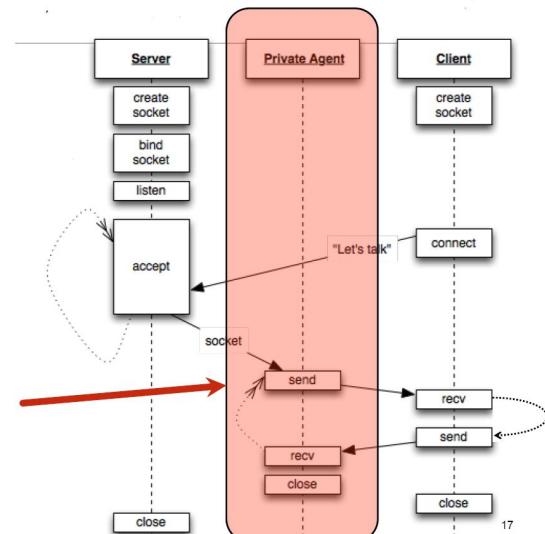
```

## Blocking and Talking to Many Clients

- On the server side
  - Accept is blocking (until a “let's talk” request comes in)
  - If serving the client directly...
    - In bound “Let's talk” are queue up (up to maximum)
  - If the server wishes to get back to accepting: fork a process! or create a

**Could be a cloned process  
(fork!) ... or a thread**

thread!



- All real comm. happens on the accepted socket
- On client side
  - Connect is blocking (until the server responds or refuses)

### Server Dealing with Many Clients

- The server cannot talk to many clients at the same time
  - Let others to deal with clients
  - It can wait for new clients

## Lecture 28 - S3: Sockets Programming: Protocol

---

Mon. Nov. 18, 2019

### Network Protocols

- Rules for communications between two parties
  - Need protocols at every layer
  - TCP and IP are protocols!
- After establishing a TCP connection, two processes talk at application layer
  - Who should talk first?
  - How much can each talk?
  - How are the messages encoded into bytes?

### Byte stream in sockets

- Applications see byte streams into and out of sockets
- It is up to applications to interpret bytes
- Two common options (both parties have to agree !)
  - Everything is in plain text. Need conversion and parsing
    - Example: XML | JSON
  - Raw data. Issues: endianness and padding

**Example:** int v = 128. Send v.

Option 1: three ASCII characters with a delimiter: '1', '2', '8', '\n'

Option 2: four bytes: 0x00, 0x00, 0x00, 0x80 (or 0x80, 0x00, 0x00, 0x00?)

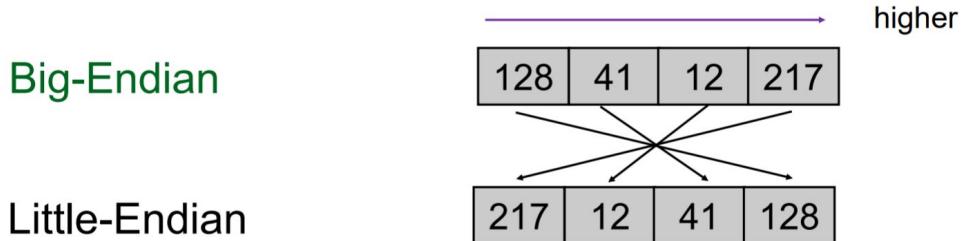
**Is the size of v really 4?**

### Endianness

- Two ways to use bytes if a data item has more than 1 byte
  - Some hosts use big-endian while some use little-endian
- Both parties need an agreement to interpret the bytes

Integer: 2150173913 or 0x80 29 0C D9

Bytes: 128(0x80), 41(0x29), 12(0x0C), 217(0xD9)



Change byte ordering for 16-bit or 32-bit integers

- Two orderings may not be the same
    - Host Byte Ordering natural for platform
    - Network Byte Ordering always big-endian
  - Convert to network byte ordering before send
  - Convert to host byte ordering after receive
    - htonl and ntohl for 32-bit. htons and ntohs for 16-bit

```
#include <arpa/inet.h>

uint32_t htonl(uint32_t hostlong);
uint16_t htons(uint16_t hostshort);
uint32_t ntohl(uint32_t netlong);
uint16_t ntohs(uint16_t netshort);
```

## Sending a string

- We are fine in this course because we assume ASCII strings
    - Each character has only one byte
  - In real word, modern applications assume Unicode internally
    - Especially important for web servers
      - Your customers/users speak many kinds of languages!

## Example:

Python 3 uses Unicode strings by default

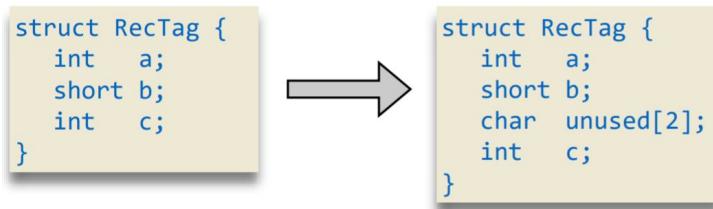
```
"♥♫".encode('utf-8') # encode a string to byte array using UTF-8  
data.decode('utf-8') # decode a byte array using UTF-8 to a string
```

## Issues with structure: alignment and padding

- Compilation can introduce padding inside structures
    - Affect the structure size and the offset of the fields
    - Important when the two endpoints have different architectures

## CSE 3100 Master Notes

- e.g. 32 bit vs. 64 bit even if both are Intel
- How to deal with the issue?
  - Use explicit padding! (or reorder fields)
  - Do not use any language/machine dependent data/structure



**Remember to change byte order for a, b, and c !**

How about arrays?

- Convert each element to bytes individually
- How about arrays of structures that have array members?
  - You can see it is becoming very tedious

**JSON: JavaScript Object Notation is an open-standard file format that uses human-readable text to transmit data objects consisting of attribute– value pairs and array data types (from Wikipedia)**

Example: SMTP(Simple Mail Transfer Protocol)

S: 220 smtp.example.com ESMTP Postfix  
C: HELO relay.example.com  
S: 250 smtp.example.com, I am glad to meet you  
C: MAIL FROM:<bob@example.com>  
S: 250 Ok  
C: RCPT TO:<alice@example.com>  
S: 250 Ok  
C: DATA  
S: 354 End data with <CR><LF>.<CR><LF>  
C: From: "Bob Example" <bob@example.com>  
C: To: Alice Example <alice@example.com>  
C: Date: Tue, 15 Jan 2008 16:02:43 -0500

Each line ends with , "\r\n"

Server may respond with many lines, for example, to EHLO command

**Based on an example on Wikipedia**

## Example: SMTP - 2

```
C: Subject: Test message
C:
C: Hello, Alice,
C: This is a test message.
C: Bob
C: .           Here is the line indicating the end of DATA
<CR><LF>.<CR><LF>
S: 250 Ok: queued as 12345
C: QUIT
S: 221 Bye
{The server closes the connection}
```

## Example: TFTP-style applications

### Trivial File Transfer Protocol (TFTP)

- Server waits for commands and performs operations accordingly
- Client sends a command to the server and waits for the response
- Three commands are supported
  - LS              List files in the directory on the server side
  - GET fn          Download a file named “fn” from the server
  - EXIT             End the session

## Commands

Client sends a command

Server responds

| Commands | Response             |
|----------|----------------------|
| LS       | A list of file names |
| Get fn   | A file               |
| EXIT     | nothing              |

## How to encode the command, messages, and files?

## Command Options: Text

- A command is an ASCII string ending with a delimiter, e.g., '\n'
  - Both parties have to agree on the delimiter

Example:

## CSE 3100 Master Notes

```
GET(fn)
    send_str("get hello.txt\n"); // no need to send NUL
```

| Pro                     | Cons                                    |
|-------------------------|-----------------------------------------|
| Easy for humans to read | Require parsing                         |
| Flexible and compatible | Variable length. Check of the delimiter |
|                         | Not compact for long commands           |

### Command Options: Binary, fixed size

- A command is represented by a structure
  - All commands have the same size (good or bad?)
- Example:

```
#define LS      0
#define GET     1
#define PUT     2
#define EXIT    3

struct Command {
    int code;
    char fileName[252];
} cmd;

cmd.code = htonl(GET);
send_all(sid, &cmd, sizeof(cmd));
```

| Pro              | Cons                      |
|------------------|---------------------------|
| Nothing to parse | Hard to read and maintain |
| Fixed size       | Fixed size                |
|                  | Endianness                |
|                  | Padding                   |

### Commands Options: Dynamic Size

- LS and EXIT need only 4 bytes (or we can use only one byte)
- GET needs to send additional (filename length + 4) bytes

```
#define LS      0
#define GET     1
#define PUT     2
#define EXIT    3

int code = htonl(GET);
int fnlen = strlen(fname);
int nlen = htonl(fnlen);

send_all(fd,&code,sizeof(int));
send_all(fd,&nlen,sizeof(int));
send_all(fd,fname,fnlen);
```

| Pro              | Cons                      |
|------------------|---------------------------|
| Nothing to parse | Hard to read and maintain |
| Dynamic sizing   | Endianness                |
|                  | More brittle              |

## Sending the response

```
typedef struct ResponseHeader_tag {
    int code;
    int length;
} ResponseHeader;
```

- code: type of the payload: text message or file
- length: the number of bytes in the text message or file

A payload is followed by length bytes of message or file contents

Remember the byte ordering of code and length!

## Recapping

- Check return value
- Validate input
- Understand strings in C
  - Always think about if you have enough space for strings
    - strcpy(), strcat(), etc.

| Don'ts                | Do's                              |
|-----------------------|-----------------------------------|
| scanf("%s", buf);     | scanf("%10s", &buf);              |
| gets(buf);            | fgets(buf, BUF_SIZE, stdin);      |
| sprint(buf, "5s", s); | snprintf(buf, BUF_SIZE, "%s", s); |

# Lecture 29 - T6: Threads - Some Loose Ends

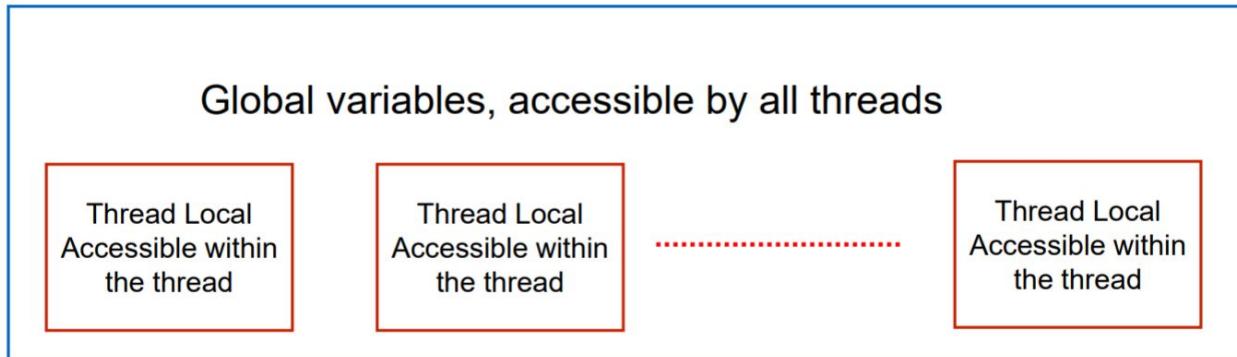
Weds. Nov. 20, 2019

## Overview

- Thread Local Storage
- Thread Cancellation
  - Purpose
  - APIs
- Priority Inversion

## Thread Local Storage (TLS)

- Purpose
  - Have “Globals” that are “Local” to each thread!
  - Can be accessed by all the functions in a thread



## TLS in C

- Quite easy to do with modern C compilers
- In C
  - MSVC, gcc and clang have a new storage keyword: `__thread`
  - C11 has `_Thread_local`  
`__thread int x = 0;`
- “x” is a global to each thread.
- Each thread regardless of when it is created gets its own copy of “x” !!!

## Cancellation use case

- Imagine that you have several threads to find a solution to a problem
  - You do not know which could find the answer the fastest

## CSE 3100 Master Notes

- Example: search a value in a large data set
- What do you do when one thread finishes?
  - It would be better to ask other threads to stop

### Our need

- Ability to ask from thread A to terminate thread B
- Requirements
  - Thread A must know B's identity
  - But B should terminate "safely"

### What do we mean by "Safely"

- The thread receiving the request may be...
  - In the middle of doing some IO (blocking in the kernel)
  - In the middle of updating some data-structure
  - In the middle of a critical section
    - Therefore holding locks!
  - Asleep in a POSIX condition (so on wake up he would hold locks!)
- Stopping "brutally" in any such scenario is a recipe for disaster

### Cancellation is "civilized termination"

- One thread does not "kill" another
- One thread asks another thread to...

**"please stop at your earliest convenience"**
- The recipient (cancellation target) can
  - Ignore it
  - Honor it immediately
  - Defer to a "safe point"
    - After a critical section
    - IO done
    - Locks released, etc....

### Two sides to the story

- The thread asking another to cancel itself
  - Quite easy, just call `pthread_cancel()`
- The thread receiving the request to self-terminate
  - A bit more involved to handle safety

### Sender: call `pthread_cancel`

```
int pthread_cancel(pthread_t thread);
```

## CSE 3100 Master Notes

- It requires the thread identifier of the “victim”
- It returns as soon as the request is logged
- Recall that it is merely a request
- Victim may keep on executing for a while

### Recipient's perspective

- Recipient can be in different “moods”
  - Mood is based on State and Type

|                       | Type<br><b>Deferred</b>                        | Type<br><b>Asynchronous</b>                |
|-----------------------|------------------------------------------------|--------------------------------------------|
| State <b>disabled</b> | cancellations remain pending                   | cancellations remain pending               |
| State <b>enabled</b>  | cancellations occur at next cancellation point | cancellations may be processes at any time |

### State and Type

- A thread can
  - Change its cancellation **state** at any time

```
int pthread_setcancelstate(int state, int *oldstate);
```

- • Change its handling **type** at any time

```
int pthread_setcanceltype(int type, int *oldtype);
```

Previous values are returned through the second argument

### Request Deferred until a Cancellation Point

- A cancellation point is a function in the code where a cancellation request can be safely honored.
- Most of the cancellation points are
  - System calls
  - C library functions

## CSE 3100 Master Notes

- Thread related
- IO related

**Remember cancellation request is honored only if the thread is in the right “mood”**

### List of Cancellation Point

- Predefined cancellation point  
man 7 pthreads,  
or <http://man7.org/linux/man-pages/man7/pthreads.7.html>

```
accept, aio_suspend, close, connect, creat, fcntl, fsync, lockf, msgrcv, msgsnd, msync,
nanosleep, open, pause, poll, pread, pselect, pthread_cond_timedwait, pthread_cond_wait,
pthread_join, pthread_testcancel, pwrite, read, readv, recv, recvfrom, recvmsg, select,
sem_wait, send, sendmsg, sendto, sigpause, sigsuspend, sigwait, sleep, system, tcdrain,
usleep, wait, waitpid, write, writev
```

- Most are system calls

### pthread\_testcancel

- What if your code does not call any of the functions?

Add a cancellation point:

```
pthread_testcancel();
```

- It does nothing if there is not cancellation request or cancelability is not enabled
- It will handle the cancellation request otherwise

### Clean up

- The cancelled thread must use the cancellation point to “cleanup”
  - Release locks
  - Close resources
  - Free memory
  - ...

Leave data in a coherent state

- How to do this?
  - With cleanup handlers

### Cleanup Handlers

- Functions that must be executed at cancellation time
  - A handler is one function
  - You get access to the function via pointers
  - Invoked handler receives a void\* to an arbitrary data structure
- Multiple handlers?
  - You have a stack of “scheduled” cleanup handlers
  - You push a handler /data pair on the stack if a cleanup is needed

- You pop a handler / data pair if the cleanup is no longer needed

### Demo!

## Asynchronous Cancellation?

### STAY AWAY FROM IT

- No need for victim to check with pthread\_cancel
- Victim killed without any chances to recover
- Rarely needed
- Difficult to use correctly
  - Safe to use when computation bound with no locks / resources

## Doing something only ONCE

- Make sure that an operation is done only once, even if called from several threads

```
#include <pthread.h>

int pthread_once(pthread_once_t *once_control,
                 void (*init_routine)(void));
```

- Use case examples
  - Initialize a mutex once
  - Set the seed of a random number generator once

## Pthread scheduling

- You can change
  - Thread scheduling policies: FIFO, round-robin, ...
  - Thread scheduling priorities: min. 32 priority levels

### See pthread\_setschedparam()

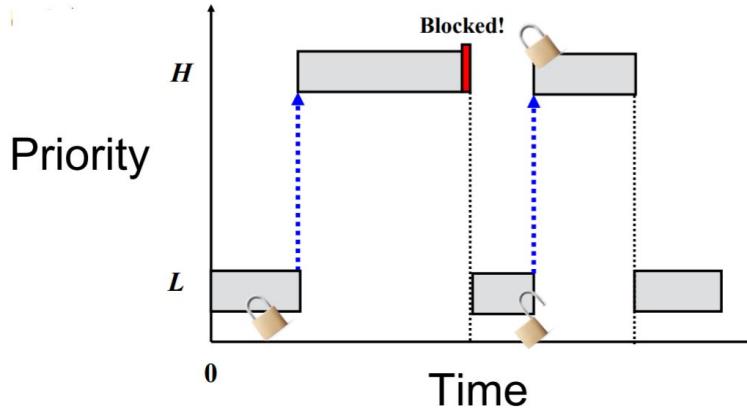
## Priority inversion

- High priority threads cannot run while low priority threads can
- The result of a conflict between
  - Resource locking
  - Priority scheduling

## Priority inversion example

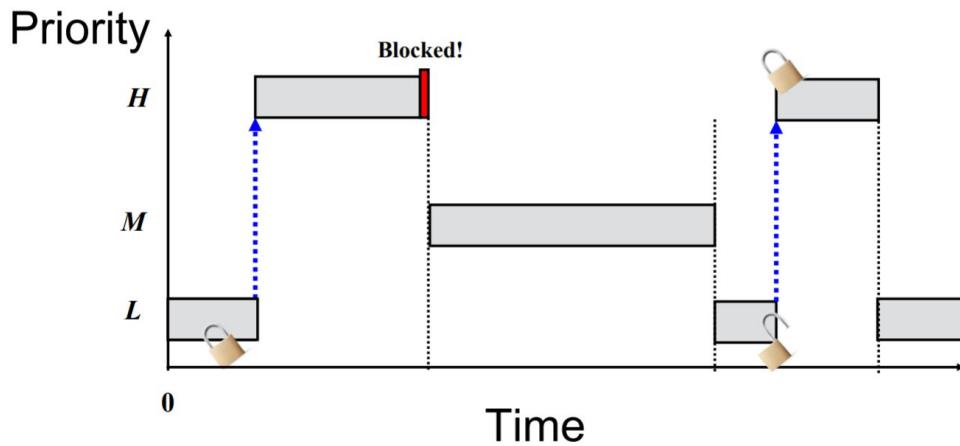
- Mutex is locked by a low priority thread L
- A high-priority thread H trying to grab a lock...

- H goes to sleep because the mutex is locked
- So the low priority thread gets to execute before the high-priority thread!



Can be worse!

- Thread M has a priority level between H and L
- M does not need the mutex. It runs before L, and before H

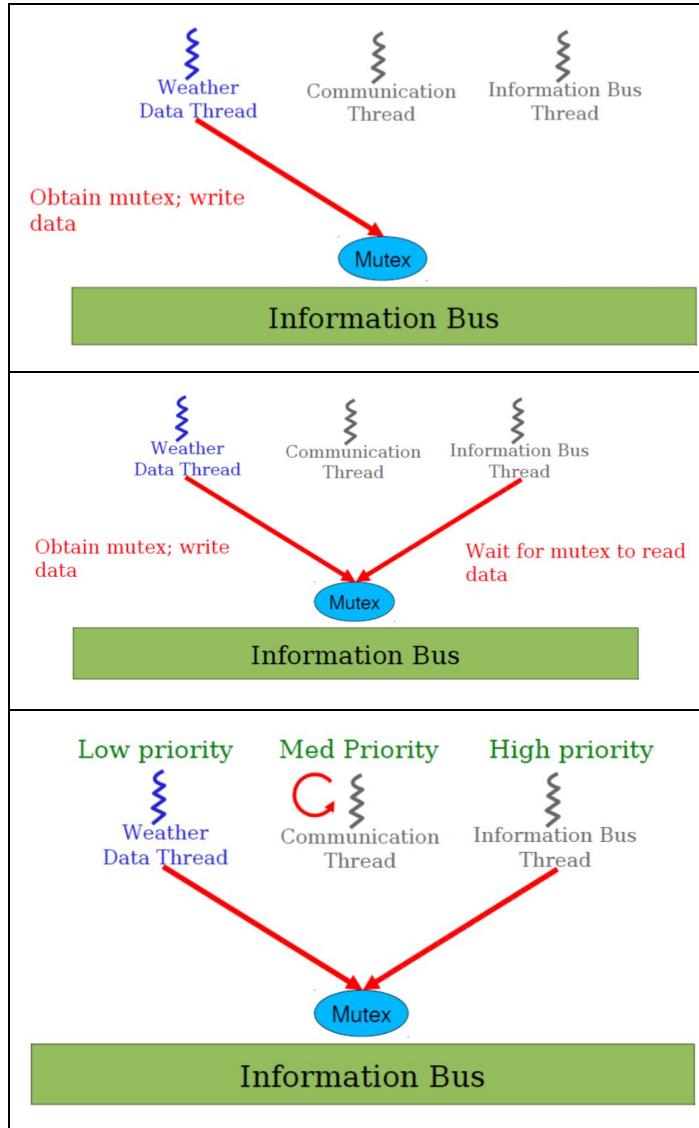


Solution? Priority boosting

- Trouble in the example
  - Low priority thread holding a lock preventing high priority thread to run
- Solution
  - Temporarily boost the priority of the thread holding the lock to match the priority of the blocked high-priority thread
  - It will release the mutex sooner and let the real high-priority thread run
- In practice
  - Handled by a priority inheritance attribute associated to the mutex

## Priority inversion on Mars

- Priority inversion caused Mars Pathfinder rover dozens of unexpected system resets



## Lecture 30 - Practice Exam 3

---

Mon. Dec. 2, 2019

## Rock, paper, and scissors

---

## CSE 3100 Master Notes

This is a game everyone has played. There are two players in the game. Each player makes a choice of a hand shape, rock, paper, or scissors and compares their choices.

In this problem, we will use two threads to simulate two players playing the game for `n` rounds. A third thread serves as a referee, a trusted party that compares players' choices. In each round, players make a decision and inform the referee their choice. The referee compares the choices and announces the result to both players.

`thread_player()` and `thread_referee()` are the main function of two types of thread. Complete the functions so two players can play the game.

Structure `shared_int_t` is used for players to send their choice to referee. The referee shares `choice1` with player 1 and `choice2` with player 2. `mutex` and `cond` are used for synchronization.

Structure `result_t` is for referee to announce the outcome to players. The referee shares `result` with both players. `barrier` is used for synchronization.

The program takes several arguments. For example, `-n` option specifies the number of rounds. Read the code for details.

Here is the sample output of the program.

```
$ ./rock-paper -q
Player 1 won 2 times, lost 5 times, and tied 3 times.
Player 2 won 5 times, lost 2 times, and tied 3 times.

$ ./rock-paper -n100 -q
Player 1 won 29 times, lost 38 times, and tied 33 times.
Player 2 won 38 times, lost 29 times, and tied 33 times.

$ ./rock-paper -n1000 -q
Player 1 won 318 times, lost 333 times, and tied 349 times.
Player 2 won 333 times, lost 318 times, and tied 349 times.
```

## Additonal challenges

---

Try different synchronization mechanisms.

### Use mutex and cond in `result_t`

---

Structure `result_t` does not have to use barrier for synchronization. It can rely on mutex and cond.

## Use barrier in `shared_int_t`

In this problem, we can use barrier in `shared_int_t` for synchronization.

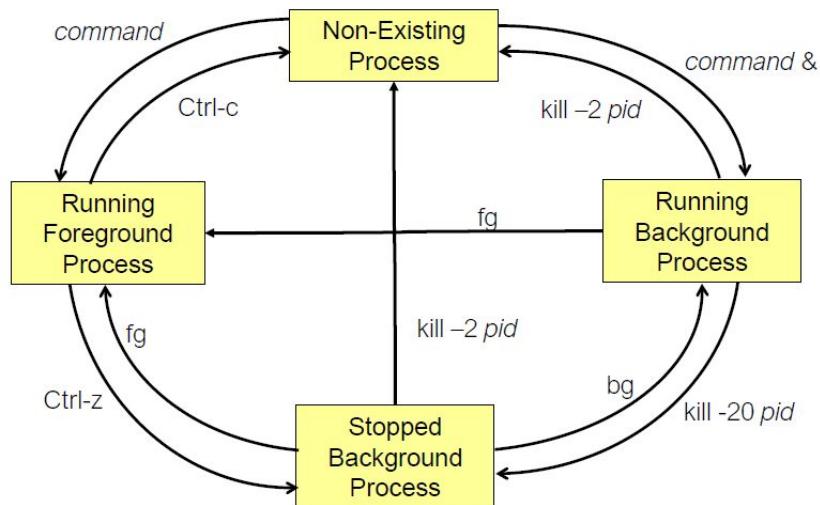
## Lecture 31 - X1: Misc Topics: Signals

Weds. Dec. 4, 2019

### Process Control & Signals

- Unix process control
- Signals (ABC 12.4)
  - From keyboard
  - Sent via function calls
  - Signal handlers
- Overview of signals
  - man -7 signals

### Unix process states



### Signals

- A signal is an **asynchronous** event that is delivered to a process
  - Asynchronous means that the event can occur at any time
  - May be unrelated to the execution of the process, e.g., user types `ctrl+C`
- What happens when an event is triggered?
  - Event gains attention of the OS

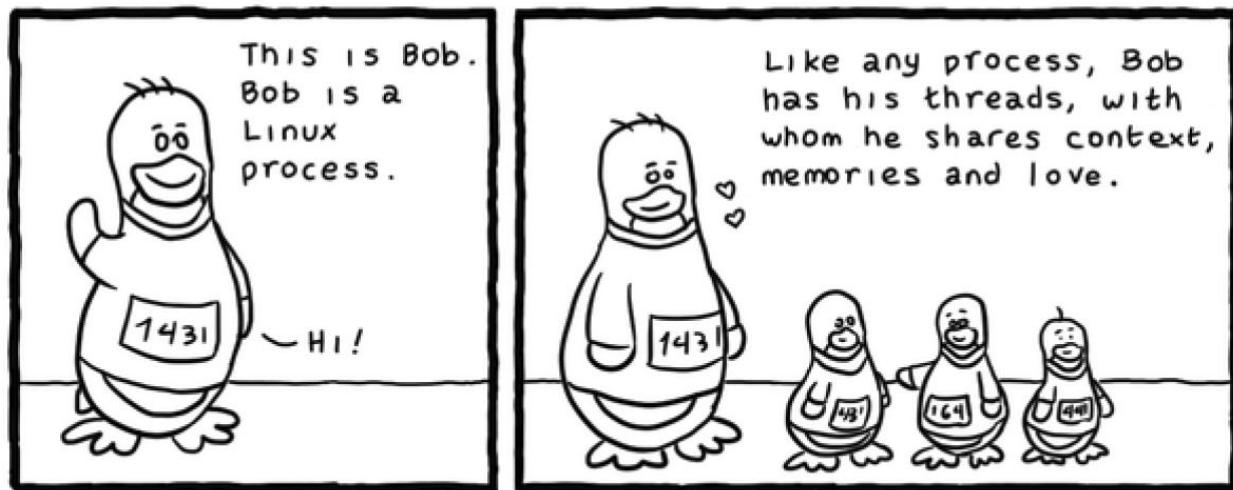
## CSE 3100 Master Notes

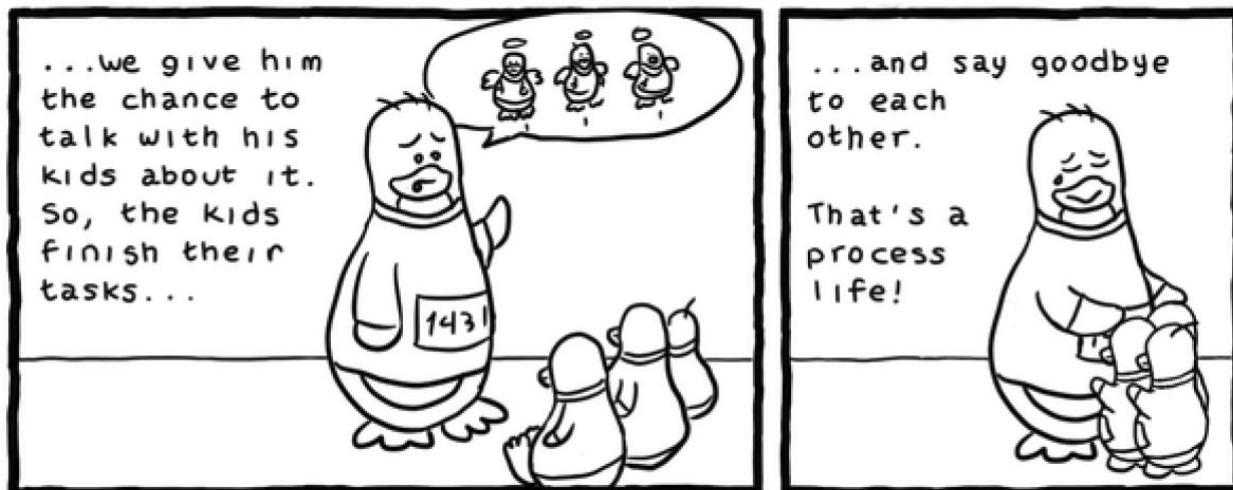
- OS stops the application process immediately
- **Signal handler** executes to completion
- Application process resumes where it left off (unless signal is to terminate it)
- Some example signals. Run “kill -l” for full list

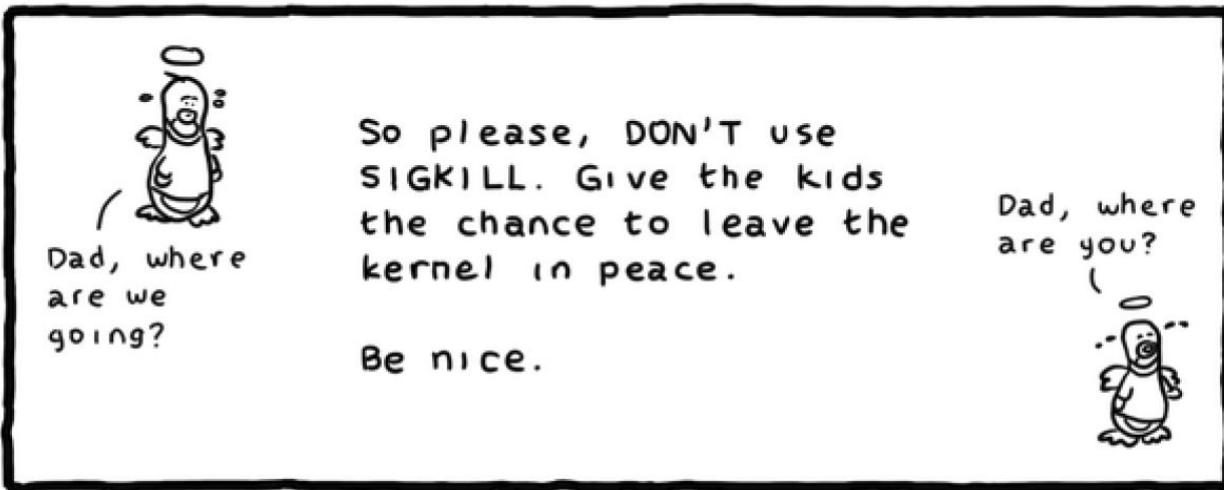
| Name    | Description                  | Default Action    |
|---------|------------------------------|-------------------|
| SIGINT  | Interrupt character (Ctrl-C) | terminate process |
| SIGTERM | Terminate process            | terminate process |
| SIGKILL | kill -9                      | terminate process |
| SIGFPE  | Floating-point exception     | create core image |
| SIGSEGV | Invalid memory reference     | create core image |
| SIGPIPE | Write to a(no reader) pipe   | terminate process |
| SIGALRM | alarm() clock ‘rings’        | terminate process |
| SIGSTOP | Suspend the process          | suspend process   |

- SIGTERM vs. SIGKILL ???
  - Graceful vs. forced termination

Don't SIGKILL (unless you have to)







## Sending Signals via Commands

- “kill command
  - **kill -signal pid**
    - Send a signal of type signal to the process with id pid
    - Can specify either signal type name (-SIGINT) or number (-2)
  - No signal type name or number specified => sends 15/SIGTERM signal
    - Default 15/SIGTERM handler exits process
  - Editorial comment: Better command name would be sendsig
- Examples  
`kill -2 1234`    or    `kill -SIGINT 1234`
  - Same as pressing Ctrl-C if process

## Sending Signals via Function Calls

```
int kill(pid_t pid, int sig);
int raise(int sig); // send to itself
```

- Return 0 on success
- Send a signal sig to process pid of itself
- `raise(sig)` is equivalent to `kill(getpid(), sig)` for single-thread apps
- Editorial comment: Better function name would be `sendsig()`

### Example

```
// Process sends itself a SIGINT signal (commits suicide?)
kill(getpid(), SIGINT);
raise(SIGINT);
```

### Sending Signals via Keystrokes

- Three signals can be sent from keyboard:
  - Ctrl-C → 2/SIGINT signal
    - Default handler exits process
  - Ctrl-Z → 20/SIGTSTP signal
    - Default handler suspends process
  - Ctrl-\ → 3/SIGQUIT signal
    - Default handler exits process
    - Core dump may be saved

### Process Control Implementation

#### Exactly what happens when you:

- Type Ctrl-C?
  - Keyboard sends hardware interrupt
  - Hardware interrupt is handled by OS
  - OS sends a 2/SIGINT signal
  - Signal handler for 2/SIGINT signal executes to completion
    - Default signal handler for 2/SIGINT signal exits process
- Type Ctrl-Z?
  - Keyboard sends hardware interrupt
  - Hardware interrupt is handled by OS
  - OS sends a 20/SIGTSTP signal
  - Signal handler for 20/SIGTSTP signal executes to completion
    - Default signal handler for 20/SIGTSTP signal suspends process
- Issue a “fg” or “bg” command?
  - OS sends an 18/SIGCONT signal (and does some other things too!)
  - OS sends an 19/SIGCONT signal

### Signals generated by abnormal events

- Example: process makes illegal memory reference
  - Event gains attention of OS
  - OS stops application process immediately, sending it a 11/SIGSEGV signal
  - Signal handler for 11/SIGSEGV executes to completion
    - Default signal handler for 11/SIGSEGV signal prints “segmentation fault” and exits process

### Responding to Signals

- A process can:
  - Carry out the default action for that signal

## CSE 3100 Master Notes

- Ignore the signal by setting a mask ( `sigprocmask()` )
- Catch a signal with signal handlers ( `signal()` or `sigaction()` )
  - And then possibly resume execution or terminate
- A process cannot ignore and/or catch SIGKILL and SIGSTOP

### sigaction

```
#include <sys/time.h>
int sigaction(int signum, const struct sigaction*act,
              struct sigaction*oldact);
```

- signum: specifies the signal to change
- act: new action for signal
- oldact: the existing action

```
struct sigaction{
    // the signal handler is either sa_handler or sa_sigaction
    void (*sa_handler)(int); // only takes one parameter
    void (*sa_sigaction)(int, siginfo_t*, void *); // 3 parameters
    sigset_tsa_mask; // mask of signals to be blocked in sig.    handler
    int sa_flags; // Flag SA_SIGINFO selects sa_sigaction
    void (*sa_restorer)(void);
};
```

### Block signals in signal handler

- During the execution of signal handler, the same signal is blocked
- Additional signals can be blocked by setting mask `sa_mask` in struct `sigaction`

#### Example:

```
sigemptyset(&sa.sa_mask); // clear the set
sigaddset(&sa.sa_mask, SIGINT); // block SIGINT
sigaddset(&sa.sa_mask, SIGQUIT); // block SIGQUIT
```

**Note:** Many functions cannot be called safely in a signal handler

“Man signal-safety” for a list of safe functions

### Example Problem

- What if the user types Ctrl-C?
  - OS sends a 2/SIGINT signal to the process
  - Default handler of 2/SIGINT exits the process
- Problem: temporary files not deleted
  - Process dies before `remove("tmp.txt")` is executed
- Challenge: Ctrl-C could happen at any time
  - Which line of code will be interrupted???
- Solution: Install a signal handler

## CSE 3100 Master Notes

- Define a “clean up” function to delete the file
- Install the function as a signal handler for 2/SIGINT

### Examples

#### Catch a signal (in demo code x1.sigaction)

- Set a timer (in demo code x1.timer)

```
alarm()  
setitimer()  
timer_create()
```

### Alarms

```
#include <unistd.h>  
unsigned int alarm(unsigned int seconds);
```

- Sends 14/SIGALRM signal after **seconds** seconds
- Cancels pending alarm if **seconds** is 0
- Uses real time, a.k.a. **wall-clock time**
  - Time spent executing other processes counts
  - Time spent waiting for user input counts
- Used to implement time-outs
- Better to use setitimer() or POSIX timers

### Interval Timers

```
#include <sys/time.h>  
int setitimer(int which, const struct itimerval *new_value,  
              struct itimerval *old_value);  
int getitimer(int which, struct itimerval *curr_value);
```

- What gets counted is specified by which
  - **ITIMER\_REAL**, decrements in real time (wall time), delivers SIGALRM
  - **ITIMER\_VIRTUAL**, decrements only when process is executing (CPU time), delivers **SIGVTALRM**
  - **ITIMER\_PROF**, decrements when process is executing or OS is executing on behalf of the process, delivers **SIGPROF**

## POSIX timer

```
#include <signal.h>
#include <time.h>
// link with -lrt
inttimer_create(clock_t clockid, struct sigevent * sevp,
                 timer_t * timerid);

inttimer_settime(timer_t timerid, int flags,
                 const struct itimerspec * new_value,
                 const struct itimerspec * old_value);
```

- Create a timer first and then settime
- clockid can be CLOCK\_REALTIME, CLOCK\_MONOTONIC, etc.
- When timer expires, deliver a signal, or create a thread
  - Or check manually