



# P4: Inter-Process Communication with Pipes (ABC 12.3)

---

Ion Mandoiu

Laurent Michel

Revised by M. Khan and J. Shi



# Inter-process communication (IPC)

---

- Files
- **Pipes**
- Named pipes
- Sockets
- Message queues
- Shared memory

## Synchronization primitives

- Semaphores, Signals, etc.

# pipe()

---

```
#include <unistd.h>

int pipe(pipefd[2]);
```


Creates a one-way pipe (a buffer to store a byte stream)

Two FDs in pipefd. **pipefd[0] is the read end**, pipefd[1] is the write end

Returns 0 if successful

Pipes allow IPC. One process writes and the other one reads


write to  
pipefd[1]



read from  
pipefd[0]

# Connecting two processes

- Parent creates a pipe and gets two FDs (e.g., 3 and 4)
- After fork(), the child has 3 and 4, too
- One process can write to FD 3, and the other can read from FD 4
  - Close unused FD!



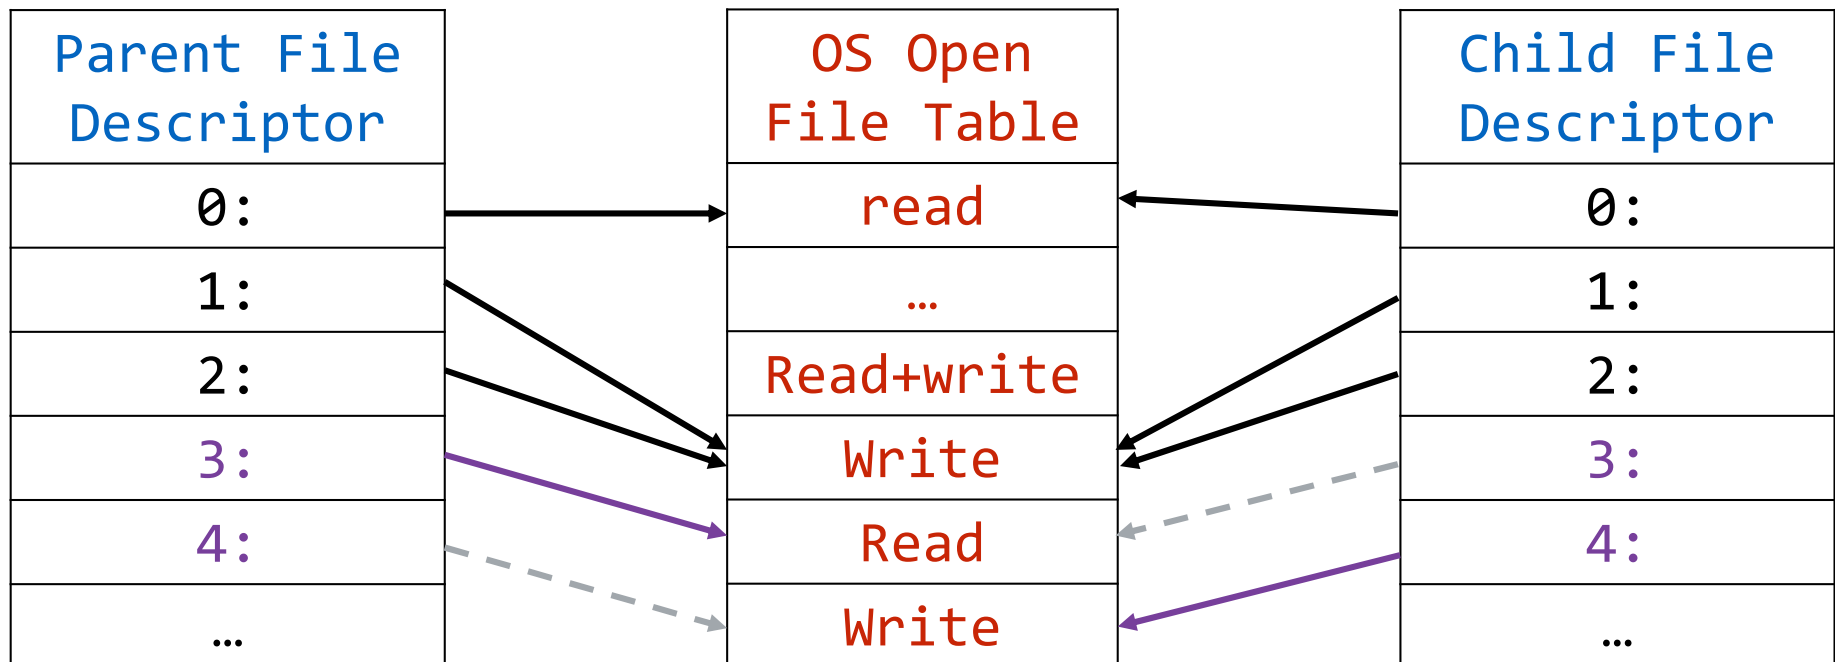
# Closing FDs not in use

If the pipe is for parent to read and for child to write,

Parent: `close(4);`

Child: `close(3);`

Then child can write to and parent can read from the pipe. See demo code!





# Questions

---

- What would you do if you need two-way communications between parent and child?
- After exec, the new program gets the file descriptors for the pipe, too
- How can the new program use the pipe?
  - A program is aware of FDs 0, 1, and 2, but not 3 or 4

# Pipeline in shell

---

- Shell supports pipelines

```
cmd1 | cmd2 arg21 arg22 | cmd3 arg31 ...
```

- stdout of a command is connected to stdin of the next command
  - Done with pipes on Linux/Unix
  - cmd1 writes to a pipe and cmd2 reads from it

Example:

```
ls | tr a-z A-Z | wc
```

# Example: connect two programs with a pipe

---

Start a pipeline in program S (aka, the shell): A | B


- High-level strategy (missing clean up !)
  - Create a pipe
  - Fork #1
    - In child process
      - Redirect stdout to the write end of the pipe
      - Start A, by calling exec
  - Fork #2
    - In child process
      - Redirect stdin to the read end of the pipe
      - Start B, by calling exec



# At the beginning


---

- S has only 0, 1, and 2 open



# Pipe creation

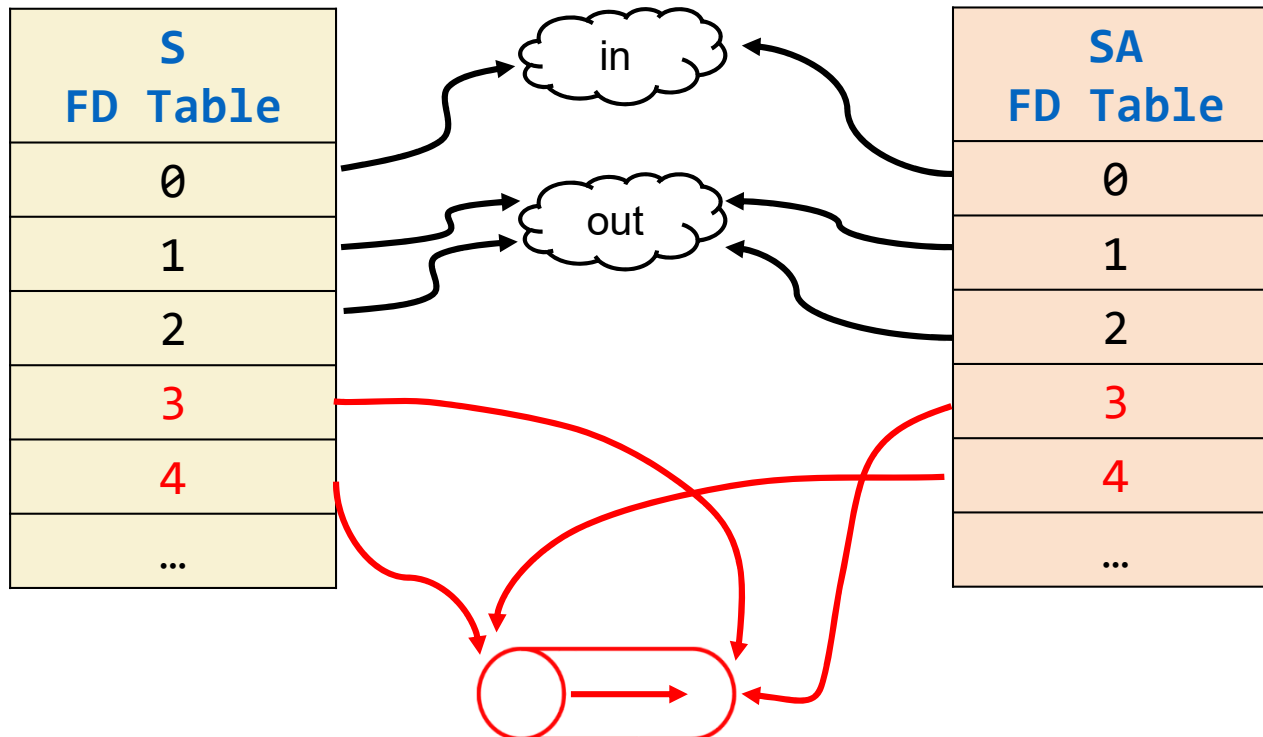
- S creates a pipe by calling pipe()
  - A pair of FDs is returned



# Fork #1

S: fork()

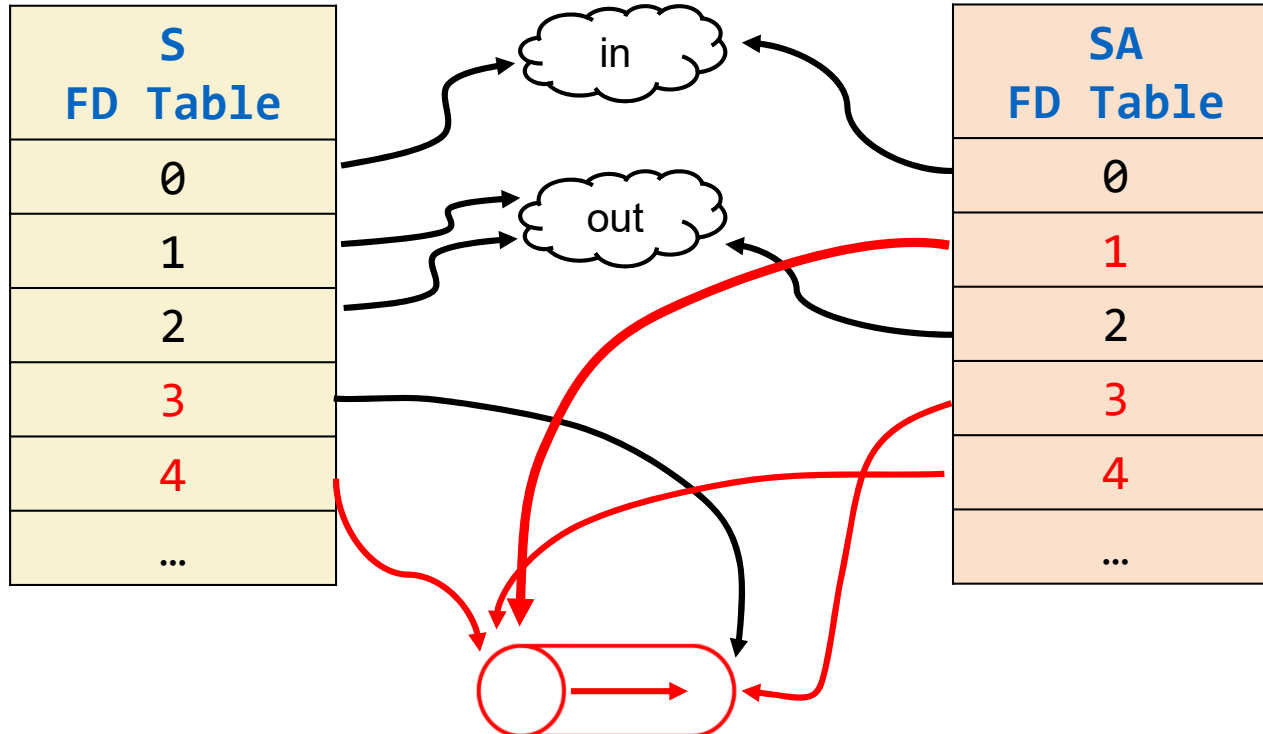
- FD table is duplicated



# Redirect in first child process

SA: `dup2(4, 1)`

- Or `close(1); dup(4);`




# Clean up #1

S: close(4)

SA: close(4); close(3)

SA can then exec into A

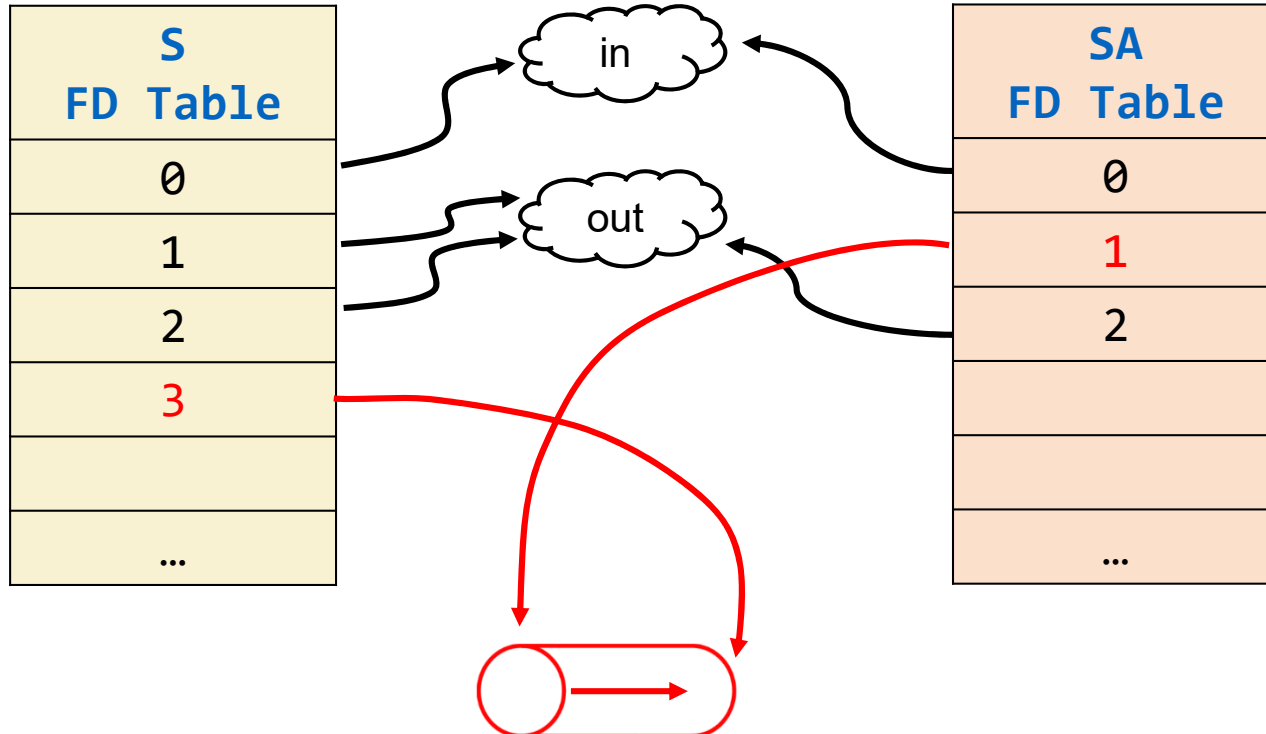


# After clean up #1

S: close(4)

SA: close(4); close(3)

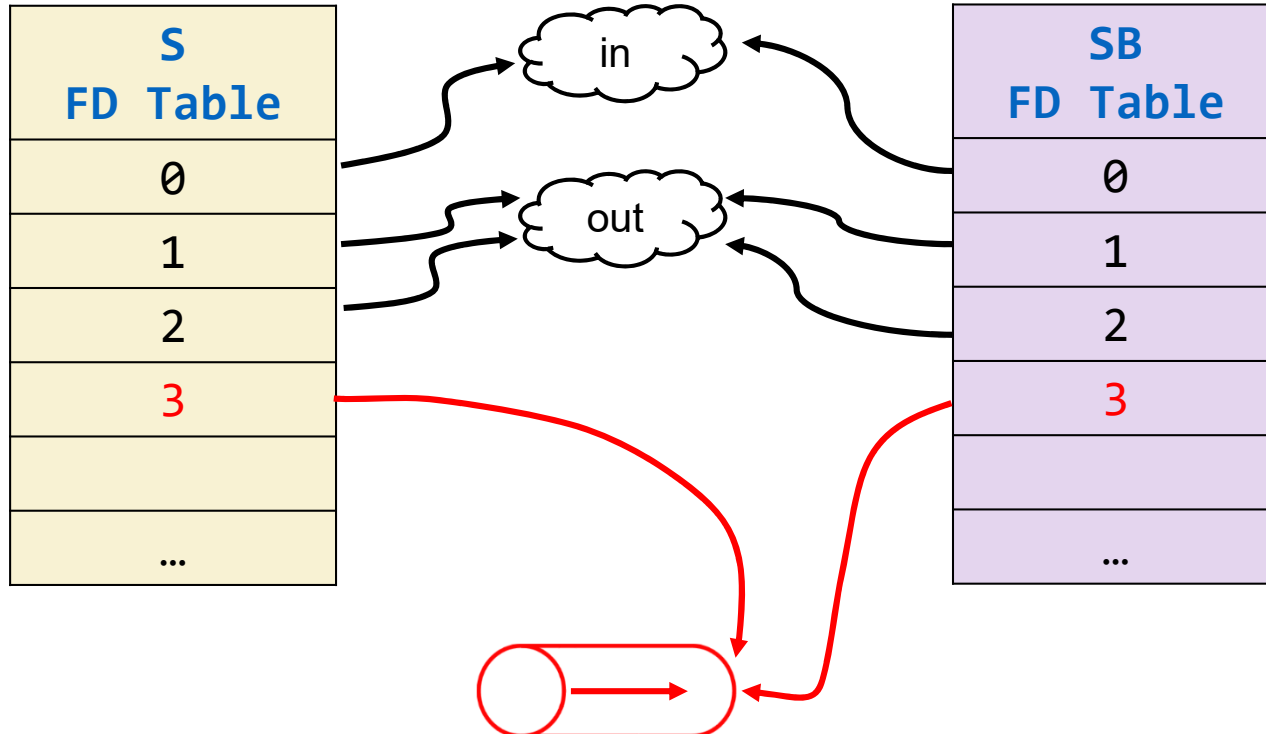
SA can then exec into A



# Fork #2

## S: fork()

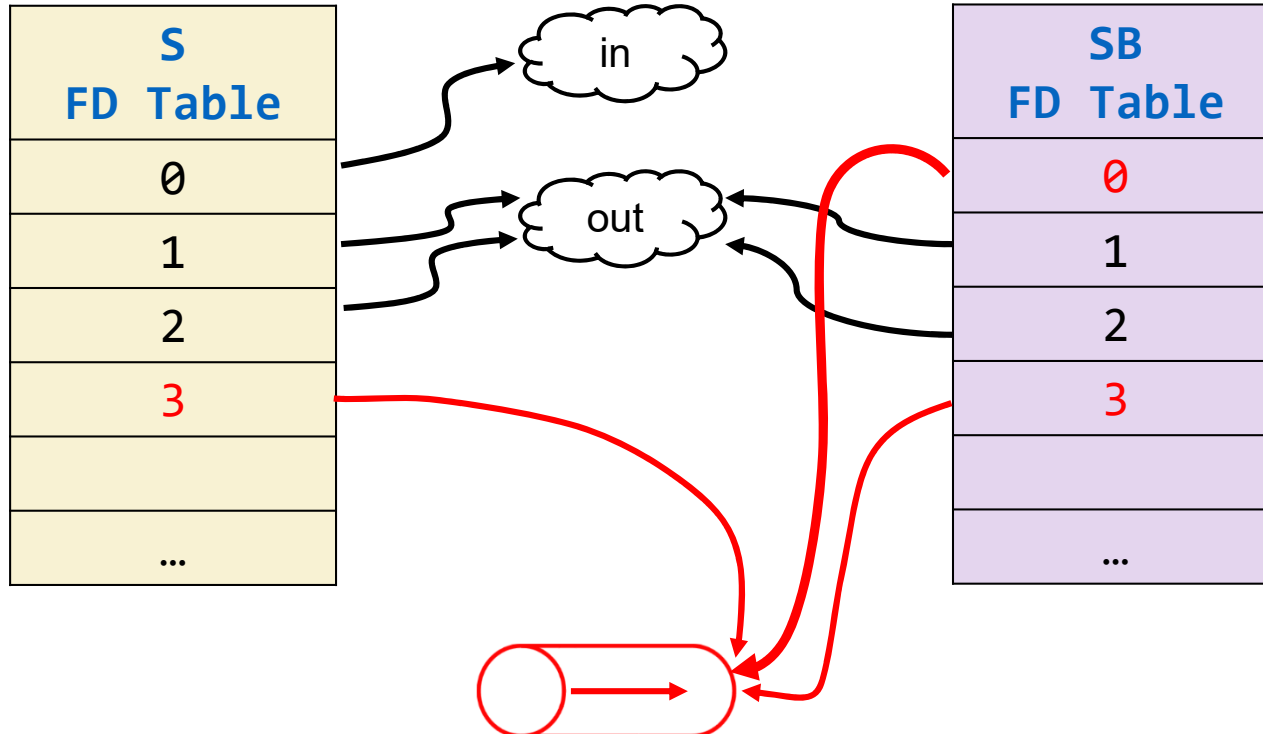
- Note that 4 has been closed in S



# Redirect in second child process

SB: `dup2(3, 0)`

- Or `close(0); dup(3);`



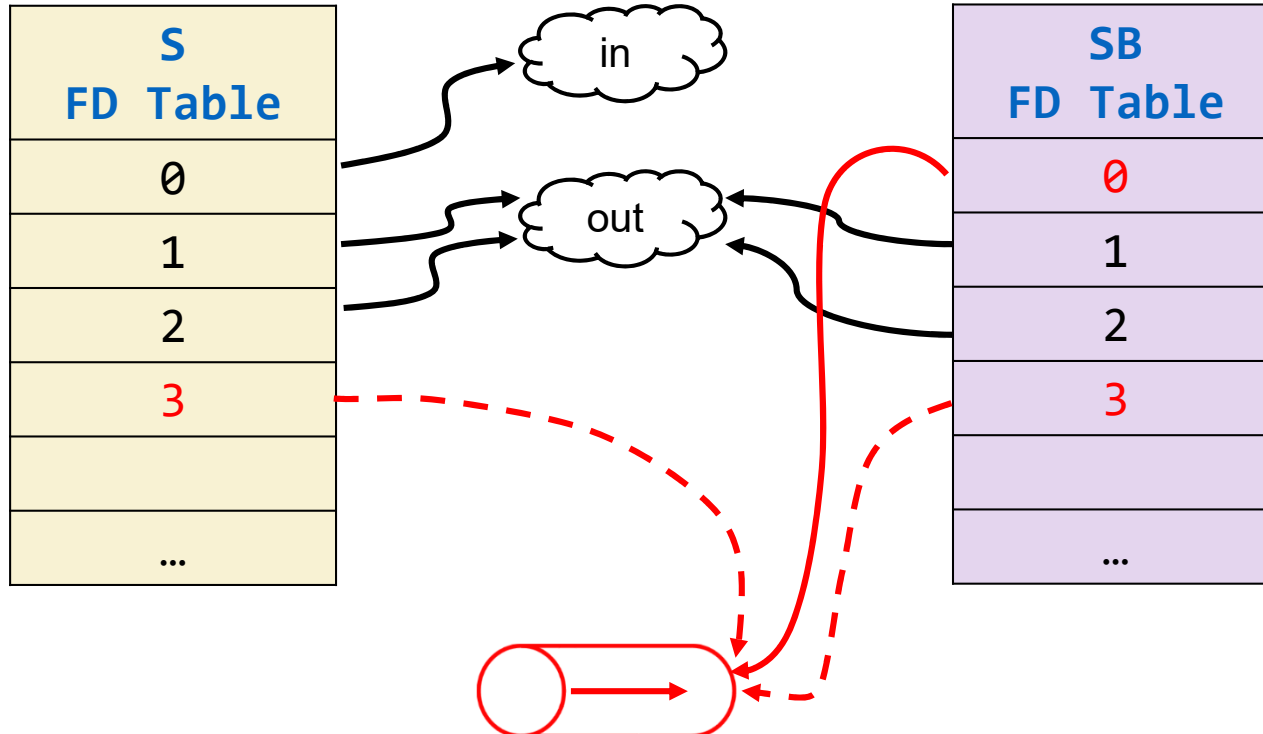


# Clean up #2

S: close(3)

SB: close(3)

SB can then exec into B

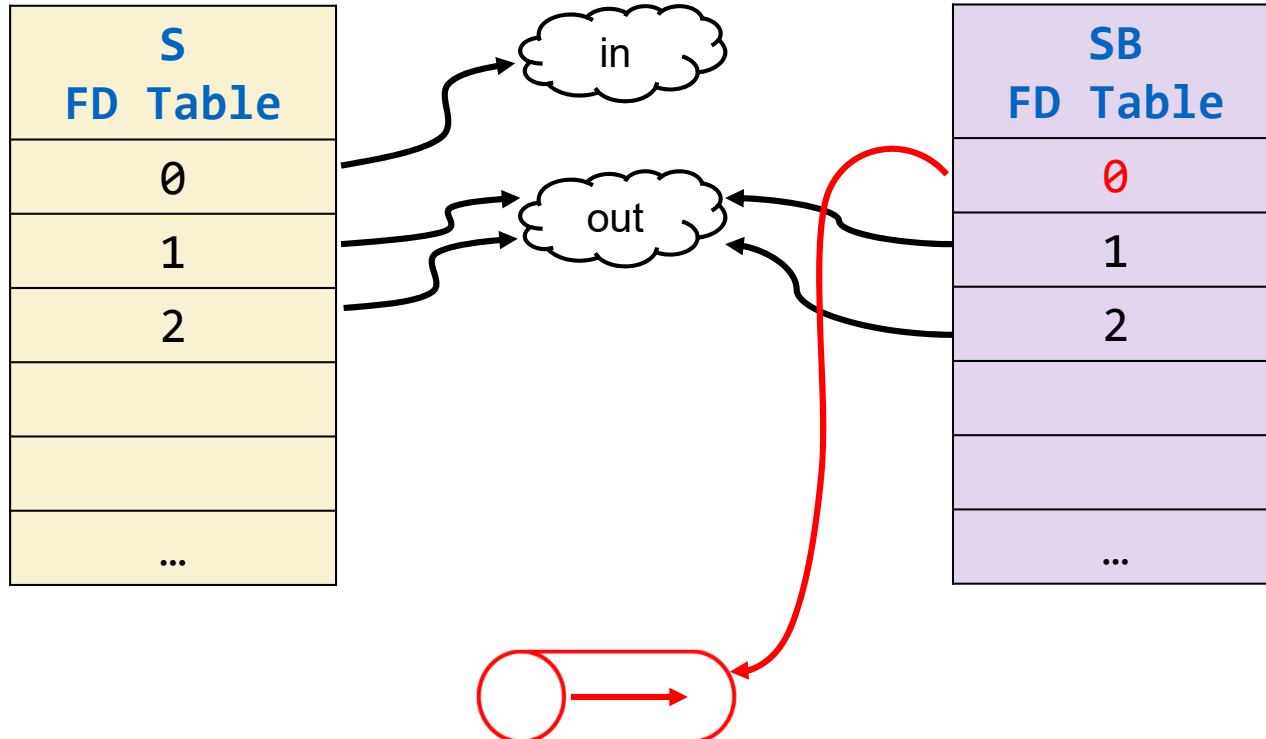


## After clean up #2

S: close(3)

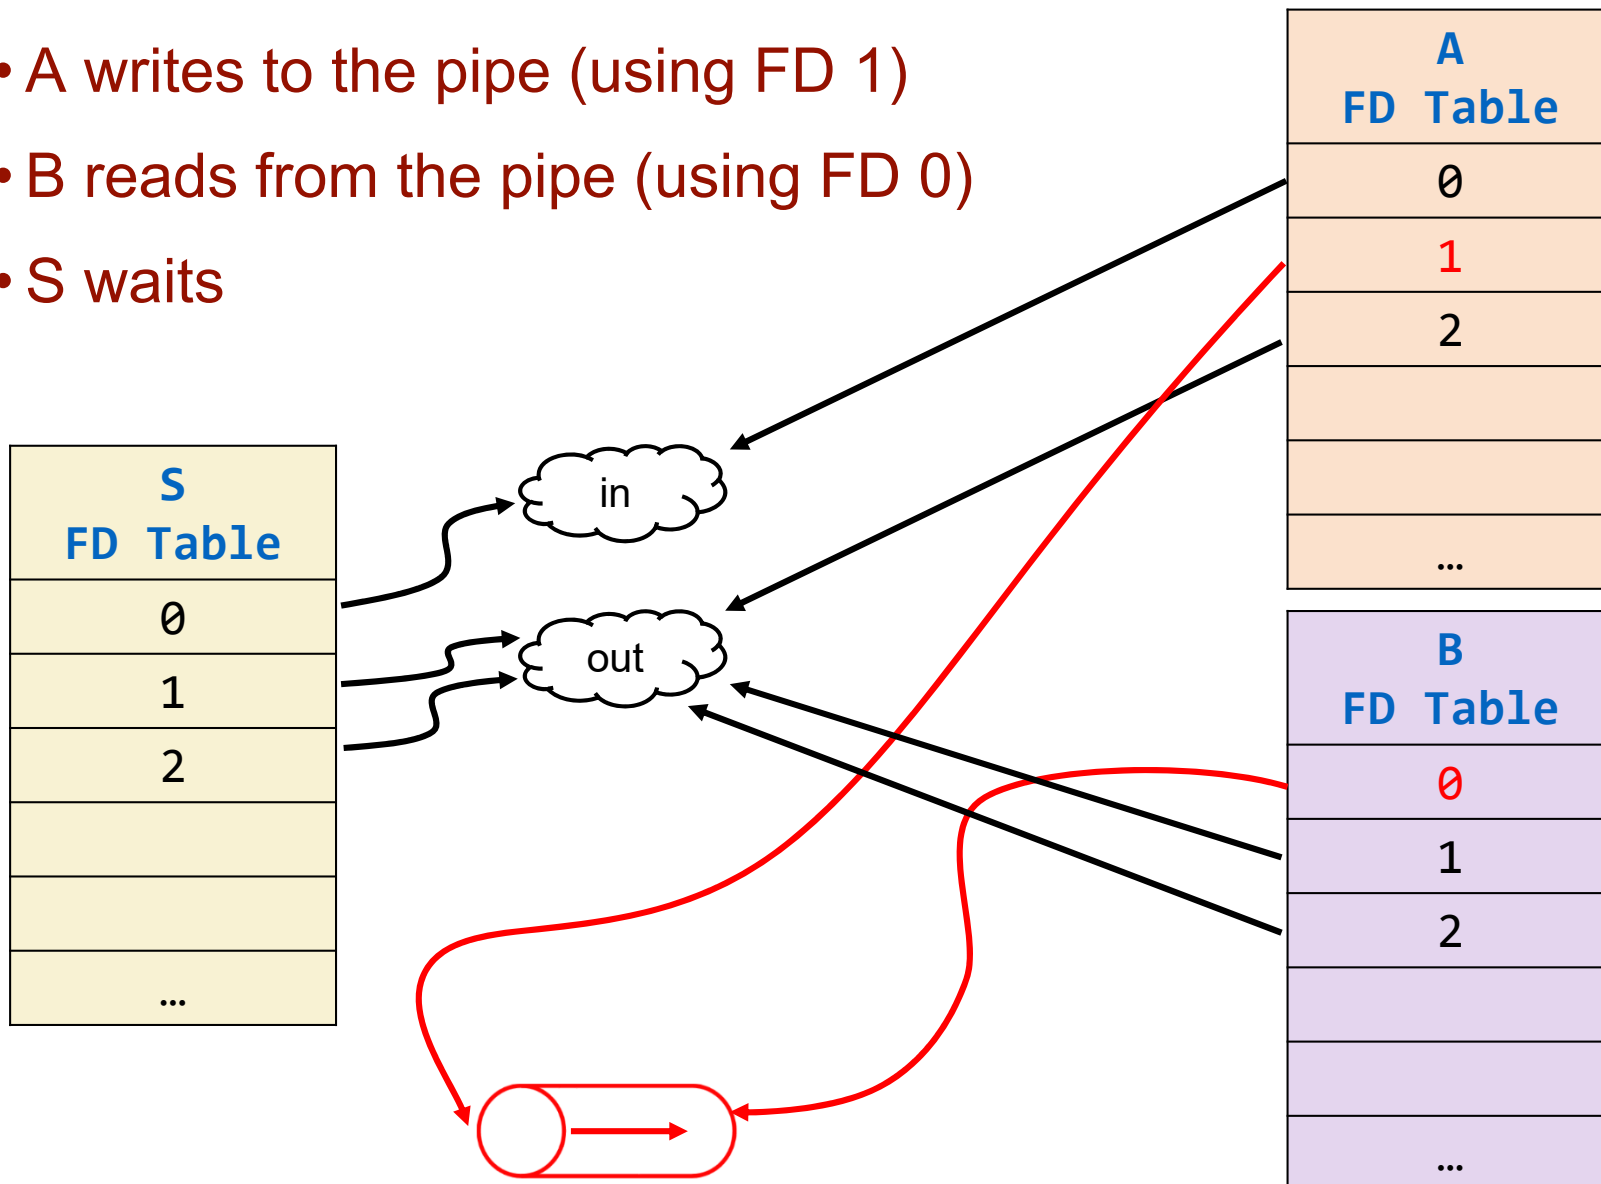
SB: close(3)

SB can then exec into B



# Final set up

- A writes to the pipe (using FD 1)
- B reads from the pipe (using FD 0)
- S waits




# FDs of a dying process

- When a process ends, all its open FDs are automatically closed
- What happens to the processes on the other end of the pipe?

Example:

Assume S does not read or write, but have FDs of the pipe

- If both A and S die, B gets EOF when all buffered data are consumed
- If A dies, B will wait for more data (assuming S may write)
- If both B and S die, A gets an error (SIGPIPE) when writing
- If B dies, A will wait if the pipe is full (assuming S will read)



Close file descriptors  
a process does not use!

# Going further...

---

- You can repeat this to create a long pipeline
  - E.g., connect B's stdout to stdin of another process C
- Draw pictures to find out how pipes are used
  - And what FDs need to be closed

## Remember

- Processes are running in parallel once they are created
  - Although we showed the operations in sequence
- All processes in the pipeline are running concurrently on Linux
  - As soon as data are sent in the pipe...
  - The next process can pick them up and start to work



# Atomicity of read() and write()

---

```
nr = read(fd, buf, N);
```

```
nw = write(fd, buf, N);
```

write() and read() returns the number of bytes actually read/written

The returned values may be less than the requested

- Atomicity of write () is guaranteed if the number of bytes is less than PIPE\_BUF
  - The bytes will be consecutive
  - The default value of PIPE\_BUF is 4096 on Linux
- For read(), it is fine if all writes and reads are of the same size
  - Otherwise, need special handling



---

Study the remaining slides yourself



# Starting a 2-stage pipeline - 1

---

// A | B

```
pipe(pipefd)    // pipefd is an array of 2 int's
pid_a = fork()  // for A
if (pid_a == 0) { // child process for A
    dup2();      // setup stdout for A
    close both FDs in pipefd
    exec to start A // remember to exit from child on error
}
close(pipefd[WR_END]); // No need to keep it open in parent
```





## Starting a 2-stage pipeline - 2

---

```
pid_b = fork()           // for B
if (pid_b == 0) {        // child process for B
    dup2();              // setup stdin for B
    close(pipefd[RD_END]);
    exec to start B // remember to exit from child on error
}
close(pipefd[RD_END]); // No need to keep it open in parent

// Add code to check return value for errors!
```



# Using Pipes to Sum Matrix Rows Concurrently

See the complete code [in the demo repo](#).

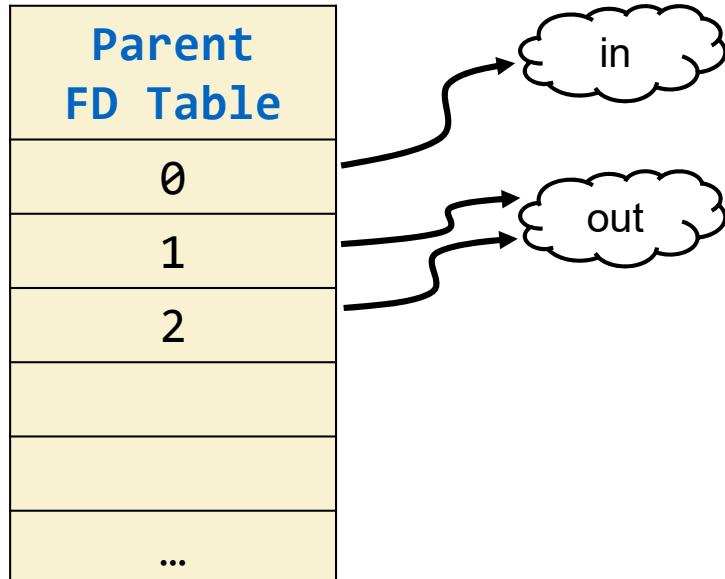

```
int main(void)
{
    int i, row_sum, sum = 0, pd[2], a[N][N] = {{1, 1, 1}, {2, 2, 2}, {3, 3, 3}};

    if (pipe(pd) == -1) error_exit("pipe() failed"); /* create pipe */

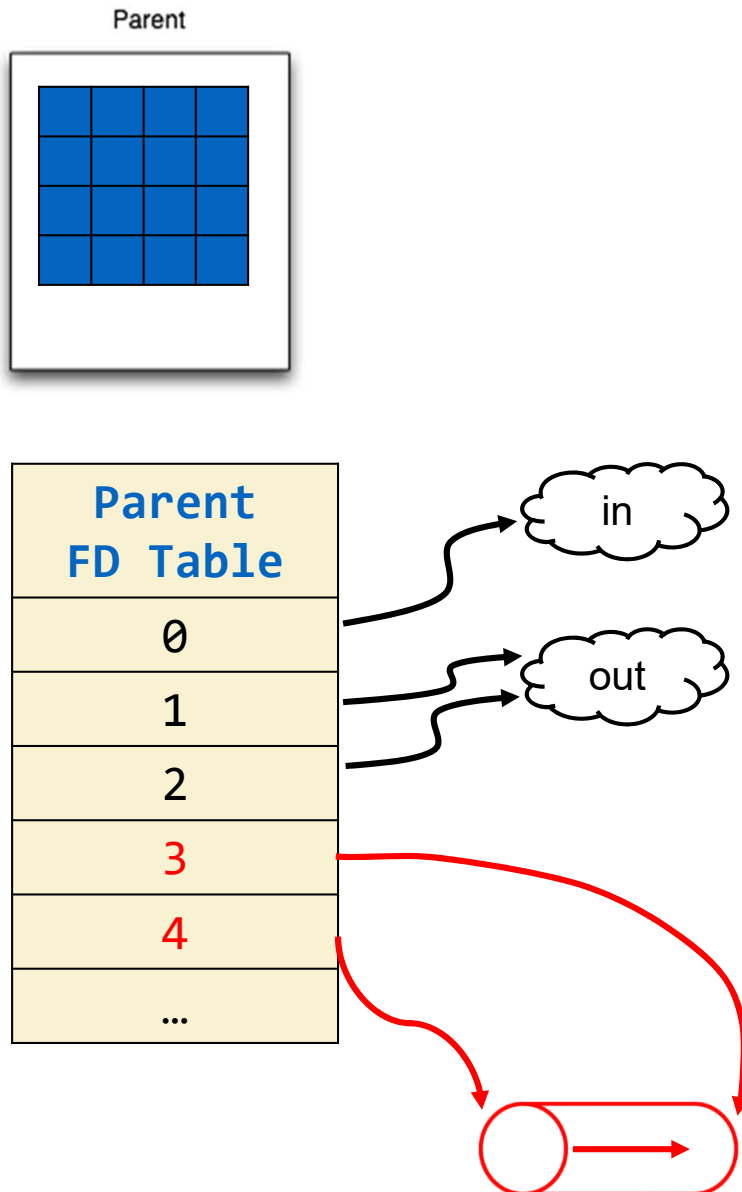
    for (i = 0; i < N; ++i)
        if (fork() == 0) { /* create a child process for each row */
            row_sum = add_vector(a[i]); /* compute the sum of a row */
            if (write(pd[1], &row_sum, sizeof(int)) == -1) /* write to pipe */
                error_exit("write() failed");
            return 0; /* exit from child */
        }
    /* better to close the write end in the parent */
    for (i = 0; i < N; ++i) {
        if (read(pd[0], &row_sum, sizeof(int)) == -1) /* read from pipe */
            error_exit("read() failed");
        sum += row_sum; /* calculate the total */
    }
    printf("Sum of the array = %d\n", sum);
    /* wait for child processes*/
}
```

# Parent process

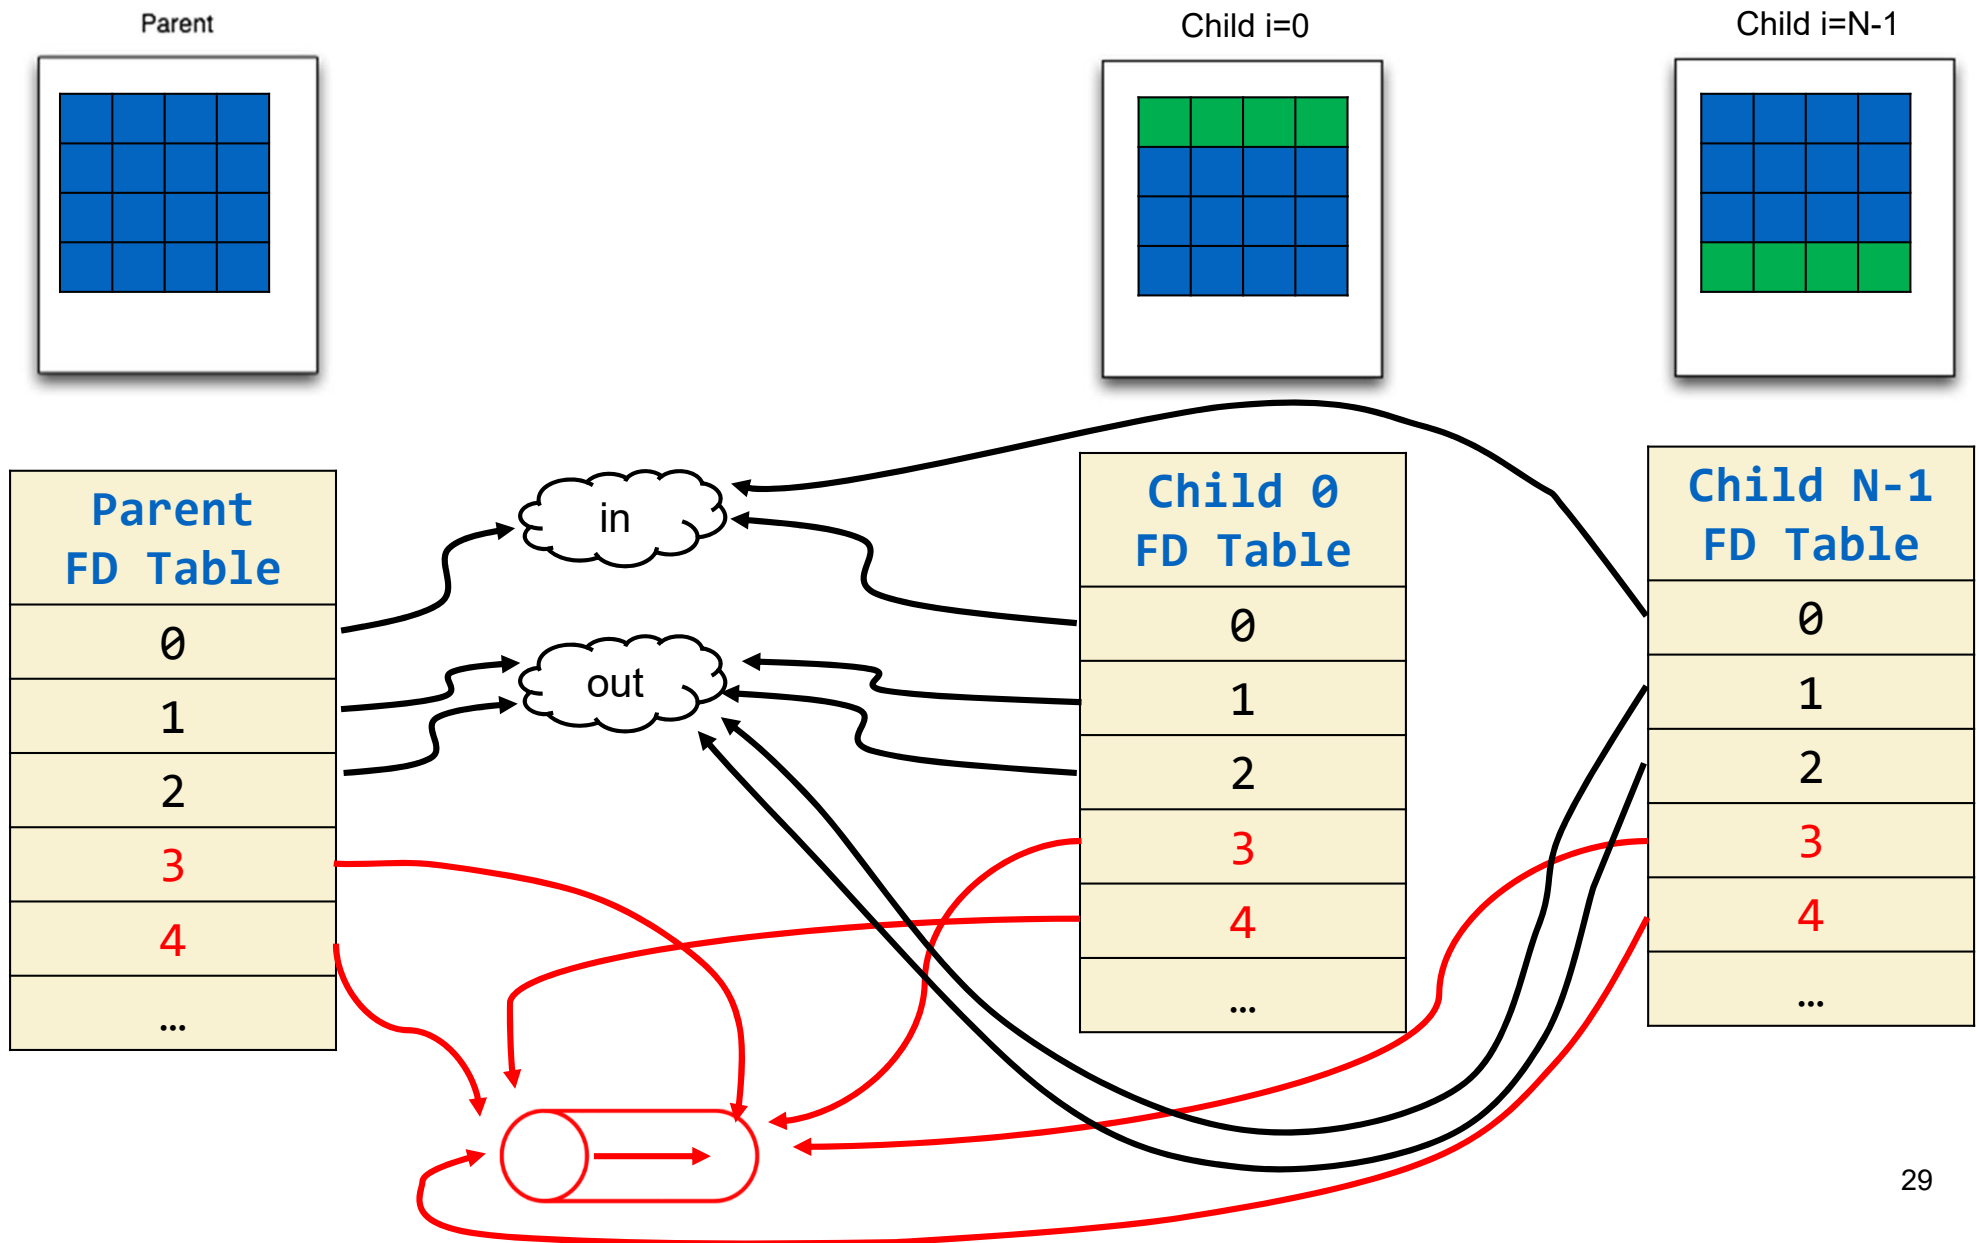
Parent



# Pipe creation



# After forking (without cleanup)



# After forking (with cleanup in parent)

