



P2: Exec and low-level I/O

Ion Mandoiu

Laurent Michel

Revised by M. Khan and J. Shi



Pitfall

- Anything might go wrong?

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1); // exit if fork() fails
} else if (pid == 0) {
    child_tasks();
} else {
    parent_tasks();
}
more_parent_tasks();
```

Pitfall

- Anything might go wrong?

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1); // exit if fork() fails
} else if (pid == 0) {
    child_tasks();
    exit(0); // terminate the child process
} else {
    parent_tasks();
}
more_parent_tasks();
```



Process upgrades

- Usually....
 - A fresh clone wants to run *different code*
- This is done by
 - Loading another executable into the process address space
 - [picked up from the file system of course]
- Note
 - **Opened files are NOT AFFECTED** by the upgrade operation

The exec family

- The act of ‘upgrading’ is done by the child with a system call
 - Many variants. "man -S3 execl" for all details

```
#include <unistd.h>
```

```
int execl(const char *path, const char *arg0, ...  
          /*, (char *) NULL */ );
```

- **The path** to the executable to load inside our own address space
- A list of arguments to be passed to the new executable
- **A final NULL pointer** to give the “end of argument list”
- If successful, `execl()` **does not return!** Started a new process

Exec example

- We will turn the child process into the following executable

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char* argv[]) {
    int i, sum=0;
    for(i=1; i<argc; i++)
        sum += atoi(argv[i]);
    printf("sum is: %d\n", sum);
    return 0;
}
```

This is a simple
“adder” program
that computes the
sum of its integer
arguments



Parent Program

```
int main() { // complete code is in demo repo
    char *cmd1 = "./adder", *cmd2 = "expr";
    pid_t child = fork();
    if (child == 0) {
        printf("In child!\n");
        execl(cmd1, cmd1, "1", "2", "3", "10", NULL);
        printf("Oops.... something went really wrong!\n");
        perror(cmd1);
        return -1;
    } else {
        printf("In parent!\n");
        execl(cmd2, cmd2, "100", "+", "300", NULL);
        printf("Oops.... something went really wrong!\n");
        perror(cmd2);
        return -1;
    }
}
```

How is executable found?

- Specify a path, like `/bin/ls`
- Specify a file, and the system searches in directories listed in `PATH`
 - `echo $PATH` in bash to see directories separated by ':'

```
int execl(const char *path, const char *arg0, ...  
          /*, (char *) NULL */ );
```

// `execvp()` searches paths for file

```
int execvp(const char *file, const char *arg0, ...  
          /*, (char *) NULL */ );
```




execv family

// If the number of arguments is unknown at compile time

```
#include <unistd.h>
```

```
int execv(const char *path, char *const argv[]);
```

```
int execvp(const char *file, char *const argv[]);
```


- The arguments in `exec` are placed in an array
 - `argv` is the `argv` you see in the `main` function!
- `execv` needs a path while `execvp` can search file in `PATH`
- Start a new process if successful. Similar to `exec`

argv to execv and execvp

- Note the NULL pointer at the end
- Why?

The last pointer is NULL !

```
$gcc -o ab ab.c
```





Question

- What might go wrong?

```
pid_t pid = fork();
if (pid < 0) {
    perror("fork()"); exit(1); // exit if fork() fails
}
else if (pid == 0) {
    // in child process
    execlp("genie", "genie", "clean the house", NULL);
}
// in parent process
online_shopping();
```

File APIs

- Remember the (C standard library) IO APIs
 - The “f” family (fopen, fclose, fread, fgetc, fscanf, fprintf,...)
 - All these use a FILE* abstraction to represent a file
 - Additional features: user-space buffering, line-ending translation, formatted I/O, etc.
- UNIX has lower-level APIs for file handling
 - Directly mapped to **system calls**
 - open, close, read, ...
 - Use *file descriptors* [which are just **integers**]
 - Deal with bytes only

Some low level file APIs

- Read the man pages (man -s2 ...) for more functions

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int open(const char *path, int oflag);
int open(const char *path, int oflag, int mode);
int close(int fd);

ssize_t read(int fd, void *buf, size_t nbyte);
ssize_t write(int fd, const void *buf, size_t nbyte);

off_t lseek(int fd, off_t offset, int whence);
```



Open a file

```
#include <fcntl.h>
#include <unistd.h>
```

```
int open(const char *path, int oflag);
```

- Parameters

- path: the path to the file to be opened/created
- oflag: read, write, or read and write, and more (on the next slide)

- The function returns a file descriptor, a small, nonnegative integer

- Return -1 on error

Flags in open()

- Must include one of the following:
O_RDONLY (read only), O_WRONLY (write only), or O_RDWR (read and write)
- And or-ed (|) with many optional flags, for example,
 - O_TRUNC: Truncate the file (remove existing contents) if opening a file for write
 - O_CREAT: Create a file if it does not exist.

Example:

```
// remember open() returns -1 on error
```

```
fd1 = open("a.txt", O_RDONLY); // open for read
```

```
fd1 = open("a.txt", O_RDWR); // open for read and write
```

```
fd1 = open("a.txt", O_RDWR|O_TRUNC); // read, write, truncate the file
```



Create a file with open()

// a mode must be provided if O_CREAT or O_TMPFILE is set

```
int open(const char *path, int oflag, int mode);
```

mode: specify permissions when a new, or temporary, file is created.

```
open("b.txt", O_WRONLY|O_TRUNC|O_CREAT, 0600);
```

// open b.txt for write. If the file exists, clear (truncate) the contents.

// if the file does not exist, create one, and set the permission so that the owner of the file can read and write, but other people cannot.

leading 0 is important
Octal numbers!



File descriptor vs stream

```
#include <stdio.h>
```

```
int fileno(FILE *stream);
```

```
// returns a file descriptor for a stream
```

FD	FILE *
0	stdin
1	stdout
2	stderr



File descriptors after fork and exec

- Opened files are **NOT AFFECTED** by the upgrade operation

```
pid_t pid = fork();
assert(pid >= 0);
if (pid == 0) {
    // Child process can access FDs 0, 1, and 2
    // if execl() is successful, gcc can access FDs 0, 1, and 2
    execlp("gcc", "gcc", "a.c", NULL);
    // If control gets here, execlp() failed.
    // Remember to terminate the child process!
    return 1;
}
```