



P3: Redirection

Ion Mandoiu

Laurent Michel

Revised by M. Khan and J. Shi



Review

- Function `open()` returns a file descriptor, a non-negative integer
- The file descriptor is used later on in functions like `read()` and `close()`
- Every opened file has a file descriptor
 - `stdin`: 0, `stdout`: 1, `stderr`: 2
- Files opened in a process remain open after `fork()` and `exec`

Shell redirections

Available when executing commands in your shell (e.g. `bash`)

- **Implemented with the close/open/dup technique**

`$ command < infile > outfile`

`< infile` : Take input from file *infile*

`> outfile` : Send output to file *outfile*

- Other variants

`>> outfile` : Append output to file *outfile*

`2> outfile` : Send errors to file *outfile*

`&> outfile` : Send both output and errors to file *outfile*

Read the manual for more variants like `2>>`, `2>&1`, etc.



Shell redirection examples

```
$ sort < file.txt > sorted.txt
```

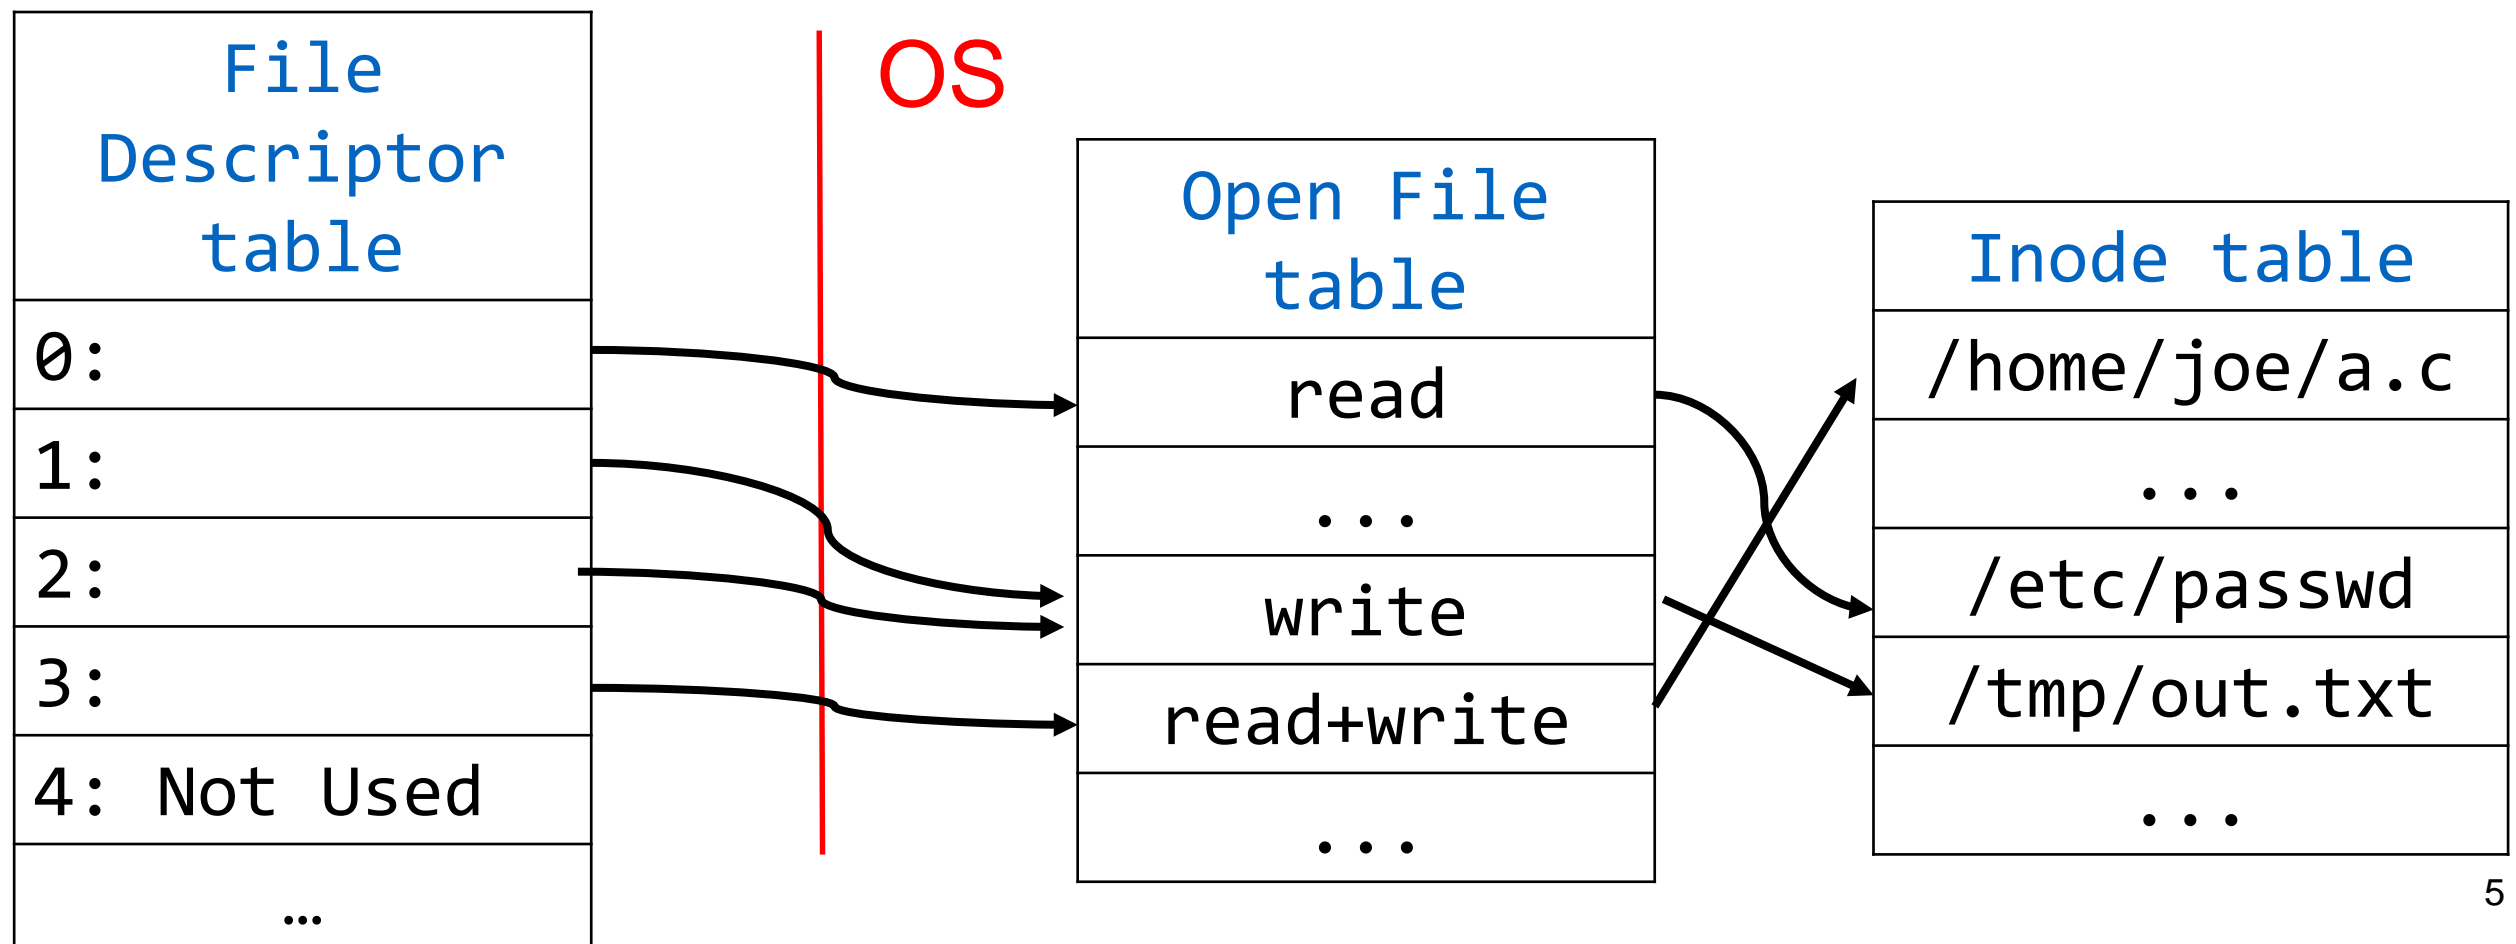
- sort will read lines from file.txt, instead of terminal
- The output of sort will be saved in sorted.txt
 - You cannot see it on screen

The statements in sort are not changed.

They read from stdin (0), and print to stdout (1)

File descriptor table

- Each process has a **file descriptor table**
 - Holds indices of entries into the **Open File Table** managed by OS
- The system-wide **Open File Table**
 - Records the *mode* of the opened files (e.g., reading, writing, appending)
 - Holds index into **the Inode Table** that has the actual file name and location on disk



Duplicating File Descriptors

- Do not change file descriptor table directly
- Used `open()` and `close()` and two new functions

```
#include <unistd.h>

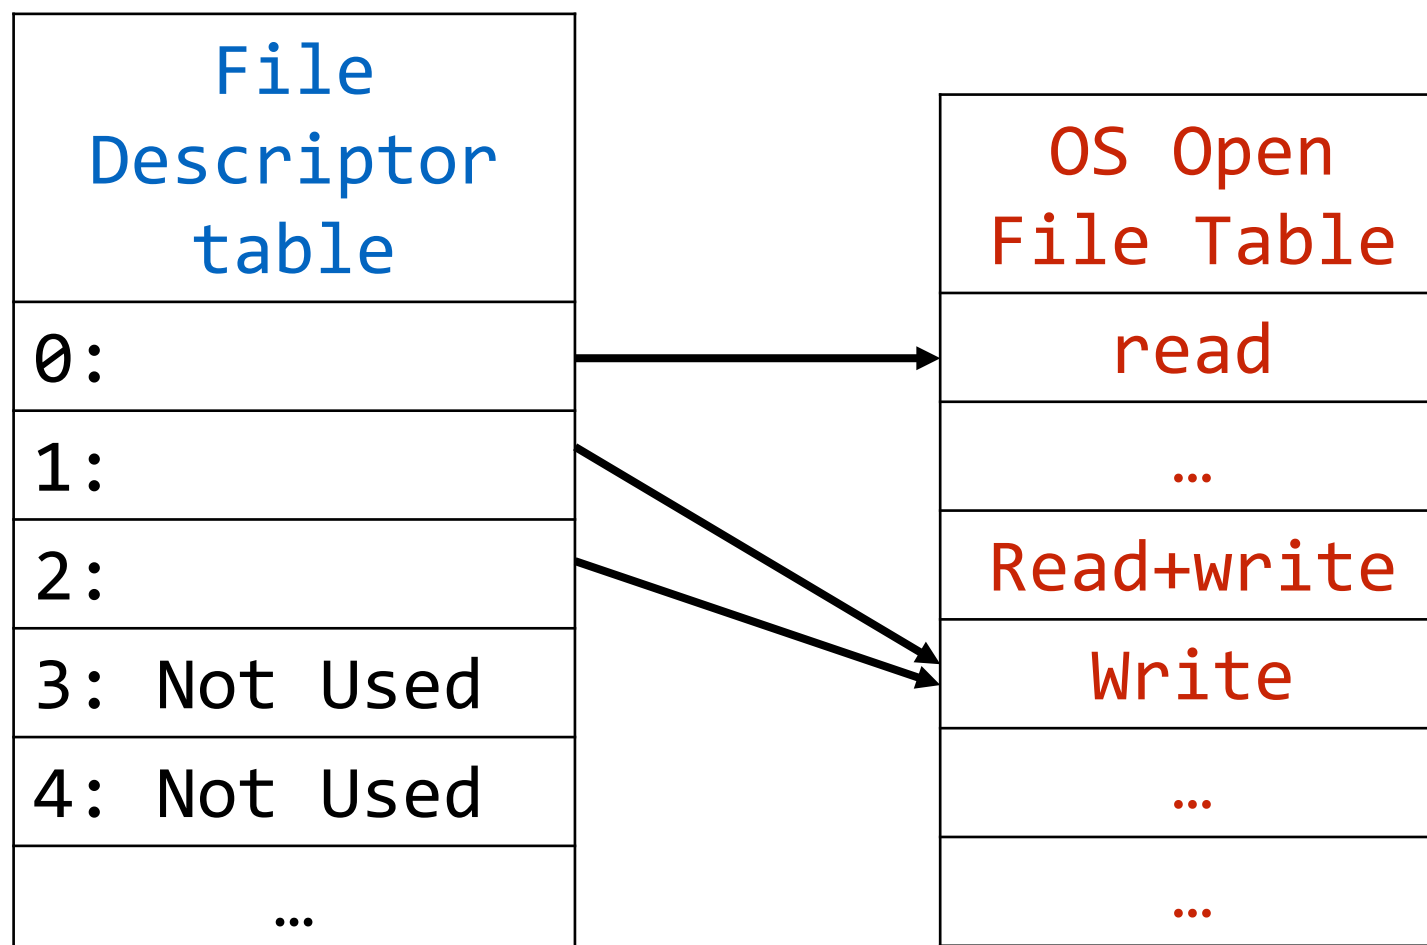
int dup(int oldfd);
int dup2(int oldfd, int newfd);
```

- `dup()` copies `oldfd` to the **first available entry** (in FD table)
- `dup2()` copies `oldfd` to `newfd`
 - Closes `newfd` first if it is in use

There is `dup3()`, but it is not in POSIX. We should not use it in this course.

Example: stdout redirect

- A program can change its standard output
- How?



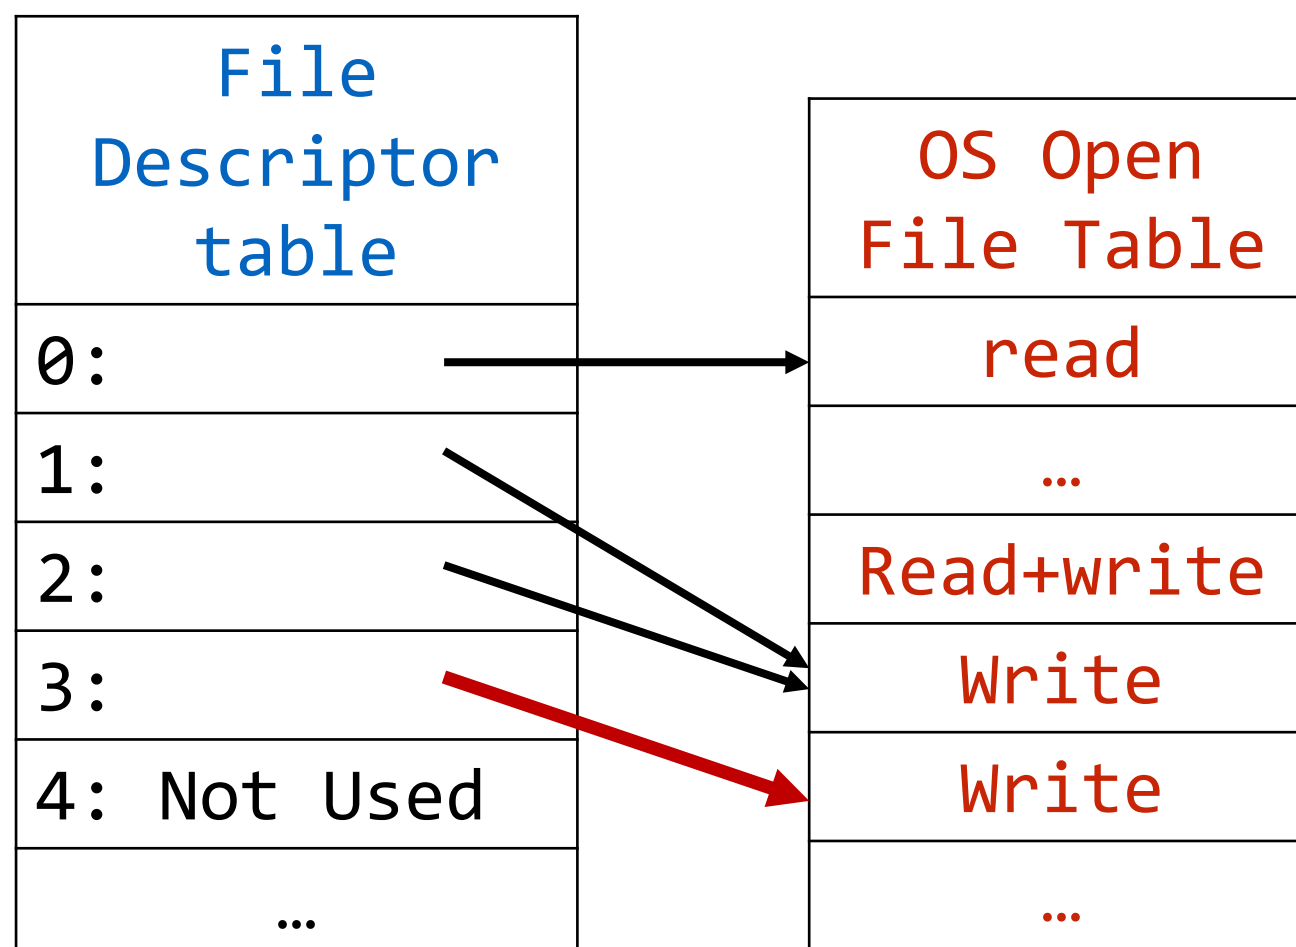


Steps for redirecting stdout

1. `open()`. open a file (and save the file descriptor in `fd`)
2. `dup2()`. copy `fd` to 1 (so that the file descriptor 1 points to the file just opened)
3. `close(fd)`

Example: stdout redirect

- Open the new file for writing; 3 is the returned fd
 - We will use 3 instead of a variable in this example



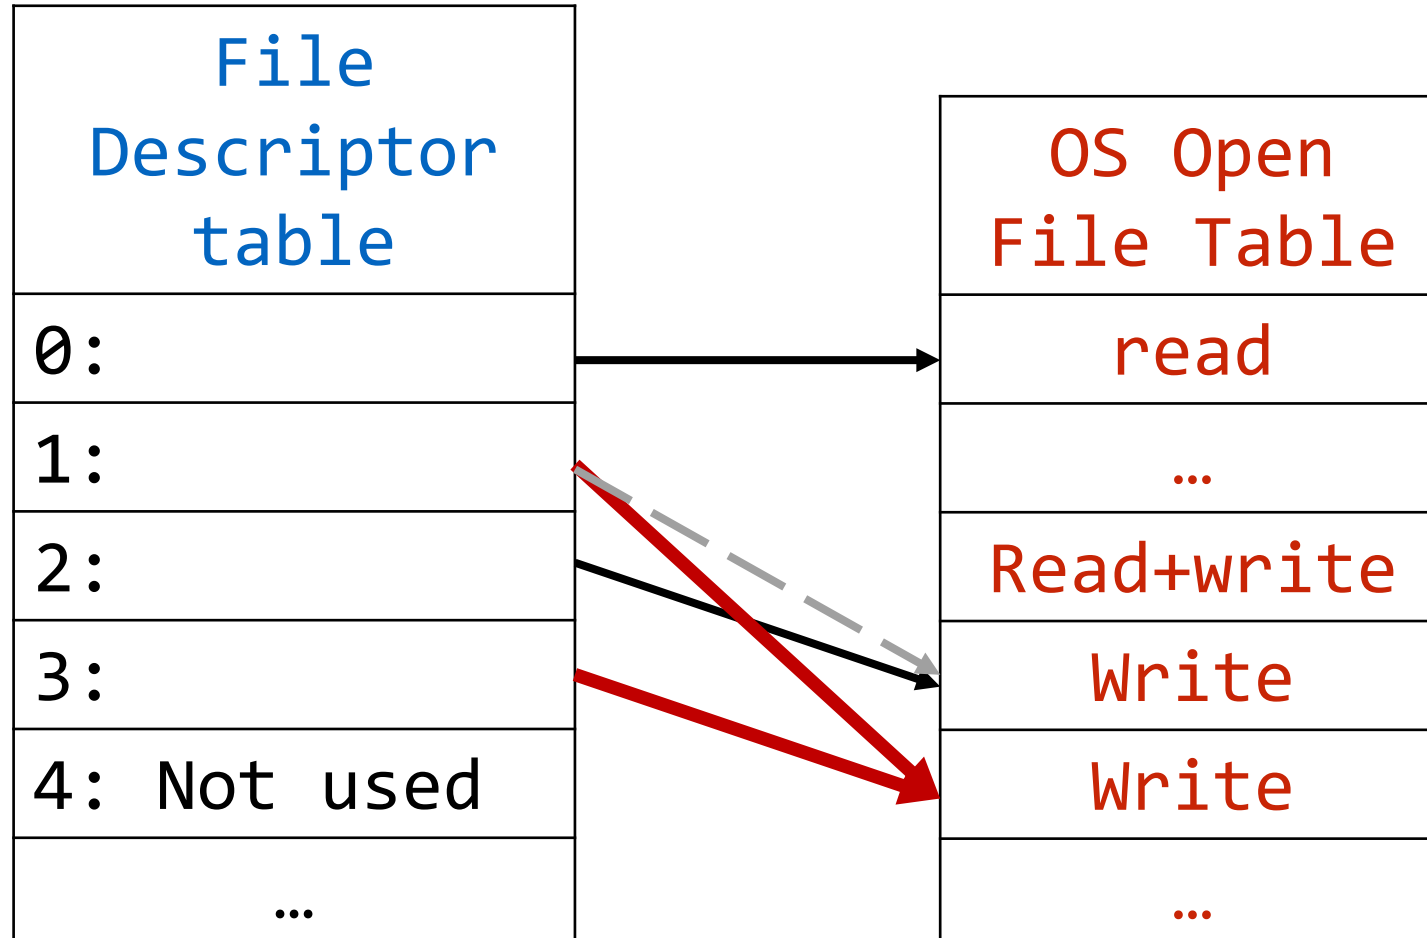
Example: stdout redirect

// Method 1: two functions. not atomic

```
close(1);      dup(3);
```

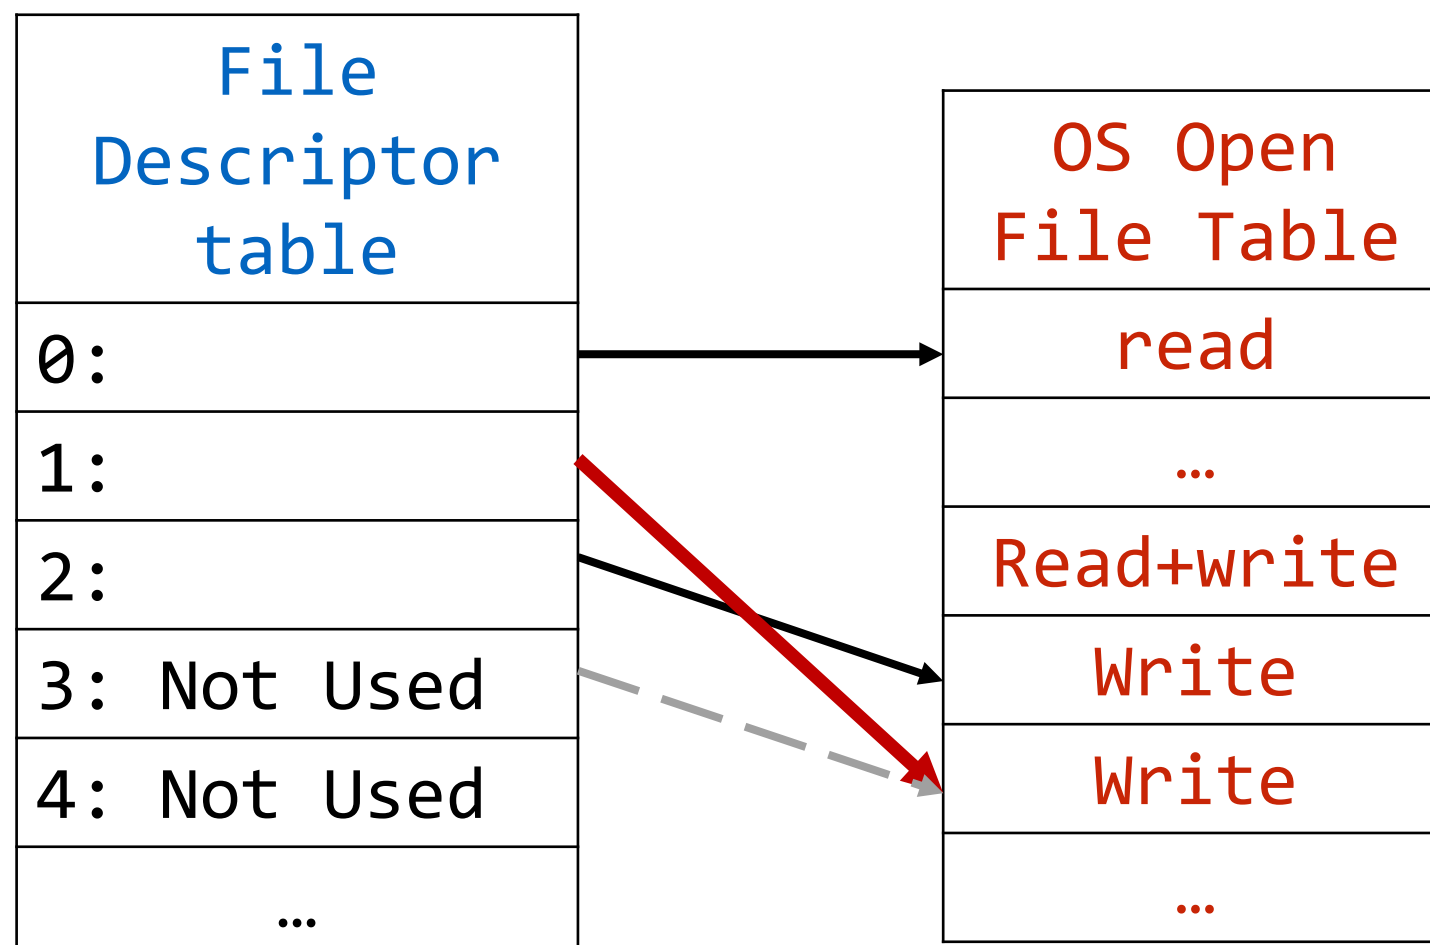
// Method 2: better. dup2() closes newfd first

```
dup2(3, 1);
```



Example: stdout redirect

`close(3);`





Implementing redirections

How does bash do redirection for other processes?

When bash starts a new process, it does `fork()` and `exec`

Recall that the file descriptor table **is preserved during fork & exec**

Idea:

In child process, set up proper file descriptors before upgrading

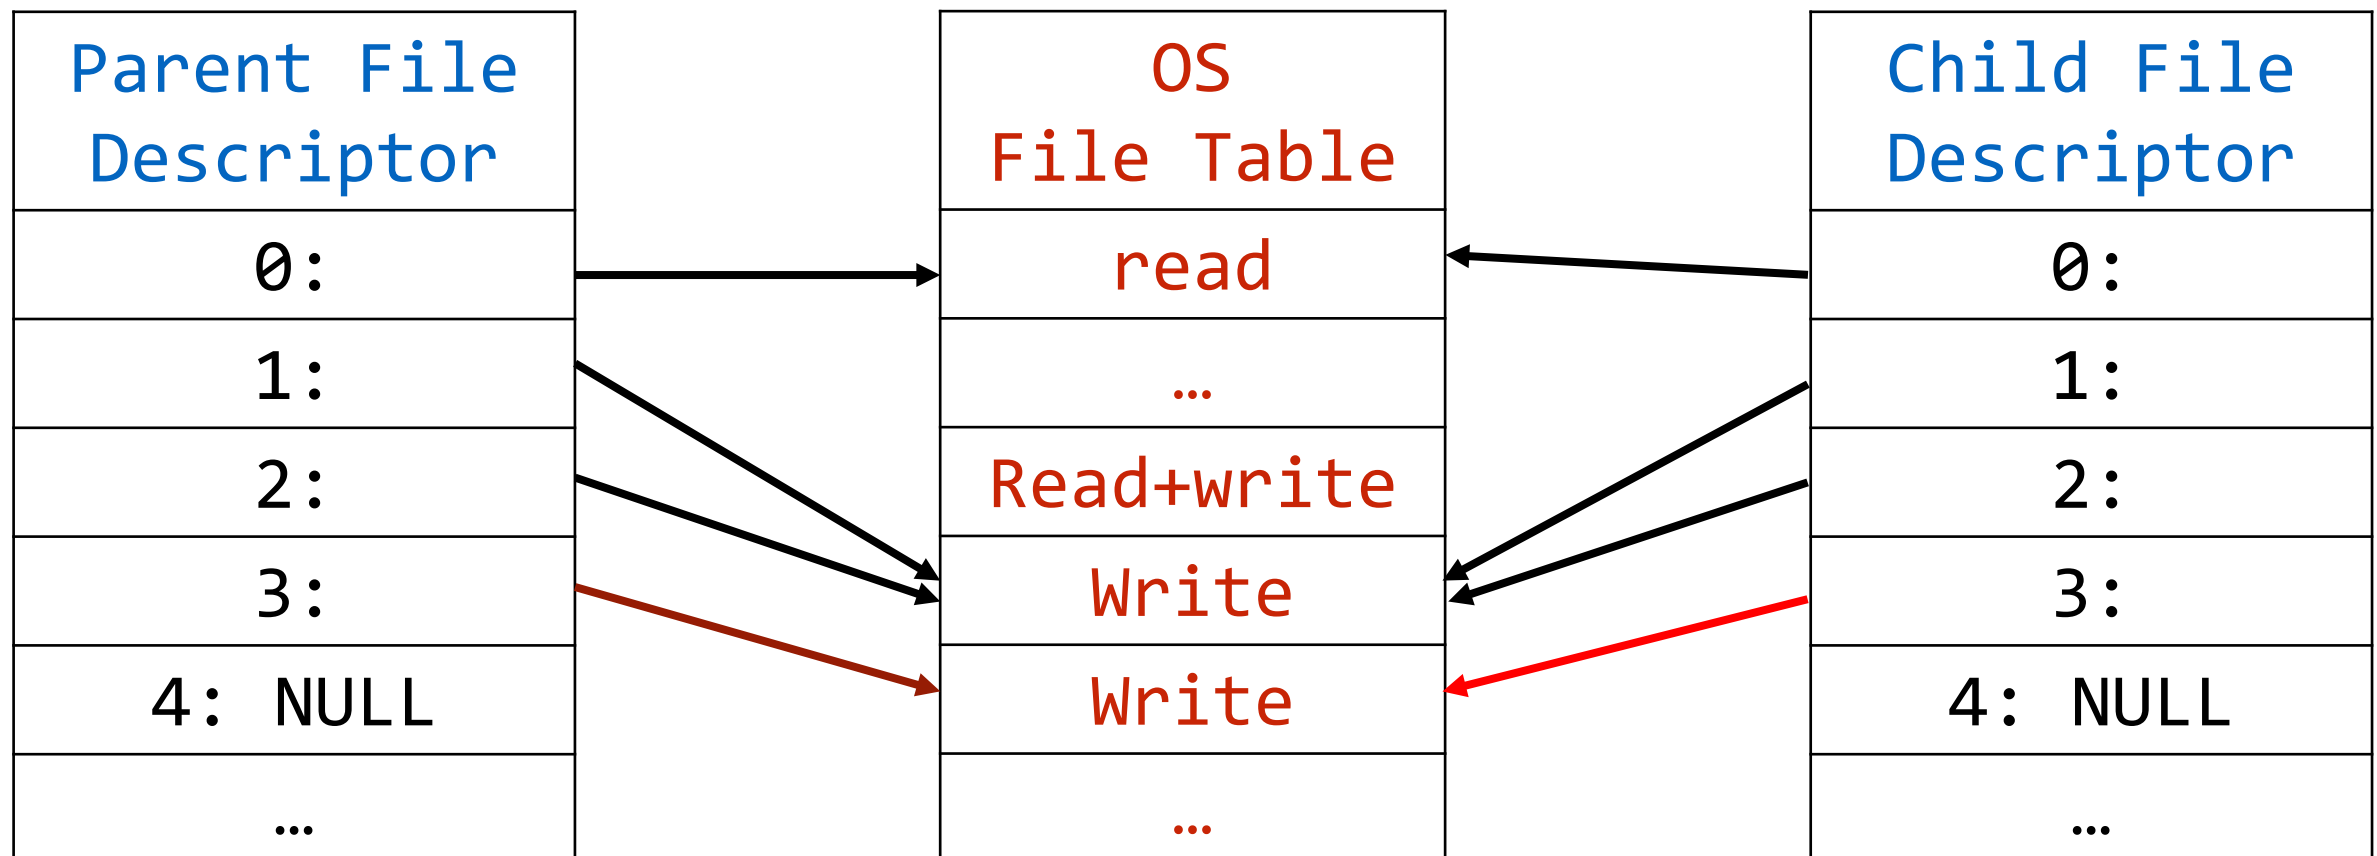
- Simply *change* the files corresponding to **stdin, stdout, or stderr**

Example: redirecting stderr for another program

Assume the parent uses FD 3.

After fork(), the child has the same file descriptors as the parent

Close FDs that are not needed !



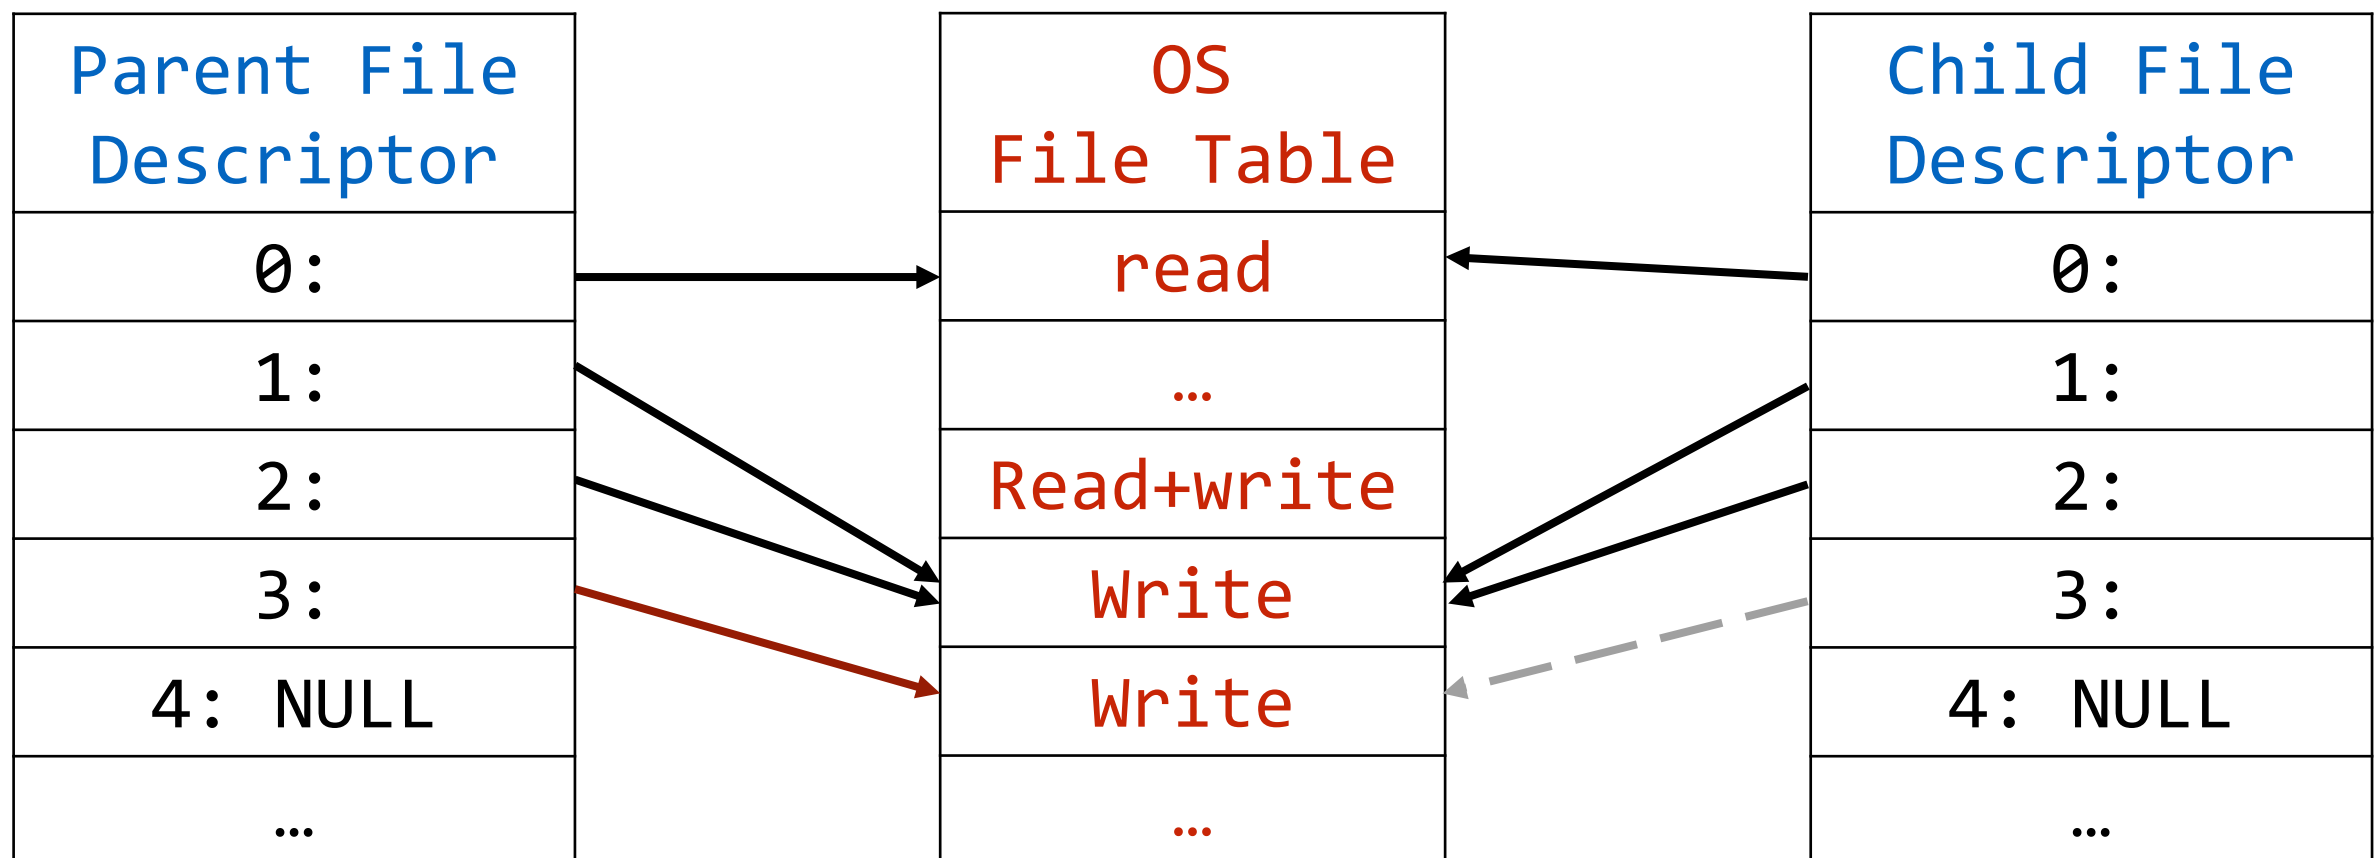
Example: redirecting stderr for another program

Assume the parent uses FD 3.

After fork(), the child has the same file descriptors as the parent

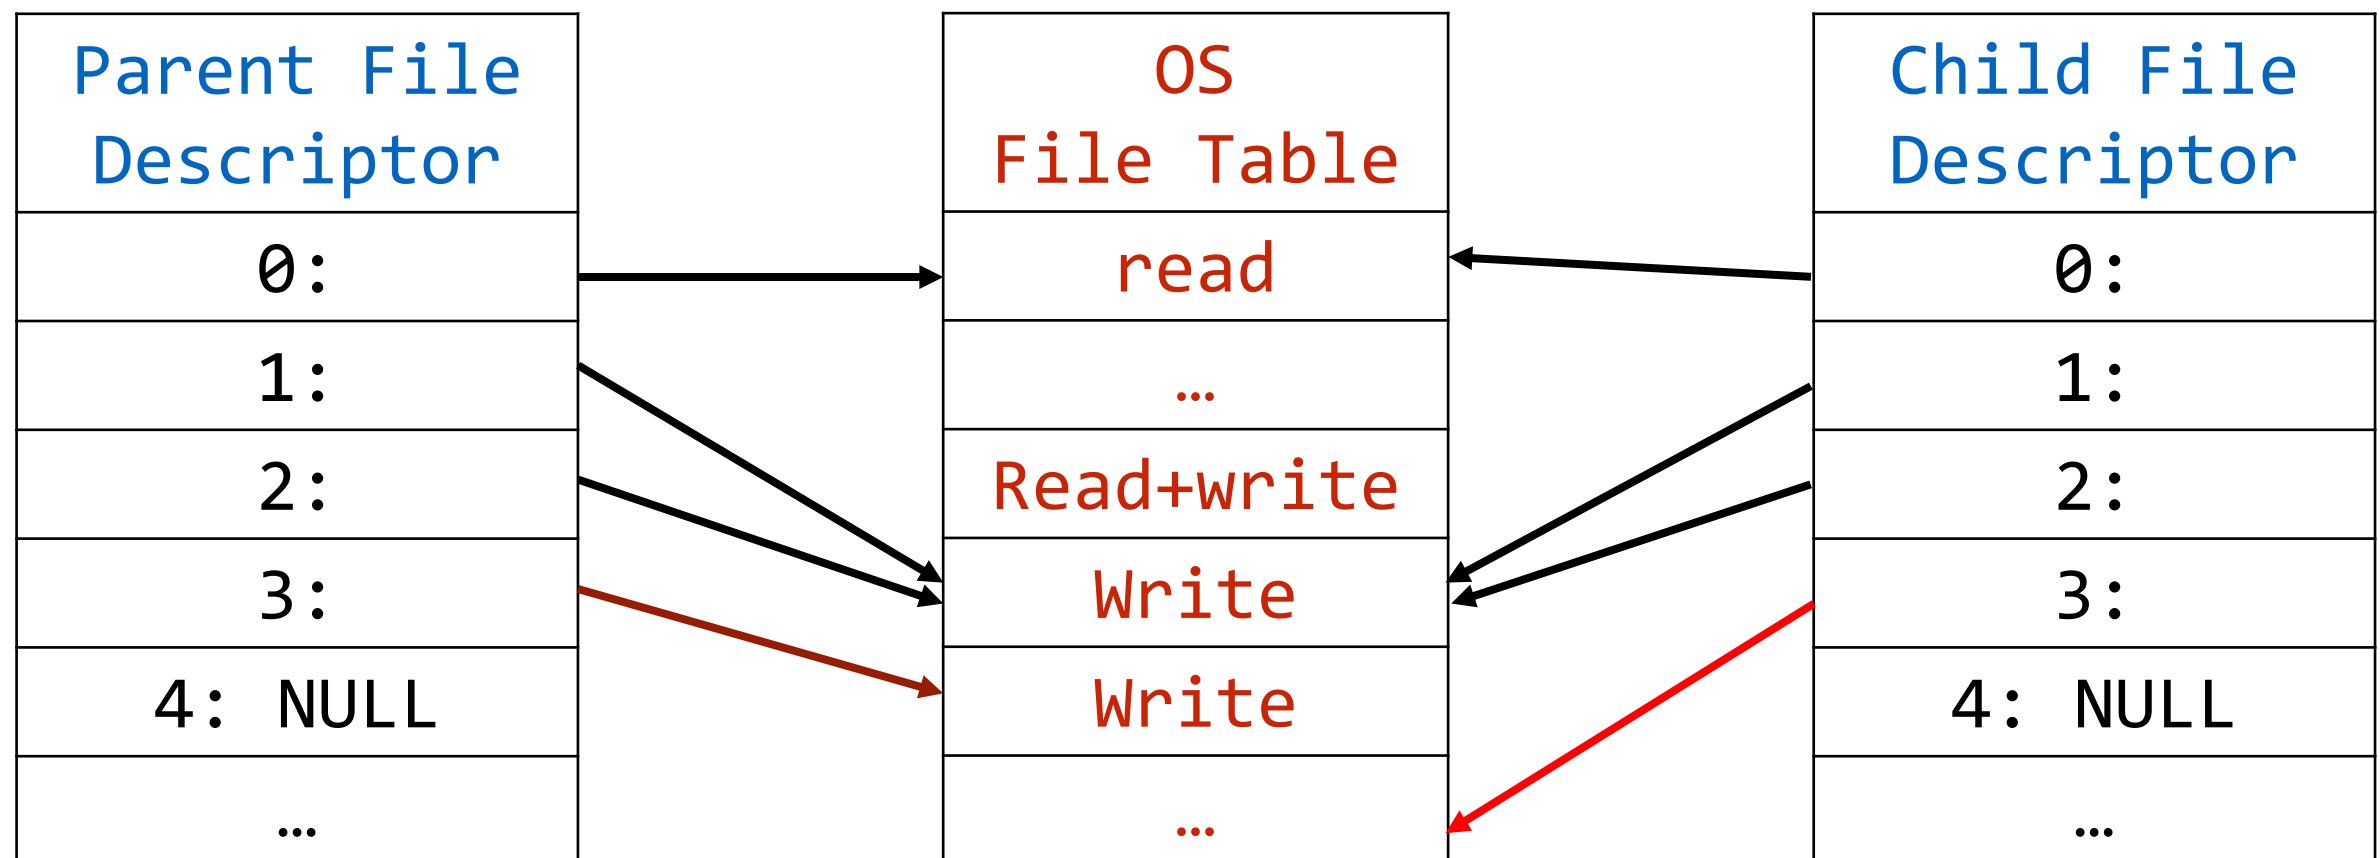
Close FDs that are not needed !

How do you close FD 3 in child process?



Example: redirecting stderr for another program

- Child: open the file (to save error output)

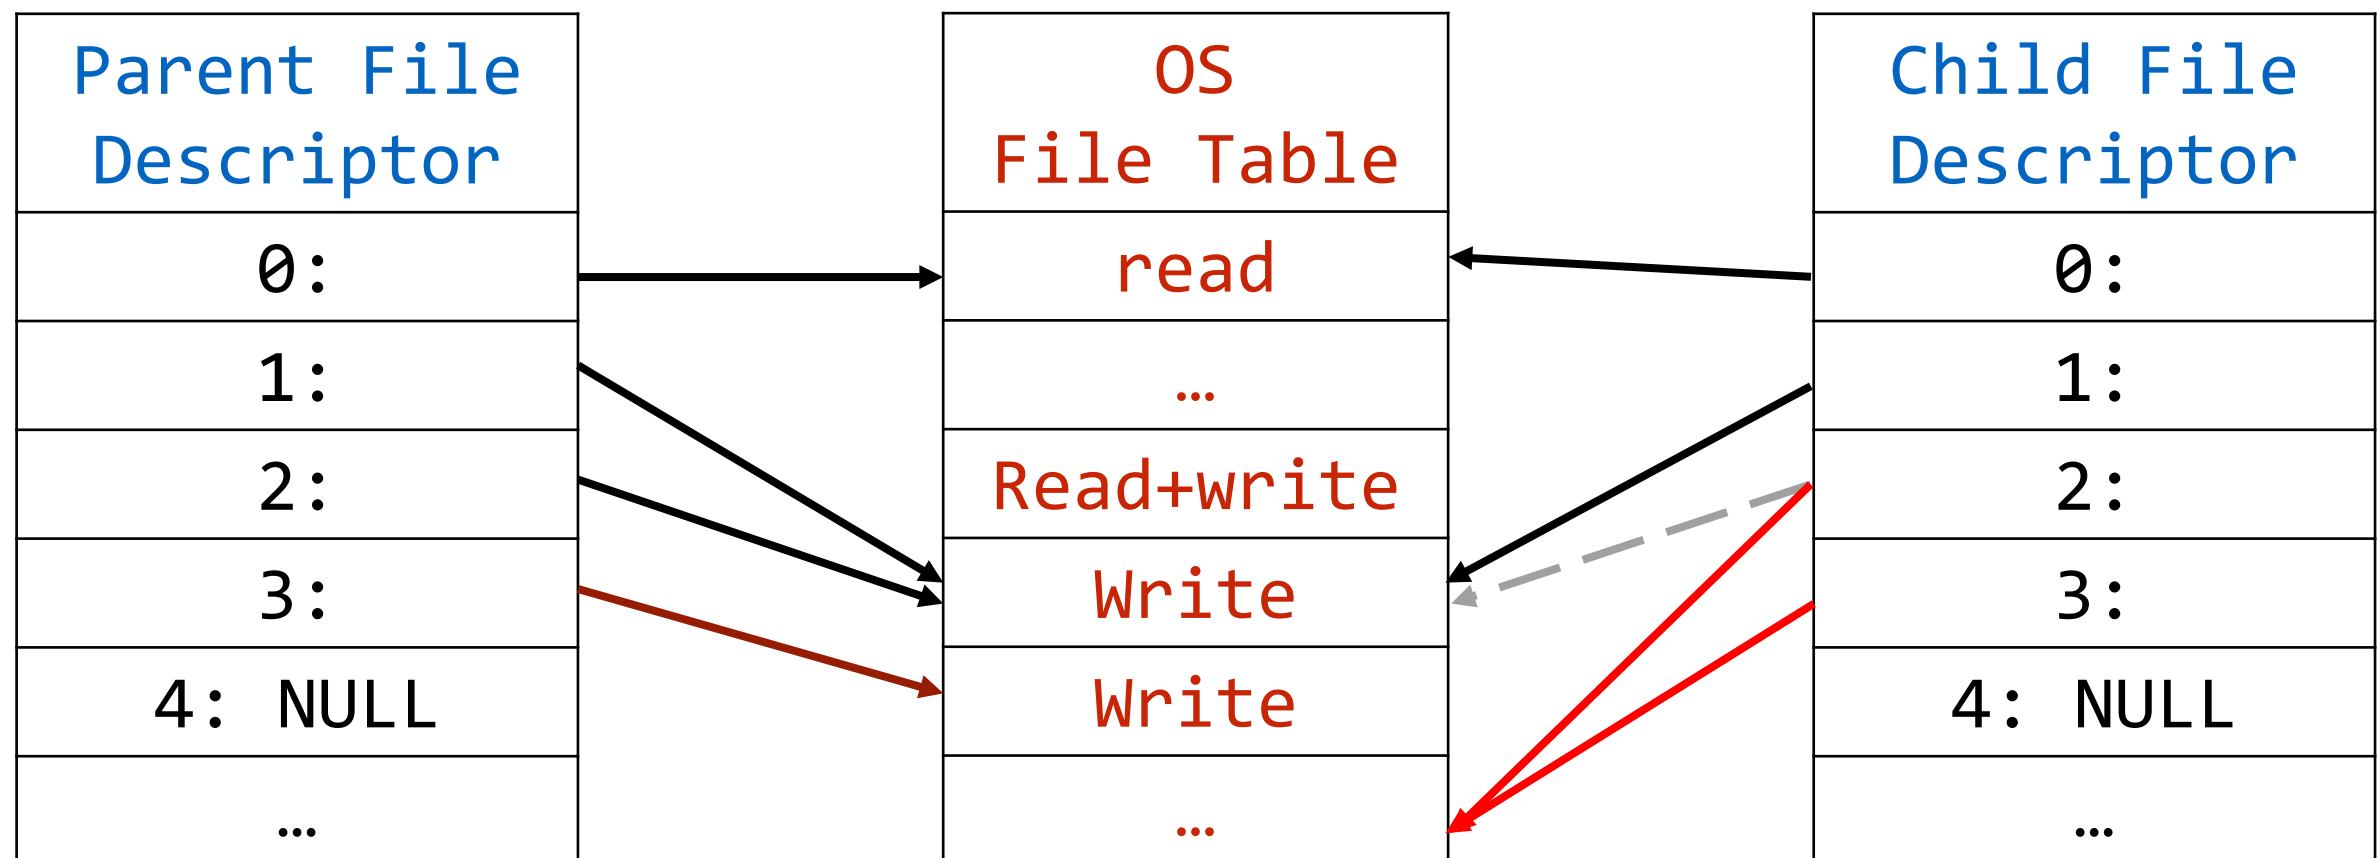


Example: redirecting stderr for another program

- Child:

```
dup2(2, 3);
```

```
close(3);
```

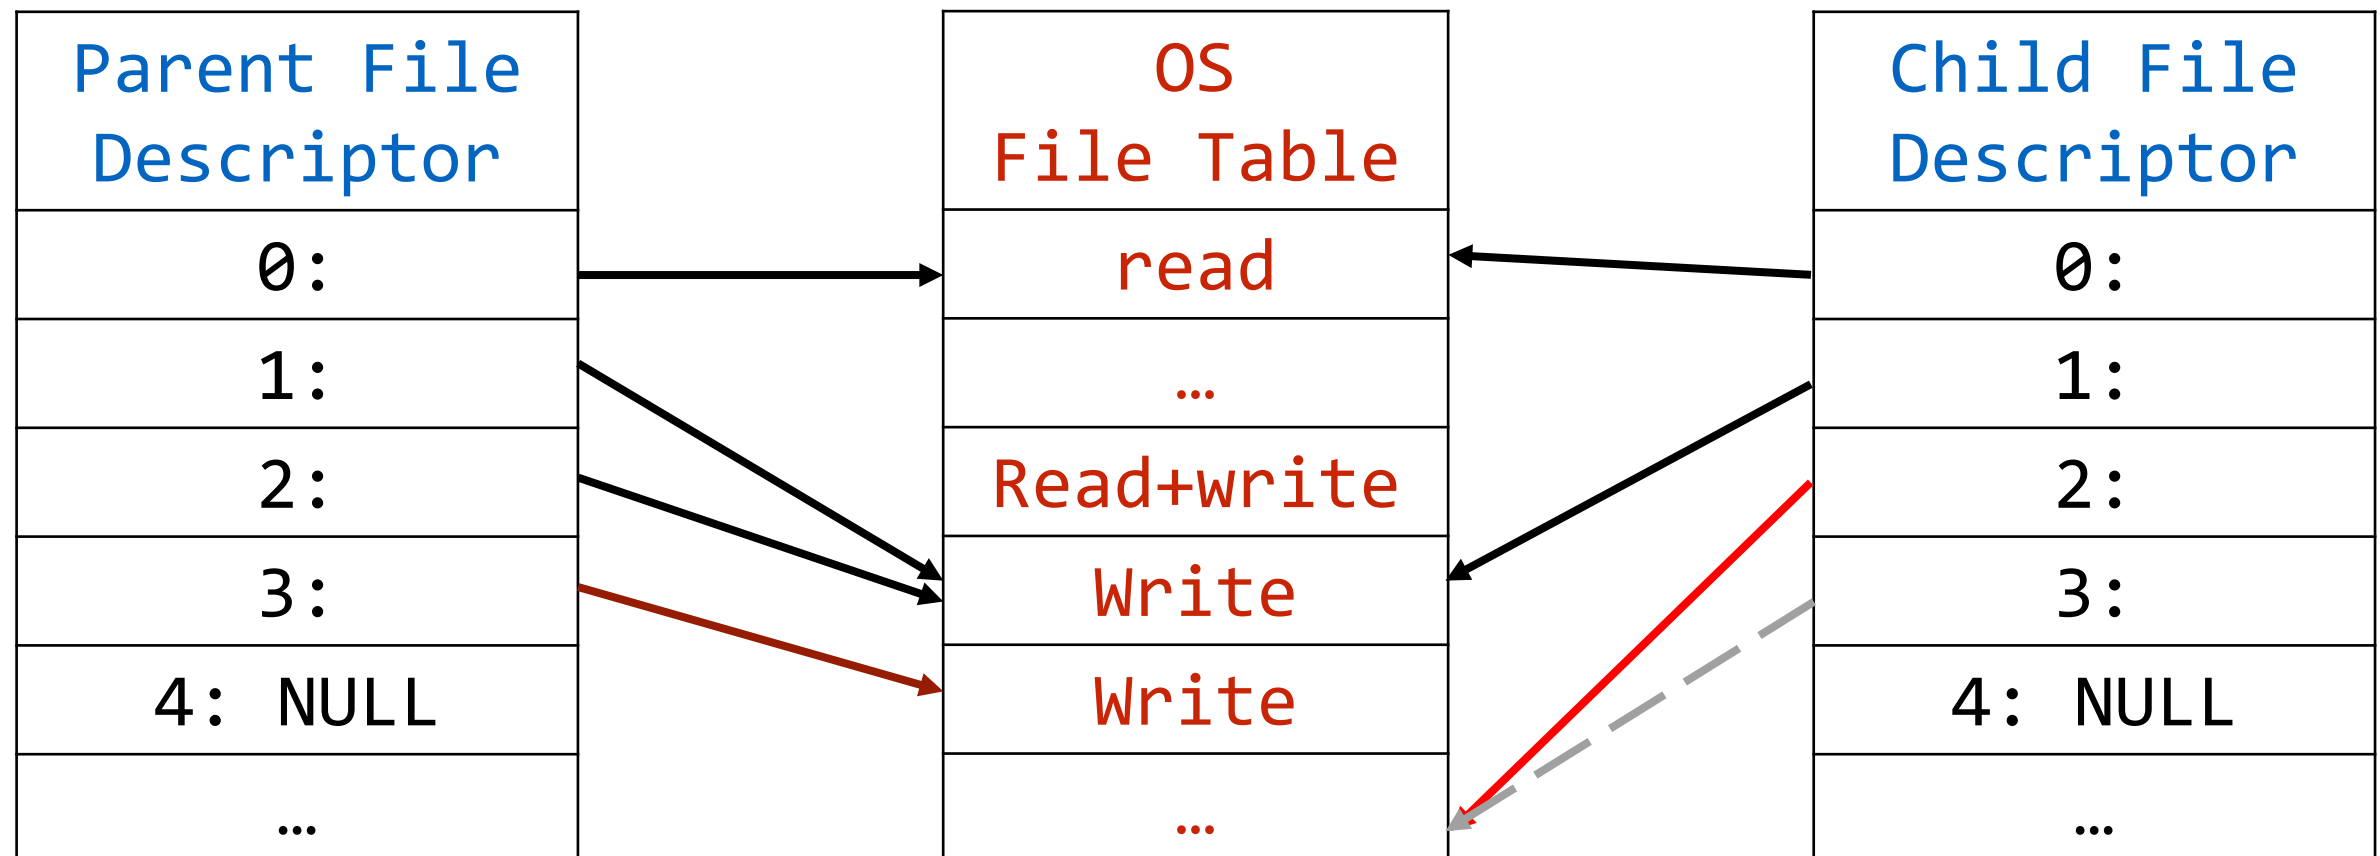


Example: redirecting stderr for another program

- Child:

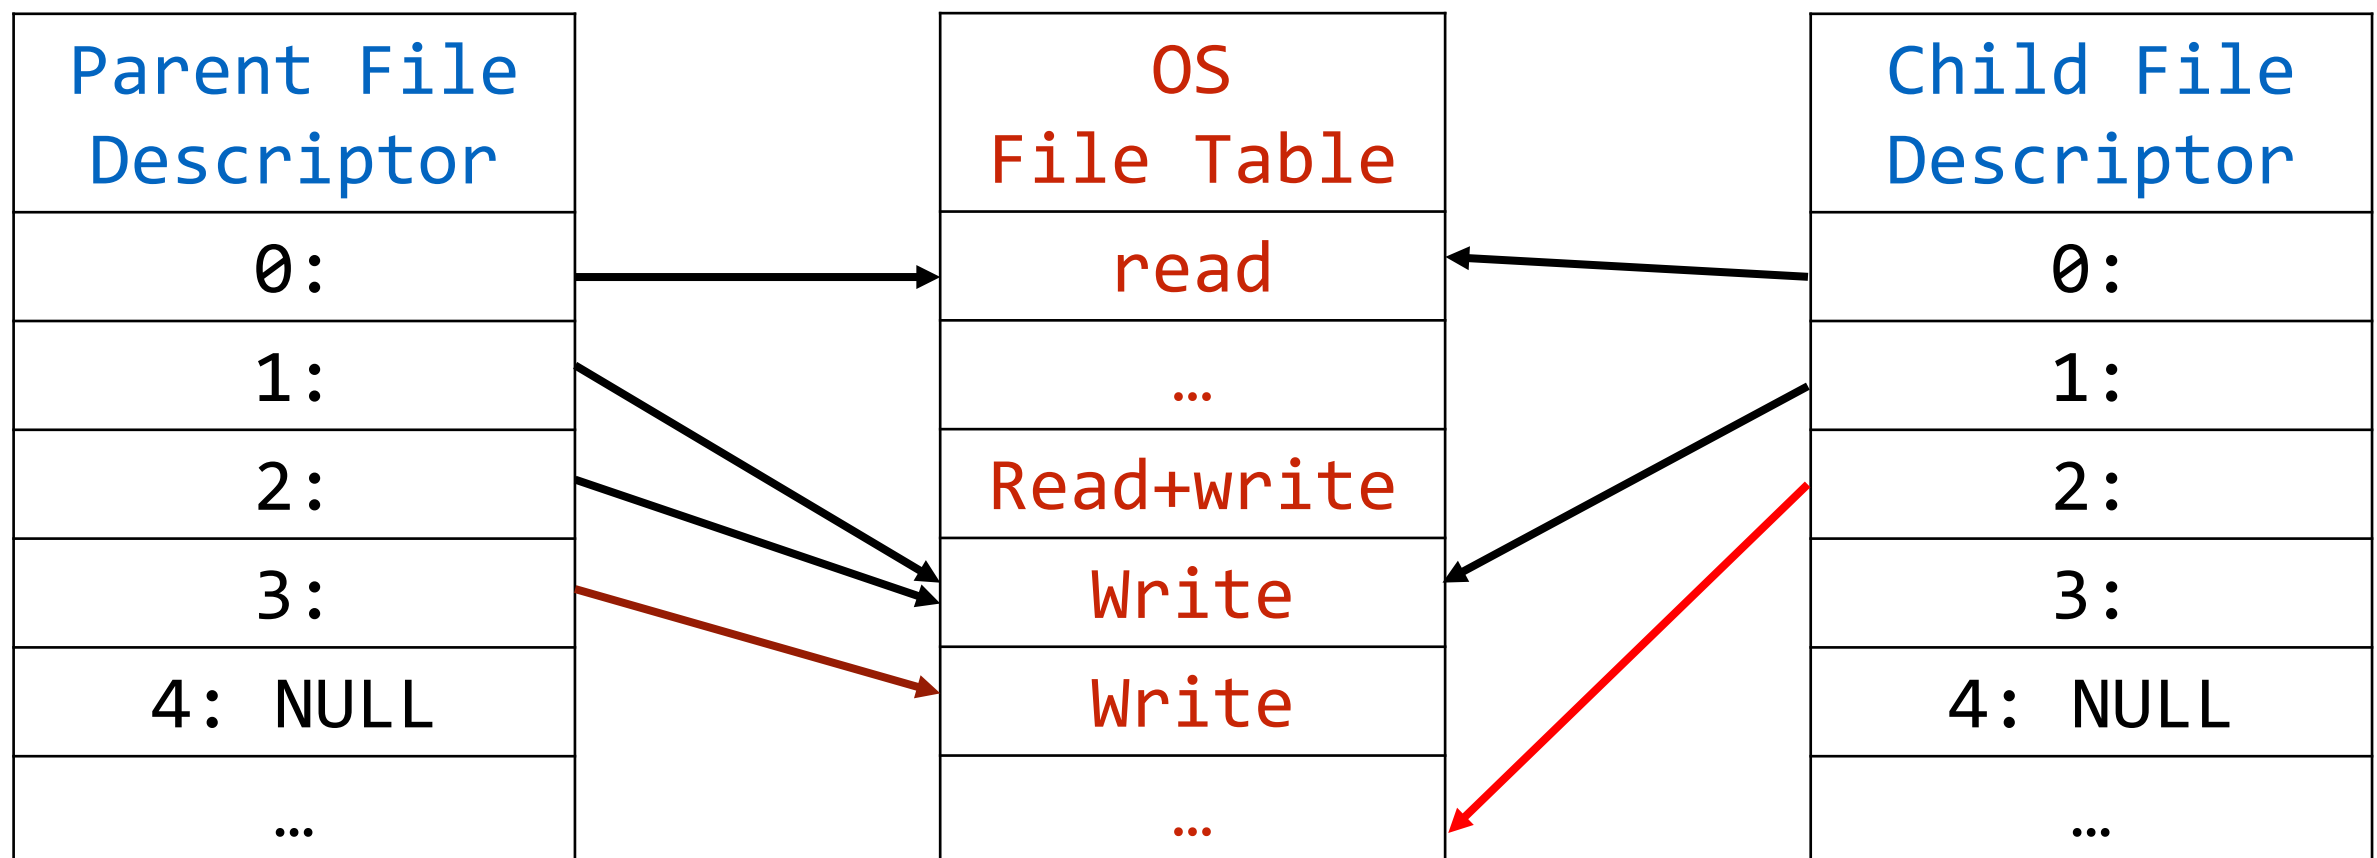
```
dup2(2, 3);
```

```
close(3);
```



Example: redirecting stderr for another program

- Child: exec and the new executable has redirected stderr!
- On exec, open file descriptors are preserved !!!



Summary of Steps in Child Process

1. Close FDs that are opened in parent and not needed in child
2. `open()`. Open a file (and save the file descriptor in `fd`)
3. `dup2()`. Copy FD to the right place
4. `close(fd)`
5. `Exec`



Redirecting stdout for a child process

A process would like to start a new program, and redirect stdout of the new process to a file.

Where & when should the file be opened?

Select the best option.

- A. Before fork(), in parent
- B. After fork() in parent
- C. Before exec in child (after fork())
- D. After exec in child (after fork())
- E. None of the above

More question on redirecting stdout

- Where & when should the new file be opened ?
- Where & when should you call `close(1)` ?
- Where & when should you call `dup` ?
- Where & when should `close(newfd)` be called ?