



A C Primer (6): Dynamic Memory Allocation

Ion Mandoiu

Laurent Michel


Revised by M. Khan and J. Shi

Recall the memory model...

- Three pools of memory
 - Static/global
 - Stack
 - Heap
- Each pool features
 - Different lifetime
 - Different allocation/deallocation policy

0xffffffff

0x00000000





Static/global memory pool

- **This is where**
 - All constants (including string literals) are held
 - Global variables
 - All variables declared “static” are held
- **Allocated when**
 - The program starts
- **Deallocated when**
 - The program terminates
- **FIXED SIZE**
 - Compiler needs to know the size and make reservations



Stack

- This is where....
 - Memory comes from for **local variables** in functions!
- Easy to manage because it is automatic!
 - Allocated automatically when entering the function
 - De-allocated automatically when you leave the function
 - Scope is that of function
 - Should not be used after the function returns
 - For example, indirectly used via a pointer

Default stack size using gcc is 2 MBytes.

Need to increase stack size for large arrays and deep recursion



Heap

- This is where...
 - Memory comes from for **manual** “on-the-fly” allocations
- Who is in charge ?
 - **The programmer** for both allocation / deallocation
- Lifetime of memory blocks ?
 - As long as they are not freed!

Requesting memory on the heap

```
#include <stdlib.h>
```

```
void* malloc(size_t size);
```

- size_t is an unsigned integer data type defined in <stdlib.h>
- used to represent sizes of objects in bytes
- **If successful, a call to malloc(n) returns a generic pointer (void *)**
 - It points to a memory block of **n** bytes on the heap
- **If not successful, NULL is returned**



```
char* p = malloc(100); // request 100 bytes
```

Generic pointers: void*

Pointer to a memory block whose content is “un-typed”

- Use for raw memory operations or in generic functions
- Automatic casting when assigned to other pointer types

```
int * pox = malloc(6 * sizeof(int));
```

- Requires casting before dereferencing for read / write

```
*(int *)pv; // use pv as an int *
```

- NULL, a special pointer value useful for initializations, error handling

```
#define NULL ((void*)0)
```



Can't always get what you want

- A call to `malloc()` may fail
 - For example, if you are out of memory
- In this case you get back a `NULL` value
 - Not much to do except report the error (and terminate nicely)
- Idiom

```
char* p = malloc(100); // request 100 bytes
if (p == NULL) {
    // report error and finish
    perror("Not enough memory");
    exit(1);
}
```




How Much Space ?

- You need to tell malloc() how many **bytes** you need
- To request space for an array, need
 - Number of elements
 - Amount of space for each array element
 - sizeof(**T**) returns number of bytes needed for a value of type **T**
- **Example**

```
int* pox = malloc(6 * sizeof(int)); // request space for 6 ints
if (pox == NULL)
    report error and finish;
```

Another way

```
#include <stdlib.h>
```

```
void* calloc(size_t nmemb, size_t size);
```

- `calloc()` is implemented in terms of `malloc()`
 - **`calloc()` also initializes the content to 0**



```
int* pox = calloc(6, sizeof(int)); // request space for 6 ints
if (pox == NULL)
    report error and finish;
```

Adjusting the size

```
#include <stdlib.h>
```

```
void* realloc (void* ptr, size_t size);
```



What if you change your mind?

- You requested 100 bytes, but now need 200!

```
char* p = malloc(100); // request 100 bytes
...
p = realloc(p, 200); // p may change!
```

- Before a call to `realloc(p, size)`, `p` must be
 - A pointer returned by a previous `malloc/calloc/realloc`
 - Or `NULL`, in which case the call is equivalent to `malloc(size)`

Deallocation

```
#include <stdlib.h>  
void free(void *ptr);
```

- Straightforward
 - Simply call the library function “free”
 - Takes a pointer to the block to free

```
free(ptr);
```

- Do not free a pointer twice!
 - After freeing a pointer, set it to NULL!



Two key rules


Rule 1: Everything you requested should be freed, eventually

Rule 2: Only free what is allocated via malloc/calloc/realloc

- Consequences of not following the rules
 - **Memory “leaks”**
 - Your program will eventually run out of memory
 - **Undefined behavior and horrible crashes**
 - Freeing unallocated memory or already freed memory
 - May cause a memory error and a program crash
 - Worse, may corrupt the heap and cause a crash later
 - Even worse, the program may keep running, totally corrupting your data, and writing it to disk without you realising

Pointers and arrays

After `pox = malloc(6 * sizeof(int));`



- That looks like an array!!!
 - And you **can** use pox as if it is an array



Example: use pointer as array

```
// programmers have to manage memory themselves before C99
void doSomething(int n) {
    int * pox;
    pox = malloc(sizeof(int)*n); // request mem from heap
    *pox = 0;    // set the int at the address pox to 0
    pox[0] = 0;  // same thing
    pox[1] = 1;  // the int after pox[0]
    // more lines here ...
    free(pox);   // remember to free!
}
```

Array of pointers

- Pointers, like integers, can be placed in an array

```
int a0;  
int a1;  
int a2;
```

```
char *p0 = malloc(10);  
char *p1 = malloc(10);  
char *p2 = malloc(10);
```

```
// array of int's  
int a[3];
```

```
// array of pointers  
char * p[3];
```

```
// Can also do with a loop  
p[0] = malloc(10);  
p[1] = malloc(10);  
p[2] = malloc(10);
```



Example: command line arguments

```
int main (int argc, char *argv[]);
```

- `argc`: the number of arguments on the command line
- `argv`: array of pointers to characters
 - The number of elements is `argc`
 - Each element in an array points to null-terminated strings

- Example:

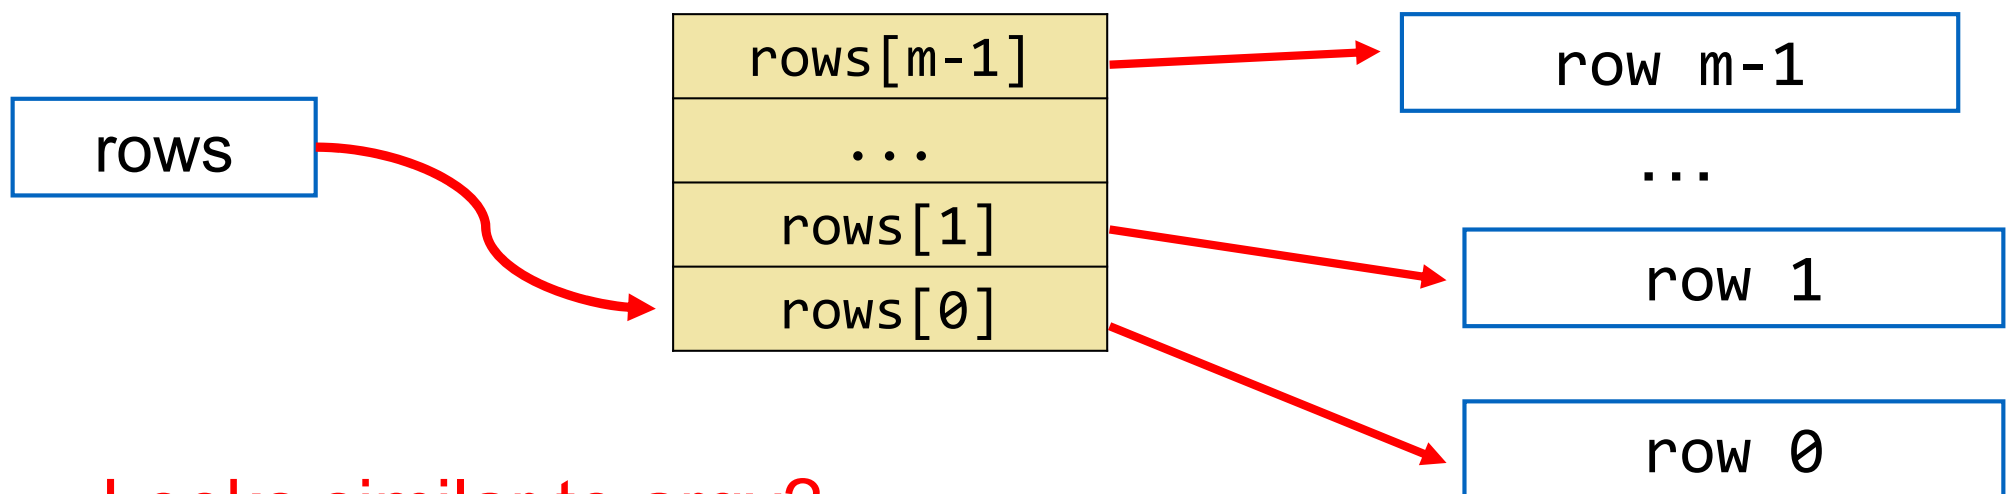
```
$gcc -o ab ab.c
```



What is `argv[3][1]`?

Example: allocating 2d dynamical array

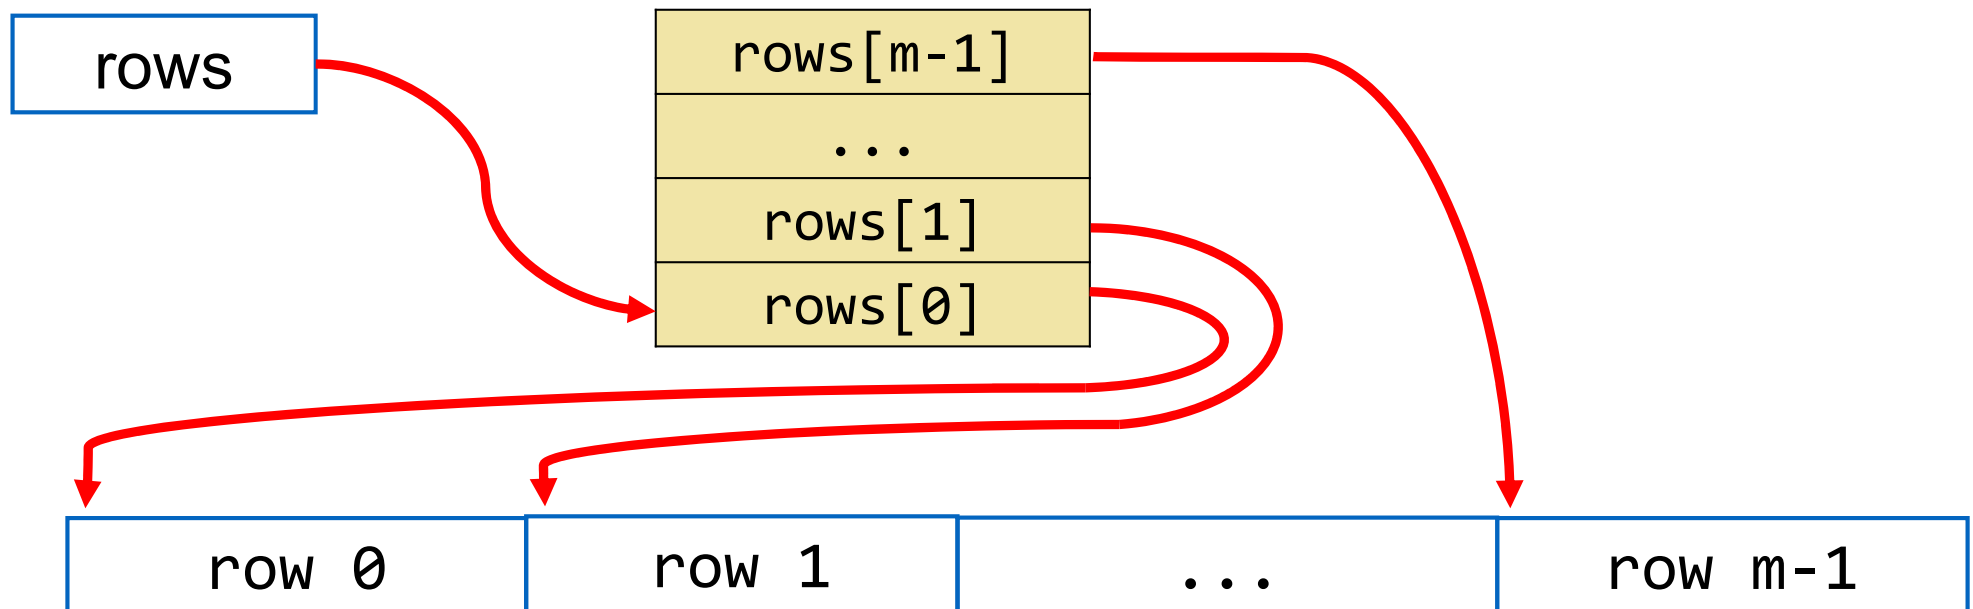
```
void doSomething(int m, int n)
{ int **rows;
  rows = malloc(sizeof(int *) * m); // array of pointers
  for (int i = 0; i < m; i++)
    rows[i] = malloc(sizeof(int)*n); // one int array for each row
  ...
  for (int i = 0; i < m; i++)
    free(rows[i]);
  free(rows);
}
```



Looks similar to argv?

Example: allocating 2d dynamical array (take 2)

- Request all the memory space needed by data with one `malloc()` call
 - Instead of call `malloc()` m times, for each row
- Calculate `rows[1]`, `rows[2]`, and so on (pointer arithmetic! next lecture)
 - You can even calculate the address of each element (e.g. `rows[10][5]`)






Example: allocating 2d dynamical array (take 3)

```
void doSomething(int m, int n)
{ int (* arr)[n];    // arr is a pointer to an array of n int's
  arr = malloc(m * n * sizeof(int)); // one malloc() for data
  // to access elements
  arr[1][2] = 1; arr[2][3] = arr[1][2] + 10;
  ...
  free(arr); // free memory
}
```

n must be constant if Variable Length Array(VLA) is not supported





Study the remaining slides yourself

Pointers taking the address of ...

- A static ?

- The address is never going to go “bad”
- The static lives as long as the program!

- A stack [automatic] variable ?

- The address is valid as long as the variable is!
- When the function returns.... The address is bogus



- A heap variable ?

- The address is valid as long as the variable is!
- The variable disappears when explicitly de-allocated (freed)