

CSE 3400 - Lecture set 8: Public Key Infrastructure (PKI)

Last updated: 4/21/20

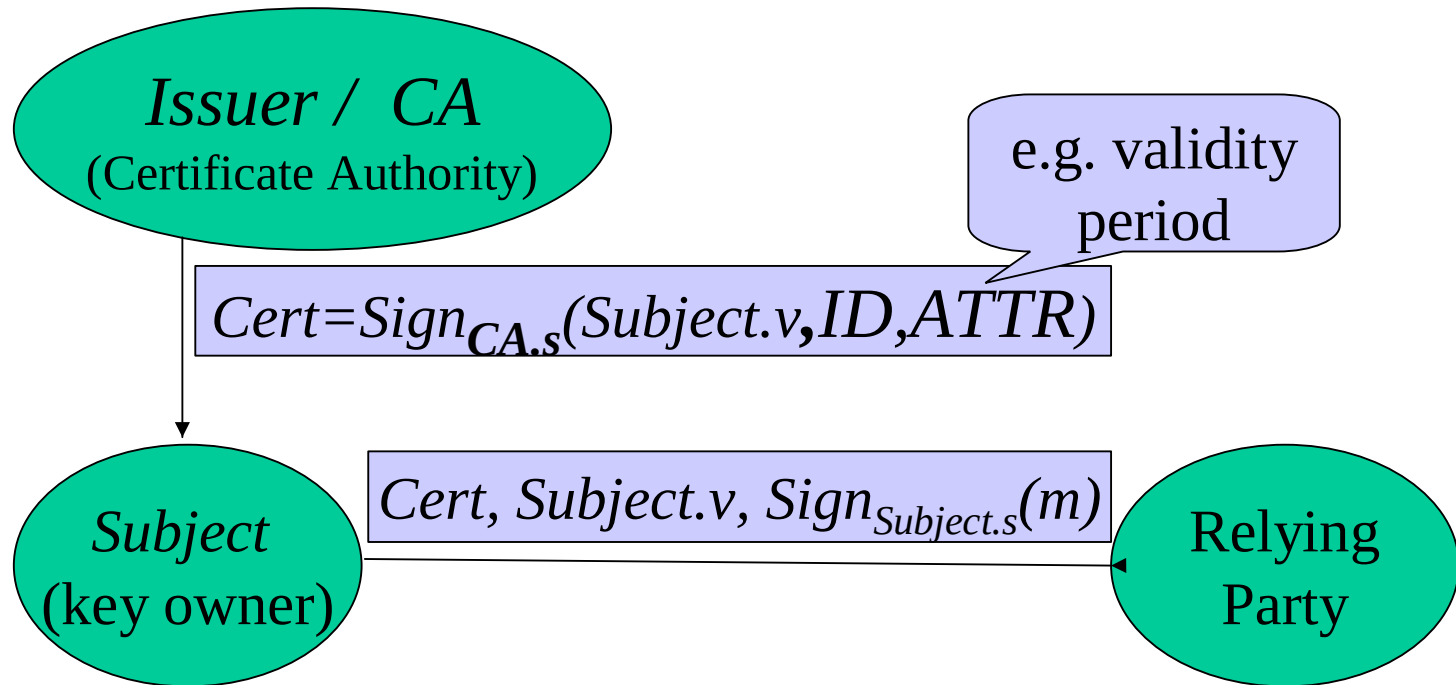
© Prof. Amir Herzberg

Public keys are very useful...

- Secure web connections
- Software signing (against malware)
- Secure messaging, email
- Crypto-currency, blockchains, financial crypto...
- But to use a PK, it must be authenticated
 - Mainly: signed by a **trusted Certificate Authority**
 - E.g., in TLS, browsers maintain list of 'root CAs'
 - So far, not much use of **personal** certificates
 - Secure email: not widely deployed
 - Secure messaging: auth by provider (& user ??)

Public Key Certificates & Authorities

- **Certificate**: signature by **Issuer / Certificate Authority (CA)** over subject's public key and **attributes**
- **Attributes**: identity (ID) and others...
 - Validated by CA (liability?)
 - Used by **relying party** for decisions (e.g., use this website?)



Can we trust a certificate /

■ CA ? Browser default: trust ~100 'root CAs'

- Each can define 'intermediate CAs'

■ But failures have happened... E.g.:

- 2001, VeriSign: attacker gets code-signing certs
- 2008, Thawte: email-validation (attackers' mailbox)
- 2008,2011: Comodo not performing domain validation
- 2011: DigiNotar compromised, >500 rogue certs
 - Incl. wildcard cert for Google for mass-interception in Iran
- 2011: TurkTrust issued intermediate-CA certs to users
- 2014: India CCA / NIC compromised (rogue certs)
- 2015: Root CA CNNIC (China) issued CA-cert to MCS (Egypt) who issued rogue certs
- 2015,17: Symantec issued unauthorized certs for >176 domains

PKI Lecture: Topics

- **X.509 Certificates**
- Revoking certificates
- Dealing with CA failures: CT (Certificate Transparency) and other defenses
- Conclusions and challenges...

The X.500 Global Directory

■ X.500: an ITU standard, first issued 1988

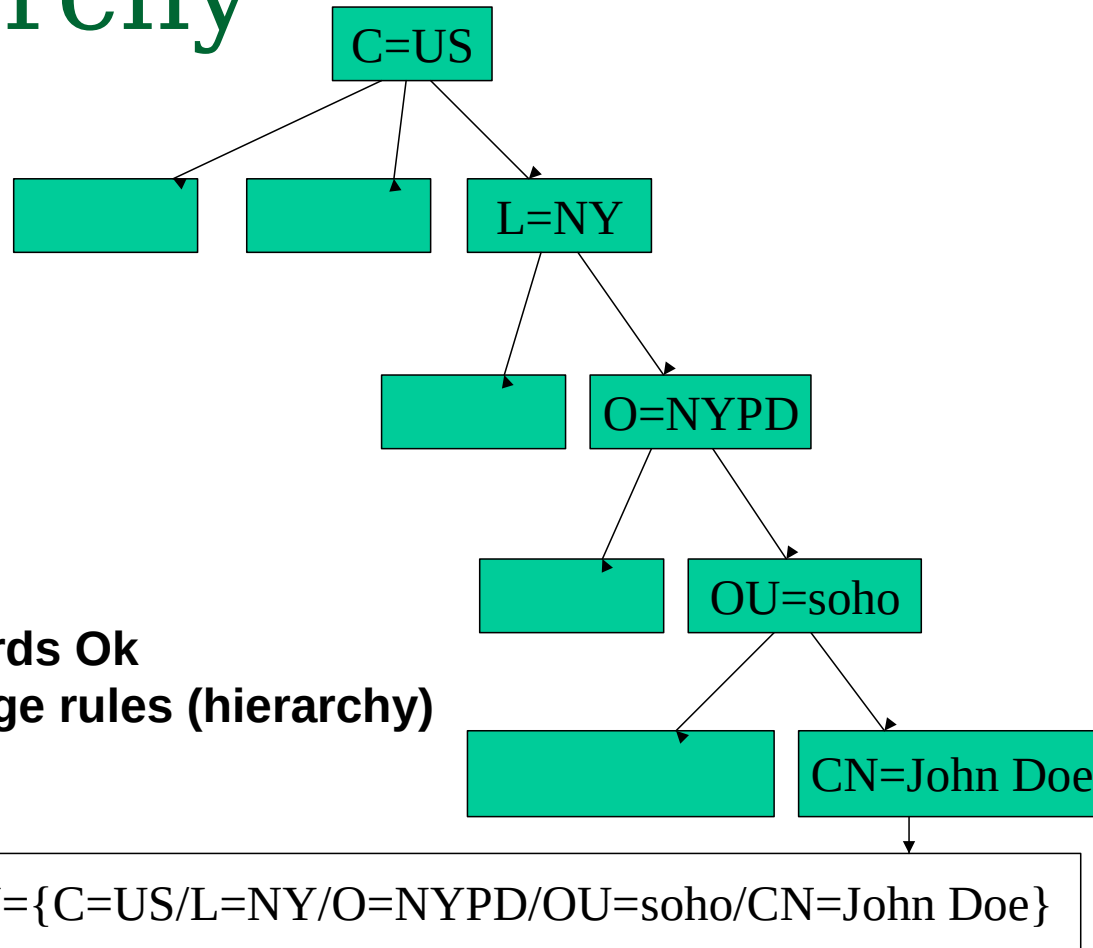
- ITU: International Telcos Union
- Idea: trusted global directory
 - Operated by hierarchy of trustworthy telcos
 - Never happened
 - Too complex, too revealing, too trusting of telcos
- Directory bind identifiers to attributes
 - Standard attributes (incl. public key)
 - Standard identifiers: Distinguished Names
 - Goal: unique, meaningful, decentralized identifiers

X.500 Distinguished Names

- Goal: meaningful, unique and decentralized identifiers
- Sequence of keywords, a string value for each of them
- Distributed directory, responsibility \square *hierarchical DN*

Keyword	Meaning
C	Country
L	Locality name
O	Organization name
OU	Organization Unit name
CN	Common Name

Distinguished Name (DN) Hierarchy



Comments:

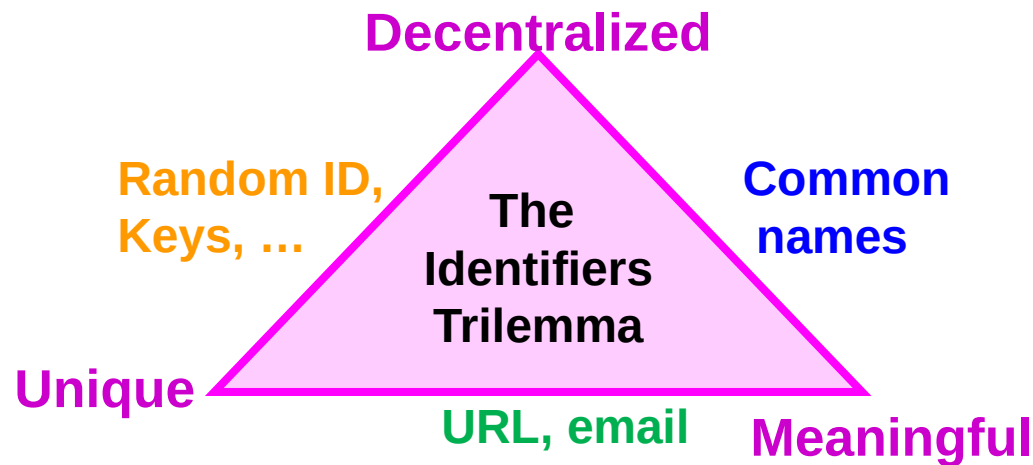
1. Other keywords Ok
2. No strict usage rules (hierarchy)

Goals for Identifiers in Certificates

- Meaningful (to humans)
 - Memorable, reputation, off-net, legal
- Unique identification of entity (owner)
- Decentralized - with Accountability:
assigned by a trusted (certificate) authority
 - Accountability: identification of assigning authority

The Identifiers Trilemma

- Three goals: Meaningful, Unique, Decentralized
- Examples of achieving any two of the goals:
 - Unique + Meaningful: URL, email
 - Meaningful + Decentralized: common name
 - Unique + Decentralized: hash of key
- **Ensuring all three goals seems challenging!**



Distinguished Names -

Evaluation

■ Decentralized?



- Separate name spaces

■ Unique ?



- Could be, if each name space has one issuer
- TLS reality: browsers trust 100s of CAs for **every name**

■ Meaningful?



- Usually: Julian Jones/UK/IBM
- But not always: Julian Jones2/UK/IBM
 - Added 'counter' to distinguish □ mistakes, loss of meaning

Distinguished Names – More

- **Problems** Distinguished Name fields may expose
 - Organizational sensitive information (e.g. unit)
 - Privacy
- Handling changes in position, organizations
- Multiple, related hierarchies:
 - International organizations, divisions...
 - Julian Jones/UK/IBM or Julian Jones/IBM/UK ?
 - Julian Jones/Research/UK/IBM or Julian Jones/UK/Research/IBM ?

And: DNs aren't usable

identifiers

- Relying parties (users) don't know the DN

DN = {C=US/L=corp/O=Bank}



https://bank.com




- Hopefully, they know the domain (in URL)
- Naming extensions: alternative names
 - DNS name: `cert.SubjectAltName.dNSname`
 - Wildcard domain names: `*.bank.com`

X.509 public key certificates

- X.509: authentication mechanisms of X.500
- Initially: Authenticate to Directory (PW, Pub key)
 - To maintain entity's record
- Later (and now): **X.509 certificate**
 - Signature binds **public key** to distinguished name (DN)
 - And to other attributes
 - Some defined in X.509 standard, others in `extensions`
- Used widely
 - SSL / TLS, code-signing, PGP, S/MIME, IP-Sec, ...
 - In spite of complexity

Original (V1) X.509 Certs



Object Identifiers (OID):

- Global, unique identifiers
- Sequence of numbers, e.g.: 1.16.840.1.45.33
 - Hierarchical

X.509 Certs & Subject Identifiers

- V1: Distinguished Name (for subject & issuer)
- V2: unique identifiers (for subject & issuer)
- V3: extensions
 - PKIX standard: standard extensions
 - PKIX: Public Key Infrastructure working group of IETF
 - Widely adopted, including in SSL/TLS (& https)
 - Example: SubjectAltName extension
 - Including DNSname: identify website by domain name

X.509 Public Key

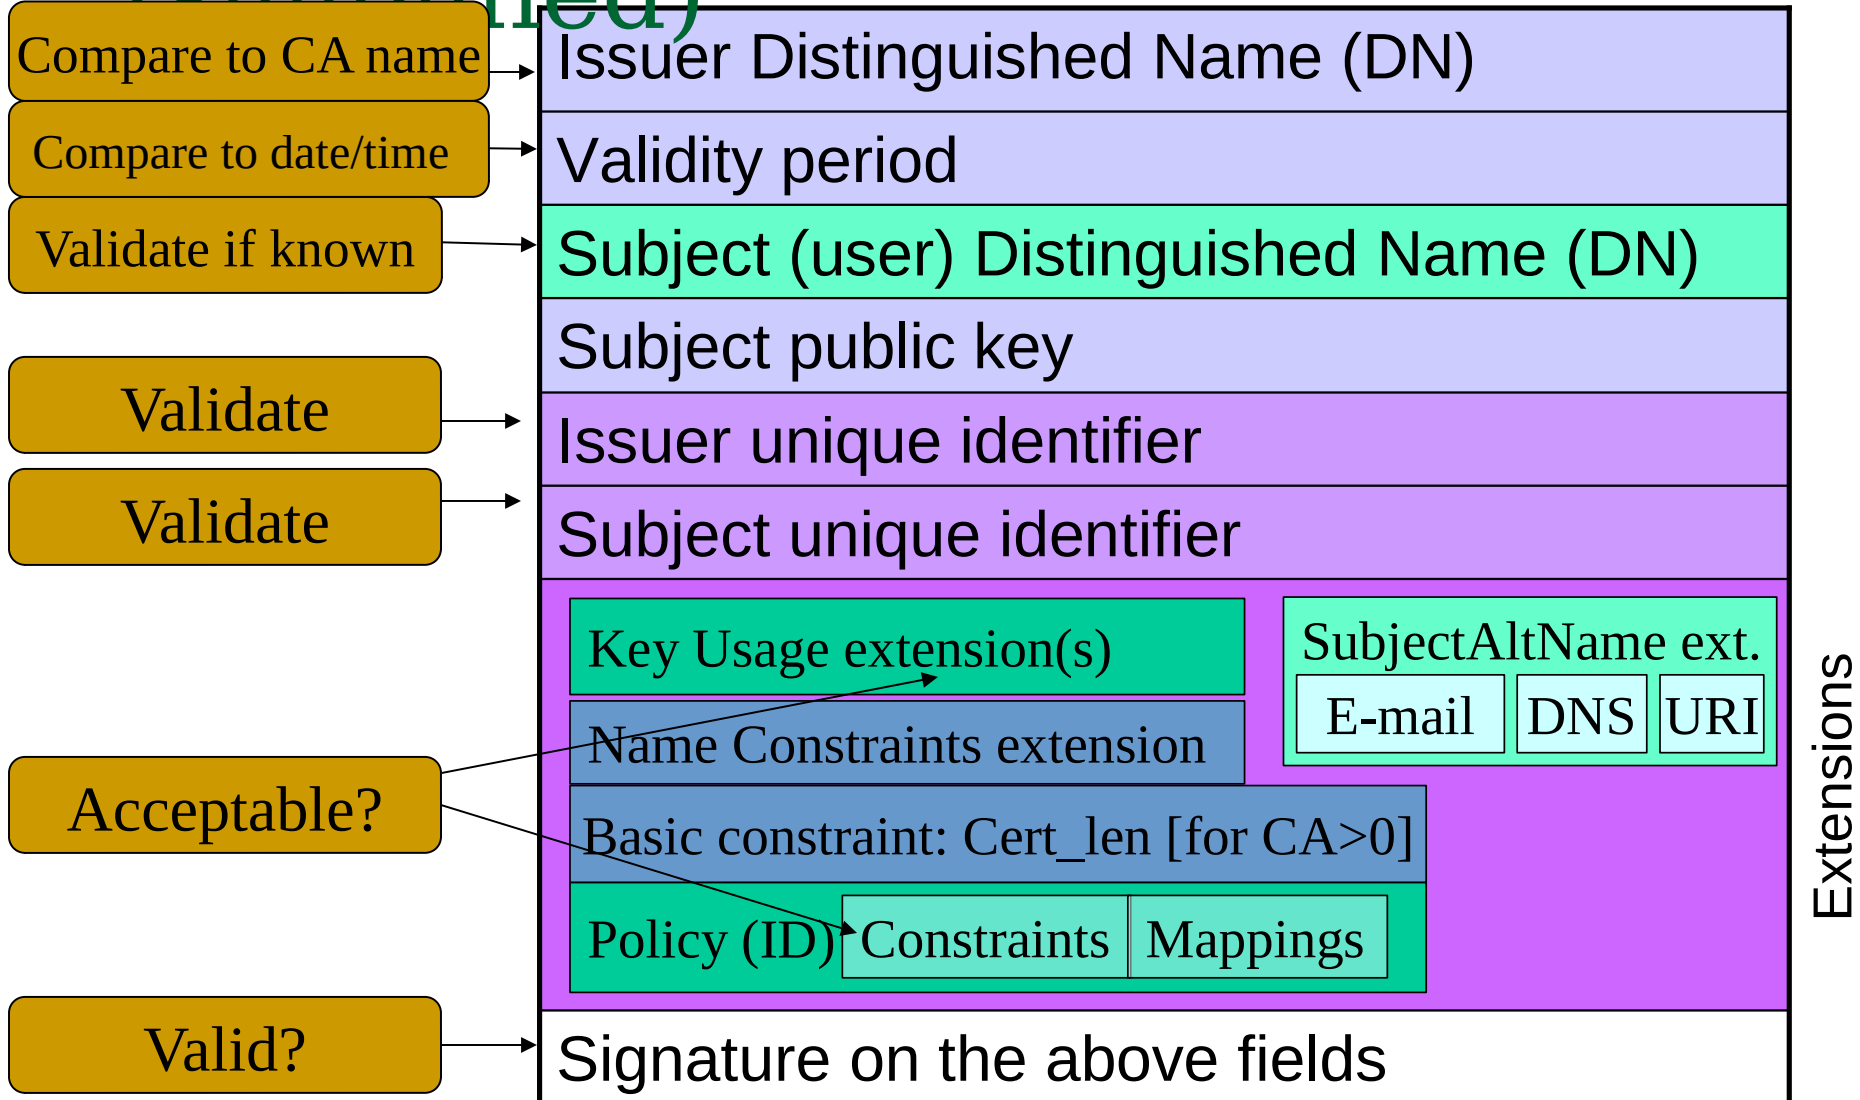
Signed fields	Version		
	Certificate serial number		
	Signature Algorithm Object Identifier (OID)		
	Issuer Distinguished Name (DN)		
	Validity period		
	Subject (user) Distinguished Name (DN)		
	Subject public key information	Public key Value	Algorithm Obj. ID (OID)
	Issuer unique identifier (from version 2)		
	Subject unique identifier (from version 2)		
	Extensions (from version 3)		
Signature on the above fields			

X.509 V3 Extensions

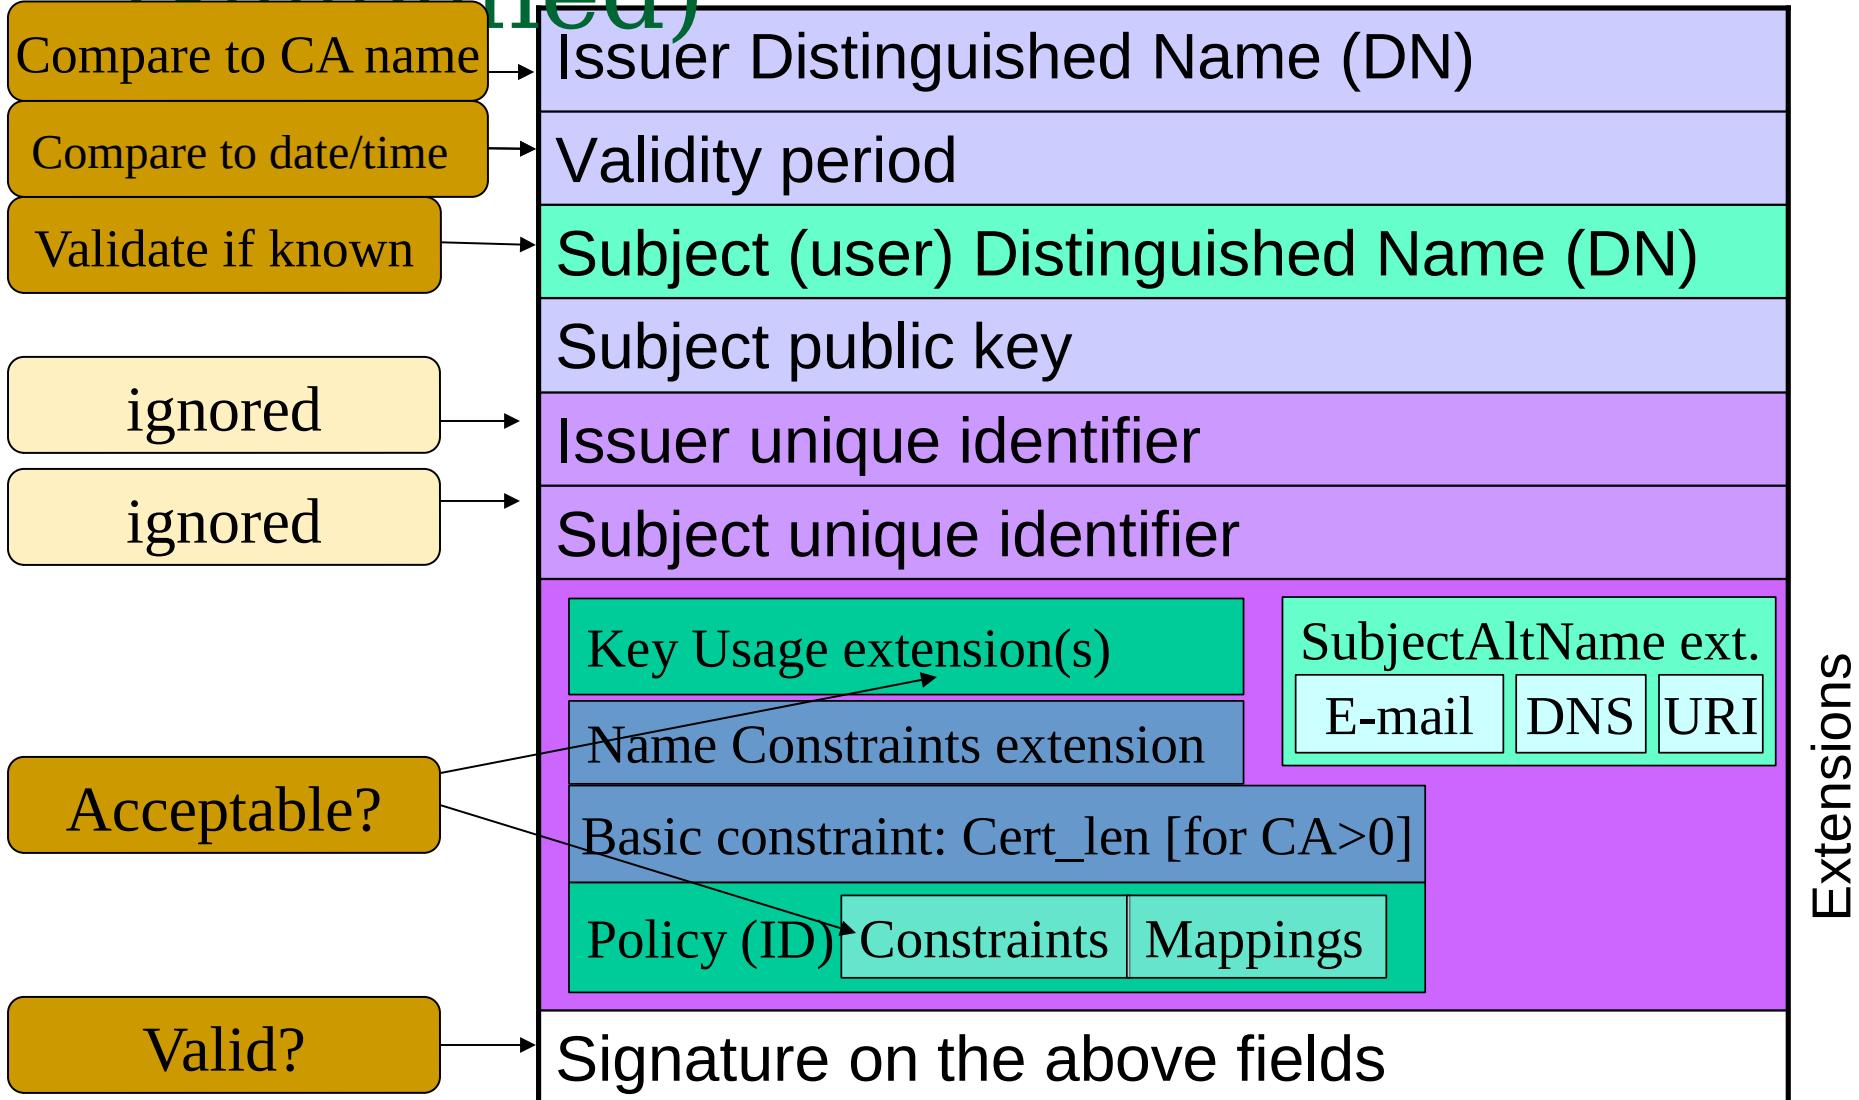
Mechanism

- Each extension contains...
 - Extension identifier
 - As an OID (Object Identifier)
 - E.g. `Naming constraints`
 - Extension value
 - E.g. `Permit C=IL`, `Exclude dNSName=IBM.COM`
 - **Criticality indicator**
 - **If critical**, relying parties **MUST NOT** use a certificate with any unknown critical extension
 - E.g. Naming constraints is `critical`
 - **If non-critical**: use certificate w/o unknown critical extensions; **ignore** unknown (non-critical) extensions

X.509 Certificate Validation (simplified)



PKIX Certificate Validation (simplified)



PKIX: Internet PK

Infrastructure

- RFC 5280: X.509 - Internet extensions
- Addressing different goals:
 - **Improved naming:**
Subject, Issuer AltName extensions
 - Establishing trust in CAs:
Basic, Name and Policy Constraints extensions
 - Other goals and extensions

SubjectAltName (SAN)

Extension

- Bound identities to the subject
 - In addition/instead of Subject Distinguished Name
 - Same extension may contain multiple SANs
- Goal: unique and meaningful names
 - Common: DNS name (dNSName), e.g., a.com
 - TLS/SSL allows wildcard domains (*.a.com)
 - Not addressed in PKIX (RFC5280); see RFC6125
 - Or: email address, IP address, URI, other
 - Don't use (ignore) X.509v2 `unique identifiers`

IssuerAltName Extension

- Bound identities to the issuer
 - In addition/instead of Issuer Distinguished Name
- Same goal, syntax as SubjectAltName (SAN)


PKIX: Internet PK

Infrastructure


- RFC 5280: X.509 + Internet extensions
- Addressing different goals:
 - Improved naming:
Subject, Issuer AltName extensions
 - **Establishing trust in CAs:**
Basic, Name and Policy Constraints extensions
 - Other goals and extensions

Different types of PKIs

- **Simple PKI:** just one trusted CA:



- **Multiple trusted CAs:**




- Both **static**: CAs known, trusted in advance
- **Dynamic PKI**: establishing trust in 'new' CA, via certificates from 'established' CAs

Dynamic PKIs


■ **Web/TLS PKI:** ‘root CAs’+‘intermediate CAs’:

- Root CA issues cert for intermediate CAs
- Supported by PKIX



■ **Or... Web-of-Trust PKIs:**

- Directed graph, not tree
- Different variants/policies
 - Some supported by PKIX



Web of Trust PKI

- PGP's friends-based Web-of-Trust:
 - Everyone is subject, CA and relying party
 - As a CA, certify (pk, name) for `friends`
 - As a subject, ask friends to sign for you
 - As a relying party, trust certificates from friends
 - Or also from friends-of-friends? Your policy....




PKIX PKI Trust Rules

- PKIX allows **path-based PKI policies**
 - A certificate path is a set c_1, \dots, c_n of certs, s.t. c_i is validated using $c_{i-1}.v$ (sig-valid-key in previous cert)
- Given cert-path c_1, \dots, c_n , **cert c_n is valid** if:
 - The first cert c_1 was issued by a trusted CA,
 - **Basic, naming and policy constraints** satisfied
 - c_1, \dots, c_{n-1} are CA certs:
 - **BasicConstraint.cA** and **keyUsage.keyCertSign** are **TRUE**
 - □ All their subjects become trusted CA ('dynamic')

Establishing trust in cert

■ Relying party A trust c_3 , given also c_1, c_2 , if:

- CA1 is a trust-anchor ('root CA')
- $\text{Valid}_{\text{CA1.v}}(c_1)$ [signed by CA1.v]
 - Let v_{i+1} be the public key signed in c_i
- $\text{Valid}_{v_1}(c_2), \text{Valid}_{v_2}(c_3)$
 - Basic constraints:
 - c_1, c_2 mark CA2, CA3 (resp.) as CAs
 - Path-Length: in c_1 at least 2, in c_2 at least 1
 - Policy, naming constraints satisfied



Example: delegation of trust

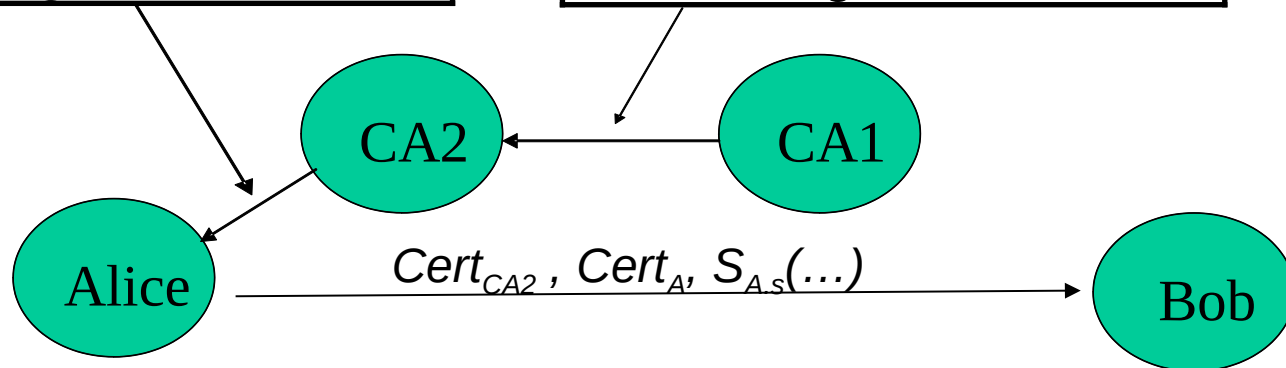
- CA1 'vouches for' CA2
- □ Interoperability (if Bob knows only CA1)

$$Cert_A = Sign_{CA2.s}(A.e, Alice)$$

Issuer DN: CA2
Subject DN: Alice
Alice's pub. key
Alice is not a CA
CA2's Signature

$$Cert_{CA2} = Sign_{CA1.s}(CA2.v, CA2)$$

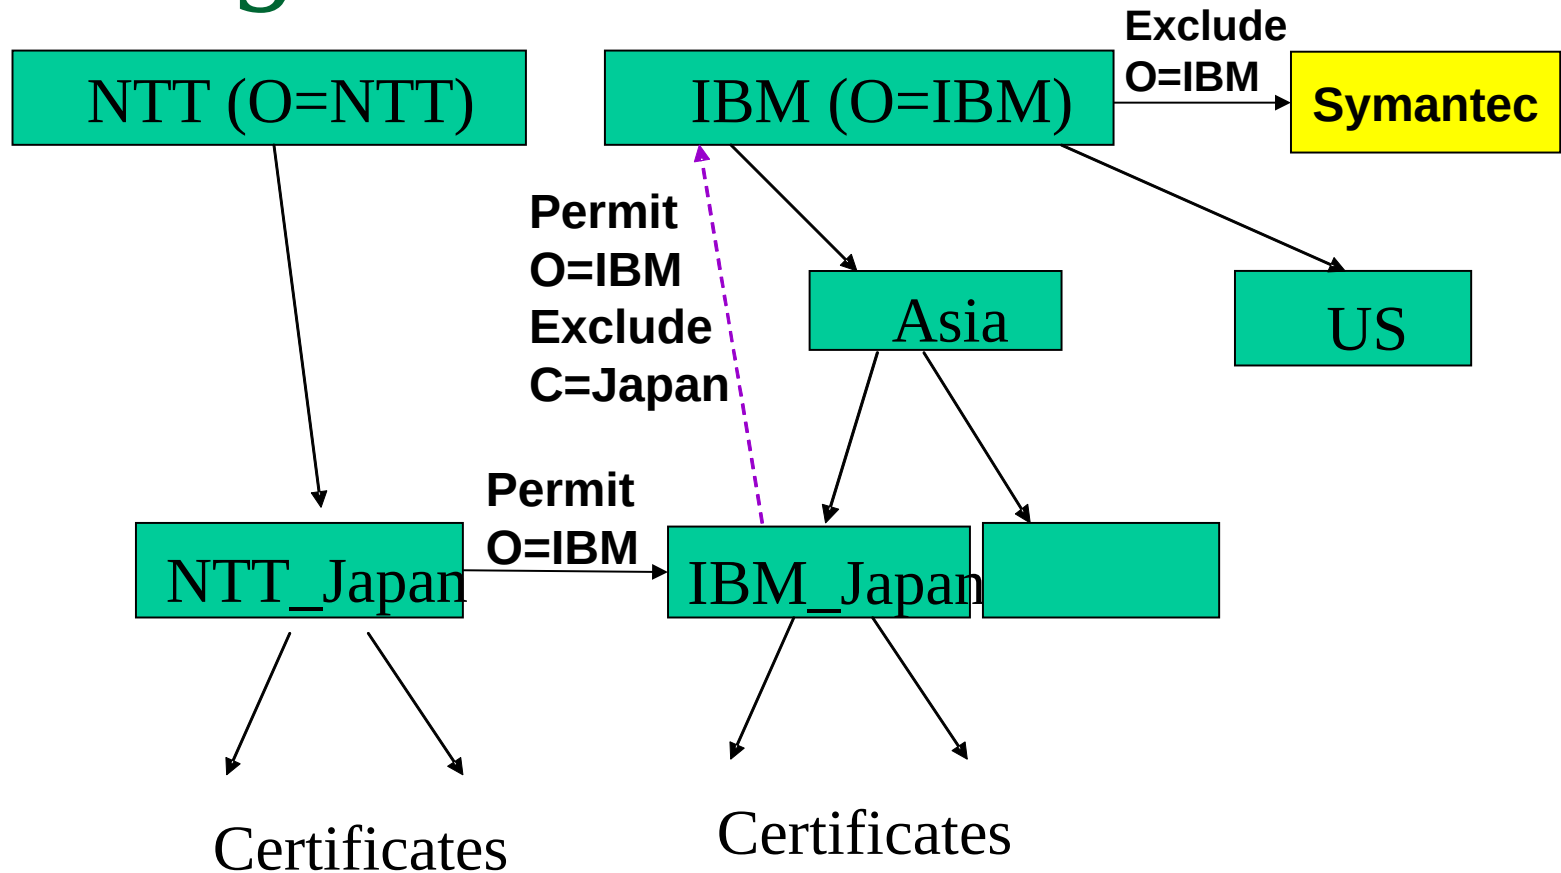
Issuer DN: CA1
Subject DN: CA2
CA2's pub. key
CA2 is a CA
CA1's Signature



Constraints

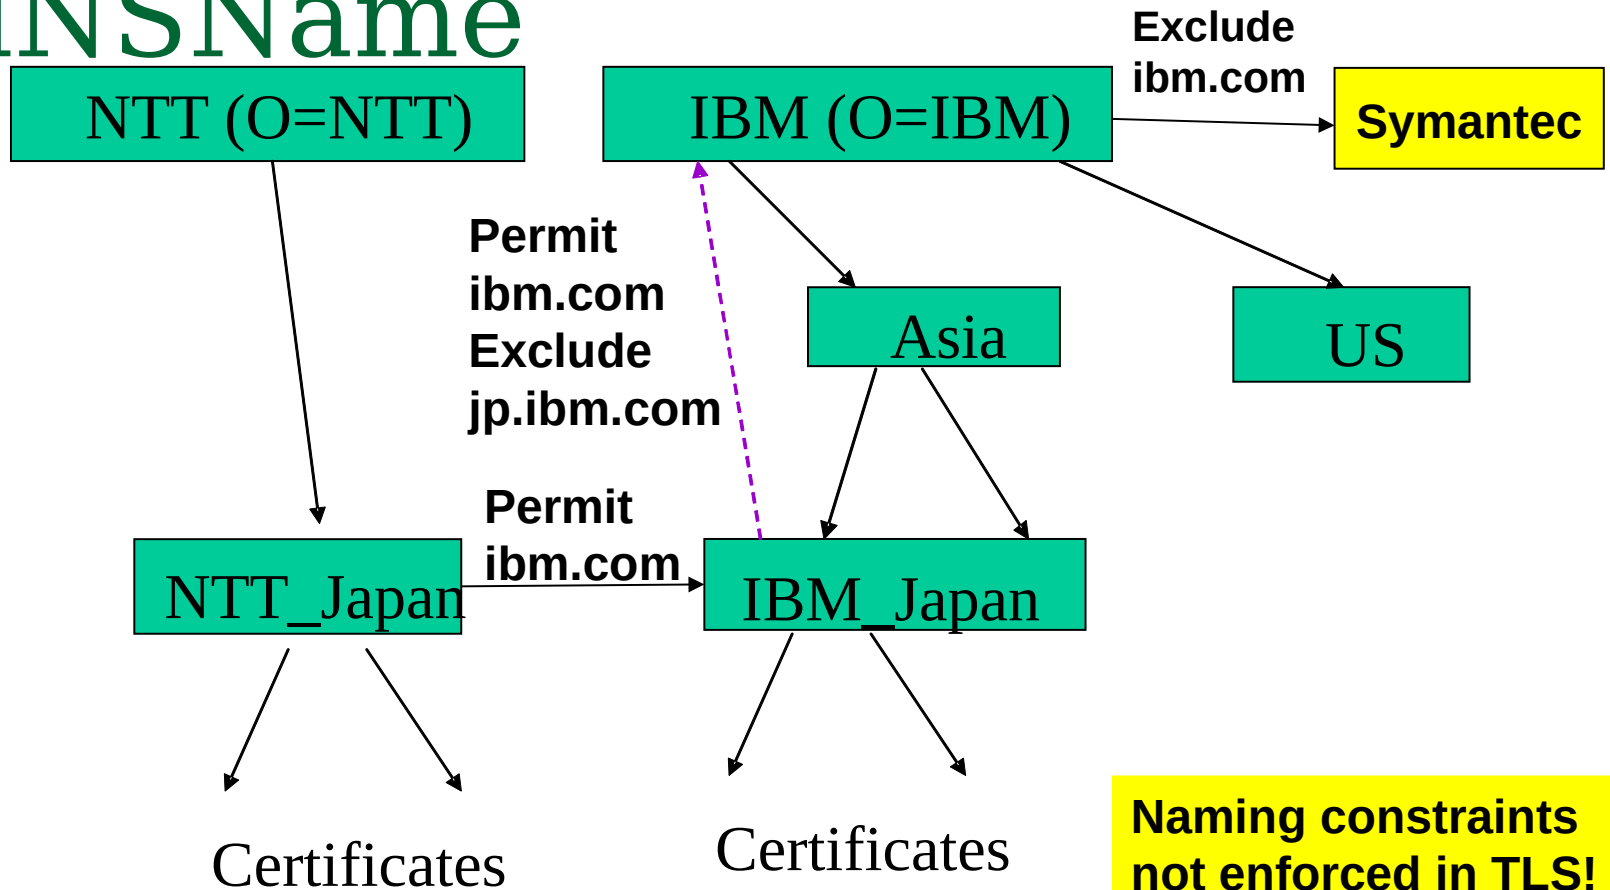
- Basic constraints:
 - Is the subject a CA? (default: FALSE)
 - Maximal depth of paths with this certificate?
 - pathLengthConstraint
- Policy constraints:
 - Interoperability of certificate-policy extensions
 - May allow/forbid 'policy mappings' (extension, too)
- Naming constraints
 - Constraints on DN and SubjectAltName
 - in certs **issued by subject**
 - Only relevant when subject is a CA !
 - 'Permit' and 'Exclude'

Naming constraints on DN



- **NTT JP permits (allows) IBM JP to certify IBMers**
- **IBM JP permits IBM to certify all IBMers, except of IBM JP**
- **IBM trusts Symantec's certificates, except for O=IBM**

Naming constraints on dNSName



PKI Lecture: Topics

- X.509 Certificates
- **Revoking certificates**
- Dealing with CA failures: CT (Certificate Transparency) and other defenses
- Conclusions and challenges...

Certificate Revocation

- Reasons for revoking certificates
 - ❑ Key compromise
 - ❑ CA compromise
 - ❑ Affiliation changed (changing DN or other attribute)
 - ❑ Superseded (replaced)
 - ❑ Cessation – not longer needed
- How to inform relying parties? Few options...
 - ❑ Do not inform – wait for end of (short?) validity period
 - ❑ Distribute ***Certificate Revocation List (CRL)***
 - ❑ Ask - **Online Certificate Status Protocol (OCSP)**
 - ❑ Skip details

X.509 CRL Format

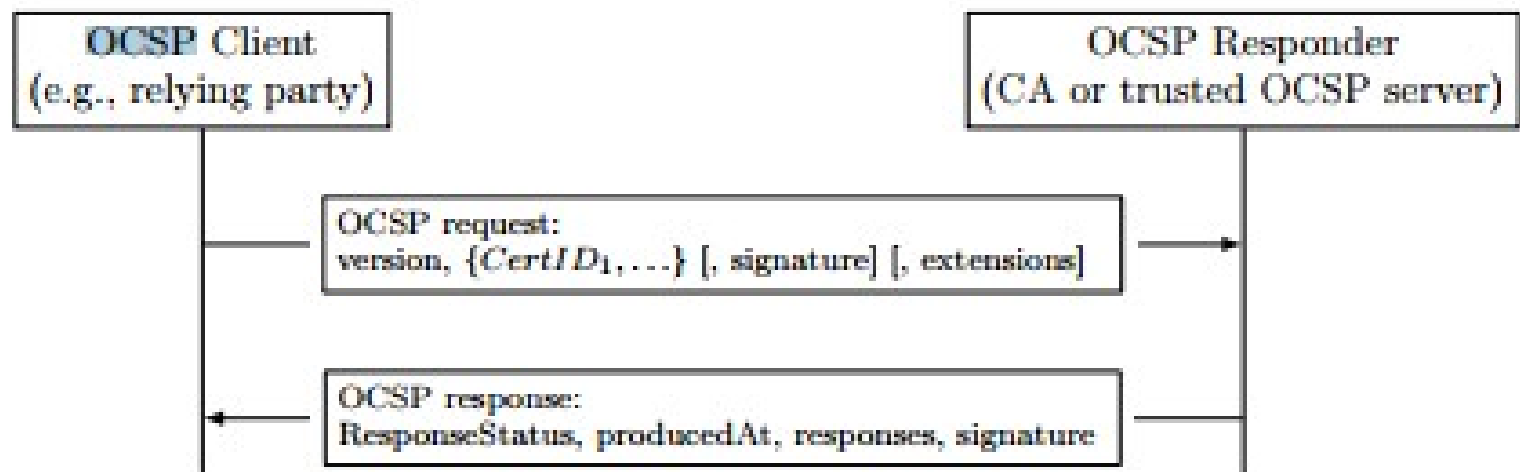
Signed fields	Version of CRL format			
	Signature Algorithm Object Identifier (OID)			
	CRL Issuer Distinguished Name (DN)			
	This update (date/time)			
	Next update (date/time) - optional			
	Subject (user) Distinguished Name (DN)			
	CRL Entry	Certificate Serial Number	Revocation Date	CRL entry extensions
	CRL Entry...	Serial...	Date...	extensions
			
	CRL Extensions			
Signature on the above fields				

Revocation with CRLs is

- If CRLs contain all revoked certificates (which did not expire)... it may be huge!
- CRLs are (also) not immediate
 - Who is responsible until CRL is distributed?
- Solutions:
 - More efficient CRL schemes
 - CRL distribution point – split certificates to several CRLs
 - Authorities Revocation List (ARL): list only revoked CAs
 - Delta CRL – only new revocations since last `base CRL`
- **Browsers mostly do not check CRLs. Instead:**
 - Online Certificate Status Protocol (OCSP)
 - Short validity for certificates

Online Certificate Status Protocol (OCSP)

- Improve efficiency, freshness cf. to CRLs
- Client asks CA about cert during handshake
- CA signs response (real-time)



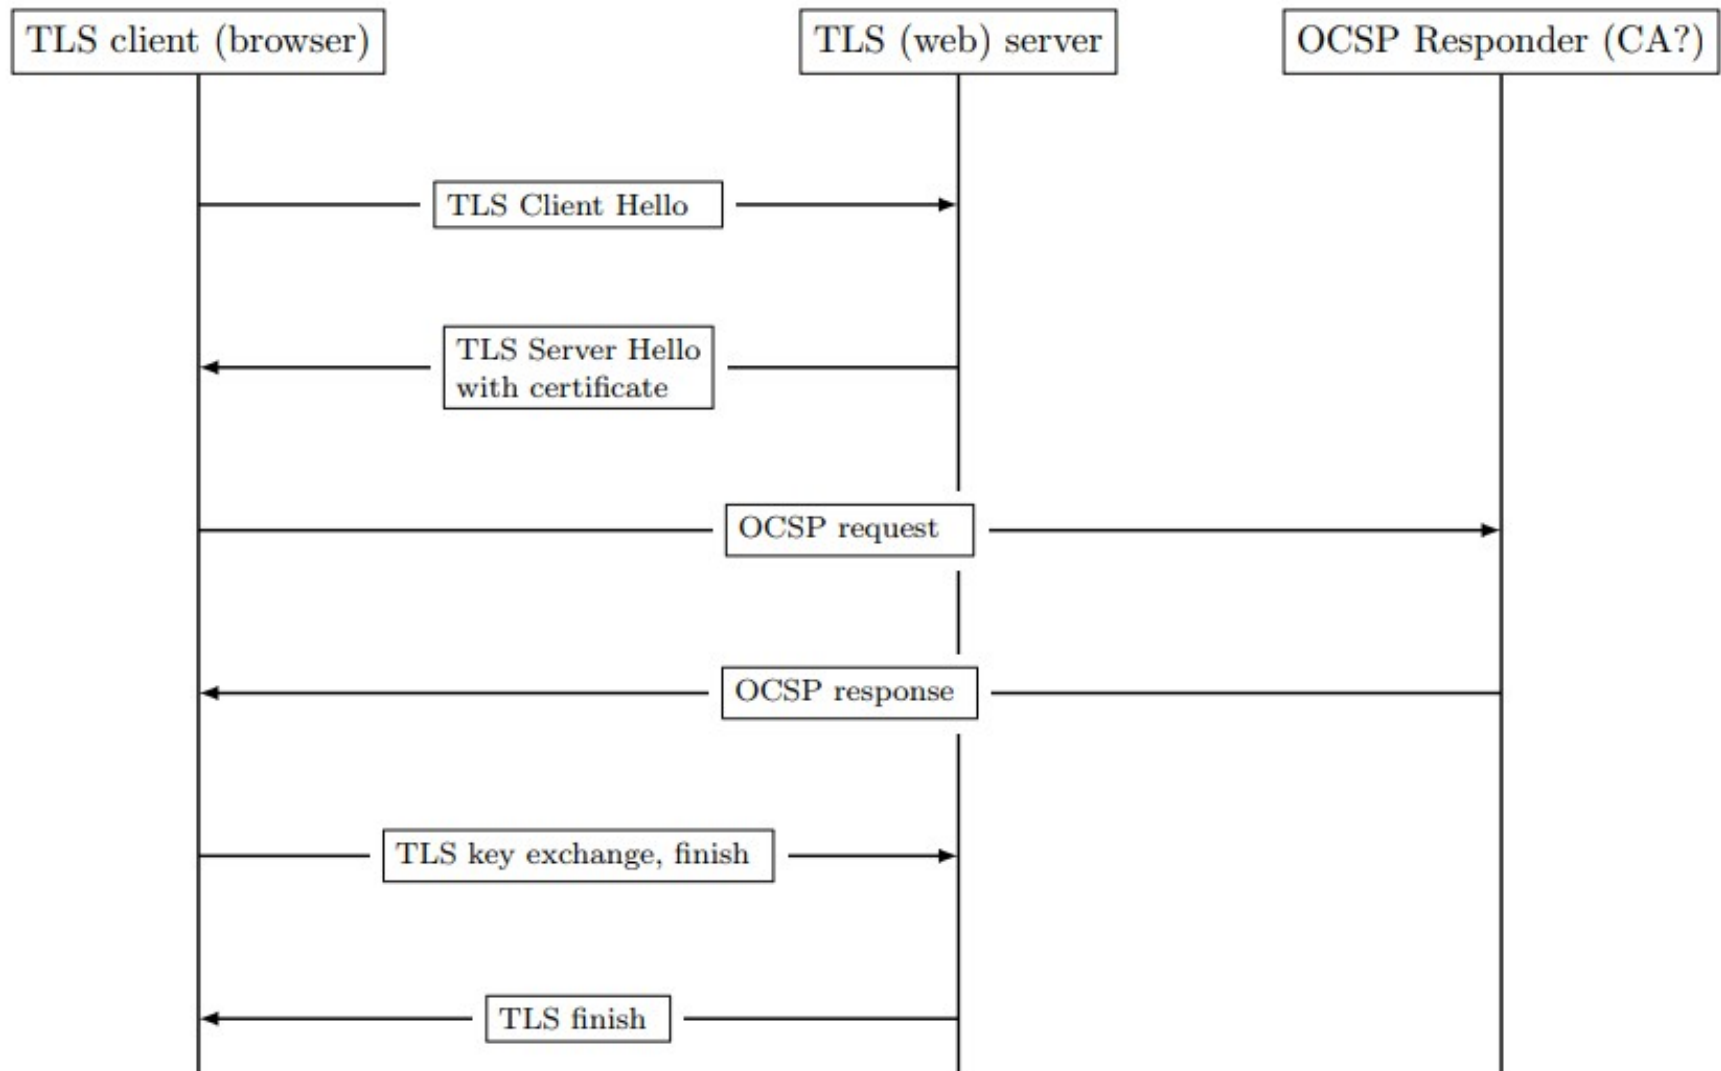


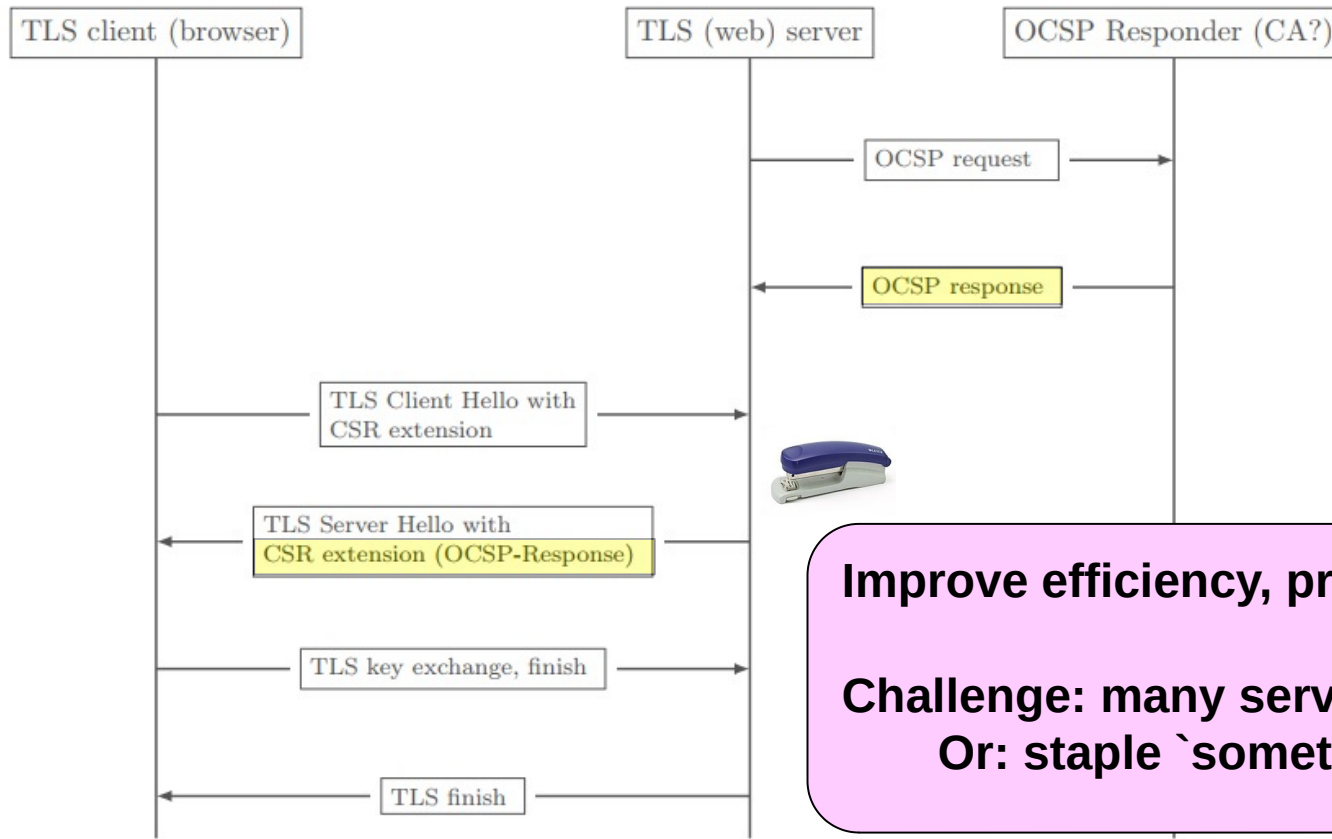
Figure 1: TLS handshake using OCSP (no stapling).

Online Certificate Status Protocol (OCSP)

- Client asks CA about cert during handshake
- CA signs response (real-time)
- Challenges:
 - Privacy: expose (domain, client) to CA
 - Load on CA
 - Delay (latency): on average, almost a second
 - Reliability: what if CA fails? No connectivity?
 - □ Most browsers skip OCSP or soft-fail: continue w/o OCSP response
- Better way to do OCSP?

Solution: OCSP-Stapling

Server runs OCSP, sends ('staples') the **CA-signed response (CSR)** during TLS handshake



Improve efficiency, privacy, reliability

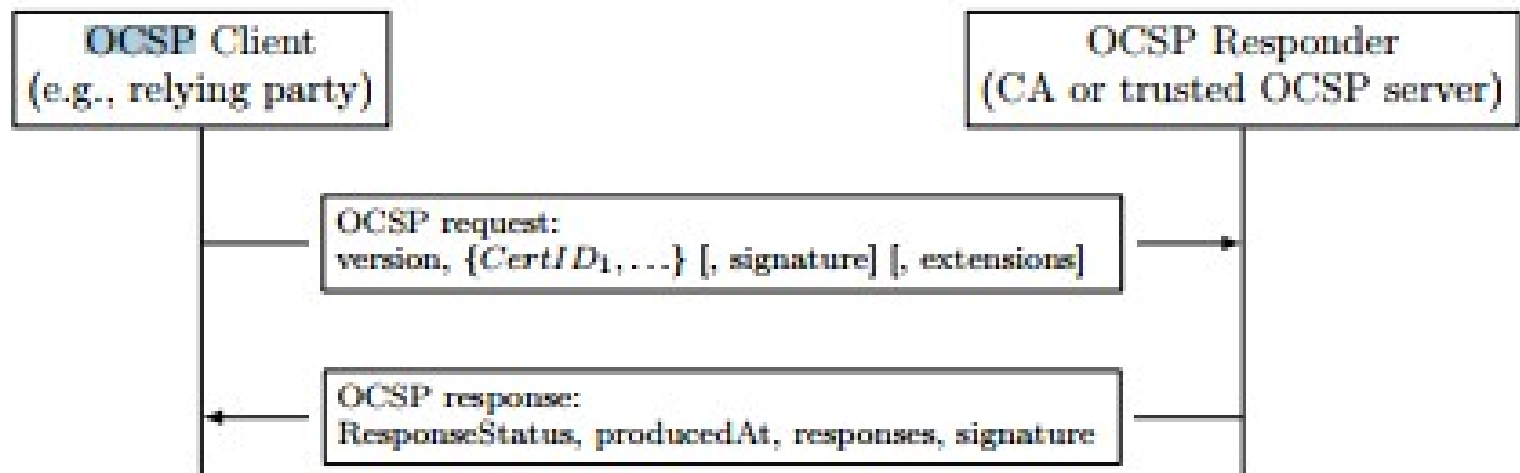
**Challenge: many servers don't staple!
Or: staple 'sometimes/usually'**

OCSP-Stapling

- Server runs OCSP, sends (`staples') the CA-signed response during TLS handshake
- Challenge: many servers don't staple!
 - Or, worse: staple `sometimes/usually'
 - So, try OCSP? Connect anyway? Disconnect?
- Solution: `Must-staple' cert. extension
 - RFC 7633
 - Mark as not critical
 - As it may not be supported by some browsers

Online Certificate Status Protocol (OCSP)

- Most browsers don't use CRLs. Why?
 - Efficiency, freshness concerns
- OCSP: check cert-status 'as needed'
- Signed responses (from trusted CA/server)



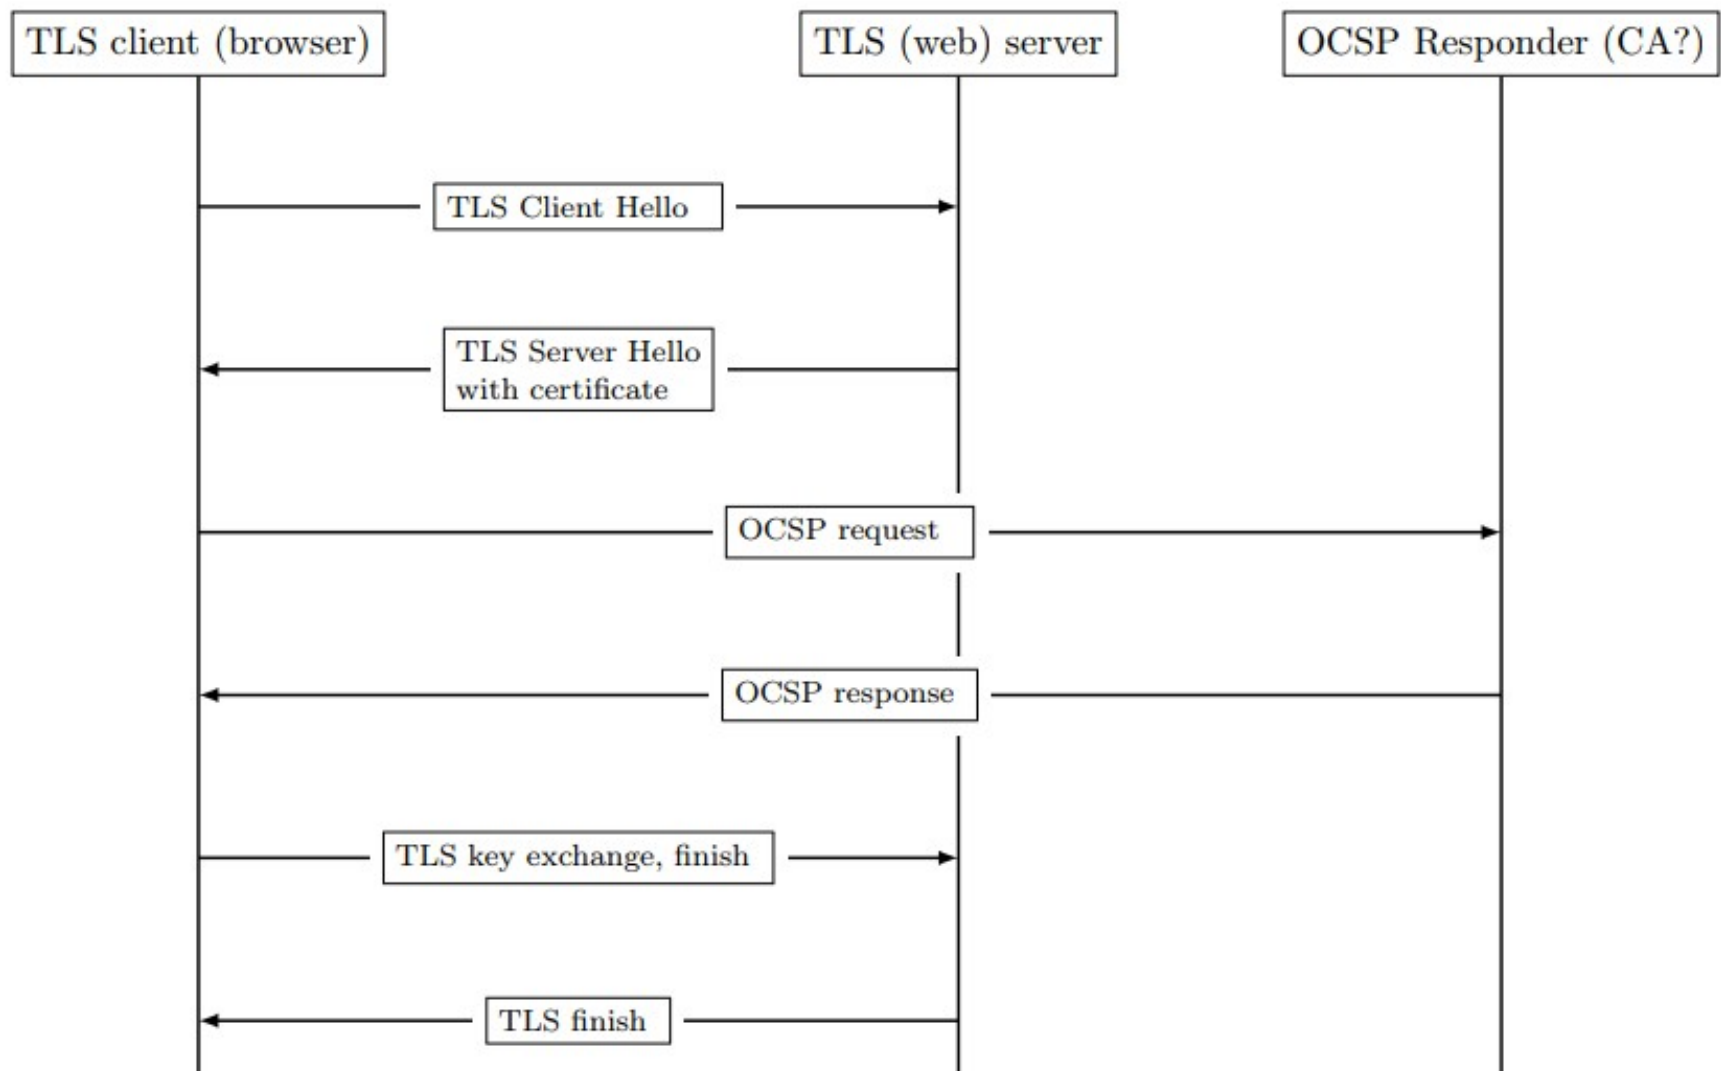
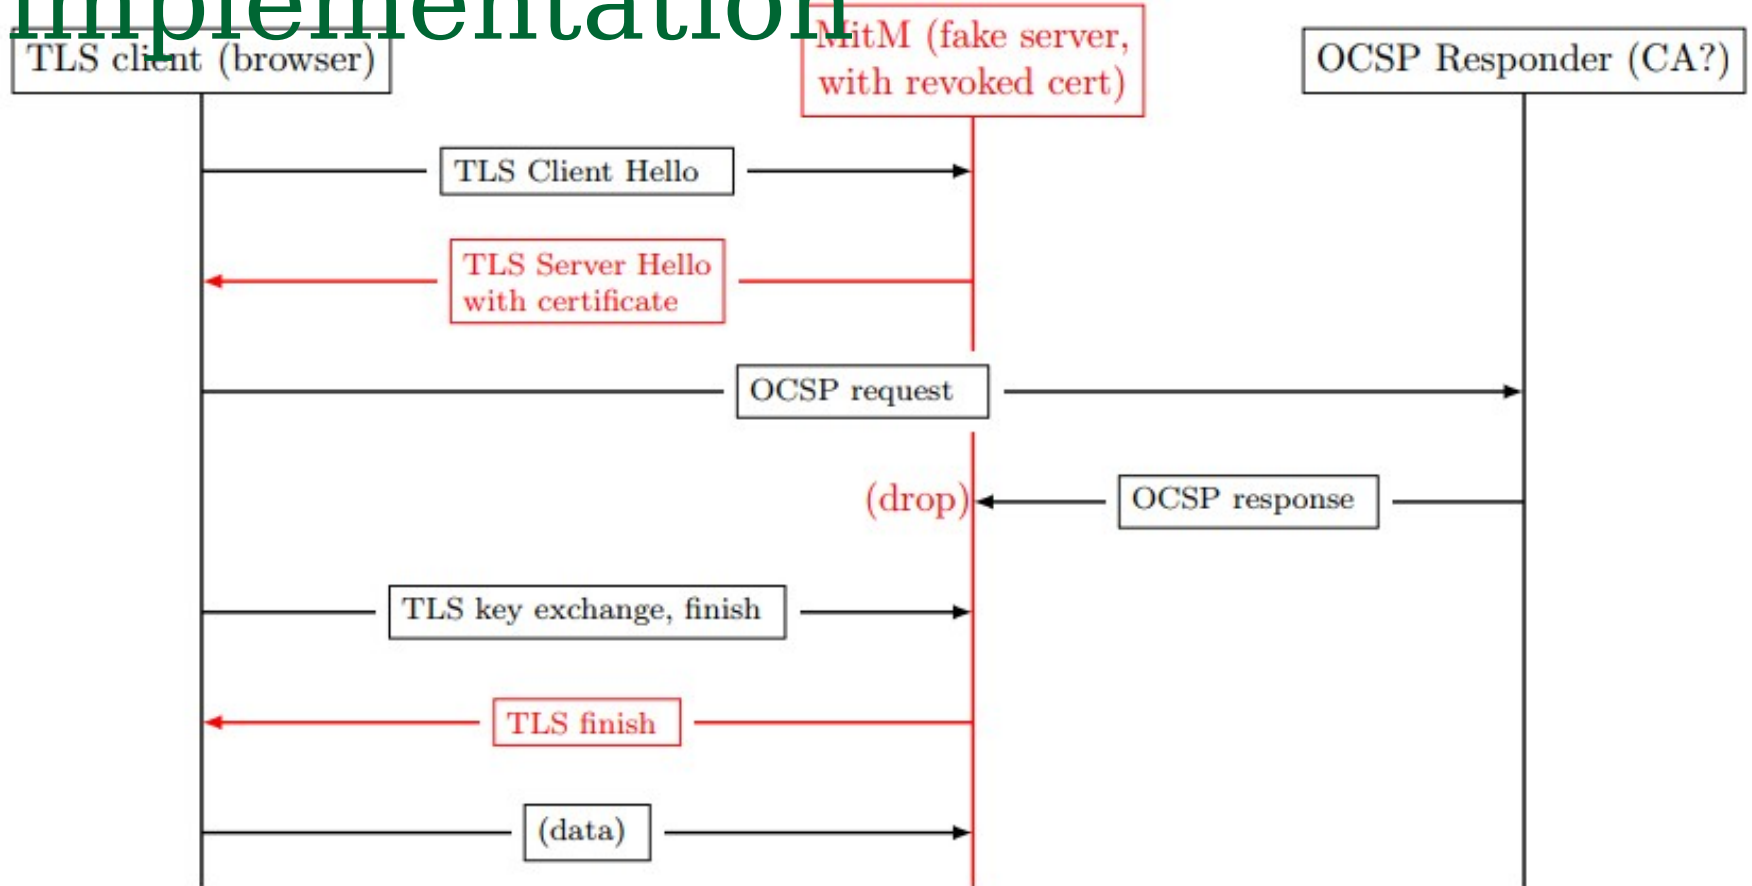


Figure 1: TLS handshake using OCSP (no stapling).

Online Certificate Status Protocol (OCSP)

- Client asks CA about cert during handshake
- CA signs response (real-time)
- Challenges:
 - Privacy: expose (domain, client) to CA
 - Load on CA
 - Delay (latency): on average, almost a second
 - Reliability: what if CA fails? No connectivity?
 - □ Most browsers skip OCSP or soft-fail: continue w/o OCSP response
- Improved mechanism: OCSP-stapling

Attacking typical OCSP implementation



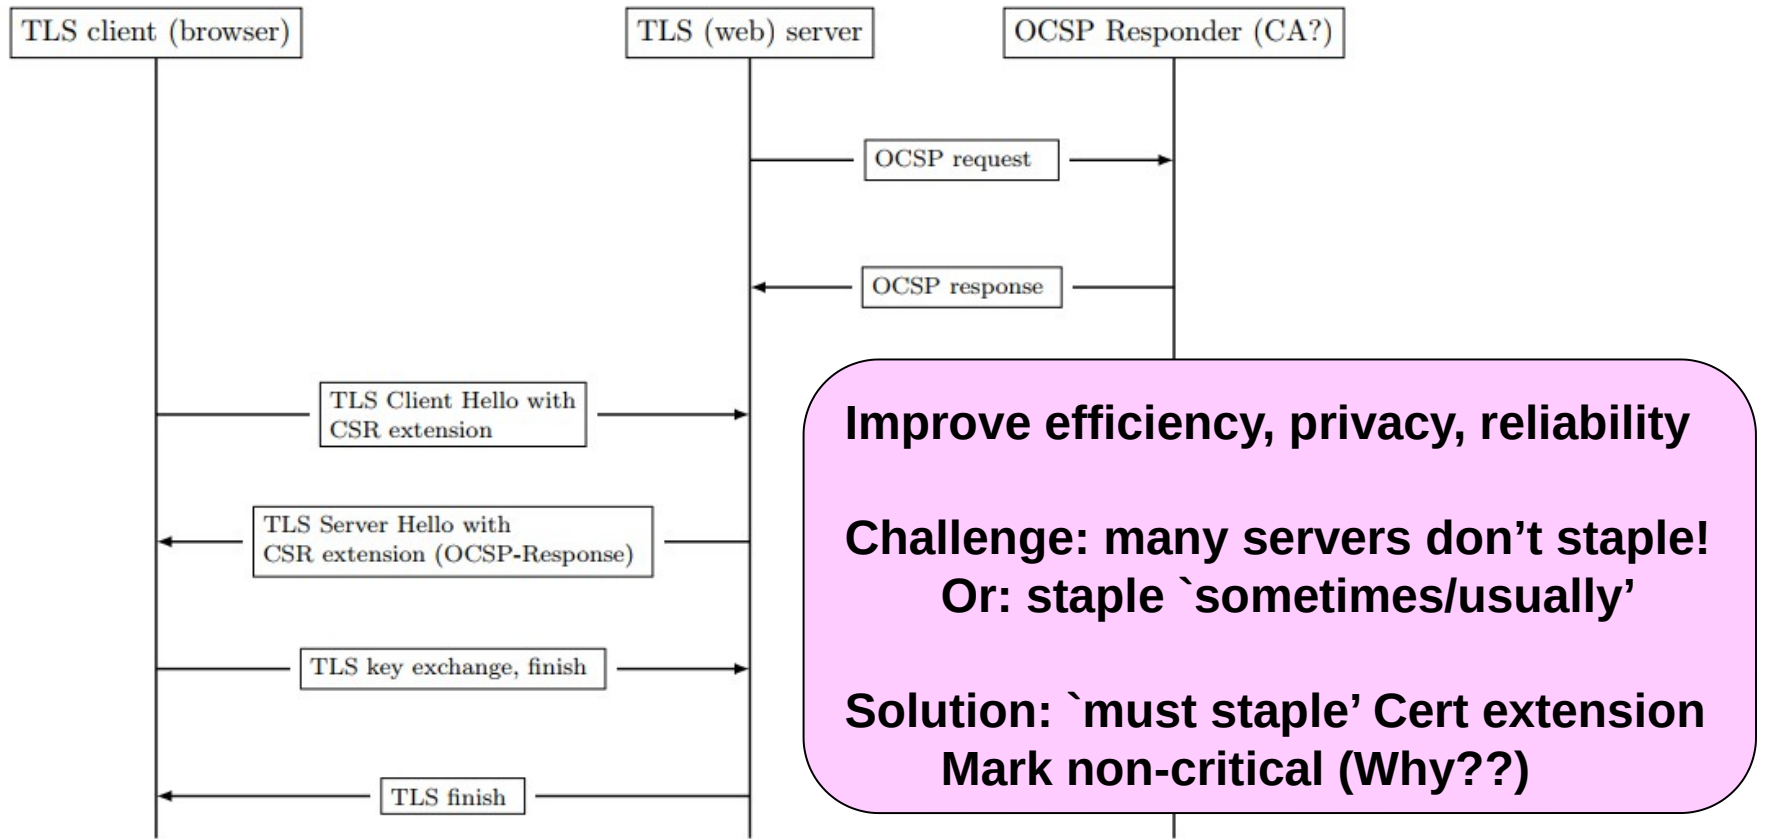
OCSP-Stapling

Server runs OCSP, sends (`staples') the CA-signed response during TLS handshake

Show sequence diagram for
TLS handshake with OCSP-
stapling

OCSP-Stapling

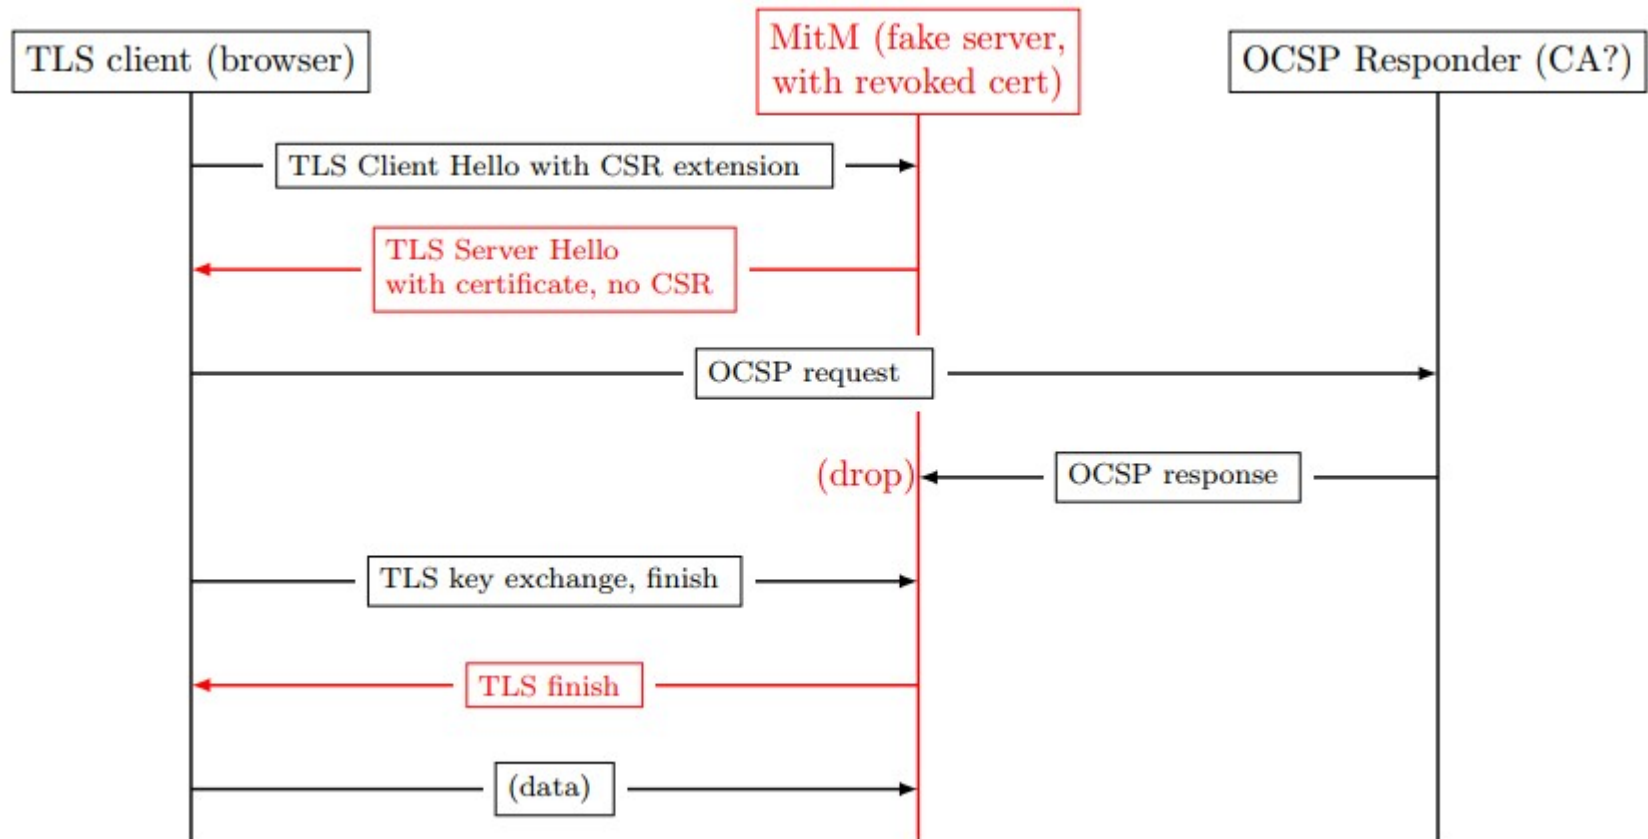
Server runs OCSP, sends (‘staples’) the CA-signed response during TLS handshake



Attacking... Stapled OCSP

- Stapled-OCSP is usually vulnerable!
 - Idea of vulnerability:
 - Most servers don't staple
 - And some staple... but not always
 - How can browser know if staple or not??
 - Present attack on (typical) browser deployment of OCSP-stapling
-

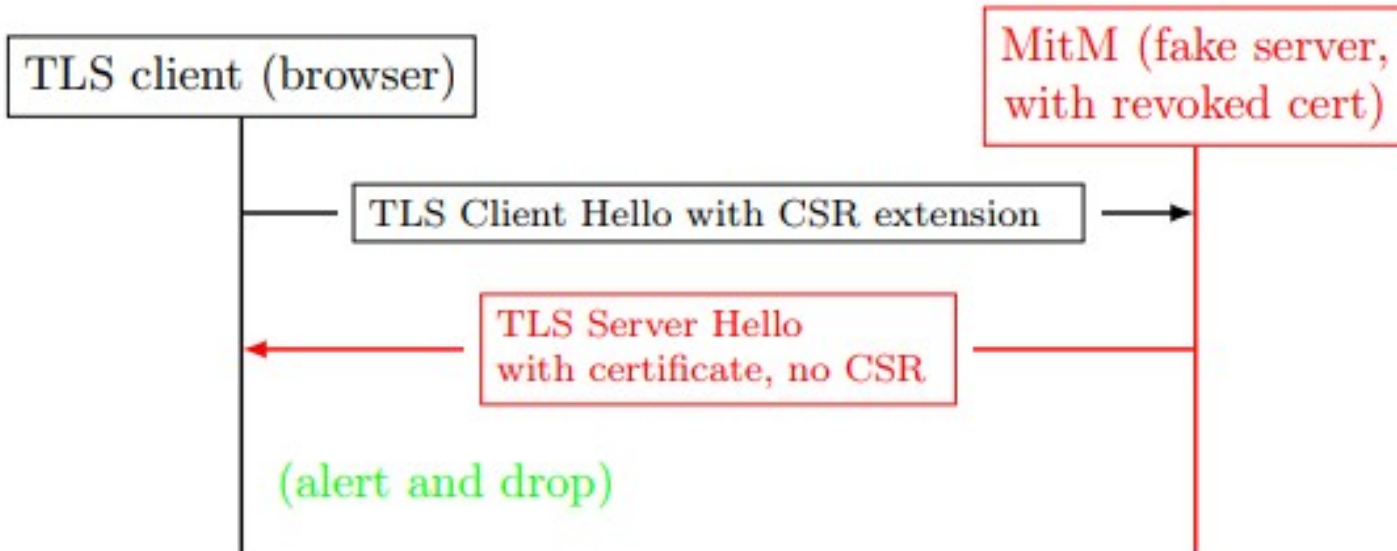
Attacking typical stapled-OCSP client



Must-Staple Certificate Extension

- Solution to attack on stapled-OCSP clients: site signals it **supports stapled OCSP**
- **How?** Add 'must staple' x.509 extension
 - A server sending this cert must staple!
- Note: extensions should be non-critical!
 - Else, a non-supporting client cannot use cert.
- How does this prevent attack?

OCSP with Must-Staple Extension



PKI Lecture: Topics

- X.509 Certificates
- Revoking certificates
- **Dealing with CA failures:
CT (Certificate Transparency) and defenses**
- Conclusions and challenges...

Why and How CAs fail?

- Many CAs `trusted' in browsers (as root)
- 'Domain-Validated' certificates
- Several well-known failures
 - DigiNotar, Comodo, Stuxnet, ...
- Every CA can certify any domain (name)
 - Naming constraints NOT applied (esp. to roots)
 - Equivocation: two certificates for same name (from same or different CAs)
 - To detect bad-CA: must find bad-certificate
 - Need: public, auditable log of certificates

Defenses against CA failures

- **Use naming constraints to limit risk**
 - who can issue global TLDs (.com, etc.)?
- **‘Burned-in’ public keys (e.g., for Google)**
 - Detected MitM in Iran, using DigiNotar CA
- **Certificate / public-key pinning (HPKP)**
 - Server: I always use this PK / Cert / Chain
 - Client: remember and implement!
- **Pre-load lists of (‘important’) revocations**
 - Chrome (‘CRLset’), FF (‘OneCRL’)
- **Certificate Transparency (CT): Accountability**

Certificate Transparency (CT)

■ X.509, PKIX: CAs sign certificate

[RFC6962]

- Given **rogue certificates**, we can know which CA issued it
 - A **rogue certificate**: same or too-similar name to that of some victim
 - Used to impersonate website, site malware, etc.
- Then: ask CA to revoke it, or **stop trusting non-trustworthy CA**
- All this – only **after** rogue certificate is detected !

■ CT: early detection of rogue certificates

- Soon after issued – possibly even before use!
- Also: detect **at what time** the rogue cert was issued
 - Helps forensics, prevents rogue-CA from ‘backdating’ a certificate
- Proposed by Google, after Chrome detected rogue Google cert

■ Main idea and service: Certificate loggers

- Maintain public log of all certificates issued (by subscribing CAs)
- To begin, let’s assume **trustworthy logger**

Simplified CT: Assume

Trusted Logger

- Goal: early detection of rogue certs

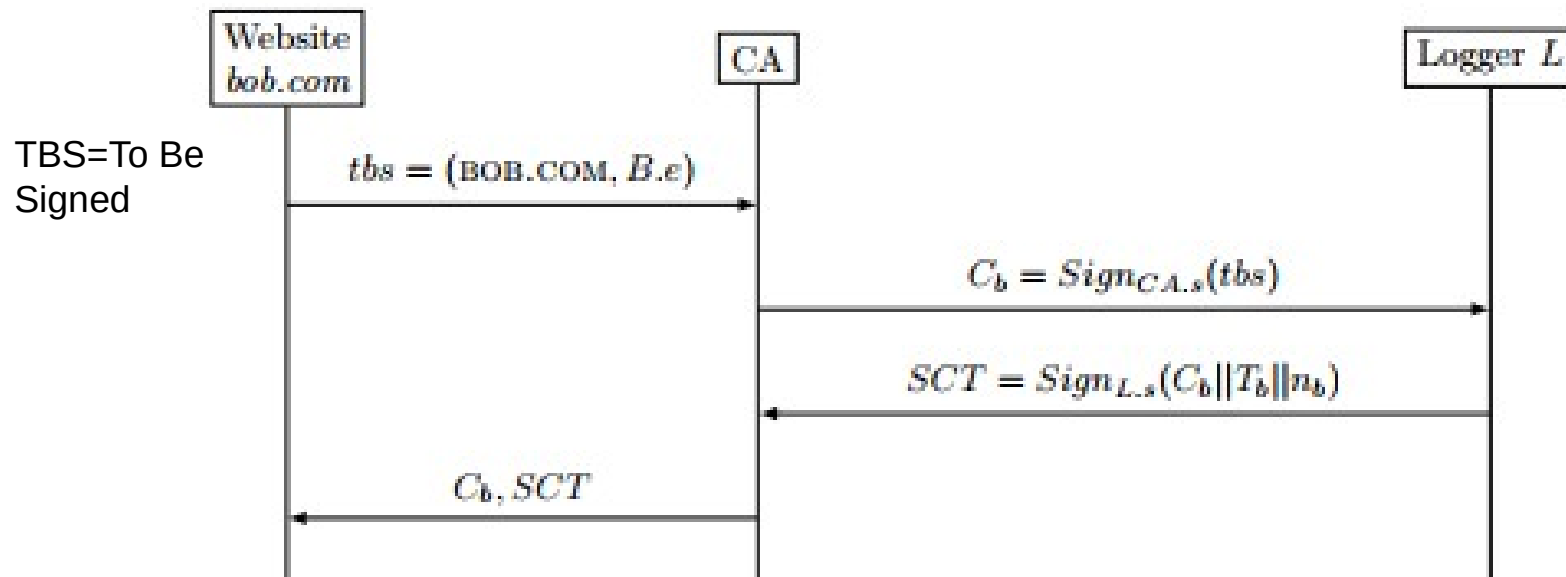
- Rogue cert: same or 'similar' to 'owned' name

■ Simplified CT solution:


- Issuer (CA) must send every cert to logger
- Logger returns Signed Certificate Timestamp (SCT)
 - Validate that the cert was logged at given time
- CA gives cert, SCT to subject
- Subject sends SCT (with cert) to relying party
- Relying party 'knows' cert was logged (and when)

Issuing a CT Certificate

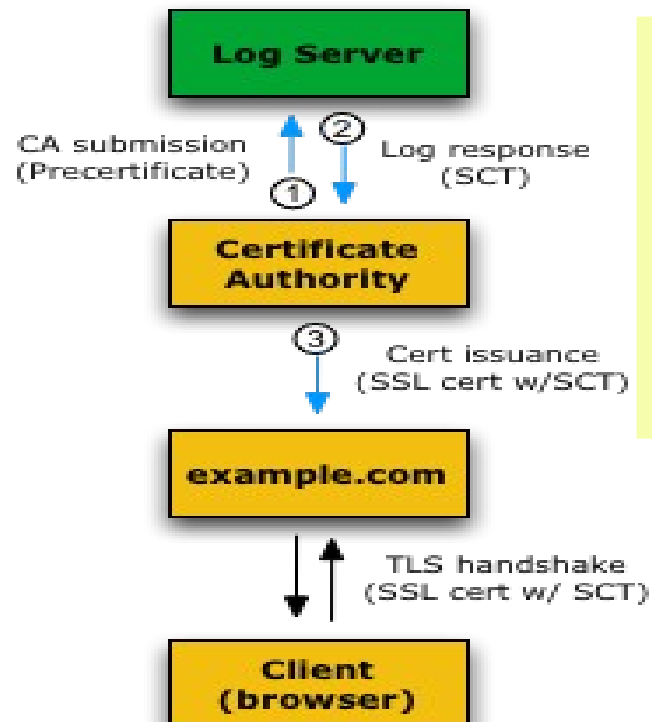
- Subject, e.g. website, send request
 - Request contains 'To Be Signed' fields: name, public-key
- CA validates request, issues 'regular' cert
- Logger adds to log, signs and returns (signed) SCT
- CA sends cert + SCT to subject (e.g., website)



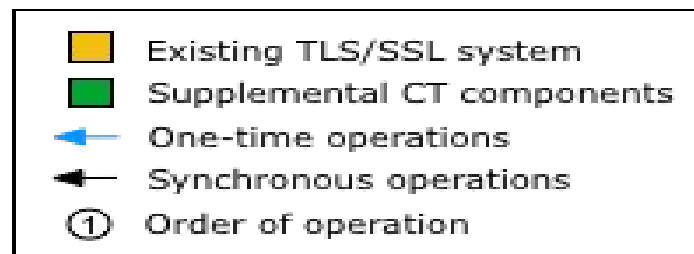
Current TLS/SSL System



TLS/SSL System with Certificate Transparency (X.509v3 Extension)



SCT: Signed Certificate Timestamp
(time that the certificate was added to log, serial number)



Simplified CT: Assume

Trusted Logger

- Goal: early detection of rogue certs

- Rogue cert: same or 'similar' to 'owned' name

- Simplified CT solution:

- Issuer (CA) must send every cert to logger
- Logger returns Signed Certificate Timestamp (SCT)
 - Validate that the cert was logged at given time
- CA gives cert, SCT to subject
- Subject sends SCT (with cert) to relying party
- Relying party 'knows' cert was logged (and when)

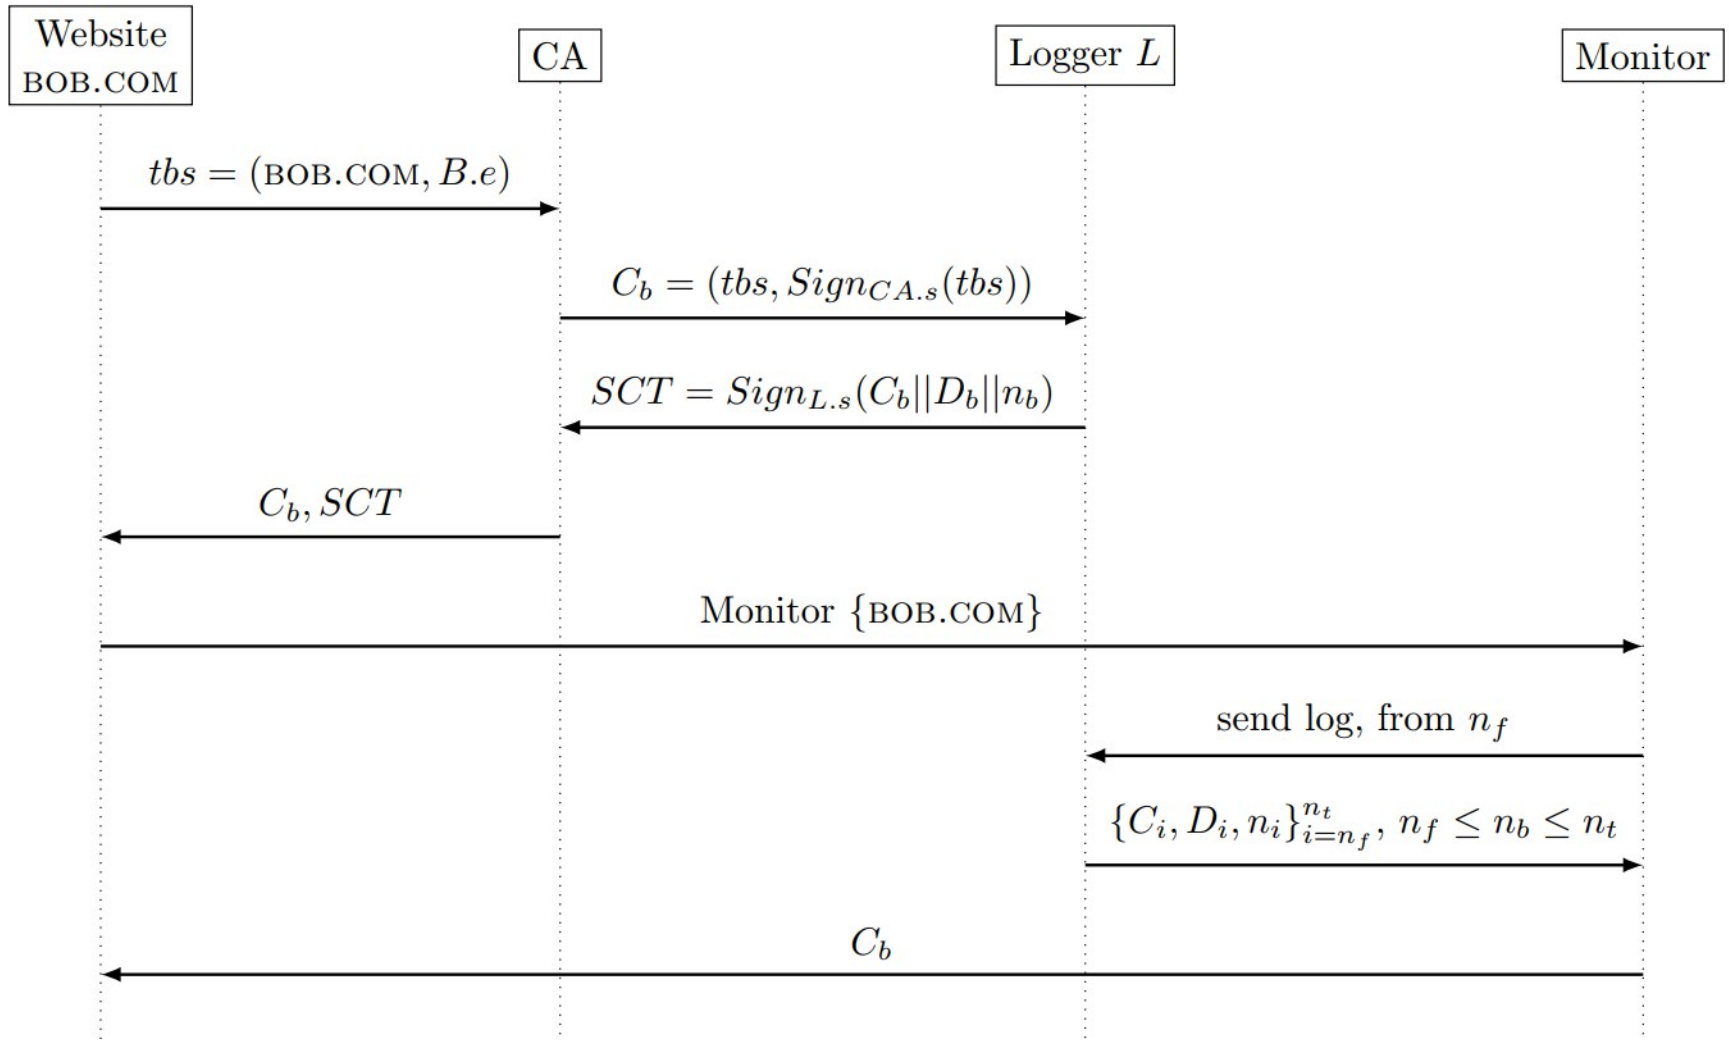
- How do we use logs to detect rogue certs?

Detecting rogue certs in log:


■ Goal: early detection of rogue certs in log

- Logs should be publicly available
- Name-owners can monitor the log
 - Download, check log for relevant names
 - □ high overhead!
- Instead: **monitors** do this (for many names)
 - Several such monitors, loggers already operate
 - Download only new certificates
 - And: ask log for seq# and/or date of last logged cert
 - Ask log to send range of certs: <from-to>
 - Optionally: maintain all certs (to check new names)

Monitoring a Log (trusted)



Detecting bad certificate (honest log)



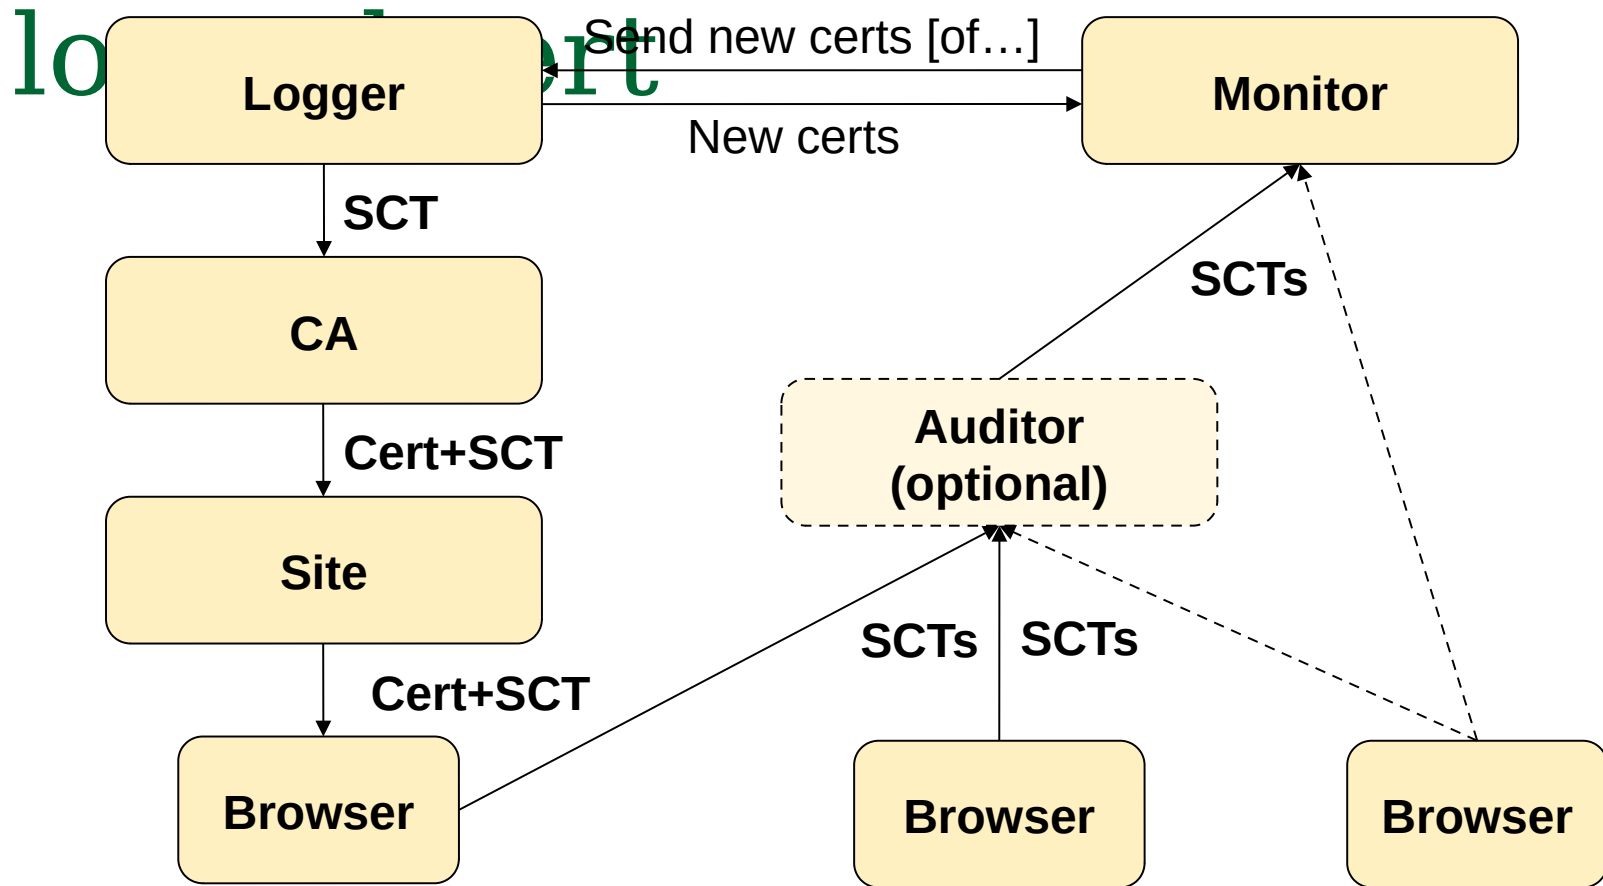
CT: Detecting Rogue

Loggers

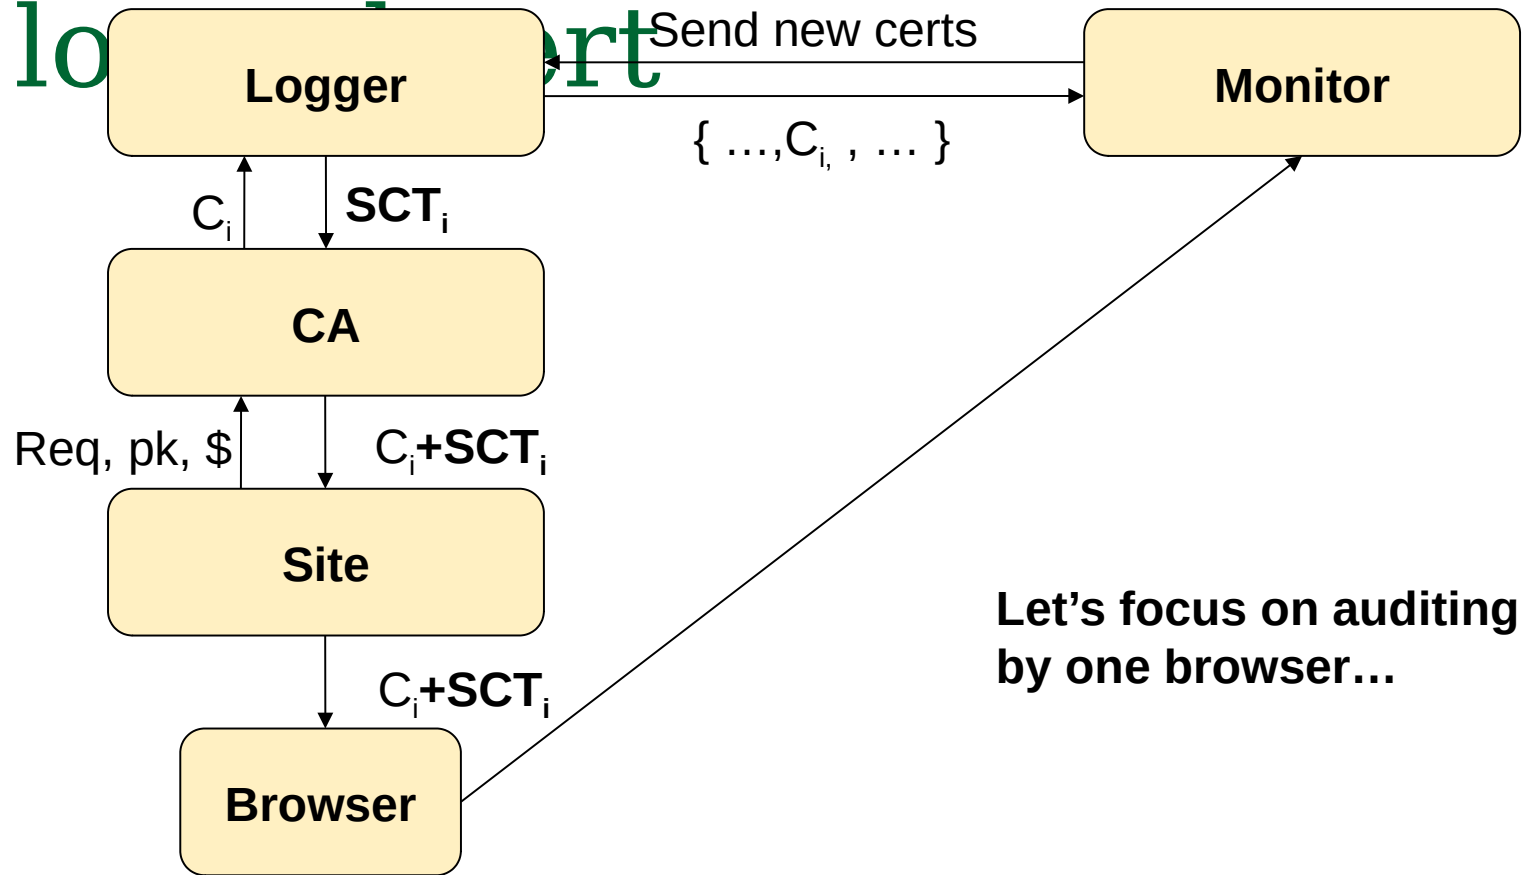
■ Can we detect a rogue logger??

- No; instead, send SCTs signed by few loggers
 - Overhead
 - How many and which loggers?
 - Current approach in Google's Chrome
 - Requires two SCTs, one of them by Google's logger
 - 'In Google we Trust' ?
- **Yes: Audit to detect rogue loggers**
 - Relying-parties (clients) validate SCT is 'auditable'
 - Using 'stapled' Proof-of-Inclusion
 - And [Randomly? Always?] audit the SCT
 - Using (a trusted) monitor

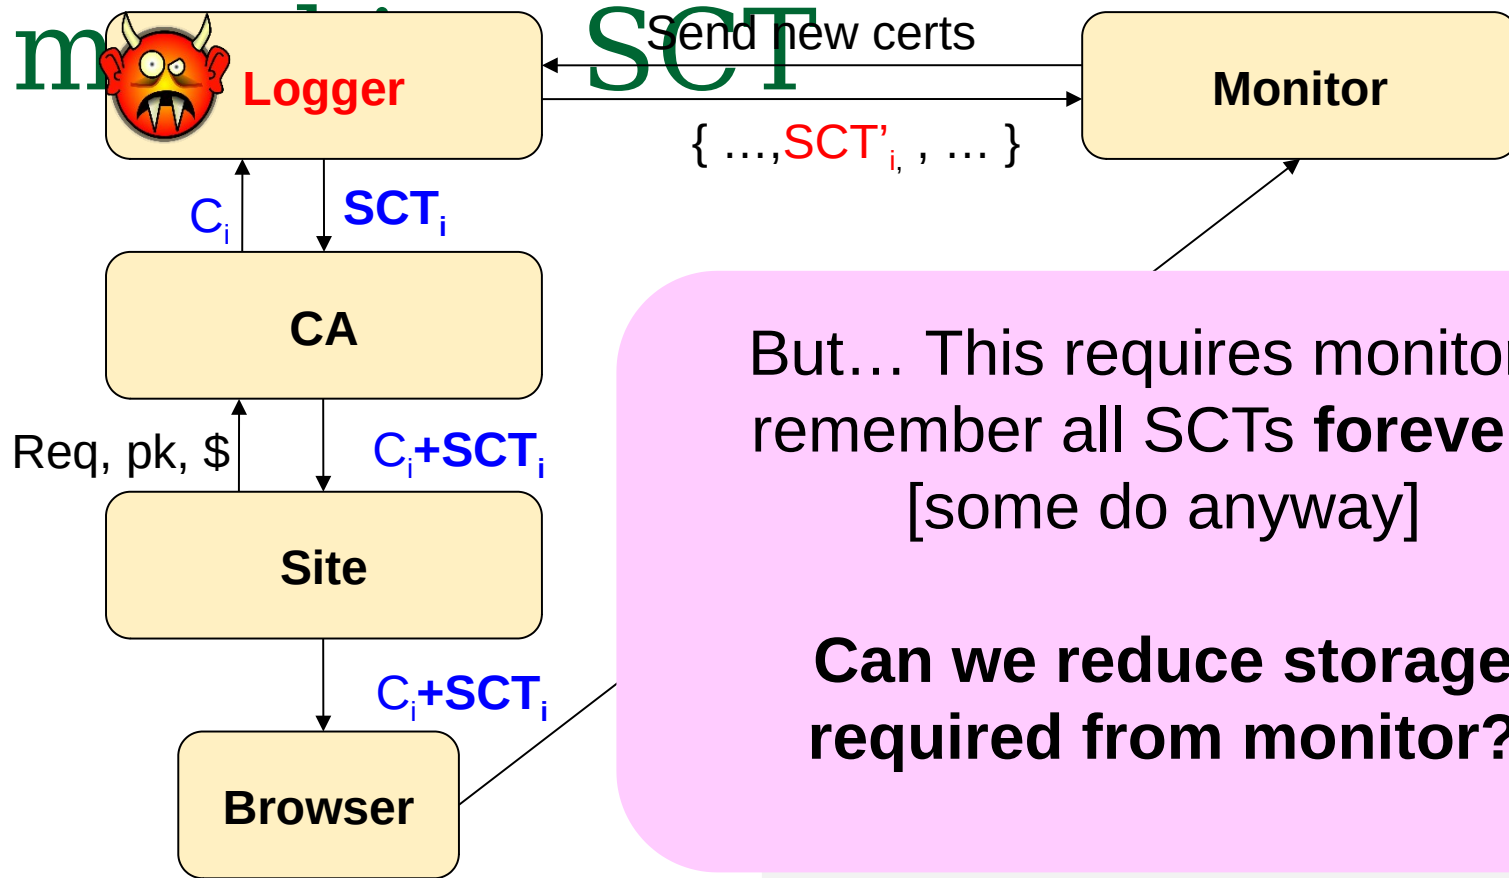
CT: Audit to detect non-



CT: Audit to detect non-logged cert



Detect and 'prove' non-

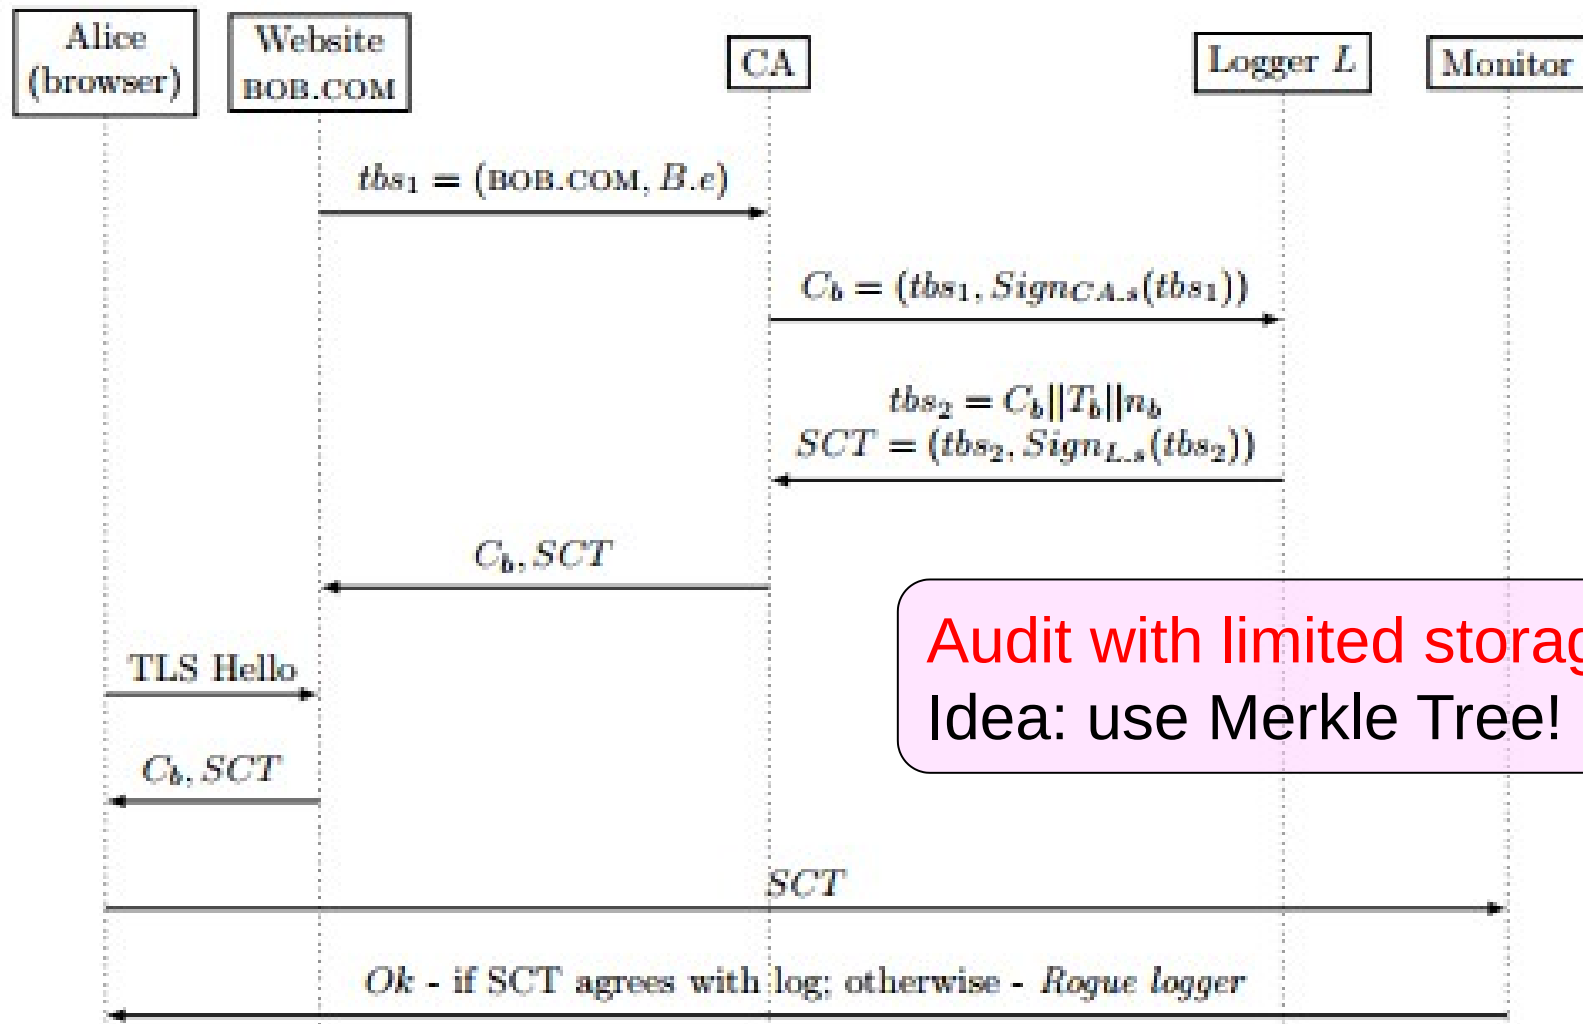


$\text{SCT}'_i = \text{Sign}_{\text{Log.s}}(i, \text{time}, C'_i)$

$\text{SCT}_i = \text{Sign}_{\text{Log.s}}(i, \text{time}, C_i)$


Logger signed both SCT_i and SCT'_i , 'proving' its 'guilt'

Auditing, simplified (monitor




Limiting Monitor Storage: Merkle Tree

- Logs use Merkle hash Tree, send **Signed Tree Hash (STH)**
- Every node is hash of its children nodes.
- d0, d1, d2, d3 – certificates.

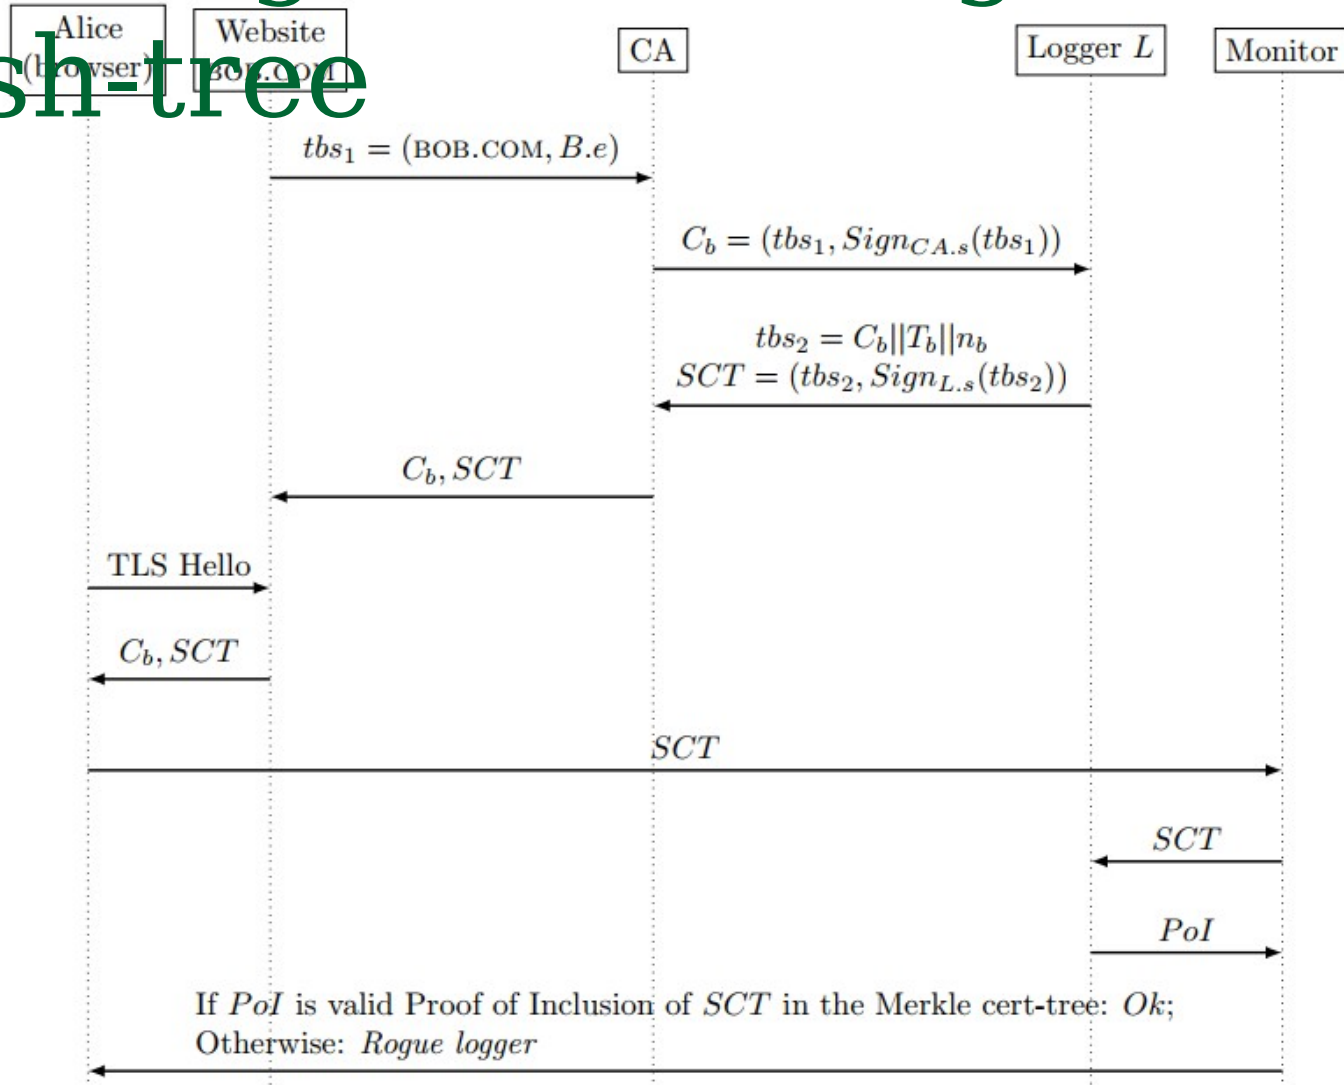


Limiting Monitor Storage: Merkle Tree

- Logs use Merkle hash Tree, send **Signed Tree Hash (STH)**
- Monitor has STH (from logger)
- Monitor can ask for **Proof of Inclusion (PoI)**
- **PoI for d1: {a,j}**



Auditing in CT, using Merkle hash-tree

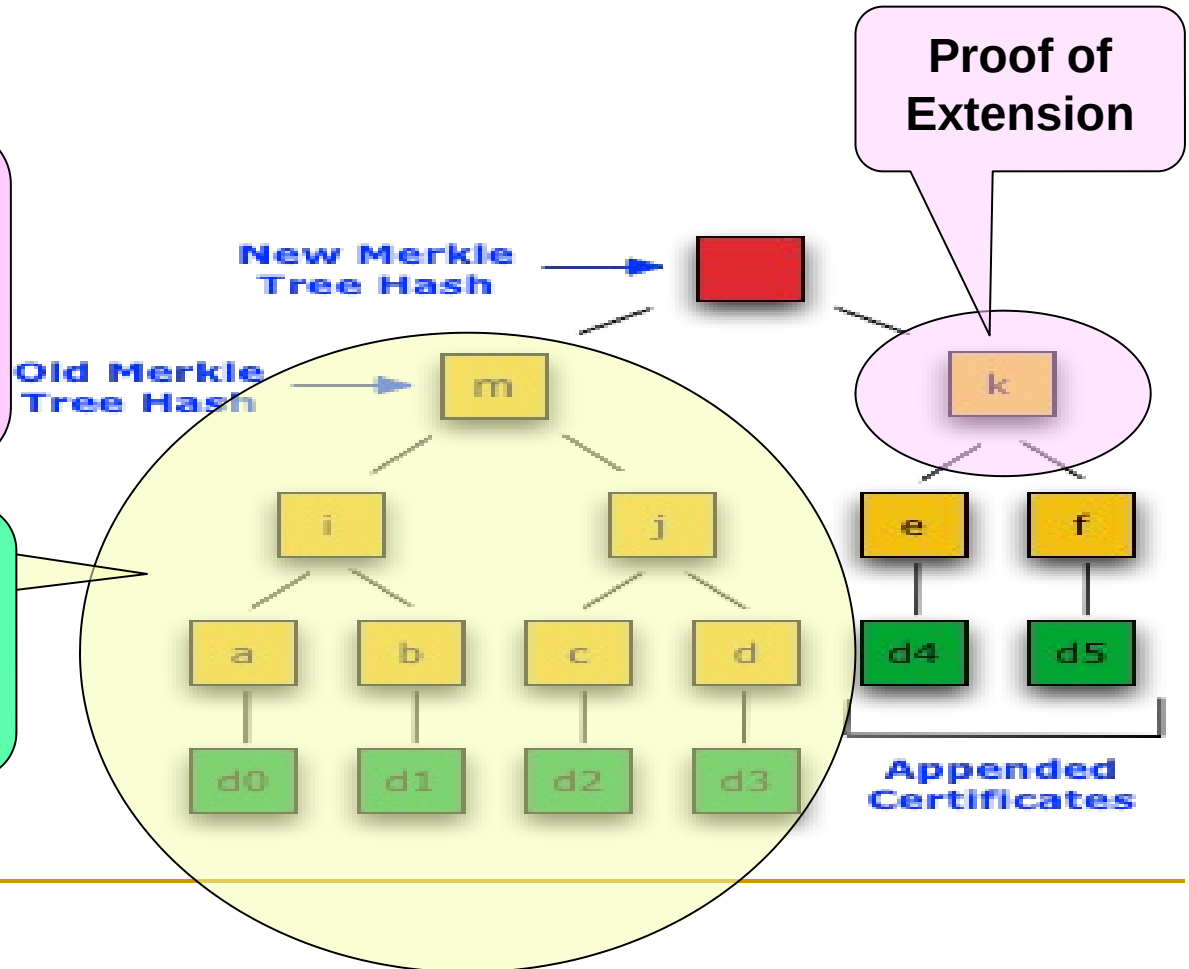


Monitors may store only tree head!

- Tree-head signed by logger
- Logger sends new tree-head, with **proof of extension** of old tree
- Append only

Risk: rogue logger sends different tree to different monitors
Must we check all monitors?

No; monitors gossip □ detect such rogue logger (details omitted... even from spec!)



So what are the benefits of CT?

■ Benefits for websites:

- Detect rogue certs, for domain
 - Or for `misleading' domain, often abused for phishing
- Once detected, owners can mitigate risk
 - Demand revocation, removal of CA from browsers,...

■ Benefit to users: less likely to fall victim...

■ Benefit to trustworthy CAs: reduced competition from shady CAs

■ Overall: more secure PKI !

■ Concerns:

- Not yet fully detecting rogue CAs [e.g., gossip not yet fixed]
- Revocation is not transparent
- Audit: concerns about overhead and privacy

PKI : Conclusions and Challenges

- ❑ Many PKI challenges remain...
- ❑ Revocation is still quite broken
 - Although OCSP-stapling is great step forward
 - E.g., incorrect 'status Ok' OCSP response from rogue CA
- ❑ Privacy: OCSP, CT-auditing exposes visited sites!
 - Must-Staple may be a solution (for OCSP; and CT too??)
 - Many other privacy concerns...
- ❑ Define, analyze PKI properties
- ❑ Client certificates...
- ❑ □ PKI still very active area of research