

Homework 2

- 1.) Consider a disk with block size $B = 512$ bytes. A block pointer $p = 6$ bytes long, and a record pointer is $P_R = 7$ bytes long. A file has $r = 100,000$ EMPLOYEE records of fixed length. Each record has the following fields: NAME(25 bytes), SSN(9 bytes), DEPARTMENTCODE(9 bytes), ADDRESS(40 bytes), PHONE(9 bytes), BIRTHDATE(8 bytes), SEX(1 byte), JOBCODE(5 bytes), and SALARY(4 bytes).

- a.) Calculate the record size R in bytes.

The record size R in bytes would be the size of all the things that make up R . This means that $25 + 9 + 9 + 40 + 9 + 8 + 1 + 5 + 4 = 110$ bytes.

- b.) Calculate the blocking factor bfr and the number of file blocks b , assuming an unspanned organization.

The blocking factor bfr is the number of records per block which means that with a block size of 512 bytes and a record size of 110 bytes there are 4.655 records per block. Since this disk is unspanned, and the not an integer then the bfr should be 4.

- c.) Suppose that the file is order by the key field SSN and we want to construct a primary index on SSN. Calculate

- i.) The index block factor $bfri$

The field size of SSN is 9 bytes, the record pointer is 7 bytes, 512 divided by 16 = 32 entries per block.

- ii.) The number of first-level index entries and the number of first-level index blocks

Since the first-level index entries in a primary index, number of file blocks = $100,000$ records / 4 bfr = 25,000 blocks. The number of first-level index blocks = $25,000$ blocks / 32 entries per block = 781 blocks.

- iii.) The number of levels needed if we make it into a multilevel index

Number of second level entries = number of first level blocks, number of second level blocks = $781 / 32 = 24$. Which makes 24 third level entries hence we would have 3 levels because $24/32$ is not a whole number and thus = 1.

- iv.) The total number of blocks required by the multilevel index

$781 + 24 + 1 = 806$

- v.) The number of block accesses needed to search for a retrieve a record form the file--- given its SSN value---using primary index.

4 block accesses because there are three levels.

- 2.) Given the same specifications of Problem 1, consider this time you are building a primary index on SSN using B-tree. Calculate

- (i) the order p for the B-tree

Given:

Block Size = 512 bytes

Block Pointer = 6 bytes

Record Pointer = 7 bytes

Key Field of SSN = 9 bytes

We are using the block pointer because it is not dense.

Since a B-tree can have up to p block pointers, $p - 1$ record pointers and $p - 1$ search field values that all must fit into one block:

$$(p-1) * (15) + 6p \leq 512$$

$$15p - 15 + 6p \leq 512$$

$$21p - 15 \leq 512$$

$$P = 25$$

So, the order of b-tree is 25.

- (ii) *The number of levels needed if blocks are approximately 69% full (round up for convenience)*

New order P will be $0.69 * p = 0.69 * 25 = 18$ rounded up from 17.25.

Root: 1 node, 17 entries, 18 tree pointers

L1: 18 nodes, 18 * 17 entries, 18 * 18 tree pointers

L2: 18 * 18 nodes, 18 * 18 * 17 entries, 18 * 18 * 18 tree pointers

L3: 18 * 18 * 18 nodes, 18 * 18 * 18 * 17 entries = 99,144 entries

Number of levels = 4 ($\log_{18}(25,000) = 3.5 = 4$)

- (iii) *The worst-case number of blocks needed to search for and retrieve a record from the file—given its SSN value—using the B-tree you are estimating.*

Worst case = # of levels + 1 so 5 block searches worst case.

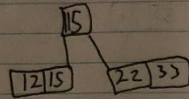
- 3.) A PARTS file with Part# as key field includes records with the following Part# values: 12, 33, 22, 15, 19, 55, 25, 11, 16, 13, 29, 44, 71, 21, 4, 42, 68, 78. Suppose that the search field values are inserted in the given order in a B+-tree of order $p = 4$ and pleaf = 3.

- (i) *Show how the tree will expand (show all steps as in Fig 18.12 (6thed)) and what the final tree will look like.*

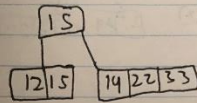
Insert first 3

12 3 3 2 2 6

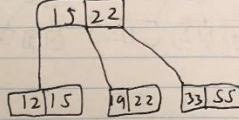
insert 15:



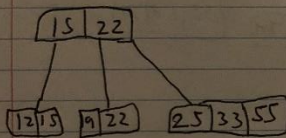
insert 19:



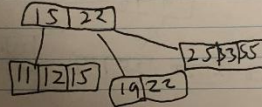
insert 55:



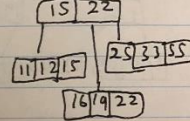
insert 25:



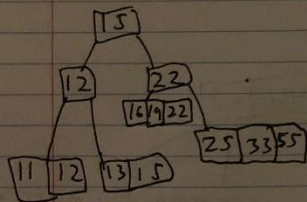
insert 11:



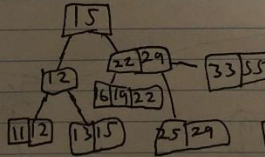
insert 16:



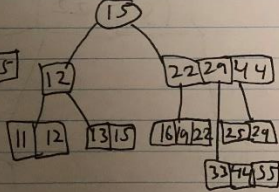
insert 13:



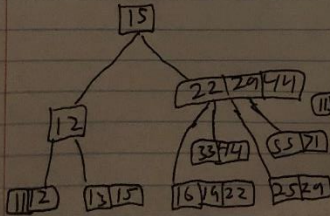
insert 29:



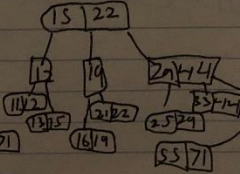
insert 44:



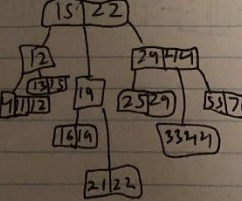
insert 71:

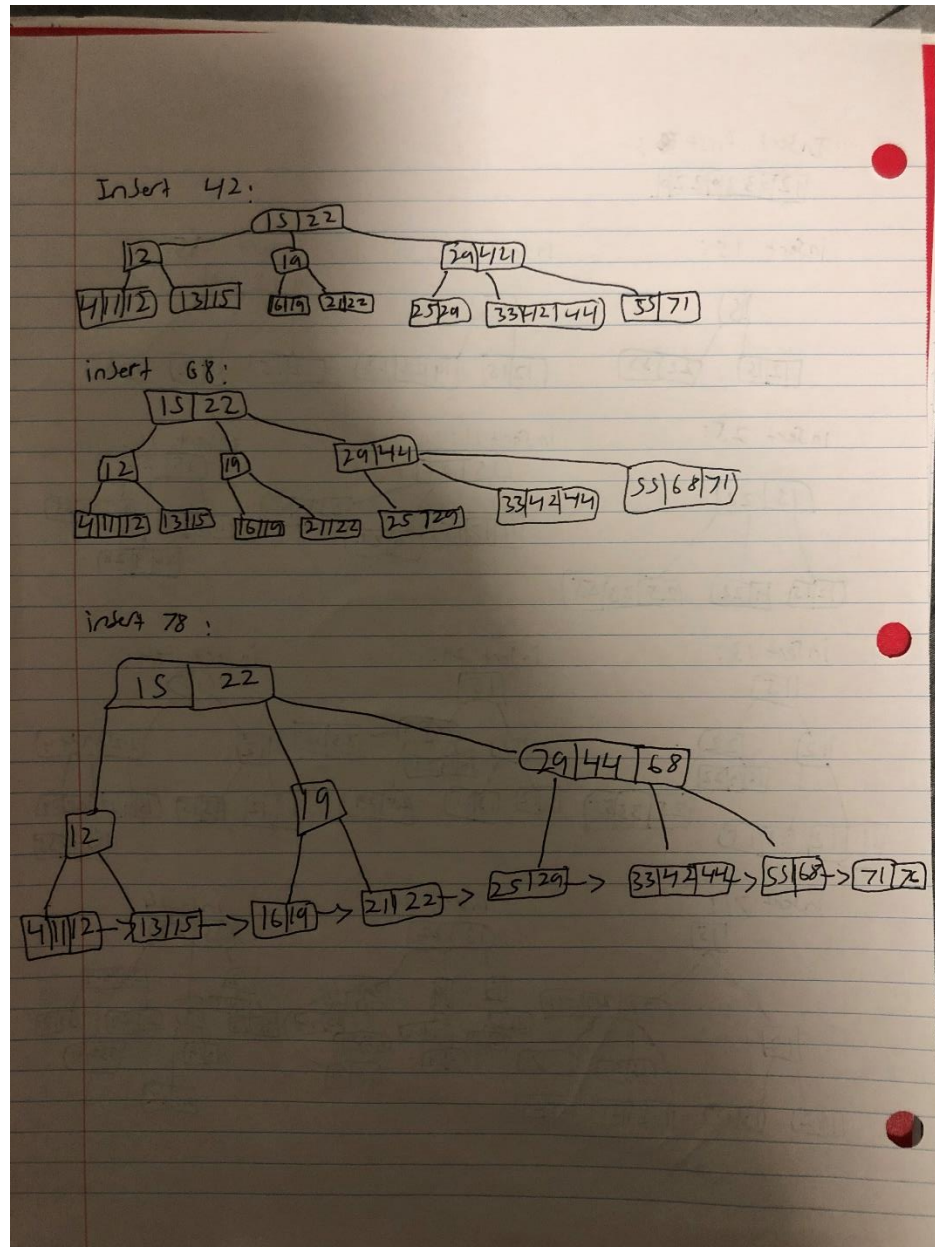


insert 21:



insert 4:





- (ii) (ii) What is the fill ratio of the final B+-tree you created? (Note: we learned 69% is the average fill ratio in class.) Here the fill ratio is defined as (# of filled key values) / (# of maximally allowed key values) regardless each key value resides in index node or sequence set node.

Since there are 25 nodes in the tree and there are 36 possible nodes we get a fill ratio of 69.44 percent.